


More Virtual Memory and JOS memory management


ECE 469, Jan 28

Aravind Machiry



Recap – Page Table & Addr Translation

Mem access #1
CR3[0x20]




	Page Directory Entry
0	Addr PT1
...	Addr PT2
0x20	Addr PT10
0x3ff	Addr PT500

PT 10	Page Table Entry
0	NO MAPPING
0x48	0x10000
0x49	0x11000
0x4a	0x50000

Virtual	Physical
0x8048000	0x10000
0x8049000	0x11000
0x804a000	0x50000

Recap – Page Table & Addr Translation


Mem access #1
CR3[0x20]



	Page Directory Entry	PT 10	Page Table Entry
0	Addr PT1	0	NO MAPPING
...	Addr PT2	0x48	0x10000
0x20	Addr PT10	0x49	0x11000
0x3ff	Addr PT500	0x4a	0x50000


Virtual	Physical
0x8048000	0x10000
0x8049000	0x11000
0x804a000	0x50000

Recap – Page Table & Addr Translation




Virtual	Physical
0x8048000	0x10000
0x8049000	0x11000
0x804a000	0x50000

Recap – Page Table & Addr Translation



Virtual	Physical
0x8048000	0x10000
0x8049000	0x11000
0x804a000	0x50000

Recap – Page Table & Addr Translation




More Virtual Memory

- Multi-level page tables
- Page Table Permissions
- Physical Memory Management


Page Directory / Table

- In x86 (32-bit), CPU uses 2-level page table
- 10-bit directory index
- 10-bit page table index
- 12-bit offset
- **2-level paging**



Size of Page Directory!

- Page Size = 4 KB



$$4096 / 4 = 1024 \text{ entries}$$

$$1024 == 2^{10}$$


10-bit index for PD

Each entry is 4-byte (32 bits)

Size of Page Table!

- Page Size = 4 KB

One page,
4KB



Each entry is 4-byte (32 bits)

$$4096 / 4 = 1024 \text{ entries}$$

$$1024 == 2^{10}$$

10-bit index for PT

Increasing Virtual Address Space


- 32-bit address:
 - 2^{32} == total 4GB
 - We ignore lower 12 bits = 2^{52} total pages
 - Page directory: 4KB.
 - Page table chunk or Second level page table: 4KB.
 - PTE/PDE entry size = 4 bytes
- **All metadata is of size 4KB!!**

Increasing Virtual Address Space

- 64-bit addresses:
 - $2^{64} == 16 \text{ EB} == 16,384 \text{ PB} == 16,777,216 \text{ TB}$
 - We ignore lower 12 bits = 2^{52} total pages
- **How should we handle paging?**


Choices of paging in x64

- 2-level paging
 - 26-bit directory index
 - 26-bit page table index
 - 12-bit offset
 - Each PTE/PDE size = 8 bytes




Choices of paging in x64

- 2-level paging
 - 26-bit directory index
 - 26-bit page table index
 - 12-bit offset
 - Each PTE/PDE size = 8 bytes
- **Page directory size or Page table size:**
 - **$(2^{26}) * 8 = 512 \text{ MB}$**




Choices of paging in x64

- 4-level paging
 - 13-bit directory indexes - 3 level
 - 13-bit page table index
 - 12-bit offset
 - Each PTE/PDE size = 8 bytes
- **Page directories size or Page table size:**
 - **$(2^{13}) * 8 = 64 \text{ KB}$**



Choices of paging in x64


- All metadata should be of 4KB (contiguous memory) chunks for efficient memory usage
- What should be the ideal level of paging for 64-bit address?
 - Each entry (PTE/PDE) size = 8 bytes
 - Number of entries in 4KB = $(4\text{KB}/8) = 512 = (2^9) = 9\text{-bit index}$
 - **64-bit address:**
 - **12-bit offset**
 - **52-bit page number:**
 - $52/9 = 5.777 = 6$
 - **6-level paging**



x86_64: 48-bit address space


- Initial amd64 processors use only 48-bit virtual address space
 - $2^{48} = 256 \text{ TB}$
- Each level of table tree can process 9 bits (512 entries in table)
- We ignore lower 12 bits
 - $48 - 12 = 36$ (total 2^{36} pages)
 - $36 / 9 = 4$ (each table can process 9 bits of address space)
- **We use 4-level page table**

Two-level paging (32-bit)



*) 32 bits aligned to a 4-KByte boundary


Four-level paging (64-bit)



More memory with increasing address space...


- 256 TB might not be enough for big data processing
 - E.g., analyzing online social network of users in Facebook
 - More than 1 billion users, more than 1 trillion edges among users
 - 1 byte per edge = 1TB
- 4 levels, 48 bit
- 5 levels? Yes we can, $48 + 9 = 57$ bits = 128PB
- 6 levels? Maybe, but $57 + 9 = 66$ bits, out of 64 bits...

Virtual Machines - Detour



Virtual Machines - Nested Page Tables (NPE/EPE)

- A case of 48-bit 4-level page table
- Suppose you run Windows
 - Install VMWARE
 - Install a Linux VM
- Host: Windows
- Guest: Linux
- Physical address of Linux
 - This is just a virtual address of Windows




JOS Memory Management

- Handling VA $\leftrightarrow$ PA mapping
- Managing Physical pages

Creating a Virtual Memory Space

- A page directory manages the entire virtual memory space
 - Of a process
- CR3 points to the Page directory, and each PDE entry points to PTs..



Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
 - 0x800000 (RW from user)
- Allocation procedure
 - Check page directory entry (PDE)
 - If not set with PTE_P, **allocate a physical page** for a new page table

PDE 0: EMPTY
PDE 1: EMPTY
PDE 2: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE 1022: EMPTY
PDE 1023: EMPTY

Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
 - 0x800000 (RW from user)
- Allocation procedure
 - Check page directory entry (PDE)
 - If not set with PTE_P, **allocate a physical page** for a new page table


PDE 0: EMPTY
PDE 1: 0x11223
PDE 2: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE 1022: EMPTY
PDE 1023: EMPTY



PTE 0: EMPTY
PTE 1: EMPTY
PTE 2: EMPTY
PTE ...: EMPTY
PTE ...: EMPTY
PTE ...: EMPTY
PTE 1022: EMPTY
PTE 1023: EMPTY

Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
 - 0x800000 (RW from user)
- Allocation procedure
 - Check page directory entry (PDE)
 - If not set with PTE_P, **allocate a physical page** for a new page table
 - Check page table entry (PTE)
 - If not set with PTE_P, **allocate a physical page** to enable access



Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
 - 0x800000 (RW from user)
- Allocation procedure
 - Check page directory entry (PDE)
 - If not set with PTE_P, **allocate a physical page** for a new page table
 - Check page table entry (PTE)
 - If not set with PTE_P, **allocate a physical page** to enable access
- We need to keep track of **'free'** physical pages...

Managing Physical memory in JOS

- Struct PageInfo
 - A metadata type that counts number of 'references' of the page
 - NOT IN USE : pp_ref == 0
- Struct PageInfo * page_free_list
 - A linked list that contains free physical pages
- We will create Struct PageInfo per each Physical Page and then
 - Create a linked list of free pages...

Struct PageInfo in JOS

- A **one-to-one** mapping from a **struct PageInfo** to a physical page
 - An 8 byte struct per each physical memory page
 - If we support 128MB memory, then we will create
 - Total number of physical pages: $128 * 1048576 / 4096 = 32768$
 - Total size = **32768** * **8** = 262,144 = 256KB
- A linked-list for managing free physical pages
 - Starting from `page_free_list->pp_link...`
- `pp_ref`
 - Count references
 - Non-zero – in-use
 - Zero – free

```
struct PageInfo {
    // Next page on the free list.
    struct PageInfo *pp_link;

    // pp_ref is the count of pointers (usually in page table entries)
    // to this page, for pages allocated using page_alloc.
    // Pages allocated at boot time using pmap.c's
    // boot_alloc do not have valid reference count fields.


    uint16_t pp_ref;
};
```

Struct PageInfo

```
struct PageInfo *pages;    // Physical page state array
```

Struct PageInfo * pages (array)

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	



Struct PageInfo (128MB)


$128 * 1048576 / 4096 = 32768$ Pages

8 byte per each entry = $32K * 8 = 256KB$

Struct PageInfo * pages (array)

idx	pp_ref	pp_link
32K	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

Physical memory (128MB)



We can put this array into our physical memory


Free Physical Memory (init)

In kern/pmap.c, boot_alloc

```
static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did *not* assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }
}
```

nextfree will point to the end of the kernel code/data




Free Physical Memory (init)

In kern/pmap.c, boot_alloc

```
static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did *not* assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }
}
```

nextfree will point to the end of the kernel code/data



Allocating struct PageInfo

```
// These variables are set in mem_init()
pde_t *kern_pgdir;    // Kernel's initial page directory
struct PageInfo *pages; // Physical page state array
static struct PageInfo *page_free_list; // Free list of physical pages
```

```
////////////////////////////////////
// Allocate an array of npages 'struct PageInfo's and store it in 'pages'.
// The kernel uses this array to keep track of physical pages: for
// each physical page, there is a corresponding struct PageInfo in this
// array. 'npages' is the number of physical pages in memory. Use memset
// to initialize all fields of each struct PageInfo to 0.
// Your code goes here:
```

```
pages =
boot_alloc(npages * sizeof(struct PageInfo));
```

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

Physical page N

Physical page 2

Physical page 1

Physical page 0

nextfree

end

0x100000

0x7fff000

Free
Physical
Memory

struct PageInfo * pages

Kernel Code


0x2000
0x1000
0x0000

Physical memory

free pages!

- in `page_init()`

```
// The example code here marks all physical pages as free.  
// However this is not truly the case. What memory is free?  
// 1) Mark physical page 0 as in use.  
// This way we preserve the real-mode IDT and BIOS structures  
// in case we ever need them. (Currently we don't, but...)  
// 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)  
// is free.  
// 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must  
// never be allocated.  
// 4) Then extended memory [EXTPHYSMEM, ...).  
// Some of it is in use, some is free. Where is the kernel  
// in physical memory? Which pages are already in use for  
// page tables and other data structures?  
//  
// Change the code to reflect this.  
// NB: DO NOT actually touch the physical memory corresponding to  
// free pages!
```




free pages!

- in `page_init()`


```
// The example code here marks all physical pages as free.
// However this is not truly the case. What memory is free?
// 1) Mark physical page 0 as in use.
//    This way we preserve the real-mode IDT and BIOS structures
//    in case we ever need them. (Currently we don't, but...)
// 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)
//    is free.
// 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must
//    never be allocated.
// 4) Then extended memory [EXTPHYSMEM, ...).
//    Some of it is in use, some is free. Where is the kernel
//    in physical memory? Which pages are already in use for
//    page tables and other data structures?
//
// Change the code to reflect this.
// NB: DO NOT actually touch the physical memory corresponding to
// free pages!
```

36



free pages!


- Page 0 is in-use
- Pages in [IOPHYSMEM ~ EXTPHYSMEM] are in-use
- Pages for the kernel code are in-use
- Pages for struct PageInfo *pages are in-use
- How can you point this?
 - pages + npages ?
 - boot_alloc(0)?



free pages!

- Page 0 is in-use
- Pages in [IOPHYSMEM ~ EXTPHYSMEM] are in-use
- Pages for the kernel code are in-use
- Pages for struct PageInfo *pages are in-use
- How can you point this?
 - pages + npages ?
 - boot_alloc(0)?

**boot_alloc(0) is
better...**



Reference counting: Knowing when a page is free

- A typical mechanism for tracking free memory blocks
- Mechanism
 - Count up the value (`pp_ref++`) if the page is referenced by others (in use!)
 - Count down the value (`pp_ref--`) if not used for one of usages anymore
 - Free if `pp_ref == 0`
- In C++, `shared_ptr<T>`
 - When a pointer is assigned to a variable, count up!
 - When the variable no longer uses the variable, count down!
 - Free the memory when the count become 0

Reference counting with pp_ref

- For in-use memory
 - Set `pp_ref = 1`
- For not-in-use memory
 - Invariant: `pp_ref == 0`
 - **Must be linked with `pages_free_list`**
- When assigning the page to a virtual address
 - `pp_ref++`
- When releasing the page from a virtual address
 - `pp_ref--`

Allocating struct PageInfo

```
// These variables are set in mem_init()
pde_t *kern_pgdir;    // Kernel's initial page directory
struct PageInfo *pages; // Physical page state array
static struct PageInfo *page_free_list; // Free list of physical pages
```

```
////////////////////////////////////
// Allocate an array of npages 'struct PageInfo's and store it in 'pages'.
// The kernel uses this array to keep track of physical pages: for
// each physical page, there is a corresponding struct PageInfo in this
// array. 'npages' is the number of physical pages in memory. Use memset
// to initialize all fields of each struct PageInfo to 0.
// Your code goes here:
```

```
pages =
boot_alloc(npages * sizeof(struct PageInfo));
```

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

Physical page N

Physical page 2

Physical page 1

Physical page 0

nextfree

end

0x100000

0x7fff000

Free
Physical
Memory

struct PageInfo * pages

Kernel Code

0x2000
0x1000
0x0000

Physical memory

Linked list of free pages

- Start with NULL at the head
 - `page_free_list = NULL;`
- After set `pp_ref` of all pages, do something like the following..

This will build a linked list...

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

page_free_list → NULL

Building free list


```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

page_free_list → NULL

Free page 0


Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    → if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



Free page 1

Building free list


```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



Free page 1


Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```




Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        ➡ page_free_list = &pages[i];  
    }  
}
```



Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```




Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



page2pa(struct PageInfo *pp)

- Changes a pointer to **struct PageInfo** to a physical address
- $idx = (pp - pages)$
 - Gets the index of pp in pages
 - E.g., $\&pages[idx] == pp$
- idx here is a physical page number



	idx	pp_ref	pp_link
pp →	4	0	
	3	0	
	2	0	
	1	0	
pages →	0	0	

```
static inline physaddr_t
page2pa(struct PageInfo *pp)
{
    return (pp - pages) << PGSHIFT;
}
```

$pp - pages = 4$
 $0x4000$ □ physical page
address!

pa2page(physaddr_t pa)

- PGNUM(pa)
 - Returns page number
- &pages[PGNUM(pa)]
 - Returns struct PageInfo * of that pa..

```
static inline struct PageInfo*  
pa2page(physaddr_t pa)  
{  
    if (PGNUM(pa) >= npages)  
        panic("pa2page called with invalid pa");  
    return &pages[PGNUM(pa)];  
}
```

idx	pp_ref	pp_link
4	0	
3	0	
2	0	
1	0	
0	0	

Sample Qs

- Which one of the following is not a job that JOS Bootloader does?
 - A. Enable protected mode
 - B. Enable paging
 - C. Load kernel image from disk
 - D. Enable A20

Sample Qs

- Which one of the following is not a job that JOS Bootloader does?
 - A. Enable protected mode
 - B. Enable paging (is done in kernel, in kern/entry.S)
 - C. Load kernel image from disk
 - D. Enable A20

Sample Qs

- In the x86 real mode, which address the following segment:offset pair points to?
- 0x8000:0x3131
 - A. 0xb131
 - B. 0x3131
 - C. 0x83131
 - D. 0x103131
 - E. 0x11131

Sample Qs

- In the x86 real mode, which address the following segment:offset pair points to?
- 0x8000:0x3131
 - A. 0xb131
 - B. 0x3131
 - C. 0x83131 ($0x8000 * 16 + 0x3131 = 0x80000 + 0x3131 = 0x83131$)
 - D. 0x103131
 - E. 0x11131

Sample Qs

- Which of the following x86 register stores the current privilege level?
 - A. ds
 - B. eip
 - C. ebp
 - D. esp
 - E. cs

Sample Qs

- Which of the following x86 register stores the end of the current stack frame (and moves if the CPU runs push/pop) ?
 - A. ds
 - B. eip
 - C. ebp
 - D. esp
 - E. cs

Sample Qs

- Which of the following x86 register stores the start of the current stack frame (also points to the address that stores previous frame's stack base pointer) ?
 - A. ds
 - B. eip
 - C. ebp
 - D. esp
 - E. cs

Sample Qs

- What kind of benefit can we enjoy by enabling virtual memory?
- Choose all (no partial credits)
 - A. Performs faster execution than when using physical memory
 - B. Suffers less memory fragmentation than when using physical memory
 - C. Provides a better isolation / protection than when using physical memory
 - D. Provides memory transparency
 - E. Enables virtual reality