



(Even) More Virtual Memory

ECE 469, Jan 30

Aravind Machiry



A gap among Architecture, Compiler and OS courses



Example



Main.c:

```
extern float sin( );
main( )
{
    static float x, val;

    printf("Type number: ");
    scanf("%f", &x);
    val = sin(x);
    printf("Sine is %f", val);
}
```

Math.c:

```
float sin(float x)
{
    static float temp1, temp2, result;

    – Calculate Sine –

    return result;
}
```



Example (cont)

- Main.c uses externally defined `sin()` and C library function calls
 - `printf()`
 - `scanf()`
- How does this program get compiled and linked?

Compiler



- Compiler: generates object file
 - Information is incomplete
 - Each file may refer to symbols defined in other files

Components of Object File



- Header
- Two segments
 - Code segment and data segment
 - OS adds empty heap/stack segment while loading
- Size and address of each segment
 - Address of a segment is the address where the segment begins.

Components of Object File (cont)




- Symbol table
 - Information about stuff defined in this module
 - Used for getting from the name of a thing (subroutine/variable) to the thing itself
- Relocation information
 - Information about addresses in this module linker should fix
 - External references (e.g. lib call)
 - Internal references (e.g. absolute jumps)
- Additional information for debugger

What could the compiler not do?




- Compiler does not know final memory layout
 - It assumes everything in .o starts at address zero
 - For each .o file, compiler puts information in the symbol table to tell the linker how to rearrange outside references safely/efficiently
 - For exported functions, absolute jumps, etc

Compiler: main.c



Compiler: math.c



Linker functionality



- Three functions of a linker
 - Collect all the pieces of a program
 - Figure out new memory organization
 - Combine like segments
 - Does the ordering matter? (spatial locality for cache)
 - Touch-up addresses
- The result is a runnable object file (e.g. a.out)

Linker – a closer look



- Linker can shuffle segments around at will, but cannot rearrange information within a segment

Linker requires at least two passes



- Pass 1: decide how to arrange memory
- Pass 2: address touch-up

Pass 1 – Segment Relocation



- Pass 1 assigns input segment locations to fill-up output segments
 - Read and adjust symbol table information
 - Read relocation info to see what additional stuff from libraries is required

Pass 2 – Address translation




- In pass 2, linker reads segment and relocation information from files, fixes up addresses, and writes a new object file
- Relocation information is crucial for this part



Putting It Together

- Pass 1:
 - Read symbol table, relocation table
 - Rearrange segments, adjust symbol table
- Pass 2:
 - Read segments and relocation information
 - Touch-up addresses
 - Write new object file

Linker





Dynamic linking

- Static linking – each lib copied into each binary
- Dynamic linking:
 - Instead of system call wrapper code, a stub that finds lib code in memory, or loads it if it is not present
- Pros:
 - all procs can share copy (shared libraries)
 - Standard C library
 - live updates



Dynamic loading

- Program can call dynamic linker via
 - `dlopen()`
 - library is loaded at running time
- Pros:
 - More flexibility -- A running program can
 - create a new program
 - invoke the compiler
 - invoke the linker
 - load it!

Memory Usage Classification



- Memory required by a program can be used in various ways
- Some possible classifications
 - Role in programming language
 - Changeability
 - Address vs. data
 - Binding time

Role in Programming Language



- Instructions
 - Specify the operations to be performed on the operands
- Variables
 - Store the information that changes as program runs
- Constants
 - Used as operands but never change

Changeability



- Read-only
 - Example: code, constants
- Read and write
 - Example: Variables

Address vs. Data



- Need to distinguish between addresses and data
- Why?
 - Addresses need to be modified if the memory is re-arranged

Binding Time



- When is the space allocated?
 - Compile-time, link-time, or load-time
 - Static: arrangement determined once and for all
 - Dynamic: arrangement cannot be determined until runtime, and may change
 - `malloc( )`, `free( )`

Classification – summary



- Classifications overlap
 - Variables may be static or dynamic
 - Code may be read-only or read and write
 - Read-only: Solaris
 - Read and write: DOS
- So what is this all about?
- What does memory look like when a process is running?



Memory Layout

- Memory divided into segments
 - Code (called text in Unix terminology)
 - Data
 - Stack
- Why different segments?
 - To enforce classification
 - e.g. code and data treated differently at hardware level

The big picture




- a.out needs address space for
 - text seg, data seg, and (hypothetical) heap, stack
- A running process needs phy. memory for
 - text seg, data seg, heap, stack
- But no way of knowing where in phy mem at
 - Programming time, compile time, linking time
- Best way out?
 - Make agreement to divide responsibility
 - Assume address starts at 0 at prog/compile/link time
 - OS needs to work hard at loading/runing time

Big picture (cont)



- OS deals with physical memory
 - Loading
 - Sharing physical memory between processes
 - Dynamic memory allocation


Connecting the dots




Process memory map




Linux User/Kernel
Memory Split




Windows, default
memory split



Windows booted
with /3GB switch




JOS




On ubuntu (check kernel map): `sudo cat /proc/iomem`

Easier context-switch



Loading



Dynamic memory allocation during program execution





- Stack: for procedure calls
- Heap: for malloc()
- Both dynamically growing/shrinking
- Assumption for now:
 - Heap and stack are fixed size
 - OS has to worry about loading 4 segments per process:
 - Text
 - Data
 - Heap
 - stack

Process Virtual Memory Layout



- OS allocates a separate virtual memory space to each process
- Transparency
 - Do not have to worry about a system's memory usage status
- Isolation
 - Others can't access my virtual memory space





Why kernel is mapped in all user processes?

Easier context-switch

Memory Maps on x64 machine



```
machiry@machiry-home:~$ cat /proc/self/maps | tail
7fe7bcb0a000-7fe7bcb0b000 r--p 00000000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb0b000-7fe7bcb2e000 r-xp 00001000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb2e000-7fe7bcb36000 r--p 00024000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb37000-7fe7bcb38000 r--p 0002c000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb38000-7fe7bcb39000 rw-p 0002d000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb39000-7fe7bcb3a000 rw-p 00000000 00:00 0
7ffda7458000-7ffda7479000 rw-p 00000000 00:00 0 [stack]
7ffda7584000-7ffda7588000 r--p 00000000 00:00 0 [vvar]
7ffda7588000-7ffda758a000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```

Memory Maps on x64 machine



```
machiry@machiry-home:~$ cat /proc/self/maps | tail
7fe7bcb0a000-7fe7bcb0b000 r--p 00000000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb0b000-7fe7bcb2e000 r-xp 00001000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb2e000-7fe7bcb36000 r--p 00024000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb37000-7fe7bcb38000 r--p 0002c000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb38000-7fe7bcb39000 rw-p 0002d000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7fe7bcb39000-7fe7bcb3a000 rw-p 00000000 00:00 0
7ffda7458000-7ffda7479000 rw-p 00000000 00:00 0 [stack]
7ffda7584000-7ffda7588000 r--p 00000000 00:00 0 [vvar]
7ffda7588000-7ffda758a000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```

```
machiry@machiry-home:~$ cat /proc/5668/maps | tail
7f300561e000-7f3005641000 r-xp 00001000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f3005641000-7f3005649000 r--p 00024000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f3005649000-7f300564a000 r--s 00000000 00:36 204 /run/user/1000/dconf/user
7f300564a000-7f300564b000 r--p 0002c000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f300564b000-7f300564c000 rw-p 0002d000 103:02 35785028 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f300564c000-7f300564d000 rw-p 00000000 00:00 0
7ffc023b3000-7ffc023d4000 rw-p 00000000 00:00 0 [stack]
7ffc023ea000-7ffc023ee000 r--p 00000000 00:00 0 [vvar]
7ffc023ee000-7ffc023f0000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```

How does OS ensure a user process does not access kernel memory?




- OS needs to ensure that a user process cannot access (read/write) kernel (or OS memory)?
 - Why?
 - Hint: Security!
 - Remember: sudo!?

How does OS ensure a user process does not access kernel memory?




- OS needs to ensure that a user process cannot access (read/write) kernel (or OS memory)?
 - Why?
 - Hint: Security!
 - Remember: sudo!?
- Permissions bits in Page directories and Page Tables!!



Page Directory / Table Entry (PDE/PTE)

- Top 20 bits: physical page number
 - Physical page number of a page table (PDE)
 - Physical page number of the requested virtual address (PTE)
- Lower 12 bits: some flags
 - Permission
 - Etc.




Permission Flags



- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable

	Page Table Entry	
0	Addr PT	
0x48	0x10000 << 12 PTE_U PTE_W	Invalid
0x49	0x11000 << 12 PTE_P PTE_W	Kernel, writable
0x4a	0x50000 << 12 PTE_P PTE_U	User, read-only


- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)





When CPU Checks Permission Bits?

- In address translation
- 1. Virtual address
 - Checks permission bits in PDE
- 2. PDE = CR3[PDX]
 - Checks permission bits in PTE
- 3. PTE = PDE[PTX]
 - Checks permission bits in PTE





When CPU Checks Permission Bits?

- A virtual memory address is inaccessible if **PDE disallows the access**
- A virtual memory address is inaccessible if **PTE disallows the access**
- Both **PDE and PTE should allow the access...**

PDE/PTE Permission Examples 0



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W | PTE_U
- PTE: PTE_P | PTE_W | PTE_U
- **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 0



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W | PTE_U
- PTE: PTE_P | PTE_W | PTE_U
- **Can user (ring 3) access it? Is it writable?**
 - Valid, accessible by ring 3, and writable
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 1



- Virtual address 0x01020304
 - PDE: PTE_P | PTE_W | PTE_U
 - PTE: PTE_P | PTE_U
 - **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
 - PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
 - PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 1



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W | PTE_U
- PTE: PTE_P | PTE_U
- **Can user (ring 3) access it? Is it writable?**
 - Valid, accessible by ring 3, but not writable
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 2



- Virtual address 0x01020304
 - PDE: PTE_P | PTE_U
 - PTE: PTE_P | PTE_W | PTE_U
 - **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
 - PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
 - PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 2



- Virtual address 0x01020304
- PDE: PTE_P | PTE_U
- PTE: PTE_P | PTE_W | PTE_U
- **Can user (ring 3) access it? Is it writable?**
 - Valid, accessible by ring 3, but not writable
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 3



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W | PTE_U
- PTE: PTE_P
- **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 3



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W | PTE_U
- PTE: PTE_P
- **Can user (ring 3) access it? Is it writable?**
 - valid, inaccessible by ring3, not writable
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 4



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W
- PTE: PTE_P | PTE_U
- **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 4



- Virtual address 0x01020304
- PDE: PTE_P | PTE_W
- PTE: PTE_P | PTE_U
- **Can user (ring 3) access it? Is it writable?**
 - valid, inaccessible by ring3, not writable
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 5



- Virtual address 0x01020304
 - PDE: PTE_P | PTE_U
 - PTE: PTE_U
 - **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
 - PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
 - PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 5



- Virtual address 0x01020304
- PDE: PTE_P | PTE_U
- PTE: PTE_U
- **Can user (ring 3) access it? Is it writable?**
 - invalid
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 6



- Virtual address 0x01020304
 - PDE: PTE_U
 - PTE: PTE_P | PTE_U
 - **Can user (ring 3) access it? Is it writable?**
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
 - PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
 - PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

PDE/PTE Permission Examples 6




- Virtual address 0x01020304
- PDE: PTE_U
- PTE: PTE_P | PTE_U
- Can user (ring 3) access it? Is it writable?
 - invalid
- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable
- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)



Valid permission bits..

- Kernel: R, User: --
 - PTE_P
- Kernel: R, User: R
 - PTE_P | PTE_U
- Kernel: RW, User: RW
 - PTE_P | PTE_U | PTE_W



Cannot have permissions such as ...

- Kernel: RW, User: R
 - PTE_P | PTE_W | PTE_U -> User RW...
 - PTE_P | PTE_W -> User --
- Kernel: R, User: RW
 - PTE_P | PTE_U | PTE_W -> Kernel RW...
 - PTE_P | PTE_U -> User R...
- Kernel: --, User: RW
 - PTE_P | PTE_U | PTE_W -> Kernel RW...



Flexibility of virtual memory!

- Virtual to physical address mapping is in N-to-1 relation
 - N number of virtual addresses could be mapped to 1 physical address
- E.g., for a physical address 0x100000
 - JOS maps VA 0x100000 to PA 0x100000
 - JOS maps VA 0xf0100000 to PA 0x100000
- Why?
 - EIP before enabling paging: 0x100025
 - EIP after enabling paging: 0x100028
 - Then jumps to 0xf010002f

```
0x00100025 ? mov    %eax,%cr0
0x00100028 ? mov    $0xf010002f,%eax
0x0010002d ? mov    *%eax
0x0010002f ? mov    $0x0,%ebp
```



Sharing a Physical Page!


- Example: Loading of the same program

Process 0, runs /bin/bash,
loads at virt addr
0x35555000

	Page Table Entry
0	...
0x155	0x10303 FLAG

Process 1, runs /bin/bash,
loads at virt addr
0x43132000

	Page Table Entry
0	...
0x132	0x10303 FLAG



2 or more mappings to 0x10303000 is
possible!




Allocating Virtual Memory

- Static allocation is inefficient:
 - Why don't we just allocate entire virtual address space to a process?
 - **Inefficient**: The process may not access entire virtual address space
- Solution: Dynamic, Request based

Dynamic space allocation



- OS allocates space (valid PTE entries in page table) as dictated by the program binary and start running the process.





Dynamic space allocation

- When a process tries to access memory that is not allocated i.e., there is no corresponding valid PTE, then, OS kills the process.
 - E.g., Segmentation Fault!
- A process needs to **explicitly request OS** to allocate additional space (and create valid PTEs).
 - brk **system call** (We will cover this later)

Dynamic space allocation



- We use malloc and free, which actually use brk system call internally
- brk only **allocates virtual memory!** **not physical memory!**




Dynamic space allocation: Example

```
#include<stdio.h>
#include<stdlib.h>
int main() {
    char p;
    int *i = (int *)malloc(16*1024*1024*1024);
    printf("We just requested 16GB of virtual memory!!\n");
    printf("Check memory now, using htop!\n");
    printf("You will be surprised to see your memory usage.\nPress any key to exit.\n");
    scanf("%c", &p);
    return 0;
}
```



What happens when we call malloc?

- Before malloc()?
 - No PTEs




- After malloc()?
 - **PDE/PTE updated but present bit not-set**
- Upon first access?
 - Assign physical page (and page table) and set the valid bit.





Handling Page faults in Kernel

- Accessing unallocated memory: Demand Paging



- Accessing Unassigned Virtual Memory: Segmentation Fault