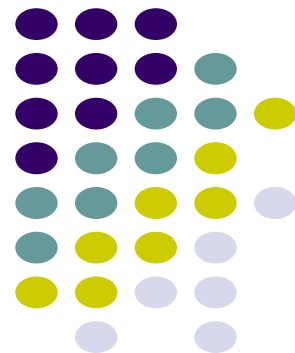


Processes

ECE 469, Jan 30

Aravind Machiry



Navigating Lab #2



- **READ COMMENTS IN THE CODE**

Free Physical Memory (init)

In kern/pmap.c, boot_alloc

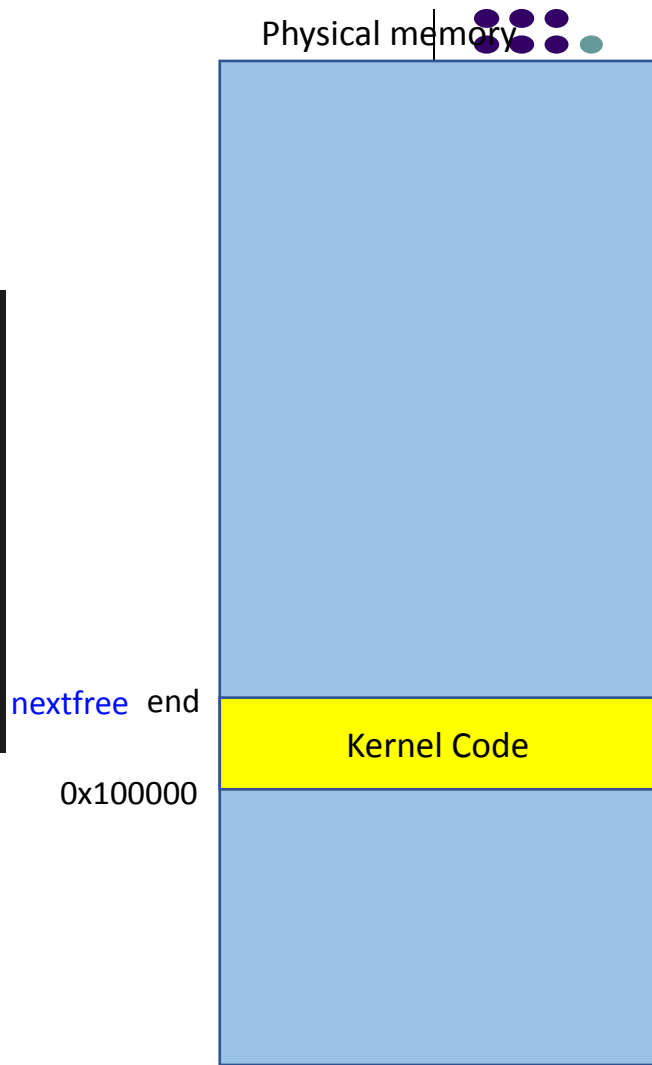
```
static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did *not* assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }
}
```

nextfree will point to the end of the kernel code/data

nextfree is virtual address.

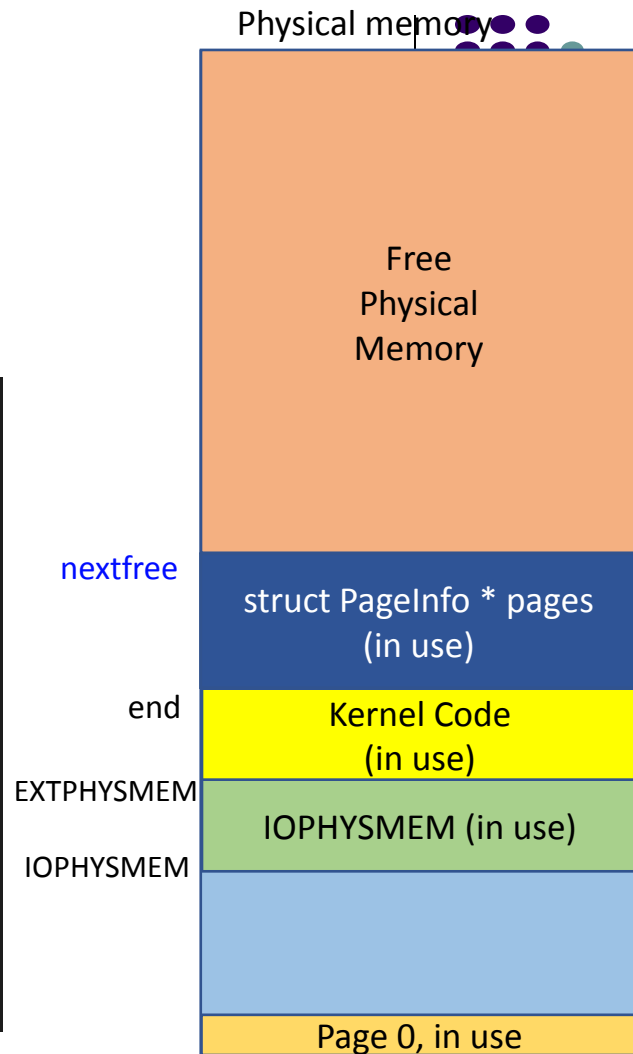
**You should allocate n bytes (rounding to PAGE boundary)
and return the old pointer and update nextfree.**



page_init: free pages!

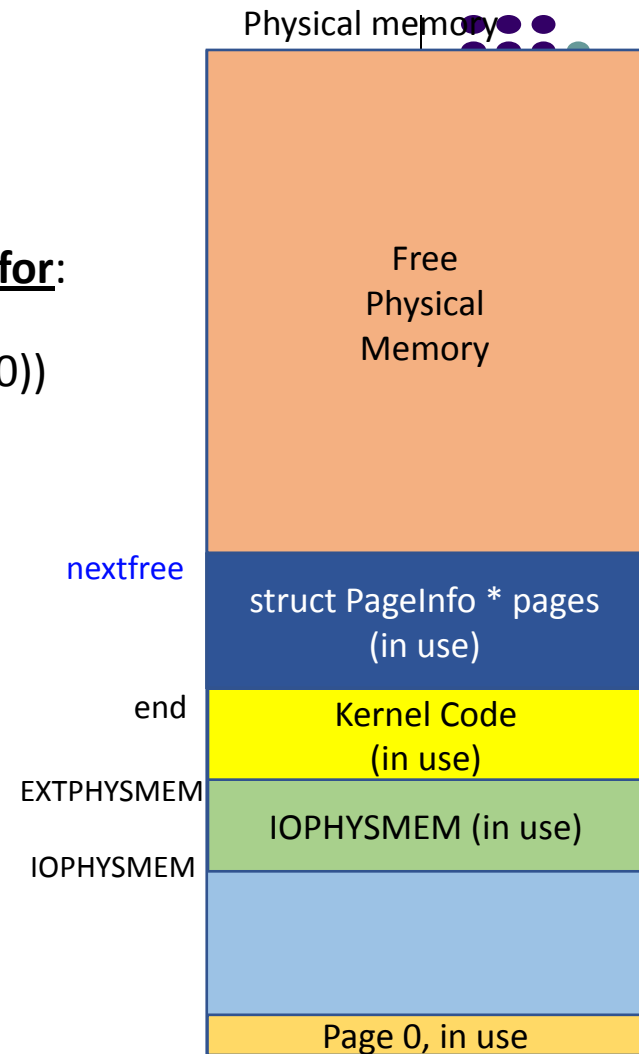
- in page_init()

```
// The example code here marks all physical pages as free.  
// However this is not truly the case. What memory is free?  
// 1) Mark physical page 0 as in use.  
// This way we preserve the real-mode IDT and BIOS structures  
// in case we ever need them. (Currently we don't, but...)  
// 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)  
// is free.  
// 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must  
// never be allocated.  
// 4) Then extended memory [EXTPHYSMEM, ...).  
// Some of it is in use, some is free. Where is the kernel  
// in physical memory? Which pages are already in use for  
// page tables and other data structures?  
//  
// Change the code to reflect this.  
// NB: DO NOT actually touch the physical memory corresponding to  
// free pages!
```



page_init: free pages!

- Iterate through the pages and mark all pages **except for**:
 - Page 0
 - Pages from IOPHYSMEM to nextfree (bootalloc(0))
- Add free pages to the list.




Reuse your code



- **boot_map_region**: Inserts mapping of given VA -> PA of given size (page aligned) with the given permission into a page directory.
- **pgdir_walk**: Gets the page table entry (or creates) corresponding to the given va in the given page directory.
- **page_lookup**: Look up PageInfo corresponding to the given VA.
- Functions can be written by re-using other functions:
 - **page_lookup**:
 - i. Can use pgdir_walk to get the pte*
 - ii. Can get the physical address of the pte*
 - iii. Convert the physical address to PageInfo using pa2page*

Today's Class

- Users, Programs and Process



Users, Programs, Processes



- Users have accounts on the system

Users, Programs, Processes



- Users have accounts on the system
- Users launch programs

Users, Programs, Processes



- Users have accounts on the system
- Users launch programs
 - Can many users launch the same program?
 - Can one user launch many instances of the same program?

Users, Programs, Processes



- Users have accounts on the system
- Users launch programs
 - Can many users launch the same program?
 - Can one user launch many instances of the same program?

□ A process is an “instance” of a program

Program vs. Process



```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Program

```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Process

Code

Data

heap

stack

main

foo

registers

PC

So What Is A Process?



- It is a running instance of a program.
- Program becomes **alive** through a process.
- Any relation between multiple instances of a program?

So What Is A Process?



- It is a running instance of a program.
- Program becomes **alive** through a process.
- Any relation between multiple instances of a program?
 - Ideally, No!
 - How is the separation maintained?

Process needs to communicate with OS



- Access system resources :
 - Network, Memory, Bluetooth, etc.
 - Maintained by OS.
- How does the communication happen between a process and OS?

Quick Detour: System Calls

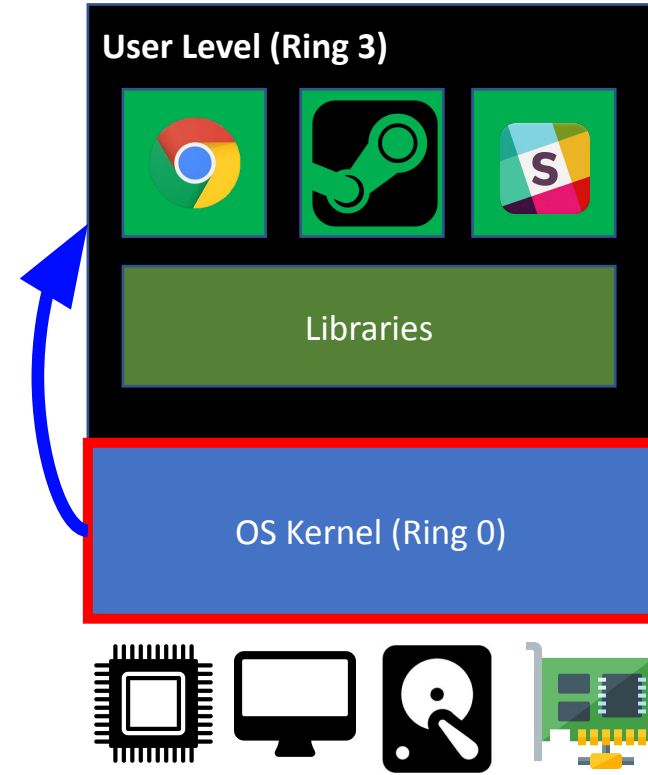



- Interface between a process and the operating system kernel
 - Kernel manages shared resources & exports interface
 - Process requests for access to shared resources
- Generally available as assembly-language instructions:
 - `syscall`

Why can't a process directly execute kernel code? Because Kernel code ...

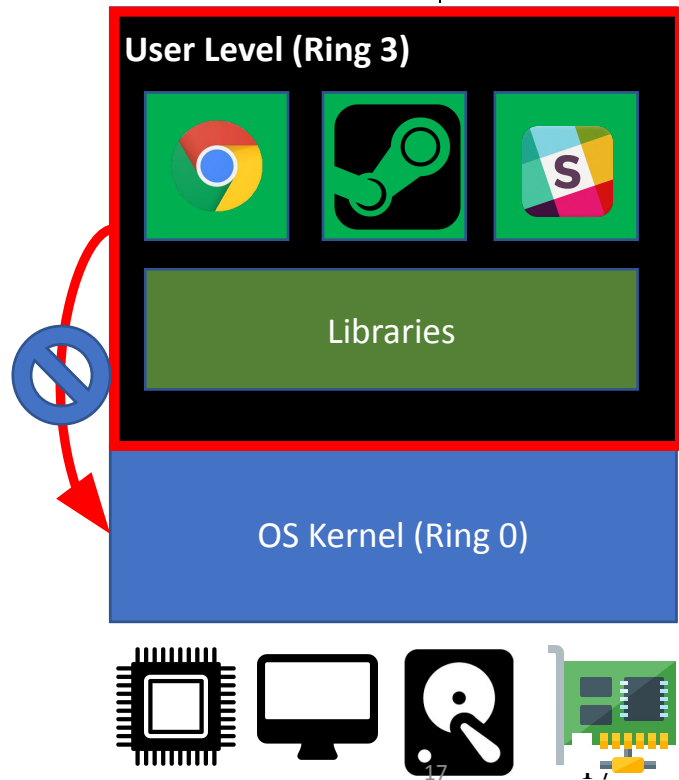



- Runs with the highest privilege level (Ring 0)
- Configures system (devices, memory, etc.)
- Manages hardware resources
 - Disk, memory, network, video, keyboard, etc.
- Manages other jobs
 - Processes and threads
- Serves as trusted computing base (TCB)
 - Set privilege
 - Restrict other jobs from doing something bad..




User mode process runs at Ring 3

- Runs with a restricted privilege (Ring 3)
 - The privilege level for running an application...
- Most of regular applications runs in this level
- Cannot access kernel memory
 - Can only access pages set with PTE_U
- Cannot talk directly to hardware devices
 - Kernel must mediate the access



Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```




Syscall: User/Kernel communication



`printf("ECE469")`

A library call in ring 3



```
int main() {  
    printf("ECE469");  
}
```



Syscall: User/Kernel communication

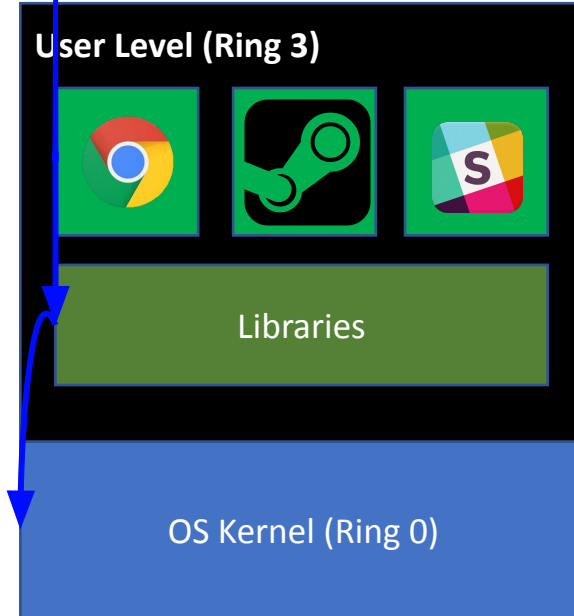


`printf("ECE469")`

A library call in ring 3

`sys_write(1, "ECE469", 6);`

A system call, **From ring 3**



```
int main() {  
    printf("ECE469");  
}
```



Syscall: User/Kernel communication



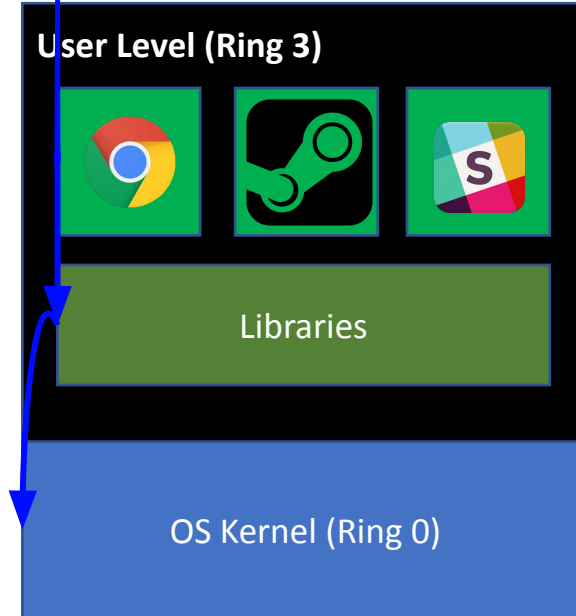
`printf("ECE469")`

A library call in ring 3

`sys_write(1, "ECE469", 6);`

A system call, **From ring 3**

```
int main() {  
    printf("ECE469");  
}
```



Syscall: User/Kernel communication



printf("ECE469")

A library call in ring 3

sys_write(1, "ECE469", 6);


A system call, **From ring 3**

```
int main() {  
    printf("ECE469");  
}
```

Interrupt!, switch from ring3 to ring0

A kernel function

do_sys_write(1, "ECE469", 6)



Syscall: User/Kernel communication



printf("ECE469")

A library call in ring 3

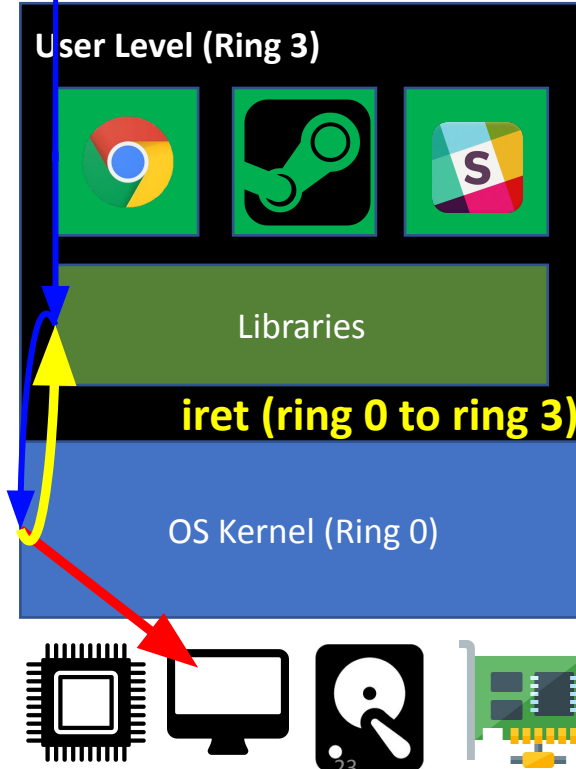
sys_write(1, "ECE469", 6);

A system call, **From ring 3**

Interrupt!, switch from ring3 to ring0

A kernel function

do_sys_write(1, "ECE469", 6)



```
int main() {  
    printf("ECE469");  
}
```


Syscall: User/Kernel communication



`printf("ECE469")`

A library call in ring 3

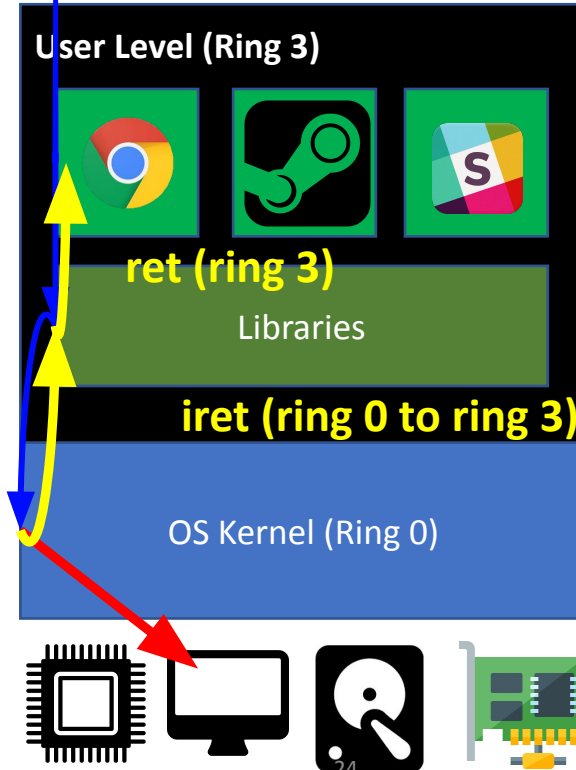
`sys_write(1, "ECE469", 6);`

A system call, **From ring 3**

Interrupt!, switch from ring3 to ring0

A kernel function

`do_sys_write(1, "ECE469", 6)`



```
int main() {  
    printf("ECE469");  
}
```

Let's get back: Process



- Who has the ability to create a process?
 - OS
- Who wants to create a process?
- How can we create a process?

Let's get back: Process



- Who has the ability to create a process?
 - OS
- Who wants to create a process?
 - User/Program
- How can we create a process?

Let's get back: Process




- Who has the ability to create a process?
 - OS
- Who wants to create a process?
 - User/Program
- How can we create a process?
 - System call
 - But, to do a syscall, we need a process!!!


Init can never die.



Process State Transition



Process Information maintained by OS



- Usually Maintained in a structure called **Process Control Block (PCB)**
- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc

Process Identifier



- Every process has an ID – process ID
- Does a program know its process ID?
- When a program is running, how does the process know its ID?

OS Support for Process



- Support to create process
- Support to wait for a process completion
- Support to terminate a process

OS Process API



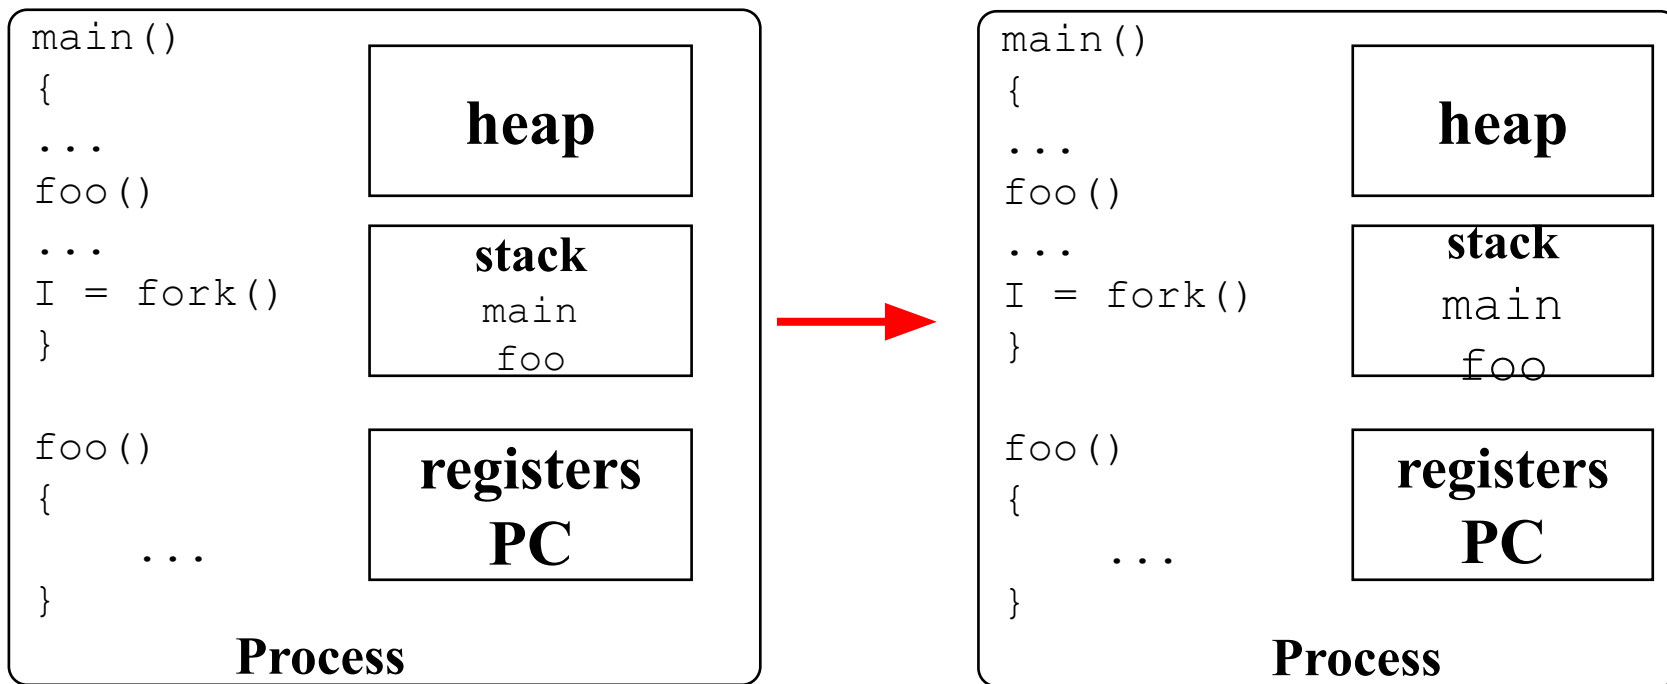
- 4 system calls related to process creation/termination:
 - Process Creation:
 - fork/clone – create a copy of this process
 - exec – replace this process with this program
 - Wait for completion:
 - wait – wait for child process to finish
 - Terminate a process:
 - kill - send a signal (to terminate) a process

fork



fork causes OS creates a copy of the calling process:

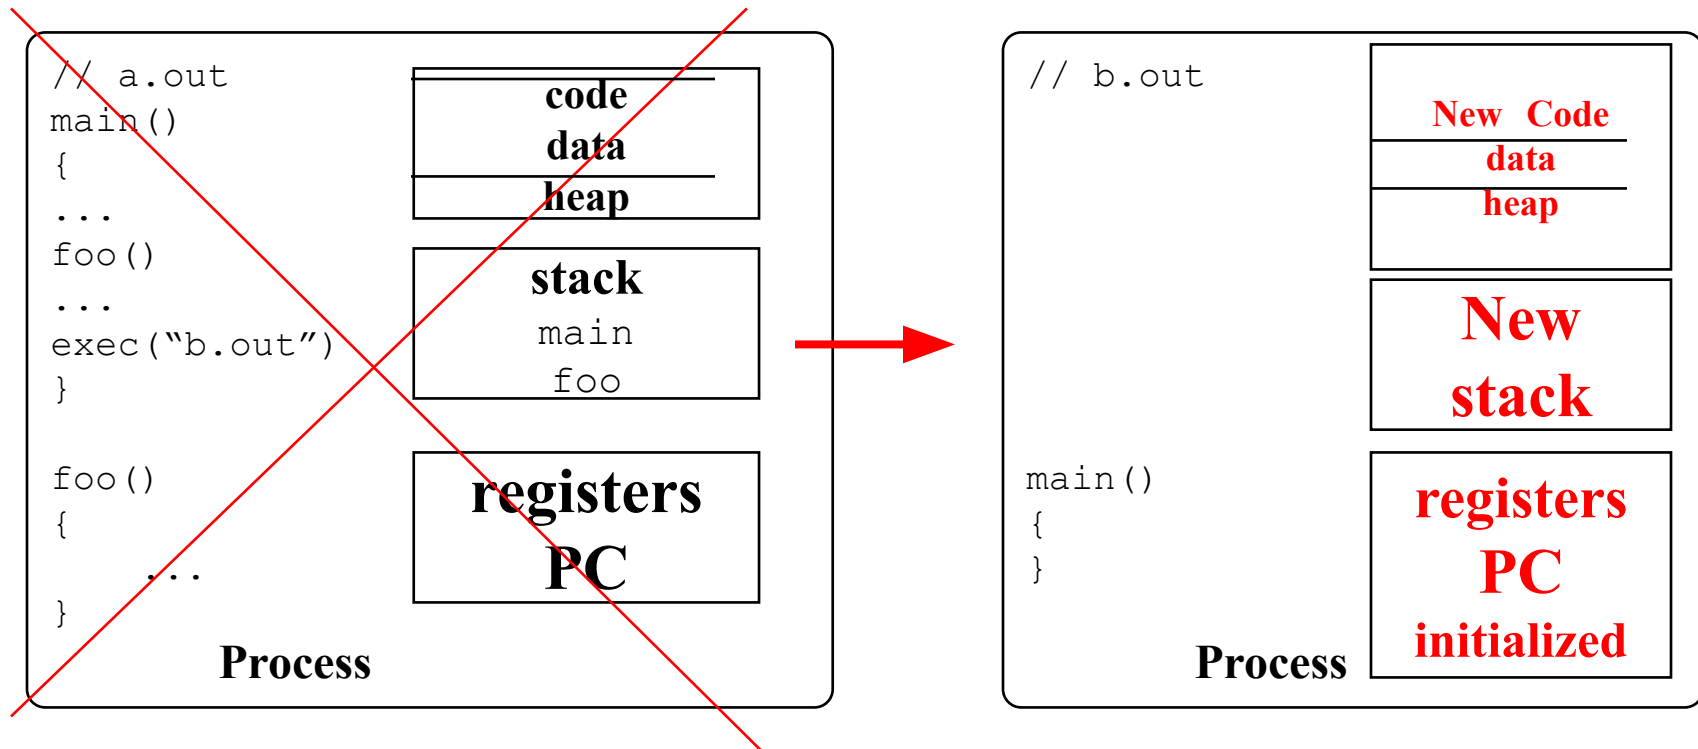
- Why?
- How can we disambiguate between new process and the calling process?



exec



Replaces current process with the content from new program.



exec

Example



wait


wait for a child process to finish



wait



What happens when the parent process dies? what happens to child process?



How our shell works?

- Fork/exec

