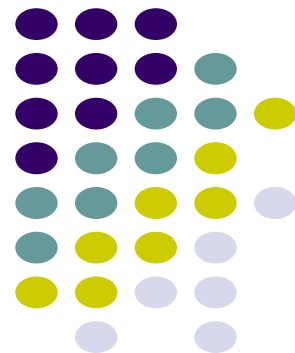


Multi Threading and Synchronization

ECE 469, Feb 22

Aravind Machiry



Web Server Example



- How does a web server handle 1 request?
- A web server needs to handle many concurrent requests
- Solution 1:
 - Have the parent process fork as many processes as needed
 - Processes communicate with each other via inter-process communication

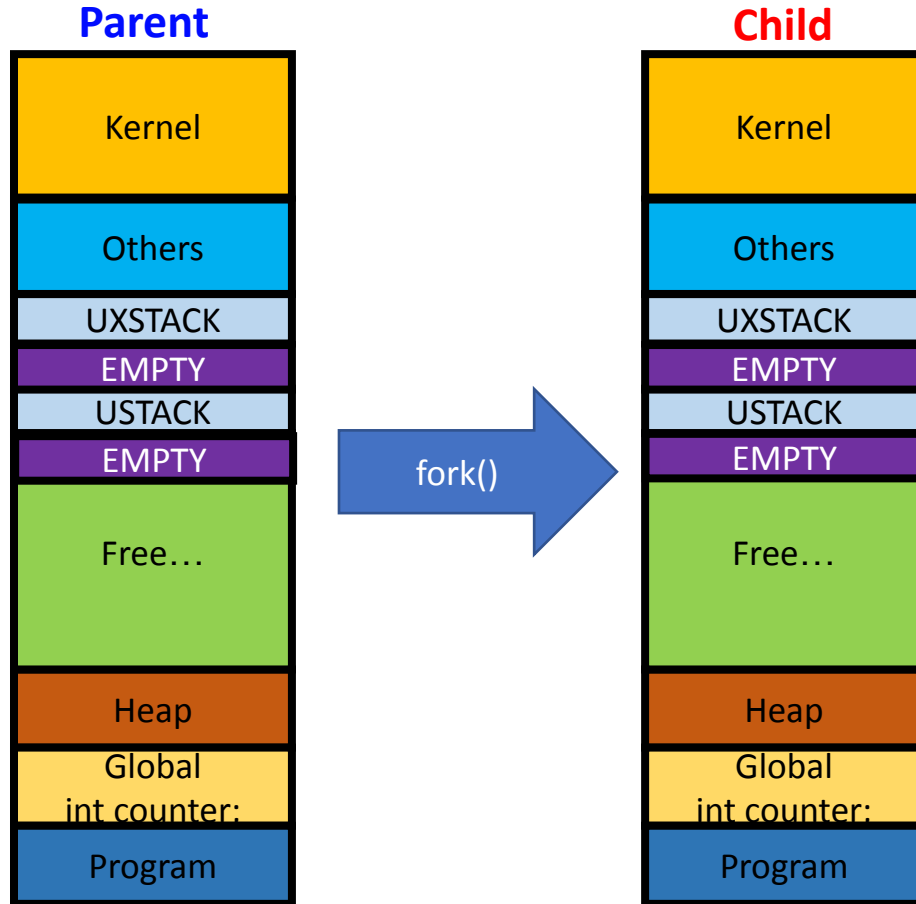
Multiple Processes



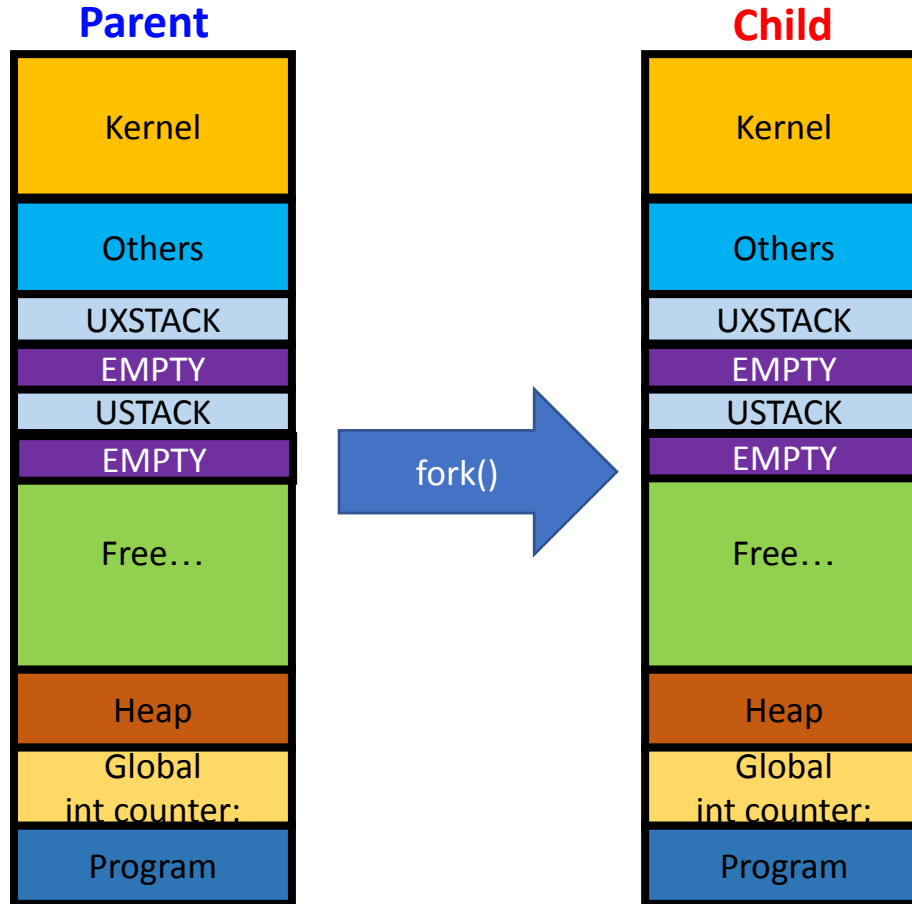
Parent



Multiple Processes



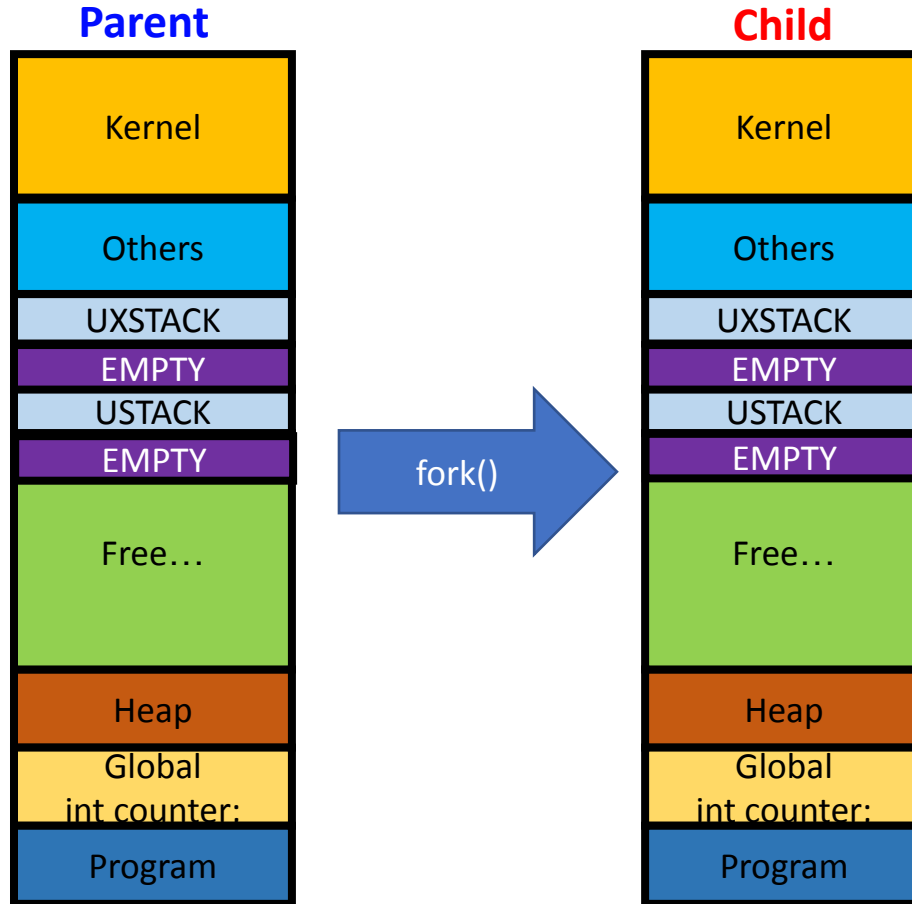
Multiple Processes



Fork() creates new process by copying memory space

Process creates a new PRIVATE memory space

Multiple Processes



Fork() creates new process by copying memory space

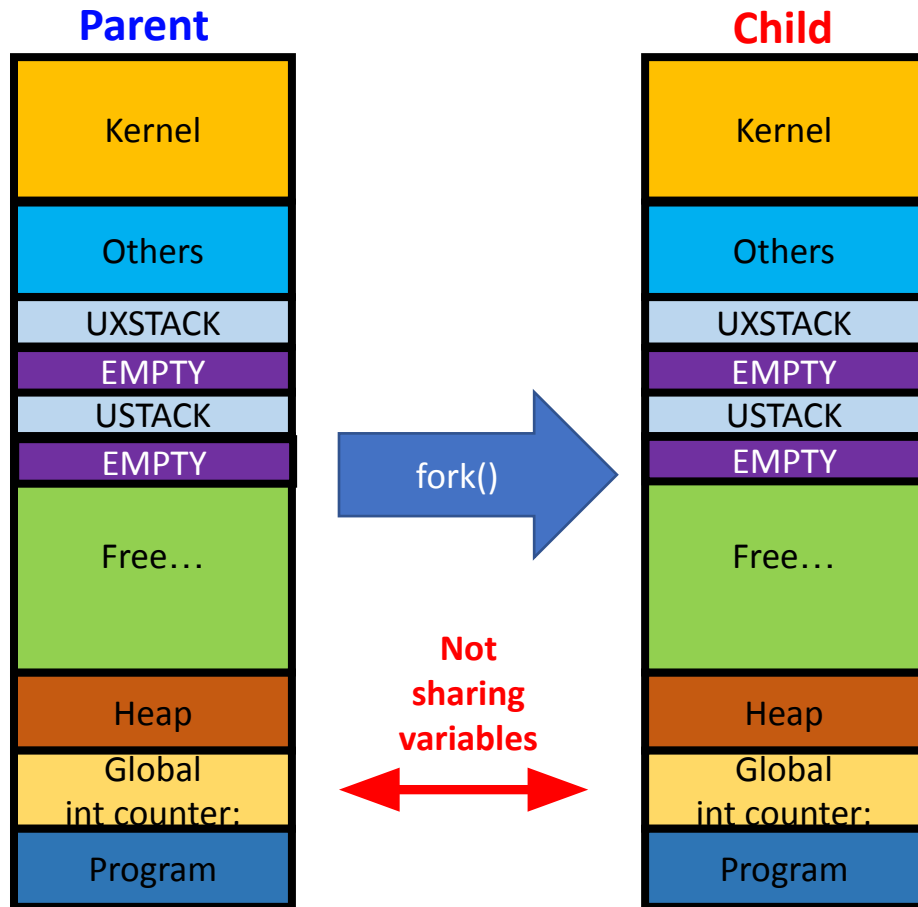
```
#include <stdio.h>
#include <unistd.h>

int counter;
volatile int value = 1;

void countup() {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
}

int main() {
    pid_t pid = fork();
    countup();
    printf("%s: %d\n", pid ? "Parent" : "Child", counter);
}
```

Multiple Processes



Fork() creates new process by copying memory space

```
#include <stdio.h>
#include <unistd.h>

int counter;
volatile int value = 1;

void countup() {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
}

int main() {
    pid_t pid = fork();
    countup();
    printf("%s: %d\n", pid ? "Parent" : "Child", counter);
}
```

Parent: 1000000
Child: 1000000

How do Process communicate?



- At process creation time
 - Parents get one chance to pass everything at `fork()`
- OS provides generic mechanisms to communicate
 - Shared Memory: multiple processes can read/write same physical portion of memory; implicit channel
 - System call to declare shared region
 - No OS mediation required once memory is mapped
 - Message Passing: explicit communication channel provided through `send()/receive()` system calls
 - A system call is required

How do Process communicate?



- IPC is, in general, expensive due to the need for system calls
 - Although many OSes have various forms of lightweight IPC

The Soul of a Process



- But all the processes in the web-server are **cooperating!**
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)
- What don't they share?
 - Each has its own execution state: PC, SP, and registers

The Soul of a Process



- Key idea: Why don't we **separate the concept of a process from its execution state**?
 - Process: address space, privileges, resources, etc.
 - Execution state: PC, SP, registers
- Exec state also called thread of control, or thread

Threads




- **Separate** the concepts of a “thread of control” (PC, SP, registers) from the rest of the process (address space, resources, accounting, etc.)
- Modern OSes support two entities:
 - the *task* (process), which defines an address space, a resource container, accounting info
 - the *thread* (lightweight process), which defines a single sequential execution stream within a task (process)

Threads vs. Process



- There can be several threads in a single address space
- Threads are the unit of scheduling; tasks are containers (address space, other shared resources) in which threads execute

Single threaded v/s multithreaded



What differs in threads of a process?



- A.K.A User Environment (JOS)
- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
 - Code, data, heap, execution stack
- I/O and file management
 - Communication ports, directories, file descriptors, etc

What differs in threads of a process?



- A.K.A User Environment (JOS)
- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
 - Code, data, heap, execution stack
- I/O and file management
 - Communication ports, directories, file descriptors, etc

Thread Control Block

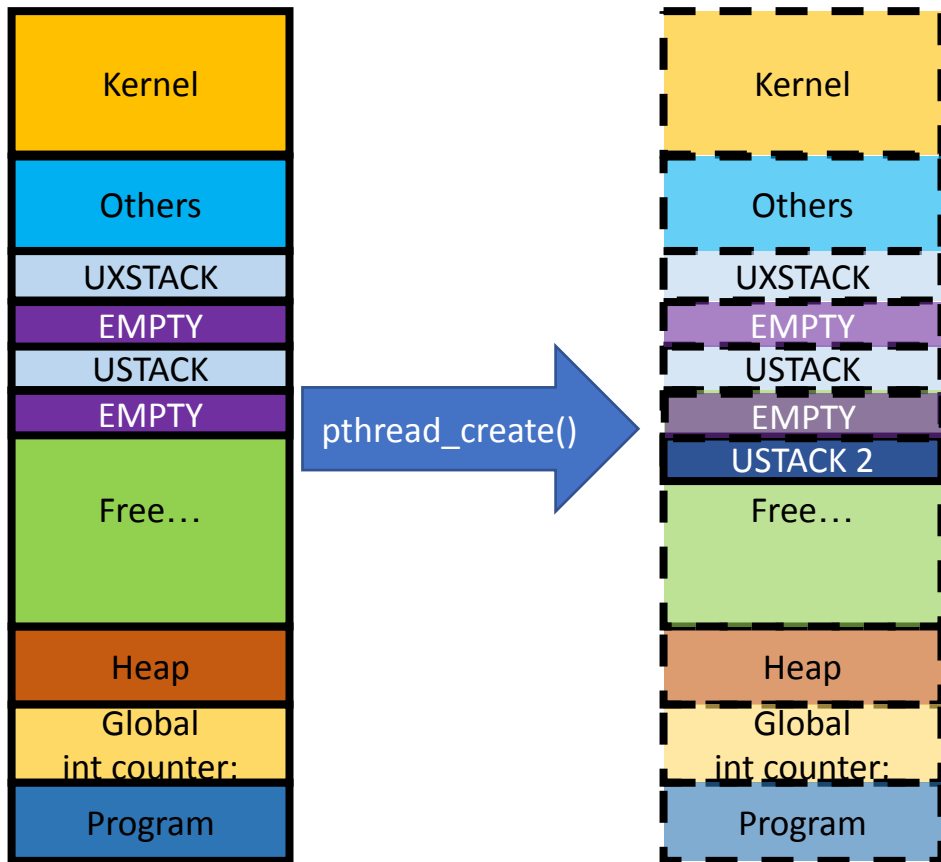


- Shared information
 - Process info: parent process
 - Memory: code/data segments, page table, and stats
 - I/O and file: comm ports, open file descriptors
- Private state
 - State (ready, running and blocked)
 - PC, Registers
 - Execution stack

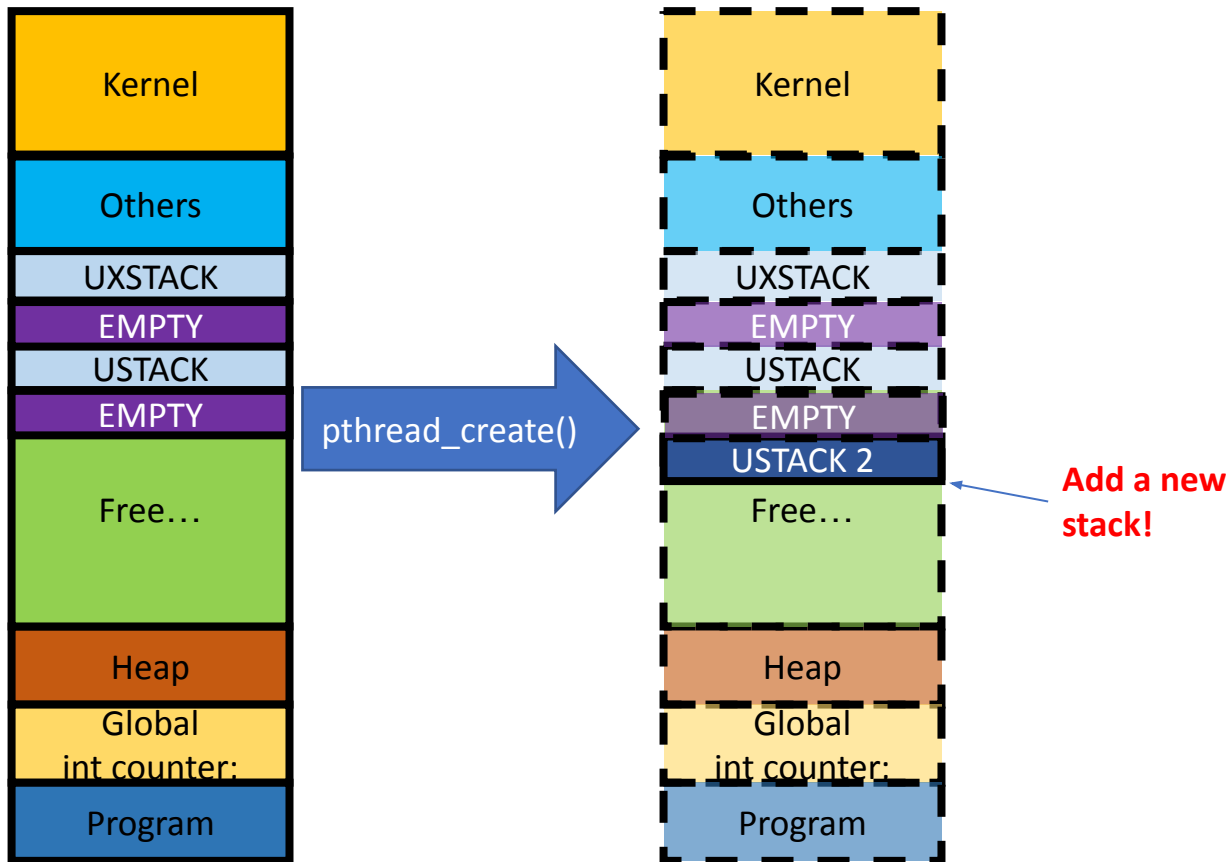
Threads



Threads



Threads



Threads



`pthread_create()`



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

**Add a new
stack!**

Threads



`pthread_create()`



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

Add a new
stack!

Adding
value..

Threads



pthread_create()



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

Add a new
stack!

Adding
value..

The same
variable..

Programming with Threads



- Flexible, but error-prone, since there no protection between threads
 - In C/C++,
 - automatic variables are private to each thread
 - global variables and dynamically allocated memory (malloc) are **shared**
- Need synchronization!

The need for synchronization!



- Cooperating processes may share data via
 - shared address space (code, data, heap) by using threads
 - Files
 - (Sending messages)
- What can happen if processes try to access shared data (address) concurrently?
 - Sharing bank account with sibling:
At 3pm: If (balance > \$10) withdraw \$10
- How hard is the solution?

“Too much milk” Problem



Person A

1. Look in fridge: out of milk
2. Leave for Walmart
5. Arrive at Walmart
6. Buy milk
7. Arrive home

Person B

3. Look in fridge: out of milk
4. Leave for Walmart
8. Arrive at Walmart
9. Buy milk
10. Arrive home

- How to put in a locking mechanism?

Possible Solution 1



Person A

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Person B

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Will this work?



Person A

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Person B

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Will this work?



Person A

```
1.if ( noMilk ) {  
  2.if (noNote) {  
    5.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

Person B

```
3.if ( noMilk ) {  
  4.if (noNote) {  
    6.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

- Process can get context switched after checking milk and note, but before leaving note

Why does this work for humans?



- Human can perform *test* (look for other person & milk) and *set* (leave note) at the same time.

Possible Solution 2



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

Will this work?



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```


Will this work?



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

- We may not have Milk: Both process can leave note and skip buying milk

Possible Solution 3



Process A

```
leave noteA
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
remove noteA
```

Process B

```
leave noteB
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

Will this work?



Process A

```
leave noteA
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
remove noteA
```

Process B

```
leave noteB
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

Works, but complicated!



Process A

leave noteA

```
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
```

remove noteA

Process B

leave noteB

```
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
```

remove noteB

- A's code is **different** from B's
- **busy waiting** is a waste

How can we solve this?





- Root cause: **Data Race**
- A thread's execution result could be inconsistent if other threads intervene its execution...

- counter += value

• edx = value;	mov	0x20087b(%rip),%edx	# 0x201010 <value>
• eax = counter;	mov	0x20087d(%rip),%eax	# 0x201018 <counter>
• eax = edx + eax;	add	%edx,%eax	
	mov	%eax,0x200875(%rip)	# 0x201018 <counter>
• counter = eax;			



Shared variable: No race

- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;




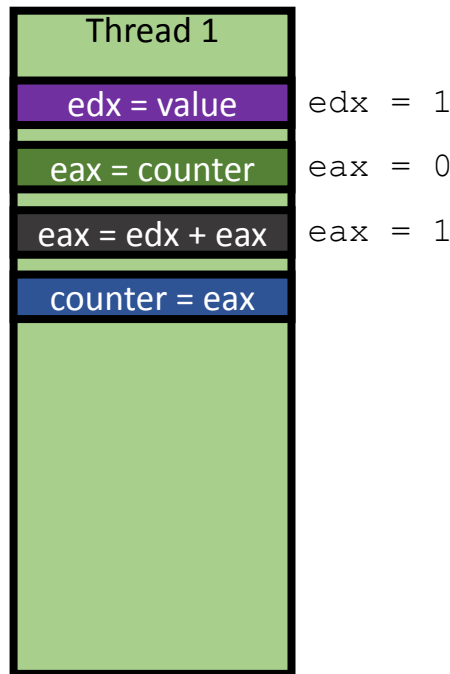
Shared variable: No race

- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;




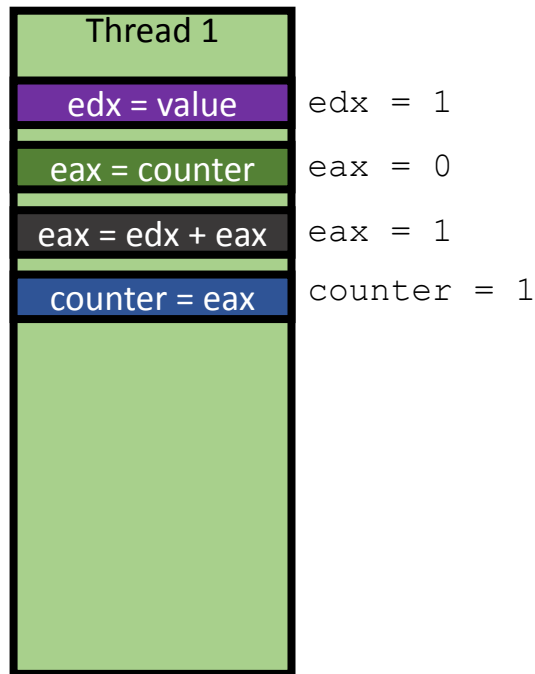
Shared variable: No race

- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: No race



- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: No race


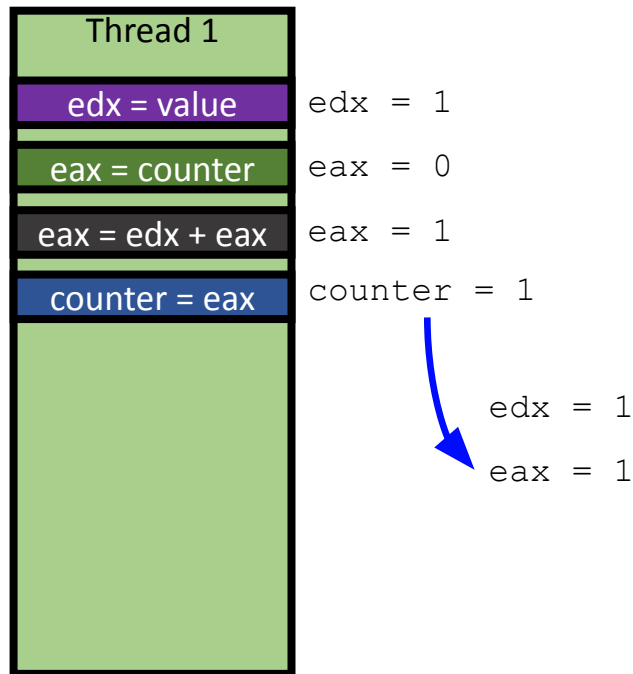


- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;




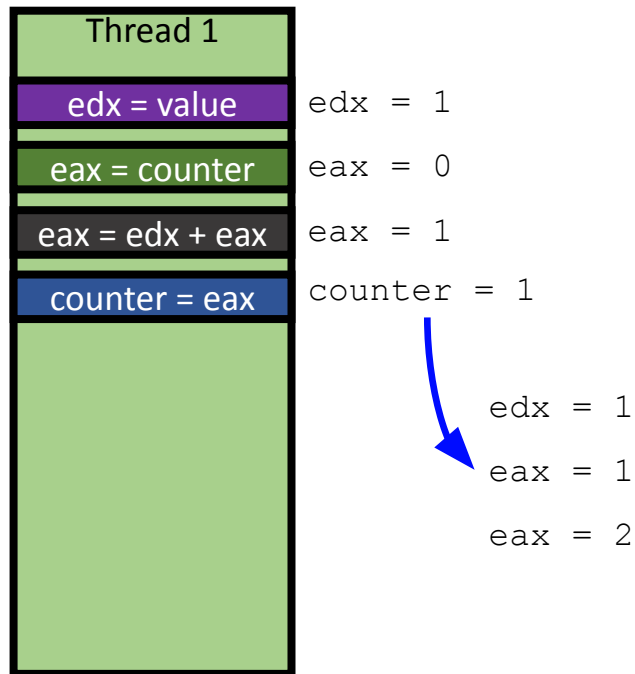
Shared variable: No race

- `counter += value`
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume `counter = 0` at start, and `value = 1`;



Shared variable: No race

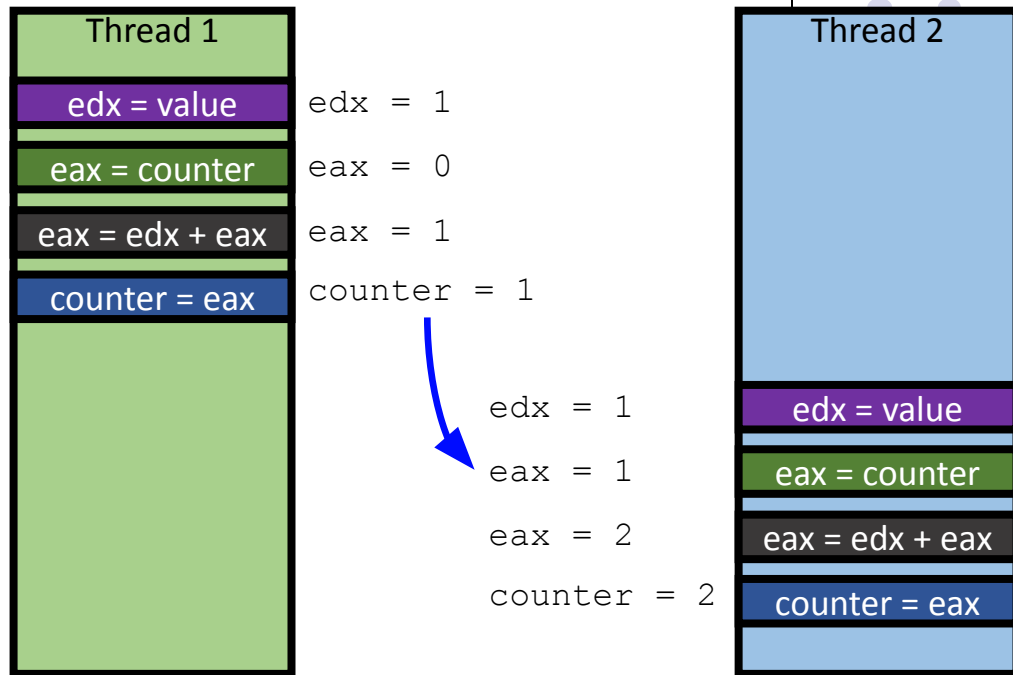
- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: No race



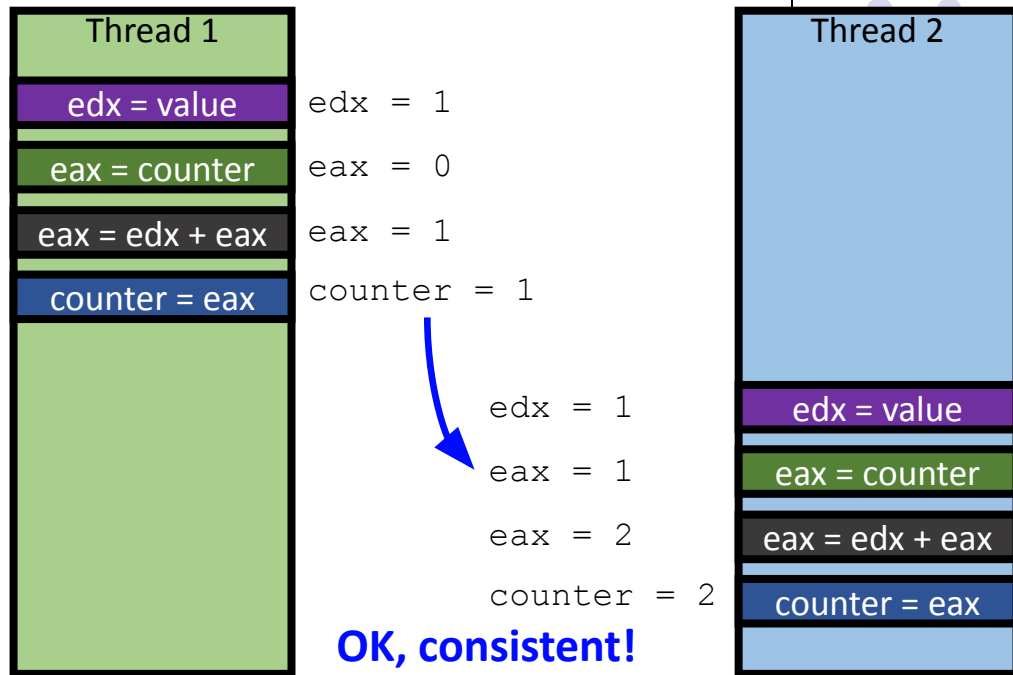
- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: No race





- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: Data race





- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: Data race

- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;

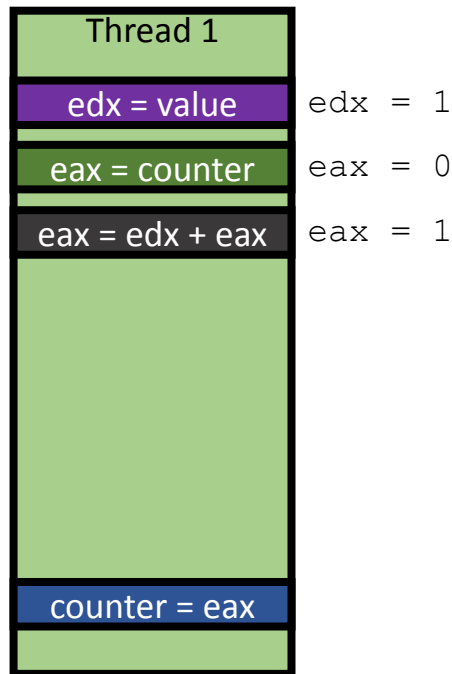


-
- ```
graph TD; T1[Thread 1] --> E1[edx = value]; E1 --> E1R[edx = 1]; E1R --> E2[eax = counter]; E2 --> E2R[eax = 0]; E2R --> E3[eax = edx + eax]; E3 --> E3R[eax = 1]; E3R --> Delay[]; Delay --> E4[counter = eax];
```
- Thread 1
- edx = value      edx = 1
- eax = counter      eax = 0
- eax = edx + eax      eax = 1
- counter = eax




# Shared variable: Data race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;




`edx = 1`




# Shared variable: Data race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



edx = 1


eax = 0




# Shared variable: Data race



- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



edx = 1  
eax = 0  
eax = 1



- 
- The diagram illustrates the execution flow of Thread 1. It consists of a vertical stack of colored rectangular blocks. From top to bottom, the blocks are: a light green block labeled 'Thread 1', a purple block labeled 'edx = value', a light green block, a green block labeled 'eax = counter', a light green block, a dark grey block labeled 'eax = edx + eax', a large light green block, a blue block labeled 'counter = eax', and a final light green block. The blocks are connected by thin black lines, indicating a sequential flow.
- ```
graph TD; A[Thread 1] --> B[edx = value]; B --> C[ ]; C --> D[eax = counter]; D --> E[ ]; E --> F[eax = edx + eax]; F --> G[ ]; G --> H[counter = eax]; H --> I[ ]
```

eax = 1

```
counter = 1
```

The diagram shows a vertical sequence of colored rectangles representing instructions in Thread 2. From top to bottom, the steps are:

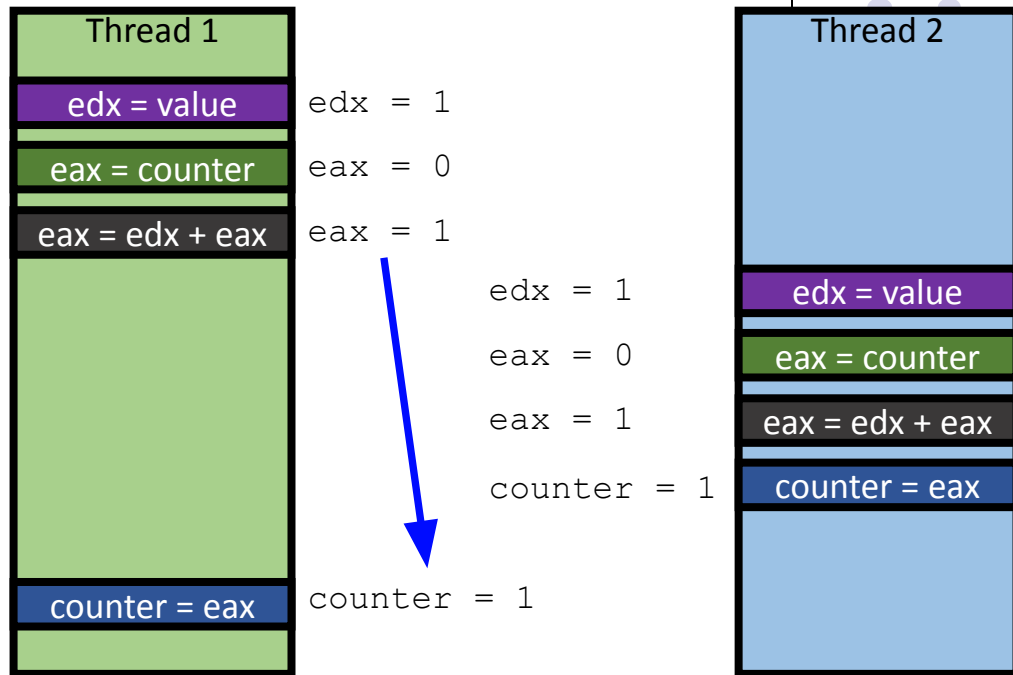
- A large light blue rectangle labeled "Thread 2".
- A purple rectangle containing the instruction `edx = value`.
- A thin white horizontal separator.
- A green rectangle containing the instruction `eax = counter`.
- A thin white horizontal separator.
- A dark grey rectangle containing the instruction `eax = edx + eax`.
- A thin white horizontal separator.
- A blue rectangle containing the instruction `counter = eax`.
- A thin white horizontal separator.
- A large light blue rectangle at the bottom.

```
counter = eax
```

Shared variable: Data race

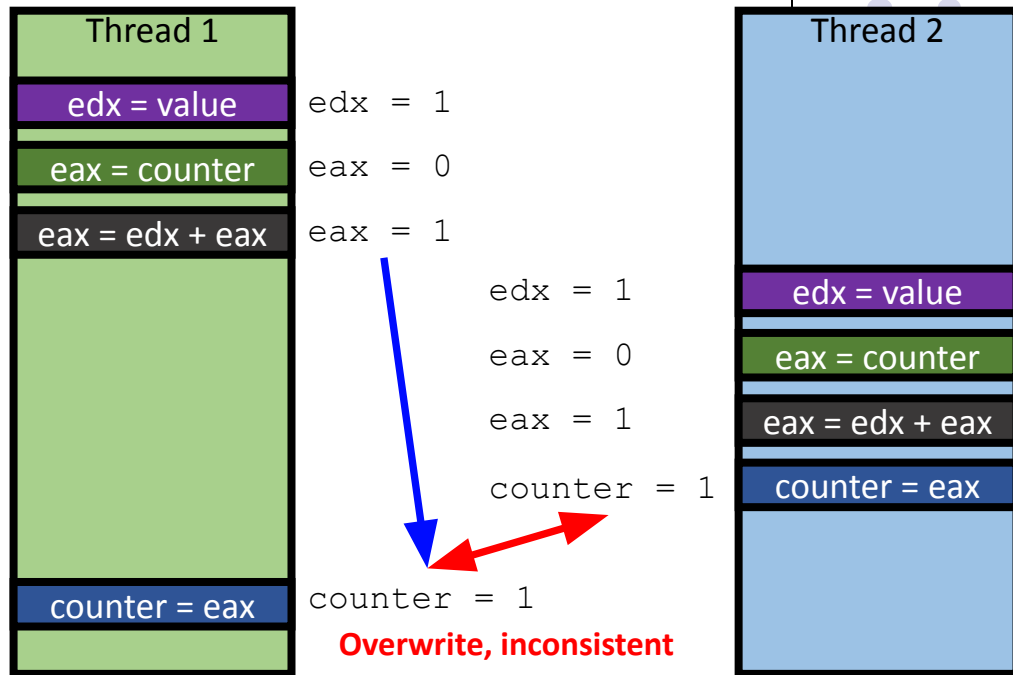


- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



Shared variable: Data race

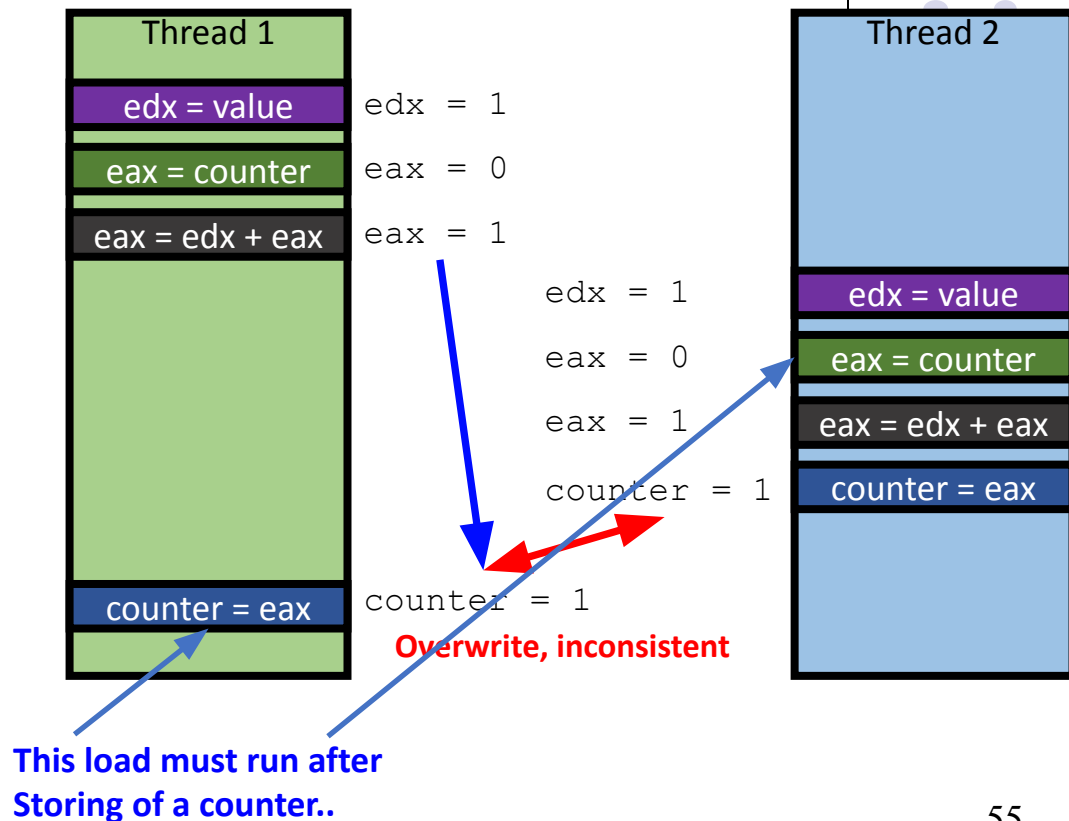
- `counter += value`
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume `counter = 0` at start, and `value = 1`;



Shared variable: Data race

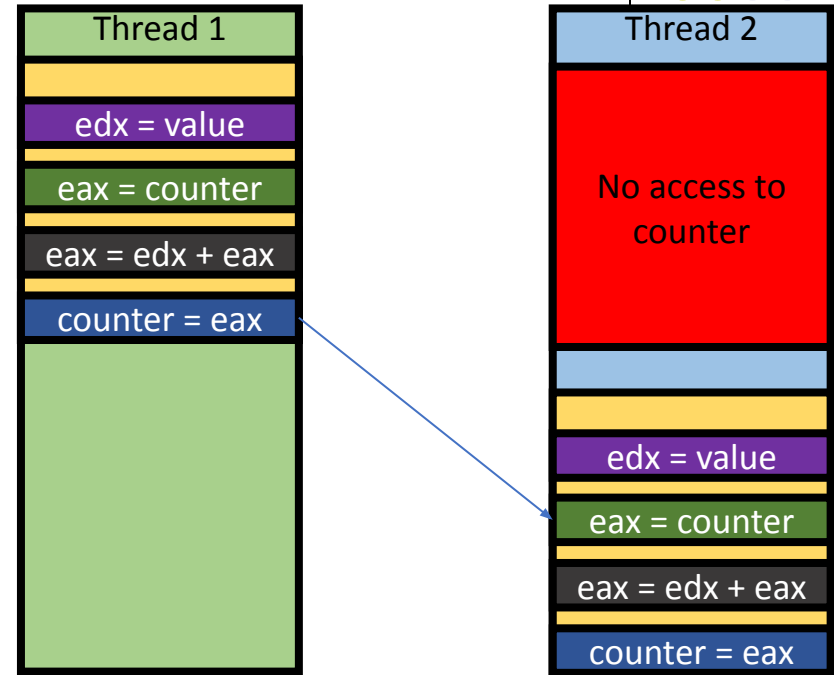


- counter += value
 - `edx = value;`
 - `eax = counter;`
 - `eax = edx + eax;`
 - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



How can we prevent data races?

- What we need?
 - **Exclusive access** to counter (shared variable)



How can we prevent data races?




- *Critical section* – a section of code, or collection of operations, in which only one process shall be executing at a given time
- *Mutual exclusion (Mutex)* - mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded)

How can we prevent data races?




- **Mutual Exclusion / Critical Section**
 - Combine multiple instructions as a chunk
 - Let only one chunk execution runs
 - Block other executions




How can we prevent data races?



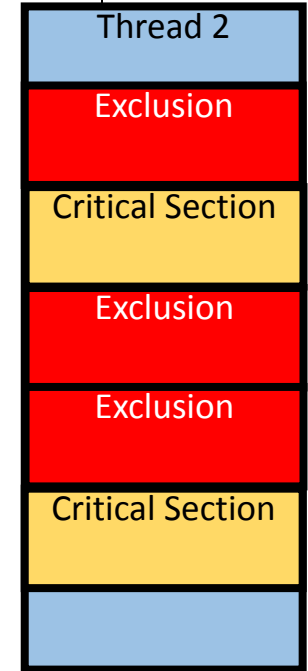
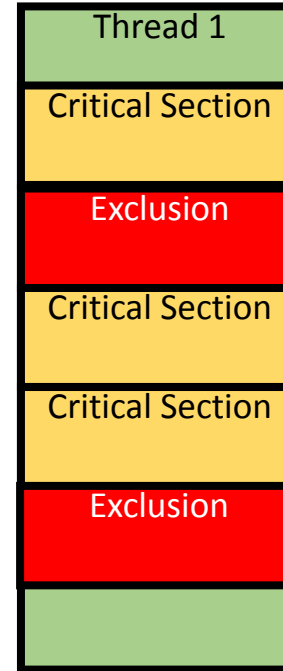
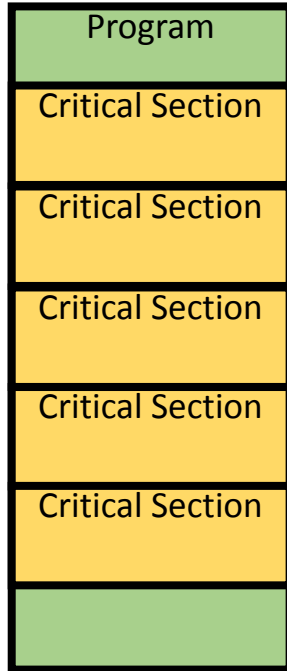
- **Mutual Exclusion / Critical Section**
 - Combine multiple instructions as a chunk
 - Let only one chunk execution runs
 - Block other executions






Does mutex renders threading useless?



Does mutex renders threading useless?



Does mutex renders threading useless?



Mutex Considerations



- Mutex can synchronize multiple threads and yield consistent result
 - No read before previous thread store the shared data
- Making the **entire program as critical section is meaningless**
 - Running time will be the same as single-threaded execution
- Apply critical section **as short as possible** to maximize benefit of having concurrency
 - Non-critical sections will run concurrently!