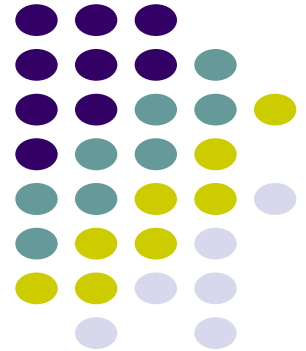


# Synchronization

---

ECE 469, Feb 27

Aravind Machiry



# Threads



- **Separate** the concepts of a “thread of control” (PC, SP, registers) from the rest of the process (address space, resources, accounting, etc.)
- Modern OSes support two entities:
  - the *task* (process), which defines an address space, a resource container, accounting info
  - the *thread* (lightweight process), which defines a single sequential execution stream within a task (process)

# Programming with Threads



- Flexible, but error-prone, since there no protection between threads
  - In C/C++,
    - automatic variables are private to each thread
    - global variables and dynamically allocated memory (malloc) are **shared**
- Need synchronization!

# The need for synchronization!



- Cooperating processes may share data via
  - shared address space (code, data, heap) by using threads
  - Files
  - (Sending messages)
- What can happen if processes try to access shared data (address) concurrently?
  - Sharing bank account with sibling:  
At 3pm: If (balance > \$10) withdraw \$10
- How hard is the solution?

# “Too much milk” Problem



## Person A

1. Look in fridge: out of milk
2. Leave for Walmart
5. Arrive at Walmart
6. Buy milk
7. Arrive home

## Person B

3. Look in fridge: out of milk
4. Leave for Walmart
8. Arrive at Walmart
9. Buy milk
10. Arrive home

- How to put in a locking mechanism?

# Possible Solution 1



Person A

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Person B

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

# Will this work?



Person A

```
1.if ( noMilk ) {  
  2.if (noNote) {  
    5.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

Person B

```
3.if ( noMilk ) {  
  4.if (noNote) {  
    6.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

- Process can get context switched after checking milk and note, but before leaving note

# Possible Solution 2



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```



# Will this work?



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

- We may not have Milk: Both process can leave note and skip buying milk

# Possible Solution 3



## Process A

```
leave noteA
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
remove noteA
```

## Process B

```
leave noteB
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

# Works, but complicated!



## Process A

leave noteA

```
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
```

remove noteA

## Process B

leave noteB

```
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
```

remove noteB

- A's code is **different** from B's
- **busy waiting** is a waste

# How can we solve this?



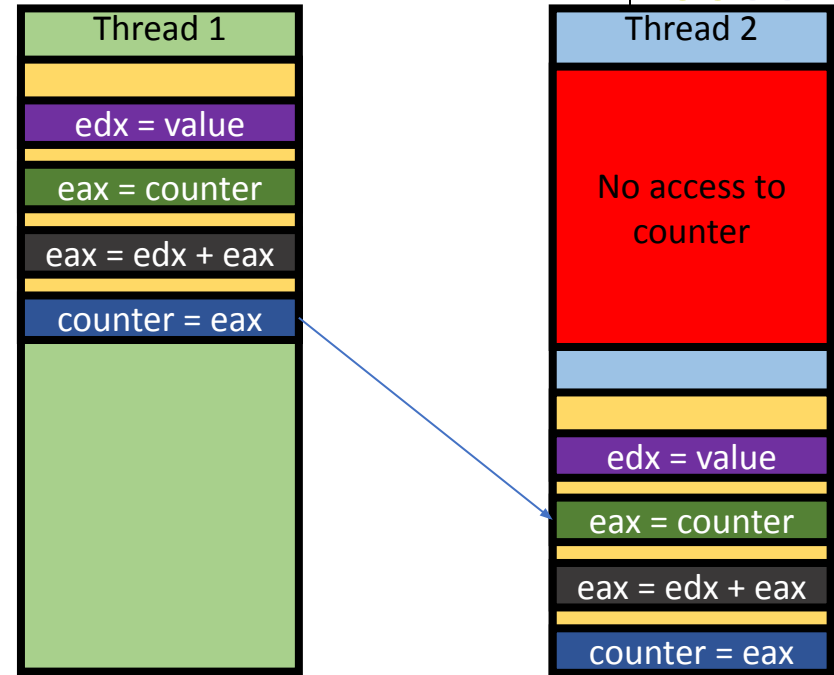
- Root cause: **Data Race**
- A thread's execution result could be inconsistent if other threads intervene its execution...

- counter += value

• <code>edx = value;</code>	<code>mov</code>	<code>0x20087b(%rip),%edx</code>	<code># 0x201010 &lt;value&gt;</code>
• <code>eax = counter;</code>	<code>mov</code>	<code>0x20087d(%rip),%eax</code>	<code># 0x201018 &lt;counter&gt;</code>
• <code>eax = edx + eax;</code>	<code>add</code>	<code>%edx,%eax</code>	
	<code>mov</code>	<code>%eax,0x200875(%rip)</code>	<code># 0x201018 &lt;counter&gt;</code>
• <code>counter = eax;</code>			

# How can we prevent data races?

- What we need?
  - **Exclusive access** to counter (shared variable)



# How can we prevent data races?




- *Critical section* – a section of code, or collection of operations, in which only one process shall be executing at a given time
- *Mutual exclusion (Mutex)* - mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded)

# How can we prevent data races?




- **Mutual Exclusion / Critical Section**
  - Combine multiple instructions as a chunk
  - Let only one chunk execution runs
  - Block other executions



# How can we prevent data races?



- **Mutual Exclusion / Critical Section**
  - Combine multiple instructions as a chunk
  - Let only one chunk execution runs
  - Block other executions





# Mutex Considerations



- Mutex can synchronize multiple threads and yield consistent result
  - No read before previous thread store the shared data
- Making the **entire program as critical section is meaningless**
  - Running time will be the same as single-threaded execution
- Apply critical section **as short as possible** to maximize benefit of having concurrency
  - Non-critical sections will run concurrently!

# Implementing Mutual exclusion



- Data races occur because of scheduler interleaving executing of different threads
- How to avoid this? Prevent interleaving

# Preventing Interleaving



- `cli`, in a single processor computer
  - Clear interrupt bit
- CPU will never get interrupt until it runs `sti`
  - Set interrupt bit
- There will be no other execution
  - **Any problems?**

- `counter += value`
  - `cli`
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
  - `sti`

# Preventing Interleaving



- `cli`, in a single processor computer
  - Clear interrupt bit
- CPU will never get interrupt until it runs `sti`
  - Set interrupt bit
- There will be no other execution
  - Any problems?
    - Multi CPU?
    - `cli/sti` available in Ring 0

- `counter += value`

- `cli`
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- `sti`

# Mutual Exclusion through locks



- Lock
  - Prevent others enter the critical section
- Unlock
  - Release the lock, let others acquire the lock


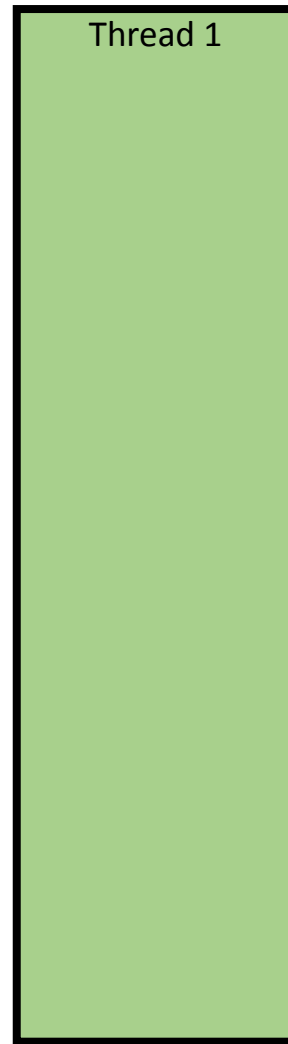
- counter += value
  - **lock()**
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
  - **unlock()**

# Mutual Exclusion through locks





- Lock
  - Prevent others enter the critical section
- How?
  - Check if any other execution in the critical section
    - If it is, wait; busy-waiting with while()
  - If not, acquire the lock!
    - Others cannot get into the critical section
  - Run critical section
  - Unlock, let other execution know that I am out!
- counter += value
  - **lock()**
    - `edx = value;`
    - `eax = counter;`
    - `eax = edx + eax;`
    - `counter = eax;`
  - **unlock()**

# Mutex Example

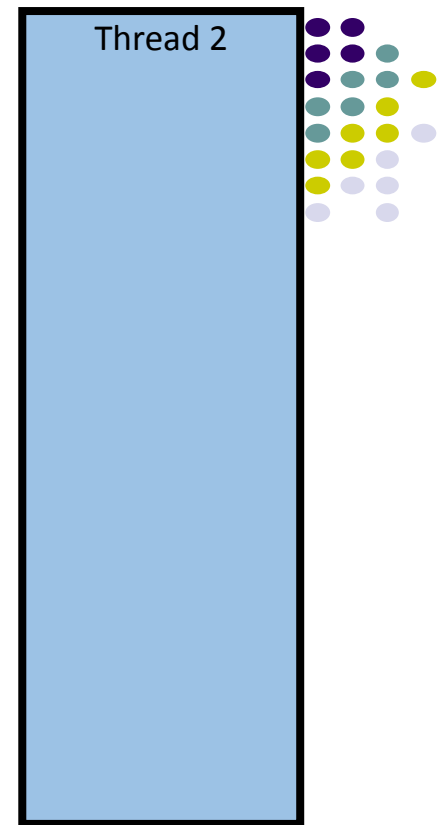
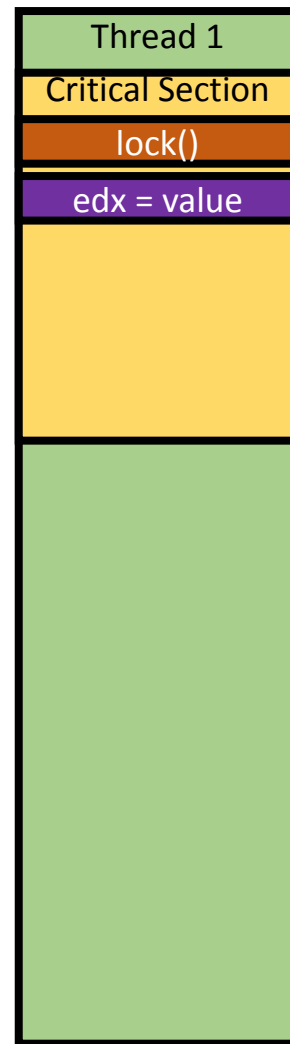


# Mutex Example


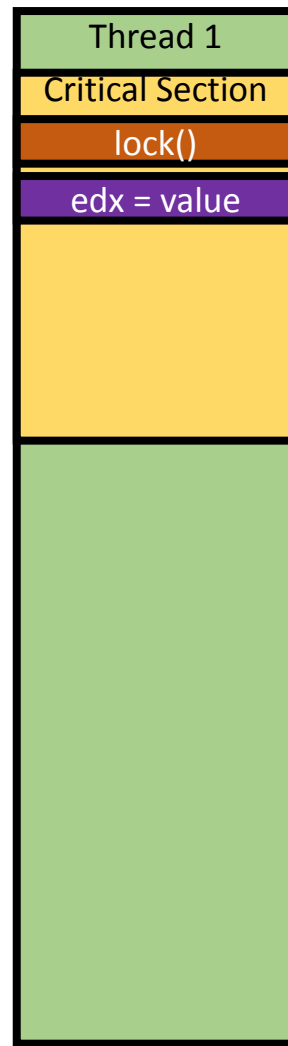





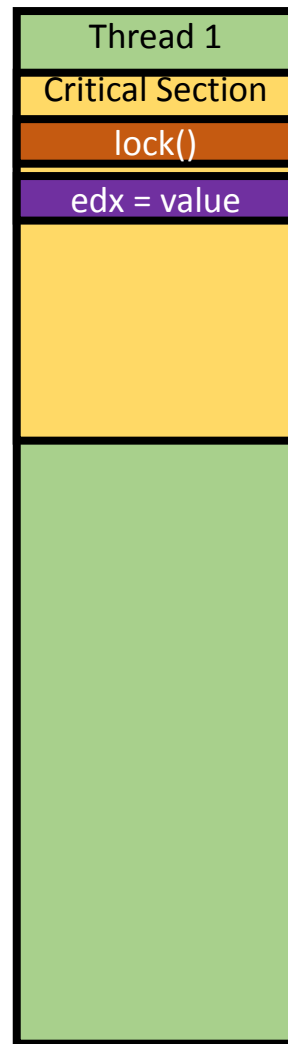
# Mutex Example




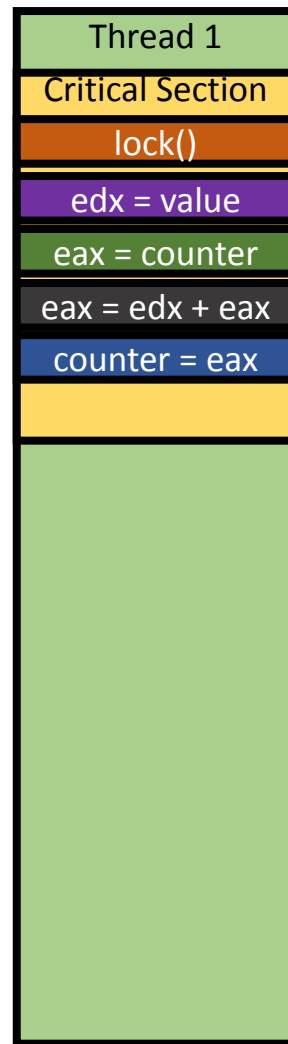
# Mutex Example





# Mutex Example



# Mutex Example

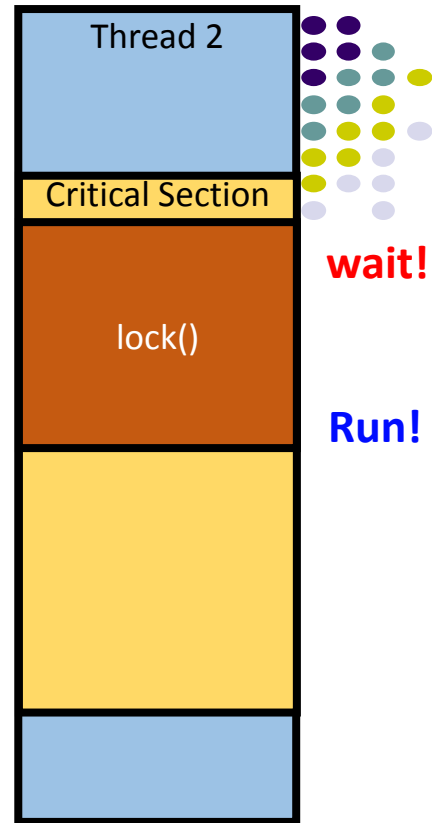
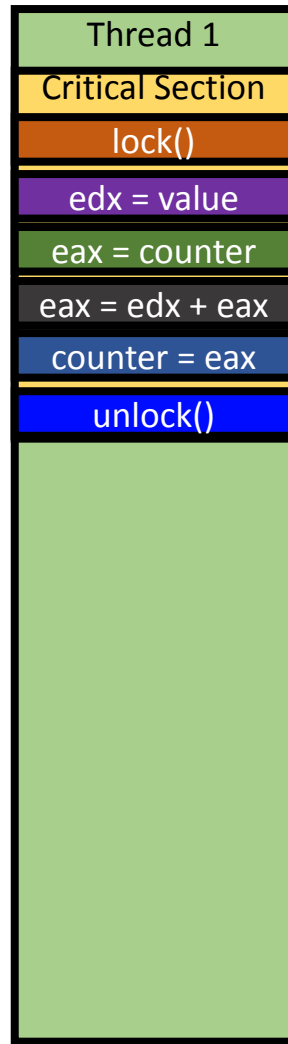


# Mutex Example


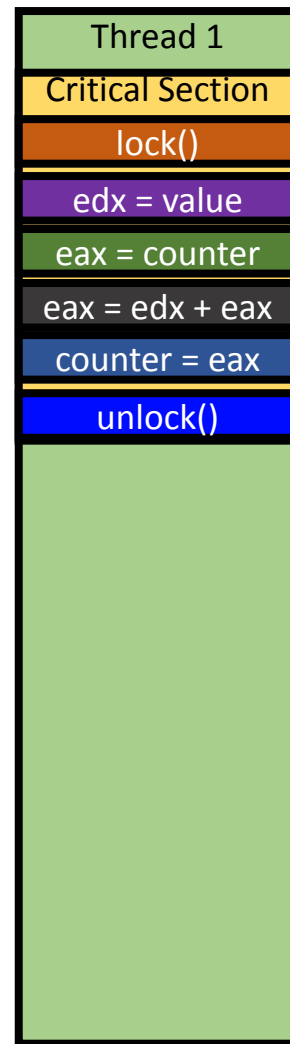


**wait!**

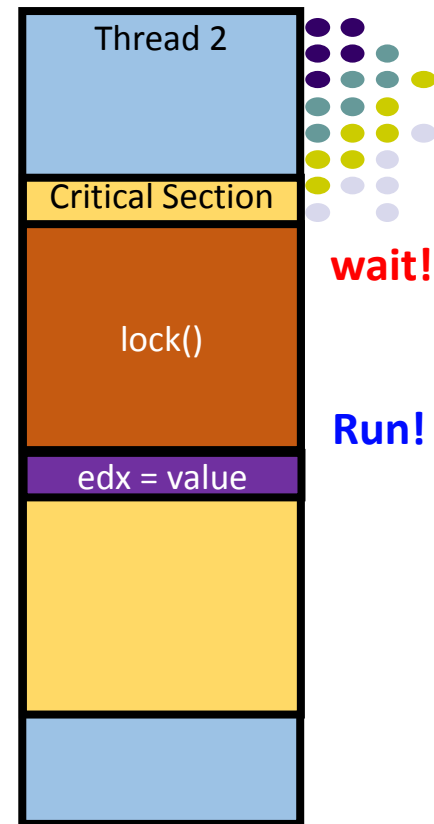
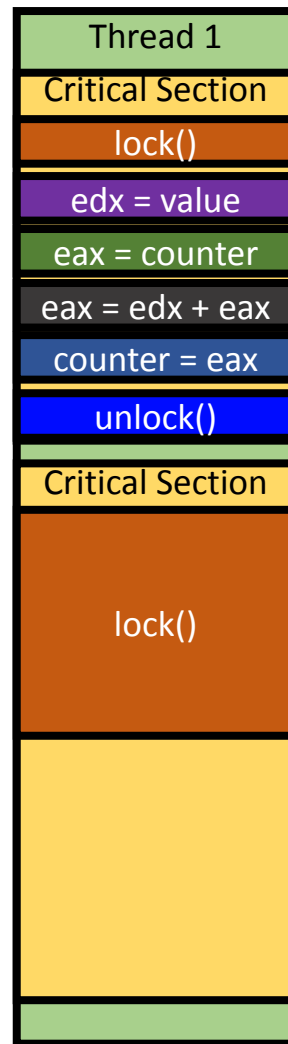
# Mutex Example



# Mutex Example

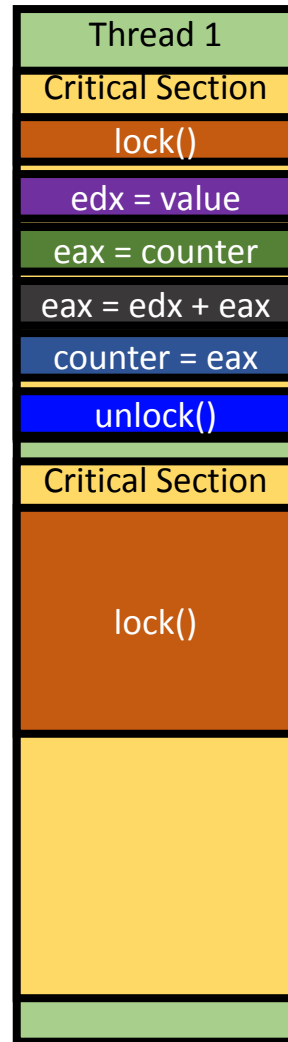


# Mutex Example



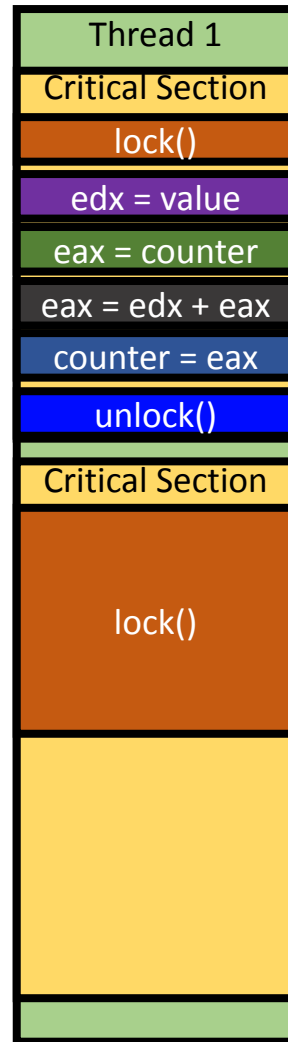


# Mutex Example

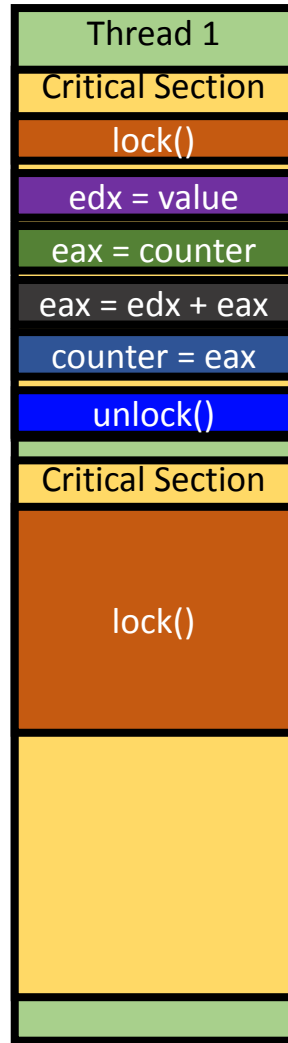


# Run!

# Mutex Example



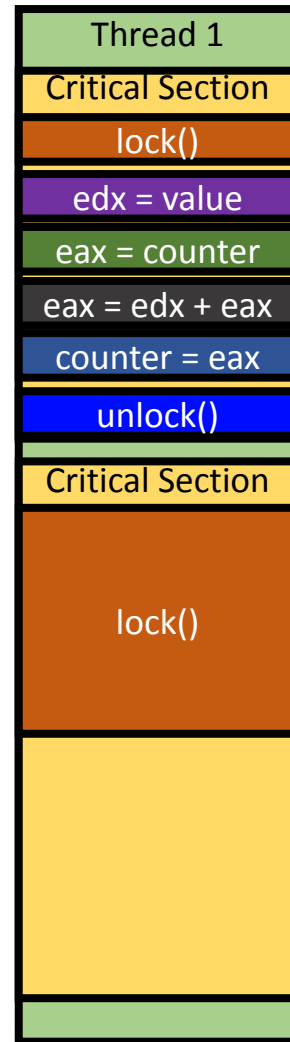
# Mutex Example



# Mutex Example


wait!

Run!



wait!

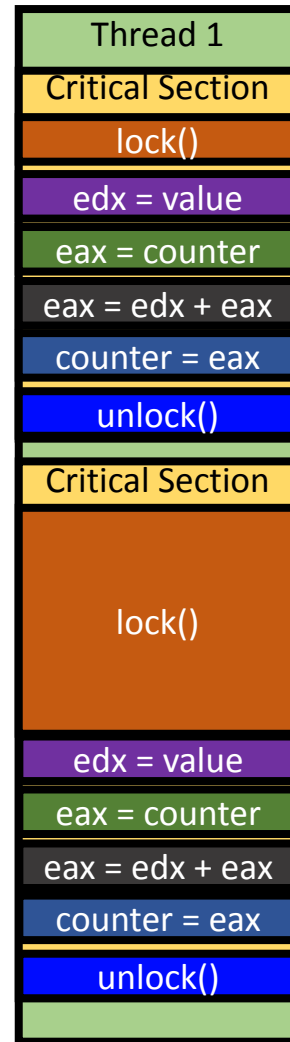
Run!



# Mutex Example


wait!

Run!



wait!

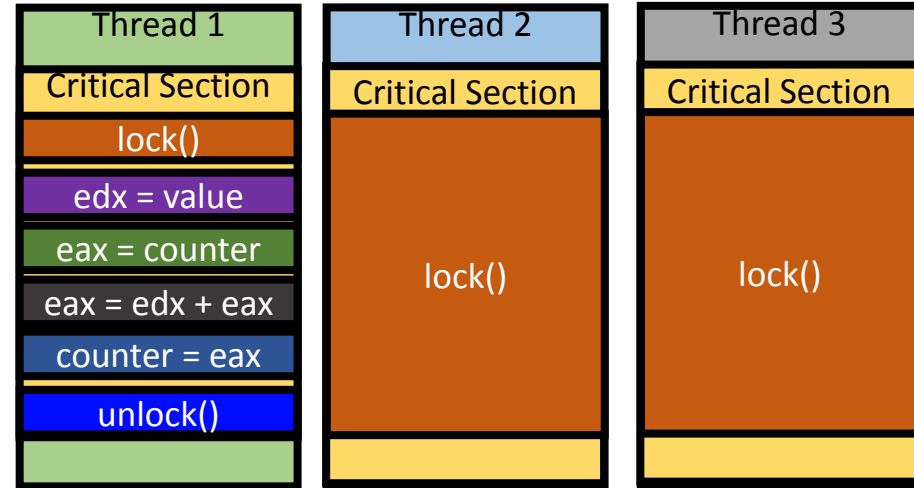
Run!



# Implementing lock



- Only one can run in critical section
- Others must wait!
  - Until nobody runs in critical section
- How can we create such
  - Lock() / Unlock() ?

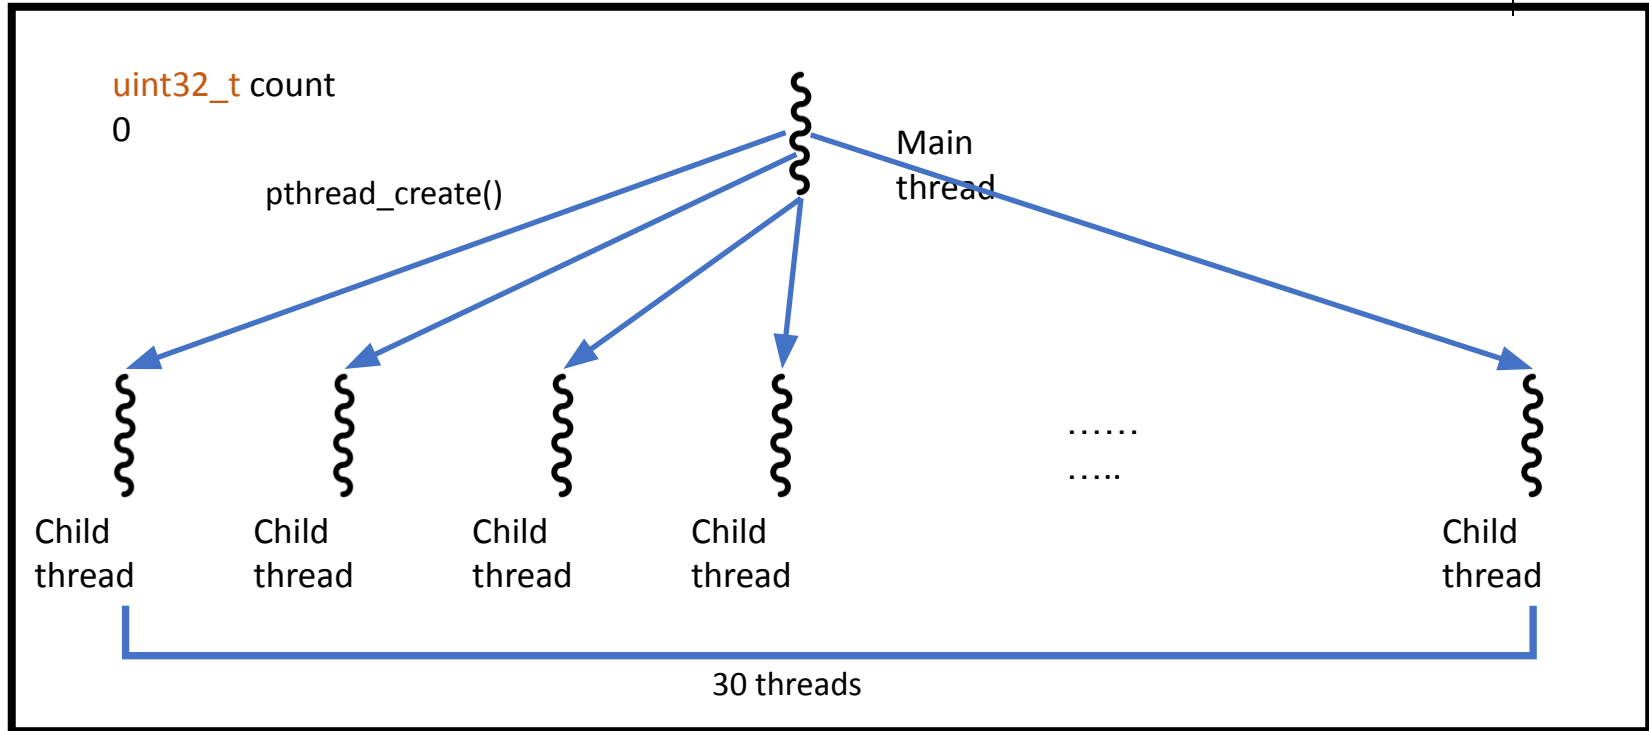


# lock example



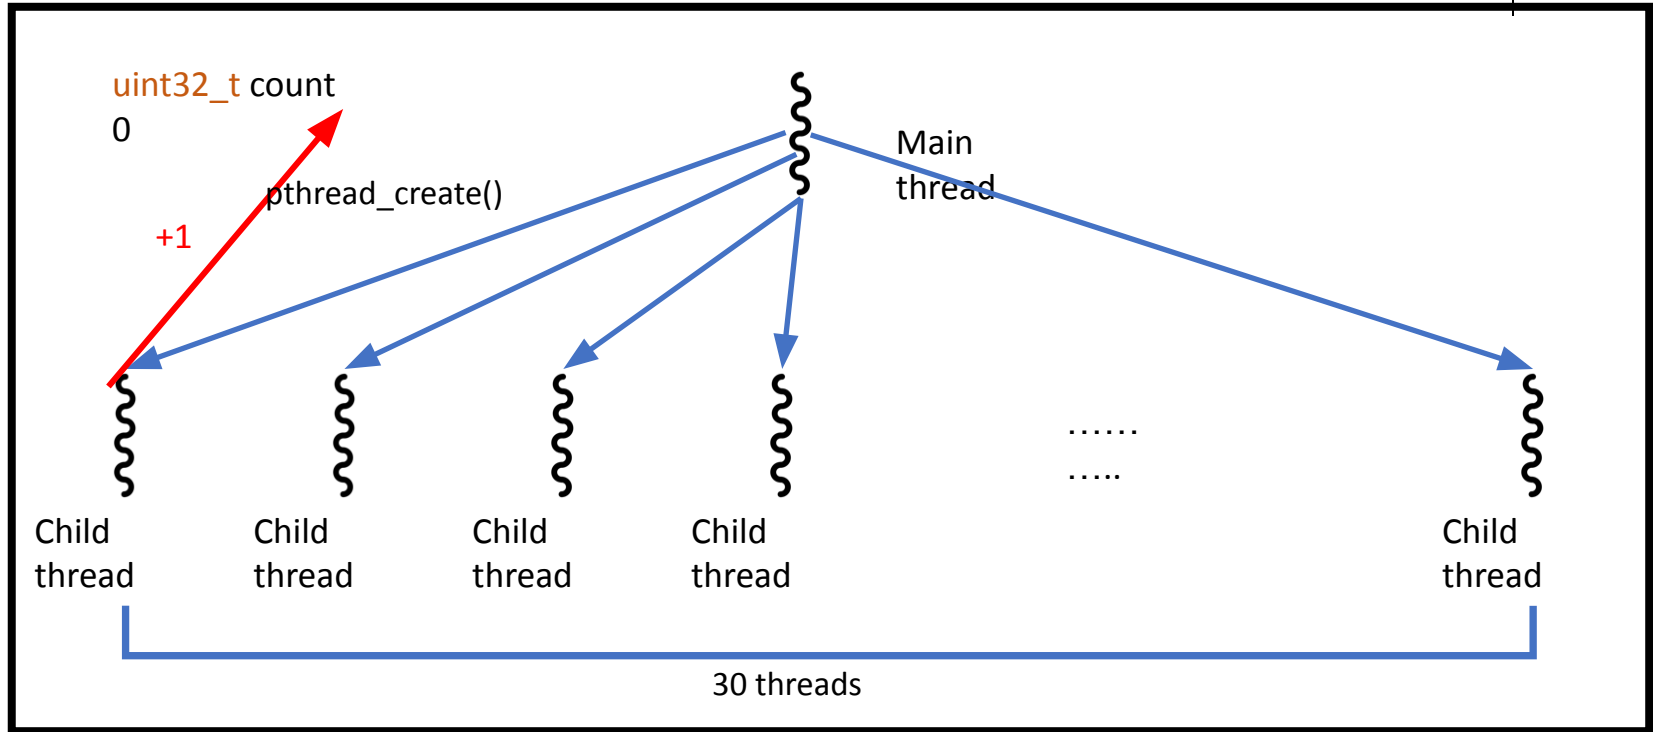
- [https://github.com/purs3lab/ee469\\_examples/tree/master/lock\\_example](https://github.com/purs3lab/ee469_examples/tree/master/lock_example)
- Run 30 threads, each count upto 50
- Build code
  - `$ make`
- Run code
  - `$ ./lock xchg` # shows the result of using xchg lock

# lock example



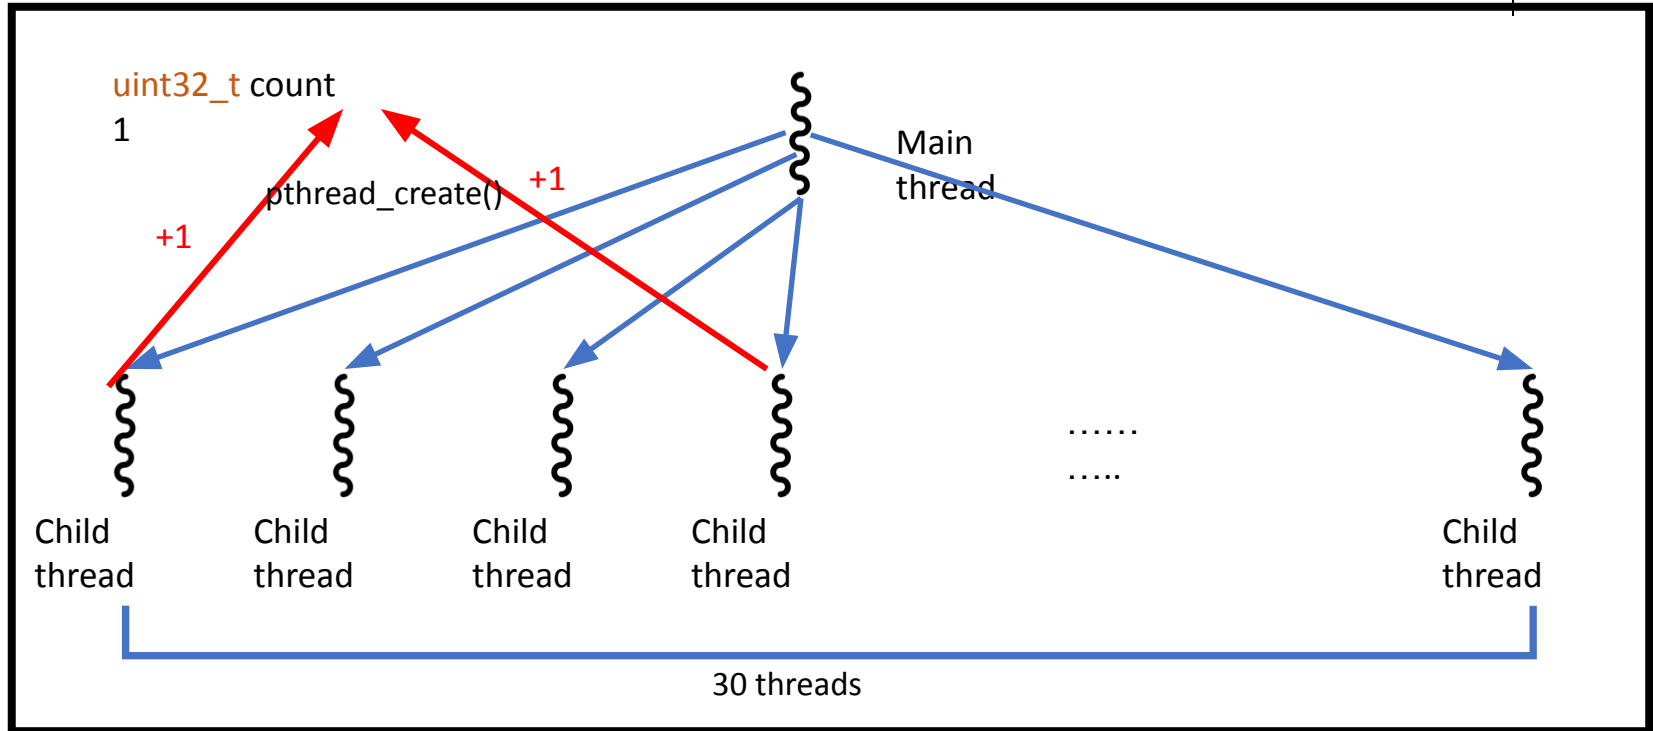


# lock example



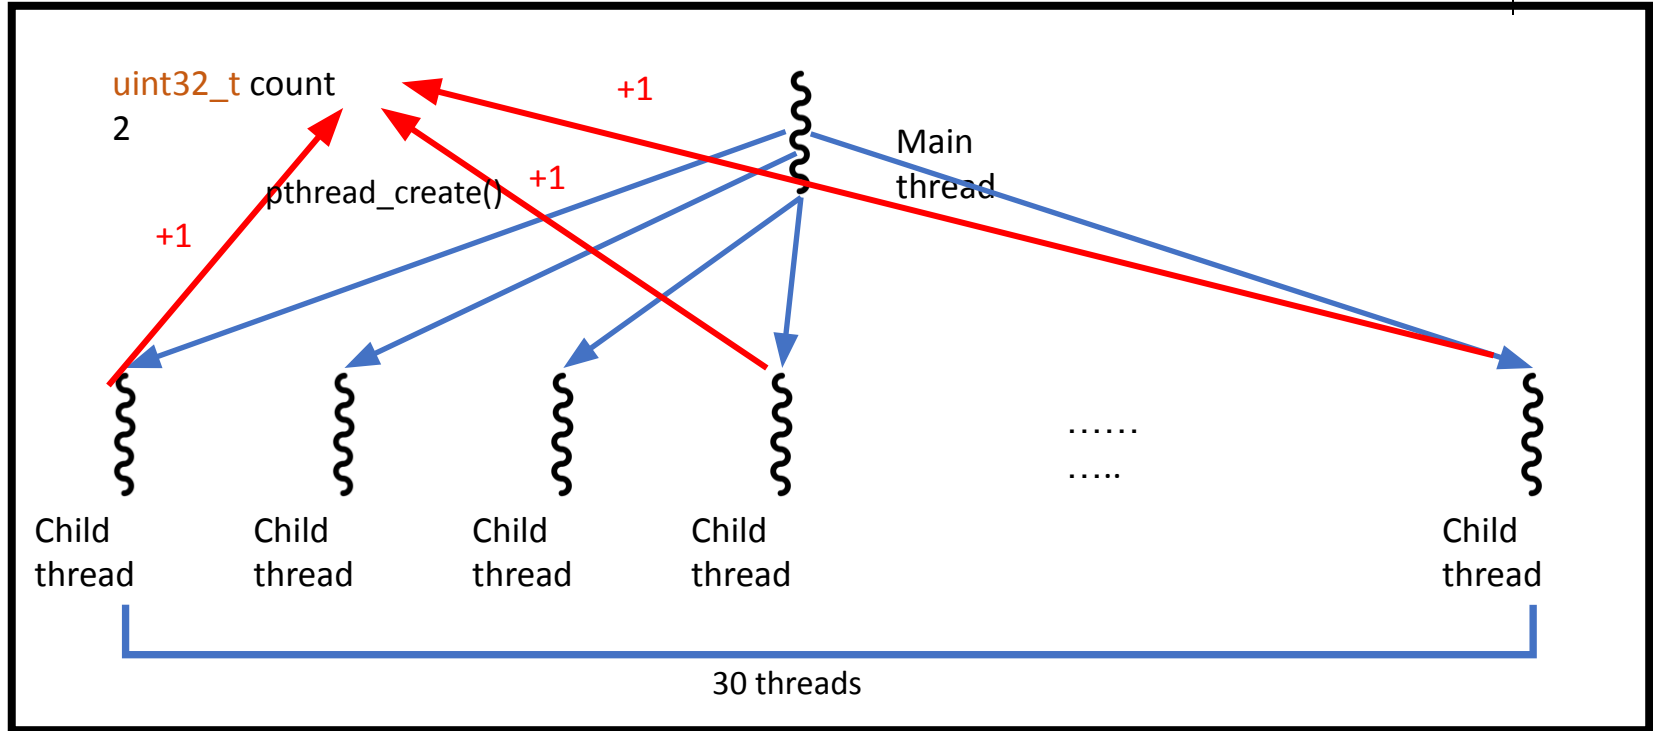
**Each thread will increase count by 1 for 50 times**

# lock example



**Each thread will increase count by 1 for 50 times**

# lock example



Each thread will increase count by 1 for 50 times



# lock example



- Running

- `$ ./lock no`                    `# using no lock at all`
- `$ ./lock bad`                    `# using a bad lock implementation`
- `$ ./lock xchg`                    `# using xchg lock`
- `$ ./lock cmpxchg`                `# using lock cmpxchg`
- `$ ./lock tts`                    `# using soft test-and-test & set with xchg`
- `$ ./lock backoff`                `# using exponential backoff cmpxchg`
- `$ ./lock mutex`                  `# using pthread mutex`

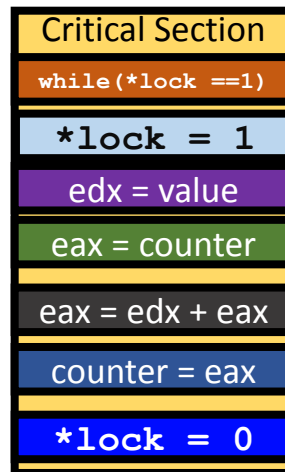
# Manual Spinlock (bad\_lock)



```
void  
bad_lock(volatile uint32_t *lock) {  
    while (*lock == 1);  
    *lock = 1;  
}
```

- Spinlock
  - Run a loop to check if critical section is empty
  - Set a lock variable, e.g., `lock`
  - Lock semantic
    - Nobody runs critical section if `*lock == 0`, so one can run the section
      - At the start of the section, set `*lock = 1`
    - Somebody runs in critical section if `*lock == 1`, so one must wait
- `lock(lock)`
  - Wait until `l` becomes 0, e.g., `while (*lock == 1);`
    - Then, nobody runs in the critical section!
  - set `*lock = 1`
- `unlock(lock)`
  - Set `*lock = 0`

Lock



Unlock

# Manual Spinlock (bad\_lock)



- What will happen if we implement lock
  - As bad\_lock / bad\_unlock?
- bad\_lock
  - Wait until lock becomes 0 (loops if 1)
  - And then, set lock as 1
    - Because it was 0, we can set it as 1
  - Others must wait!
- bad\_unlock
  - Just set \*lock as 0

```
void *  
count_bad_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        bad_lock(&lock);  
        sched_yield();  
        count += 1;  
        bad_unlock(&lock);  
    }  
}
```

```
void  
bad_lock(volatile uint32_t *lock) {  
    while (*lock == 1);  
    *lock = 1;  
}  
  
void  
bad_unlock(volatile uint32_t *lock) {  
    *lock = 0;  
}
```

# Manual Spinlock (bad\_lock)



- What will happen if we implement lock
  - As bad\_lock / bad\_unlock?
- bad\_lock
  - Wait until lock becomes 0 (loops if 1)
  - And then, set lock as 1
    - Because it was 0, we can set it as 1
  - Others must wait! **Can pass this if lock=0**
- bad\_unlock
  - Just set \*lock as 0

```
void *  
count_bad_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        bad_lock(&lock);  
        sched_yield();  
        count += 1;  
        bad_unlock(&lock);  
    }  
}
```

```
void  
bad_lock(volatile uint32_t *lock) {  
    while (*lock == 1);  
    *lock = 1;  
}  
  
void  
bad_unlock(volatile uint32_t *lock) {  
    *lock = 0;  
}
```



# Manual Spinlock (bad\_lock)



- What will happen if we implement lock
  - As bad\_lock / bad\_unlock?

- bad\_lock

- Wait until lock becomes 0 (loops if 1)
- And then, set lock as 1
  - Because it was 0, we can set it as 1
- Others must wait! **Can pass this if lock=0**  
**Sets lock=1 to block others**

- bad\_unlock

- Just set \*lock as 0

```
void *
count_bad_lock(void *args) {
    for (int i=0; i < N_COUNT; ++i) {
        bad_lock(&lock);
        sched_yield();
        count += 1;
        bad_unlock(&lock);
    }
}
```

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}

void
bad_unlock(volatile uint32_t *lock) {
    *lock = 0;
}
```

# Manual Spinlock (bad\_lock)



- What will happen if we implement lock
  - As bad\_lock / bad\_unlock?

- bad\_lock

- Wait until lock becomes 0 (loops if 1)
- And then, set lock as 1
  - Because it was 0, we can set it as 1
- Others must wait! **Can pass this if lock=0**  
**Sets lock=1 to block others**

- bad\_unlock

- Just set \*lock as 0

```
void *
count_bad_lock(void *args) {
    for (int i=0; i < N_COUNT; ++i) {
        bad_lock(&lock);
        sched_yield();
        count += 1;
        bad_unlock(&lock);
    }
}
```

**Critical Section**

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}

void
bad_unlock(volatile uint32_t *lock) {
    *lock = 0;
}
```

# Manual Spinlock (bad\_lock)



- What will happen if we implement lock
  - As bad\_lock / bad\_unlock?

- bad\_lock

- Wait until lock becomes 0 (loops if 1)
- And then, set lock as 1
  - Because it was 0, we can set it as 1
- Others must wait! **Can pass this if lock=0**  
**Sets lock=1 to block others**

- bad\_unlock

- Just set \*lock as 0  
**Sets lock=0 to release**

```
void *  
count_bad_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        bad_lock(&lock);  
        sched_yield();  
        count += 1;  
        bad_unlock(&lock);  
    }  
}
```

```
void  
bad_lock(volatile uint32_t *lock) {  
    while (*lock == 1);  
    *lock = 1;  
}  
  
void  
bad_unlock(volatile uint32_t *lock) {  
    *lock = 0;  
}
```

# Why does bad\_lock doesn't work?



```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

# Why does bad\_lock doesn't work?



```
mov    (%rdi),%eax
cmp    $0x1,%eax
je     0x400b60 <bad_lock>
movl   $0x1, (%rdi)
```

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

# Why does bad\_lock doesn't work?



```
LOAD  mov    (%rdi),%eax
      cmp    $0x1,%eax
      je     0x400b60 <bad_lock>
STORE movl    $0x1, (%rdi)
```

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

# Why does bad\_lock doesn't work?



```
LOAD  mov    (%rdi),%eax
      cmp    $0x1,%eax
      je     0x400b60 <bad_lock>
STORE movl    $0x1, (%rdi)
```

Another thread might get scheduled here.

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

# Why does bad\_lock doesn't work?



- There is a **room for race condition!**

```
LOAD  mov    (%rdi),%eax
      cmp    $0x1,%eax
      je     0x400b60 <bad_lock>
STORE movl    $0x1, (%rdi)
```

Another thread might get scheduled here.

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```



# Why does bad\_lock doesn't work?



- There is a room for race condition!

```
LOAD  mov    (%rdi),%eax
      cmp    $0x1,%eax
      je     0x400b60 <bad_lock>
STORE movl    $0x1, (%rdi)
```

Race condition may happen

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

# Why does bad\_lock doesn't work?



- There is a room for race condition!

```
LOAD  mov    (%rdi),%eax
      cmp    $0x1,%eax
      je     0x400b60 <bad_lock>
STORE movl    $0x1, (%rdi)
```

Race condition may happen

```
void
bad_lock(volatile uint32_t *lock) {
    while (*lock == 1);
    *lock = 1;
}
```

## Recall: Why does this work for humans?



- Human can perform *test* (look for other person & milk) and *set* (leave note) at the same time.



# Atomic Test and Set

- We need a way to test
  - if lock == 0
- And we would like to set
  - lock = 1
- And do this atomically
- Hardware support is required
  - `xchg` in x86 does this
  - An atomic test-and-set operation

```
mov    (%rdi),%eax
cmp     $0x1,%eax
je      0x400b60 <bad_lock>
movl    $0x1, (%rdi)
```

**Not like these four instructions...**




# xchg: Atomic Value Exchange in x86

- `xchg [memory], %reg`
  - Exchange the content in `[memory]` with the value in `%reg` atomically
- E.g.,
  - `mov $1, %eax`
  - `xchg $lock, %eax`
- This will set `%eax` as the value in `lock`
  - `%eax` will be 0 if `lock==0`, will be 1 if `lock==1`
- At the same time, this will set `lock = 1` (the value was in `%eax`)
- CPU applies 'lock' at hardware level (cache/memory) to do this
  - Hardware guarantees no data race when running `xchg`



# xchg: Atomic Value Exchange in x86

- `xchg [memory], %reg`
  - Exchange the content in `[memory]` with the value in `%reg` atomically
- E.g.,
  - `mov $1, %eax`
  - `xchg $lock, %eax` **Swap lock and eax atomically**
- This will set `%eax` as the value in `lock`
  - `%eax` will be 0 if `lock==0`, will be 1 if `lock==1`
- At the same time, this will set `lock = 1` (the value was in `%eax`)
- CPU applies 'lock' at hardware level (cache/memory) to do this
  - Hardware guarantees no data race when running `xchg`



# xchg: Atomic Value Exchange in x86

- E.g.,
  - `mov $1, %eax`
  - `xchg $lock, %eax` **Swap lock and eax atomically**
- This will set `%eax` as the value in `lock`
  - `%eax` will be 0 if `lock==0`, will be 1 if `lock==1`
- How can we determine if a thread acquired the lock?
  - if `eax == 0`
    - This means the `lock` was 0, and after running `xchg`, `lock` will be 1 (`eax` was 1)
    - We acquired the lock!!! (`lock` was 0 and now the `lock` is 1)
  - if `eax == 1`
    - This means the `lock` was 1, and after running `xchg`, `lock` will be 1
    - We did not acquired the lock (it was 1)
    - `lock == 1` means some other thread acquired this...



# Lock using xchg

- xchg\_lock
  - Use atomic 'xchg' instruction to
  - Load and store values atomically
  - Set value to 1, and compare ret
    - If 0, then you can acquire the lock
    - If 1, lock as 1, you must wait
- xchg\_unlock
  - Use atomic 'xchg'
  - Set value to 0
    - Do not need to check
    - You are the only thread that runs in the
    - Critical section..

```
void *  
count_xchg_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        xchg_lock(&lock);  
        sched_yield();  
        count += 1;  
        xchg_unlock(&lock);  
    }  
}
```

```
void  
xchg_lock(volatile uint32_t *lock) {  
    while(xchg(lock, 1));  
}  
  
void  
xchg_unlock(volatile uint32_t *lock) {  
    xchg(lock, 0);  
}
```





# Lock using xchg

- xchg\_lock
  - Use atomic 'xchg' instruction to
  - Load and store values atomically
  - Set value to 1, and compare ret
    - If 0, then you can acquire the lock
    - If 1, lock as 1, you must wait
- xchg\_unlock
  - Use atomic 'xchg'
  - Set value to 0
    - Do not need to check
    - You are the only thread that runs in the
    - Critical section..

```
void *  
count_xchg_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        xchg_lock(&lock);  
        sched_yield();  
        count += 1;  
        xchg_unlock(&lock);  
    }  
}
```

A blue arrow points from the text "Critical Section" to the loop body of the `count_xchg_lock` function.

```
void  
xchg_lock(volatile uint32_t *lock) {  
    while(xchg(lock, 1));  
}  
  
void  
xchg_unlock(volatile uint32_t *lock) {  
    xchg(lock, 0);  
}
```

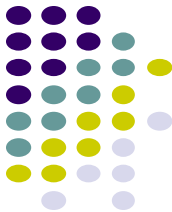


# Does xchg\_lock works?

- Yes!!

```
Running 30 threads each counting to 50 using xchg lock  
Result:1500, Time taken: 2358.591000 ms
```

- Any issues?



# Issues with xchg\_lock

- xchg will always update the value
  - If lock == 0
    - lock = 1
    - eax = 0
  - If lock == 1
    - lock = 1
    - eax = 1
- We use while() to check the value in lock
  - Will be cached into L1 cache of the CPU
- After updating a value in cache
  - We need to invalidate the cache in other CPUs...



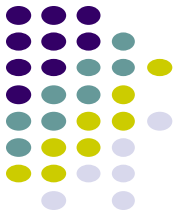
# Issues with xchg\_lock

- xchg will always update the value
  - If lock == 0
    - lock = 1      **Swap with eax == 1, update lock to 1**
    - eax = 0
  - If lock == 1
    - lock = 1
    - eax = 1      **Swap with eax == 1, update lock to 1**
- We use while() to check the value in lock
  - Will be cached into L1 cache of the CPU
- After updating a value in cache
  - We need to invalidate the cache in other CPUs...

# No need to write when lock == 1



- Let's not do that
  - xchg can't do that



# No need to write when lock == 1

- Let's not do that
  - xchg can't do that
- New method: `Test` and `test-and-set`
  - Check the value first (if lock == 0) `□ TEST`
  - If it is,
    - Do `test-and-set`
  - Otherwise (if lock == 1),
    - Do nothing
    - DO NOT UPDATE lock if lock == 1 (**No cache invalidate**)



# Lock using test and set

- `tts_xchg_lock`
- Algorithm
  - Wait until lock becomes 0
  - After `lock == 0`
    - `xchg(lock, 1)`
    - This only updates `lock = 1` if `lock` was 0
- Why **`xchg`**, why not **`*lock = 1`** directly?
  - `while` and `xchg` are not atomic
  - Load/Store must happen at
    - The same time!

```
void *  
count_tts_xchg_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        tts_xchg_lock(&lock);  
        sched_yield();  
        count += 1;  
        xchg_unlock(&lock);  
    }  
}
```

```
void  
tts_xchg_lock(volatile uint32_t *lock) {  
    while (1) {  
        while(*lock == 1);  
        if (xchg(lock, 1) == 0) {  
            break;  
        }  
    }  
}
```



# Lock using test and set

- `tts_xchg_lock`
- Algorithm
  - Wait until lock becomes 0
  - After lock == 0
    - `xchg (lock, 1)`
    - This only updates lock = 1 if lock was 0
- Why `xchg`, why not `*lock = 1` directly
  - while and `xchg` are not atomic
  - Load/Store must happen at
    - The same time!

```
void *  
count_tts_xchg_lock(void *args) {  
    for (int i=0; i < N_COUNT; ++i) {  
        tts_xchg_lock(&lock);  
        sched_yield();  
        count += 1;  
        xchg_unlock(&lock);  
    }  
}
```

```
void  
tts_xchg_lock(volatile uint32_t *lock) {  
    while (1) {  
        while(*lock == 1);  
        if (xchg(lock, 1) == 0) {  
            break;  
        }  
    }  
}
```





# Test and Set in x86

- `cmpxchg [update-value], [memory]`
  - Compare the value in `[memory]` with `%eax`
  - If matched, exchange value in `[memory]` with `[update-value]`
  - Otherwise, do not perform exchange
- `cmpxchg(lock, 0, 1)`
  - Arguments: Lock, test value, update value
  - Returns old value of lock

Test

Test-and-set



# Lock using cmpxchg\_lock

- Cmpxchg\_lock
  - Use cmpxchg to set lock = 1
    - Do not update if lock == 1
    - Only write 1 to lock if lock == 0
- Xchg\_unlock
  - Use xchg\_unlock to set lock = 0
  - Because we have 1 writer and
  - This always succeeds...

```
void *
count_cmpxchg_lock(void *args) {
    for (int i=0; i < N_COUNT; ++i) {
        cmpxchg_lock(&lock);
        sched_yield();
        count += 1;
        xchg_unlock(&lock);
    }
}
```

```
void
cmpxchg_lock(volatile uint32_t *lock) {
    while(cmpxchg(lock, 0, 1));
}

void
xchg_unlock(volatile uint32_t *lock) {
    xchg(lock, 0);
}
```



# Lock using cmpxchg\_lock

- Cmpxchg\_lock

- Use cmpxchg to set lock = 1
  - Do not update if lock == 1
  - Only write 1 to lock if lock == 0

- Xchg\_unlock

- Use xchg\_unlock to set lock = 0
- Because we have 1 writer and
- This always succeeds...

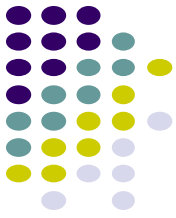
```
void *
count_cmpxchg_lock(void *args) {
    for (int i=0; i < N_COUNT; ++i) {
        cmpxchg_lock(&lock);
        sched_yield();
        count += 1;
        xchg_unlock(&lock);
    }
}
```

Critical Section

```
void
cmpxchg_lock(volatile uint32_t *lock) {
    while(cmpxchg(lock, 0, 1));
}

void
xchg_unlock(volatile uint32_t *lock) {
    xchg(lock, 0);
}
```

# Reading fine print : x86 is too COMPLEX!



This *[cmpxchg]* instruction can be used with a LOCK prefix to allow the instruction to be executed atomically. To simplify the interface to the processors bus, the destination operand receives a write cycle without regard to the result of the comparison. The destination operand is written back if the comparison fails; otherwise, the source operand is written into the destination. (The processor never produces a locked read without also producing a locked write.)

Cmpxchg designed to be Test and Test & Set instruction

However, Intel CPU gets too complex, so they decided to always update value regardless the result of comparison

# tts\_xchg\_lock v/s cmpxchg\_lock



- tts\_xchg\_lock is faster then cmpxchg\_lock

# Not everything in hardware is fast!



**Observation 2:** AddressSanitizer, despite being a software-only approach, performs on par with ICC-MPX and better than GCC-MPX. This unexpected result testifies that the HW-assisted performance improvements of MPX are offset by its complicated design. At the same time,




# Using hardware features smartly

- `backoff_cmpxchg_lock(lock)`
- Try `cmpxchg`
  - If succeeded, acquire the lock.
  - If failed
    - Wait 1 cycle (pause) for 1<sup>st</sup> trial
    - Wait 2 cycles for 2<sup>nd</sup> trial
    - Wait 4 cycles for 3<sup>rd</sup> trial
    - ...
    - Wait 65536 cycles for 17<sup>th</sup> trial..
    - Wait 65536 cycles for 18<sup>th</sup> trial..

```
void
backoff_cmpxchg_lock(volatile uint32_t *lock) {
    uint32_t backoff = 1;
    while(cmpxchg(lock, 0, 1)) {
        for (int i=0; i<backoff; ++i) {
            __asm volatile("pause");
        }
        if (backoff < 0x10000) {
            backoff <= 1;
        }
    }
}
```

- [https://en.wikipedia.org/wiki/Exponential\\_backoff](https://en.wikipedia.org/wiki/Exponential_backoff)



# Summary

- Mutex is implemented with Spinlock
  - Waits until lock == 0 with a while loop (that's why it's called spin)

- Naïve code implementation

- Load/Store must be atomic

- xchg is a “test and set” atom

- Consistent, however, many c

- Lock cmpxchg is a “test and t

- But Intel implemented this a

- We can implement test-and-


- Faster!

- We can also implement expd

- Much faster! **Faster Than p**

```
./lock no
Running 30 threads each counting to 50 using no lock
Result:1400, Time taken: 3.913000 ms
./lock bad
Running 30 threads each counting to 50 using bad lock
Result:1465, Time taken: 2.256000 ms
./lock xchg
Running 30 threads each counting to 50 using xchg lock
Result:1500, Time taken: 853.585000 ms
./lock cmpxchg
Running 30 threads each counting to 50 using cmpxchg lock
Result:1500, Time taken: 12997.561000 ms
./lock tts
Running 30 threads each counting to 50 using tts lock
Result:1500, Time taken: 1.779000 ms
./lock backoff
Running 30 threads each counting to 50 using backoff lock
Result:1500, Time taken: 0.939000 ms
./lock mutex
Running 30 threads each counting to 50 using mutex lock
Result:1500, Time taken: 5.313000 ms
```





# Other synchronization primitives

- We may want to have more than one thread/process to execute at same time

## Producer

```
while (1) {  
  
    produce an item;  
  
    lock();  
    insert(item to pool);  
    unlock();  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

# How many producers/consumers can run at a given time?



## Producer

```
while (1) {  
  
    produce an item;  
  
    lock();  
    insert(item to pool);  
    unlock();  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```



# What we want!

- To be more efficient we want to be able to allow **more than one producer/consumer, i.e., equal to the number of items that can be inserted into/removed from the pool**

## Producer

```
while (1) {  
  
    produce an item;  
  
    lock();  
    insert(item to pool);  
    unlock();  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

# Semaphore



A semaphore is like an **integer**, with three differences:

When you create the semaphore, you can initialize its value to any integer, but after that the only operations you are **allowed to perform** are **increment** (increase by one) and **decrement** (decrease by one). *You cannot read the current value of the semaphore.*

When a thread **decrements** the semaphore, if the **result is negative**, the **thread blocks itself** and cannot continue until another thread increments the semaphore.

When a thread **increments** the semaphore, if there are **other threads waiting**, **one of the waiting threads gets unblocked**.

# Semaphore operations



```
wait(S) {  
    while (S<=0) ;  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```



# Producers/consumers using Semaphores

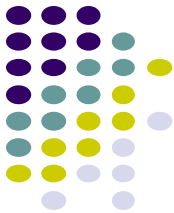
## Producer

```
while (1) {  
  
    produce an item;  
  
    lock();  
    insert(item to pool);  
    unlock();  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

Init: FULL = 0; **EMPTY = N;**



# Producers/consumers using Semaphores

## Producer

```
while (1) {  
  
    produce an item;  
  
    lock();  
    insert(item to pool);  
    unlock();  
    signal(FULL);  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

Init: FULL = 0; **EMPTY = N;**



# Producers/consumers using Semaphores

## Producer

```
while (1) {  
  
    produce an item;  
    wait(EMPTY);  
    lock();  
    insert(item to pool);  
    unlock();  
    signal(FULL);  
}
```

## Consumer

```
While (1) {  
  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

Init: FULL = 0; **EMPTY = N;**





# Producers/consumers using Semaphores

## Producer

```
while (1) {  
  
    produce an item;  
    wait(EMPTY);  
    lock();  
    insert(item to pool);  
    unlock();  
    signal(FULL);  
}
```

## Consumer

```
While (1) {  
  
    wait(FULL);  
    lock();  
    remove(item from pool);  
    unlock();  
  
    consume the item;  
}
```

Init: FULL = 0; **EMPTY = N;**



# Producers/consumers using Semaphores

## Producer

```
while (1) {  
  
    produce an item;  
    wait(EMPTY);  
    lock();  
    insert(item to pool);  
    unlock();  
    signal(FULL);  
}
```

## Consumer

```
While (1) {  
  
    wait(FULL);  
    lock();  
    remove(item from pool);  
    unlock();  
    signal(EMPTY);  
    consume the item;  
}
```

Init: FULL = 0; **EMPTY = N;**

# Is Semaphore good for producers/consumers?



Need to know the size of buffer!

How to accommodate dynamic pool size?