


File System Reliability and Journaling File System

ECE 469, April 21

Aravind Machiry



File System Reliability



- Loss of data in a file system can have catastrophic effect
 - How does it compare to hardware (DRAM) failure?
 - Need to ensure safety against data loss
- Reasons for loss of data:
 - Accidental or malicious deletion of data
 - Media (disk) failure
 - System crash during file system modifications

Handling loss of data



- Accidental or malicious deletion of data?
- Media (disk) failure?
- System crash during file system modifications?

Handling loss of data



- Accidental or malicious deletion of data?
 - **Backup**
- Media (disk) failure?
- System crash during file system modifications?

Backup



- Copy entire file system onto low-cost media (tape), at regular intervals (e.g. once a day).
 - Backup storage (cold storage)
- In the event of a disk failure, replace disk and restore from backup media
- Amount of loss is limited to modifications occurred since last backup

Handling loss of data



- Accidental or malicious deletion of data?
 - Backup
- Media (disk) failure?
 - **Data Replication (e.g., RAID)**
- System crash during file system modifications?

Data Replication



- Full replication
 - Mirroring across disks
 - Full replication to different machines (more next week)
- RAID (next lecture)
- Erasure Coding
 - Like RAID, use parity, but saves more space

Handling loss of data



- Accidental or malicious deletion of data?
 - Backup
- Media (disk) failure?
 - Data Replication (e.g., RAID)
- System crash during file system modifications?
 - **Crash Recovery**

Crash Recovery



- After a system crash in the middle of a file system operation, file system metadata may be in an *inconsistent state*
 - Can we recover from this inconsistent state!?

File Persistence under Buffer Cache/Page Cache



- Problem: file cache memory is **volatile**, but users expect disk files to be **persistent**
 - In the event of a system crash, dirty blocks in the page cache are lost !
 - Example 1: creating “/dir/a”
 - Allocate inode (from free inode list) for “a”
 - Update parent dir content – add (“a”, inode#) to “dir”

File Persistence under Buffer Cache/Page Cache



- Solution 1: use **write-through** cache
 - Modifications are written to disk immediately
 - (minimize “window of opportunities”)
 - No performance advantage for disk writes

File Persistence under Buffer Cache/Page Cache




- Solution 2: **write back** cache
 - Gather (buffer) writes in memory and then write all buffered data back to storage devices
 - e.g., write back dirty blocks after no more than 30 seconds
 - e.g., write back all dirty blocks during file close
- Problem with this?

Many “dirty” blocks in memory: What order to write to disk?



- Example: Appending a new block to existing file
 - Write data bitmap B (for new data block),
write inode I of file (to add new pointer, update time),
write new data block D



Problem




- One file operation may involve modifying multiple disk blocks (and hence multiple disk I/Os)
- After crashing, do we know which blocks were involved at the moment of crashing?



Crash after Bitmap


- Write Ordering: Bitmap (B), Inode (I), Data (D)
 - But CRASH after B has reached disk, before I or D
- Result?






Crash after inode


- Write Ordering: Inode (I), Bitmap (B), Data (D)
 - But CRASH after I has reached disk, before B or D
- Result?





Crash after bitmap and inode


- Write Ordering: Inode (I), Bitmap (B), Data (D)
 - CRASH after I AND B have reached disk, before D
- Result?






Crash after Data

- Write Ordering: Data (D) , Bitmap (B), Inode (I)
 - CRASH after D has reached disk, before I or B
- Result?






Traditional Solution: fsck


- FSCK: “file system checker”
- When system boots:
 - Make multiple passes over file system, looking for inconsistencies
 - e.g., inode pointers and bitmaps, directory entries and inode reference counts
 - Either fix automatically or punt to admin
 - How to recover?



Crash after Bitmap: Can we recover?


- Write Ordering: Bitmap (B), Inode (I), Data (D)
 - But CRASH after B has reached disk, before I or D
- Result?





Crash after inode: Can we recover?

- Write Ordering: Inode (I), Bitmap (B), Data (D)
 - But CRASH after I has reached disk, before B or D
- Result?




Crash after bitmap and inode

Can we recover?



- Write Ordering: Inode (I), Bitmap (B), Data (D)
 - CRASH after I AND B have reached disk, before D
- Result?





Traditional Solution: fsck

- Main problem with fsck: **Performance**
 - Sometimes takes hours to run on large disk volumes



Solution 2: Logging file system updates

- We need to ensure a “copy” of consistent state can always be recovered
- Either the old consistent state (before updates)
- Undo Log
 - Make a copy of the old state to a different place
 - Update the current place
- Or the new consistent state (after updates)
- Redo Log
 - Write to a new place, leave the old place intact



Redo Log

- Idea: Write something down to disk at a different location from the data location
 - Called the “write ahead log” or “journal”
- When all data is written to redo log, write it back to the data location, and then delete the data on redo log
- When crash occurs, look through the redo log and see what was going on
 - Replay complete data, discard incomplete data
 - The process is called “recovery”

Journaling File Systems



- Basic idea
 - update metadata, or all data, *transactionally*
 - “*all or nothing*”
 - *Failure atomicity*
 - if a crash occurs, you may lose a bit of work, but the disk will be in a consistent state
 - more precisely, you will be able to quickly get it to a consistent state by using the transaction log/journal – rather than scanning every disk block and checking sanity conditions

Journaling File Systems



- In file systems with page cache, the data is in two places:
 - On disk
 - In in-memory caches
- The basic idea of the solution:
 - Always leave “home copy” of data in a consistent state
 - Make updates persistent by writing them to a sequential (chronological) **journal** partition/file
 - At your leisure, push the updates (in order) to the home copies and reclaim the journal space
 - Or, make sure log is written before updates

Journal

- Journal: an append-only file containing log records
 - <start t>
 - transaction t has begun



Journal




- Journal: an append-only file containing log records
 - <start t>
 - transaction t has begun
 - <t,x,v>
 - transaction t has updated block x and its new value is v
 - Can log block “diffs” instead of full blocks
 - Can log *operations* instead of data (operations must be idempotent and undoable)

Journal



- Journal: an append-only file containing log records
 - <start t>
 - transaction t has begun
 - <t,x,v>
 - transaction t has updated block x and its new value is v
 - Can log block “diffs” instead of full blocks
 - Can log *operations* instead of data (operations must be idempotent and undoable)
 - <commit t>
 - transaction t has committed – updates will survive a crash
 - Only after the commit block is written is the transaction final
 - The commit block is a single block of data on the disk
- Committing involves writing the records – **the home data doesn't need to be updated at this time**

How does data get out of the journal?



- After a commit the new data is in the journal – it needs to be written back to its home location on the disk
- Cannot reclaim that journal space until we resync the data to disk

Journal Checkpointing



- A cleaner thread walks the journal in order, updating the home locations (on disk, not the cache!) of updates in each transaction
- Once a transaction has been reflected to the home locations, it can be deleted from the journal

Crash Recovery



- Only completed updates have been committed
 - During reboot, the recovery mechanism reapplies the committed transactions in the journal
- The old and updated data are each stored separately, until the commit block is written

If a crash occurs



- Open the log and parse
 - `<start>`, data, `<commit>` => committed transactions
 - `<start>`, ~~`no-<commit>`~~ => uncommitted transactions
- Redo committed transactions
 - Re-execute updates from all committed transactions
 - Aside: note that update (write) is *idempotent*: can be done any positive number of times with the same result.
- Undo uncommitted transactions
 - Undo updates from all uncommitted transactions
 - Write “compensating log records” to avoid work in case we crash during the undo phase

Case Study: Ext3



- Ext3: roughly ext2+journaling
- Ext3 grew out of ext2
- Exact same code base
- Completely backwards compatible (*if you have a clean reboot*)

Ext3 and Journaling



- Two separate layers
 - /fs/ext3 – just the filesystem with transactions
 - /fs/jdb – just the journaling stuff
- ext3 calls jbd as needed
 - Start/stop transaction
 - Ask for a journal recovery after unclean reboot

Ext3 and Journaling



- Journal location
 - EITHER on a separate device partition
 - OR just a “special” file within ext2
- Three separate modes of operation:
 - **Data:** All data is journaled
 - **Ordered, Writeback:** Just metadata is journaled

Data Journaling Mode



- Same example: Update Inode (I), Bitmap (B), Data (D)
- First, write to journal:
 - Transaction begin (Tx begin)
 - Transaction descriptor (info about this Tx)
 - I, B, and D blocks (in this example)
 - Transaction end (Tx end)
- Then, “checkpoint” data to fixed ext2 structures
 - Copy I, B, and D to their fixed file system locations
- Finally, free Tx in journal
 - Journal is fixed-sized circular buffer, entries must be periodically freed

When crash occurs...



- Recovery: Go through log and “redo” operations that have been successfully committed to log
- What if ...
 - Tx begin but not Tx end in log?
 - Tx begin through Tx end are in log, but I, B, and D have not yet been checkpointed?
 - What if Tx is in log, I, B, D have been checkpointed, but Tx has not been freed from log?
- Performance? (As compared to fsck?)

Problem with Data Journaling



- Data journaling: Lots of extra writes
 - All data committed to disk twice (once in journal, once to final location)
- Overkill if only goal is to keep **metadata** consistent

Metadata only journaling: Writeback mode




- **Writeback** mode
 - Just journals metadata
 - Data is not journaled. Writes data to final location directly
 - Better performance than data journaling (data written once)
 - The contents might be written **at any time** (before or after the journal is updated)
- Problems?
 - If a crash happens, metadata can point to old or even garbage data!

Metadata only journaling: Ordered mode



- **Ordered** mode
 - Only metadata is journaled, file contents are not (like writeback mode)
 - But file contents guaranteed to be written to disk before associated metadata is marked as committed in the journal
 - Default ext3 journaling mode



When crash occurs...

- Metadata will only point to correct data (no stale data can be reached after reboot).
- But there *may be data that is not pointed to by any metadata*.
- How is this better than writeback in terms of consistency guarantees?



Conclusions

- Journaling
 - Almost all modern file systems use journaling to reduce recovery time during startup (e.g., Linux ext3/ext4, ReiserFS, SGI XFS, IBM JFS, NTFS)
 - Simple idea: Use write-ahead log to record some info about what you are going to do before doing it
 - Turns multi-write update sequence into a single atomic update (“all or nothing”)
 - Some performance overhead: Extra writes to journal
 - Worth the cost?