



Distributed File Systems and Network File Systems

ECE 469, April 24

Aravind Machiry



Distributed File Systems





Client/Server Model

- **Service** – software entity running on one or more machines and providing a particular type of function to a priori unknown clients
- **Server** – which runs the service software to provide the service
- **Client** – process that can invoke a service using a set of operations that forms its *client interface*
- Traditional DFS uses client/server model

Distributed File Systems (DFS)



- You can login to any instructional machine on campus, your home dir is always there!
- but sometimes it is very slow...

DFS



- Definition: a distributed implementation of the classical time-sharing model of a file system, where multiple users share files and storage resources
- Many DFS have been proposed and developed



Motivation

- Why are distributed file systems useful?
 - Access from multiple clients
 - Same user on different machines can access same files
 - Simplifies sharing
 - Different users on different machines can read/write to same files
 - Simplifies administration
 - One shared server to maintain (and backup)
 - Improve reliability
 - Add RAID storage to server



Challenges

- Transparent access
 - User sees single, global file system regardless of location
- Scalable performance
 - Performance does not degrade as more clients are added
- Fault Tolerance
 - Client and server identify and respond appropriately when other crashes
- Consistency
 - See same directory and file contents on different clients at same time
- Security
 - Secure communication and user authentication
- Tension across these goals
 - Example: Caching helps performance, but hurts consistency

NFS (Network File System)




- First commercially successful distributed file system:
 - Developed in 1984 by Sun Microsystems for their diskless workstations
 - Designed for **robustness** and “**adequate performance**”
 - Multiple versions (**v2**, v3, v4)
 - Widely used today

NFS Overview



- Remote Procedure Calls (RPC) for communication between client and server
- Client Implementation
 - Provides transparent access to NFS file system
 - UNIX contains Virtual File system layer (VFS)
 - Vnode: interface for procedures on an individual file
 - Translates vnode operations to NFS RPCs
- Server Implementation
 - Stateless: Must not have anything only in memory
 - Implication: All modified data written to stable storage before return control to client
 - Servers often add NVRAM to improve performance

NFS Overview





NFS Design Objectives


- ***Machine and Operating System Independence***
 - Could be implemented on low-end machines of the mid-80's
- ***Transparent Access***
 - Remote files should be accessed in exactly the same way as local files
- ***Fast Crash Recovery***
 - Major reason behind stateless design
- ***“Reasonable” performance***
 - Robustness and preservation of UNIX semantics were much more important

Implementation of Transparency



- *“All computer science problems can be solved with an extra level of indirection”*
-- David Wheeler
- What were the earlier manifestations of this in this class?

Client Side





NFS Key Ideas

- NFS key idea #1: **Stateless** server
 - Server not required to remember anything (in memory)
 - Which clients are connected, which files are open, ...
 - Implication: **All client requests have all the information to complete op**
 - Example: Client specifies offset in file to write to
 - Why is this important for fast crash recovery?
- NFS Key idea #2: **Idempotent** server operations
 - Operation can be repeated with same result (no side effects)
 - Example: idempotent: $a=b+1$; Not idempotent: $a=a+1$;
 - Why is this important for crash recovery?



Consequence of Statelessness

- read and write calls must specify offset
 - Server does not keep track of current position in the file
- But user will still use conventional UNIX APIs
- Client must maintain a local offset to be used for read/write.

Advantages of Statelessness



- Crash recovery is very easy:
 - When a server crashes, client just resends request until it gets an answer from the rebooted server
 - Client cannot tell difference between a server that has crashed and recovered and a slow server
- Server state does not grow with more clients
- Simplifies the protocol
 - Client can always repeat any request



Basic NFS Protocol

- Operations at NFS layer (applications do not execute these)
 - lookup(dirfh, name) returns (fh, attributes)
 - Use mount protocol for root directory
 - create(dirfh, name, attr) returns (newfh, attr)
 - remove(dirfh, name) returns (status)
 - read(fh, offset, count) returns (attr, data)
 - write(fh, offset, count, data) returns attr
 - gettattr(fh) returns attr
 - What's missing here?
 - close. Why we do not need it?



Remote Lookup

- Returns a **file handle** instead of a file desc.
 - File handle specifies *unique location* of file
- **lookup(dirfh, name)** *returns (fh, attr)*
 - Returns file handle **fh** and attributes of named file in directory **dirfh**
 - Fails if client has no right to access directory **dirfh**



Remote Lookup

- To lookup `“/usr/joe/6360/list.txt”`

`lookup(rootfh, “usr”) returns (fh0, attr)`

`lookup(fh0, “joe”) returns (fh1, attr)`

`lookup(fh1, “6360”) returns (fh2, attr)`

`lookup(fh2, “list.txt”) returns (fh, attr)`

Mapping UNIX System calls to NFS Operations



- Unix system call: `fd = open("/dir/foo")`
 - Traverse pathname to get filehandle for foo
 - `dir_fh = lookup(root_dir_fh, "dir");`
 - `fh = lookup(dir_fh, "foo");`
 - Record mapping from fd file descriptor to fh NFS filehandle
 - Set initial file offset to 0 for fd
 - Return fd file descriptor

Mapping UNIX System calls to NFS Operations



- Unix system call: `read(fd,buffer,bytes)`
 - Get current file offset for `fd`
 - Map `fd` to `fh` NFS filehandle
 - Call `data = read(fh, offset, bytes)` and copy data into buffer
 - Increment file offset by bytes

Mapping UNIX System calls to NFS Operations



- Unix system call: close(fd)
 - Free resources associated with fd
 - No need to tell server: **stateless server**




Identifying Files in NFS


- Can we still use inode?
- NFS use File handles
- **File handle** consists of
 - ***Filesystem id*** identifying disk partition
 - ***i-node number*** identifying file within partition
 - ***i-node generation number*** changed every time i-node is reused to store a new file

Filesystem id	i-node number	i-node generation number
---------------	---------------	--------------------------


Why not only i-node number?




Why not only i-node number?



Why not only i-node number?




Why not only i-node number?



Performance



- Every read and write requires a network access
- How can we avoid this frequent network access?



Client-Side Caching


- Caching needed to improve performance
 - Reads: Check local cache before going to server
 - Writes: Only periodically write-back data to server
 - Why avoid contacting server
 - Avoid slow communication over network
 - Server becomes scalability bottleneck with more clients
- Two types of client caches
 - data blocks
 - attributes (metadata)

Cache Consistency



- Problem: Consistency across multiple copies (server and multiple clients)
 - How to keep data consistent between client and server?
 - If file is changed on server, will client see update?
 - Determining factor: Read policy on clients
 - How to keep data consistent across clients?
 - If write file on client A and read on client B, will B see update?
 - Determining factor: Write and read policy on clients

Cache Consistency Problem



Inconsistent updates



NFS Consistency: Reads

- Reads: How does client keep current with server state?
 - Attribute cache: Used to determine when file changes
 - File open: Client checks server to see if attributes have changed
 - If haven't checked in past T seconds (configurable, T=3)
 - Discard entries every N seconds (configurable, N=60)
 - Data cache
 - Discard all blocks of file if attributes of the file has been modified



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => ?
 - Client waits 70 seconds
 - Client reads block 2 => ?



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => cache
 - Client waits 70 seconds
 - Client reads block 2 => cache
 - **Block 3 is changed on server**



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => cache
 - Client waits 70 seconds
 - Client reads block 2 => cache
 - **Block 3 is changed on server**
 - Client reads block 3 => ?



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => cache
 - Client waits 70 seconds
 - Client reads block 2 => cache
 - **Block 3 is changed on server**
 - Client reads block 3 => **cache, get old value**
 - Client reads block 4 => ?



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => cache
 - Client waits 70 seconds
 - Client reads block 2 => cache
 - **Block 3 is changed on server**
 - Client reads block 3 => **cache, get old value**
 - Client reads block 4 => fetch from server
 - **Client waits 70 seconds** ————— **Attr changed, all Data blocks are discarded.**



NFS Consistency: Reads

- Eg: Client cache has file A's attributes and blocks 1, 2, 3
 - Client opens A:
 - Client reads block 1 => cache
 - Client waits 70 seconds
 - Client reads block 2 => cache
 - **Block 3 is changed on server**
 - Client reads block 3 => **cache, get old value**
 - Client reads block 4 => fetch from server
 - **Client waits 70 seconds**
 - Client reads block 3 => ?
 - Client reads block 1 => ?



NFS Consistency: Writes

- Writes: How does client update server?
 - Files
 - Write-back from client cache to server every 30 seconds
 - Also, Flush (write all dirty data) on close() (AKA *flush-on-close*)
 - Directories
 - Synchronously write to server (write through)



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2
 - Client Y reads block 1 => cache



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2
 - Client Y reads block 1 => cache
 - 30 seconds later... Data blocks from X are pushed to the server.



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2
 - Client Y reads block 1 => cache
 - 30 seconds later...
 - Client Y reads block 2 => cache, get old value



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2
 - Client Y reads block 1 => cache
 - 30 seconds later...
 - Client Y reads block 2 => cache, get old value
 - 40 seconds later... ————— Attr changed, all Data blocks are discarded.



NFS Consistency: Writes

- Example: Client X and Y have file A (blocks 1,2,3) cached
 - Clients X and Y open file A
 - Client X writes to blocks 1 and 2
 - Client Y reads block 1 => cache
 - 30 seconds later...
 - **Client Y reads block 2 => cache, get old value**
 - 40 seconds later...
 - Client Y reads block 1 => server



Conclusions

- Distributed file systems
 - Important for data sharing
 - Challenges: Fault tolerance, scalable performance, and consistency
- NFS: Popular distributed file system
 - Key features:
 - Stateless server, idempotent operations: Simplifies fault tolerance
 - Crashed server appears as slow server to clients
 - Client caches needed for scalable performance
 - Rules for invalidating cache entries and flushing data to server are not straight-forward
 - Data consistency very hard to reason about