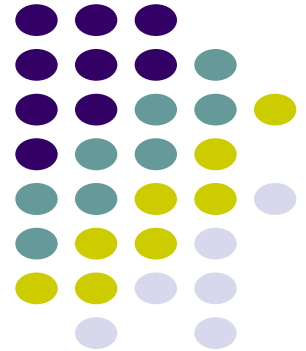


# JOS memory management

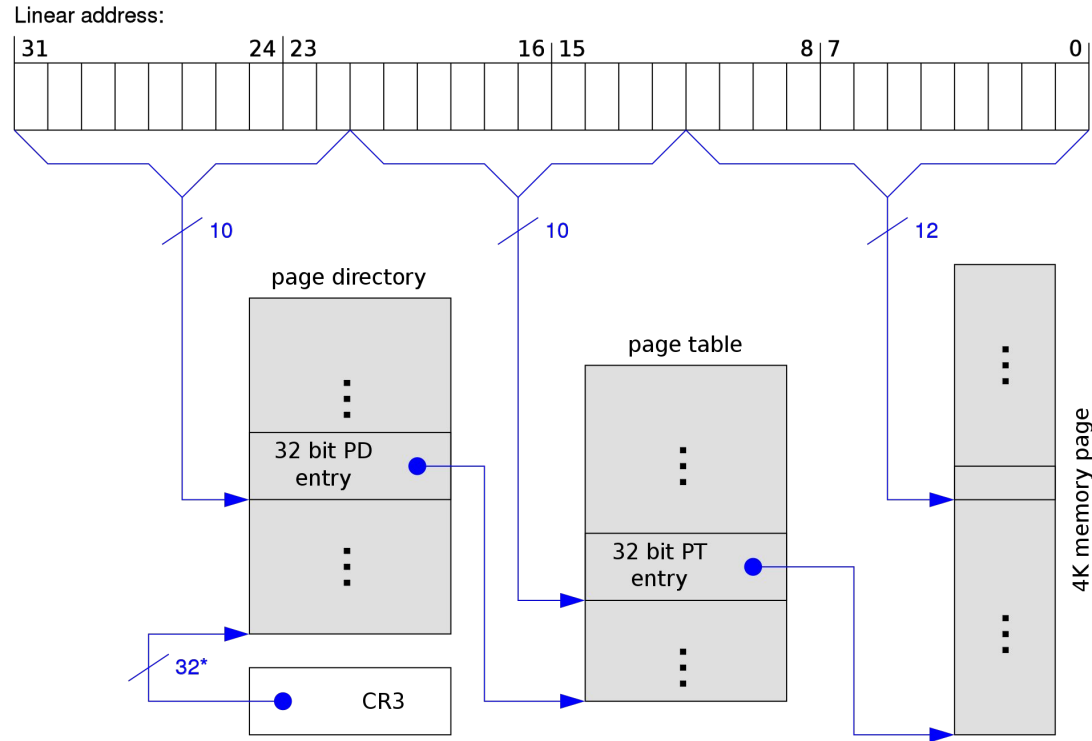
---

ECE 469, Jan 27

Aravind Machiry



# Recap: Two-level paging (32-bit)



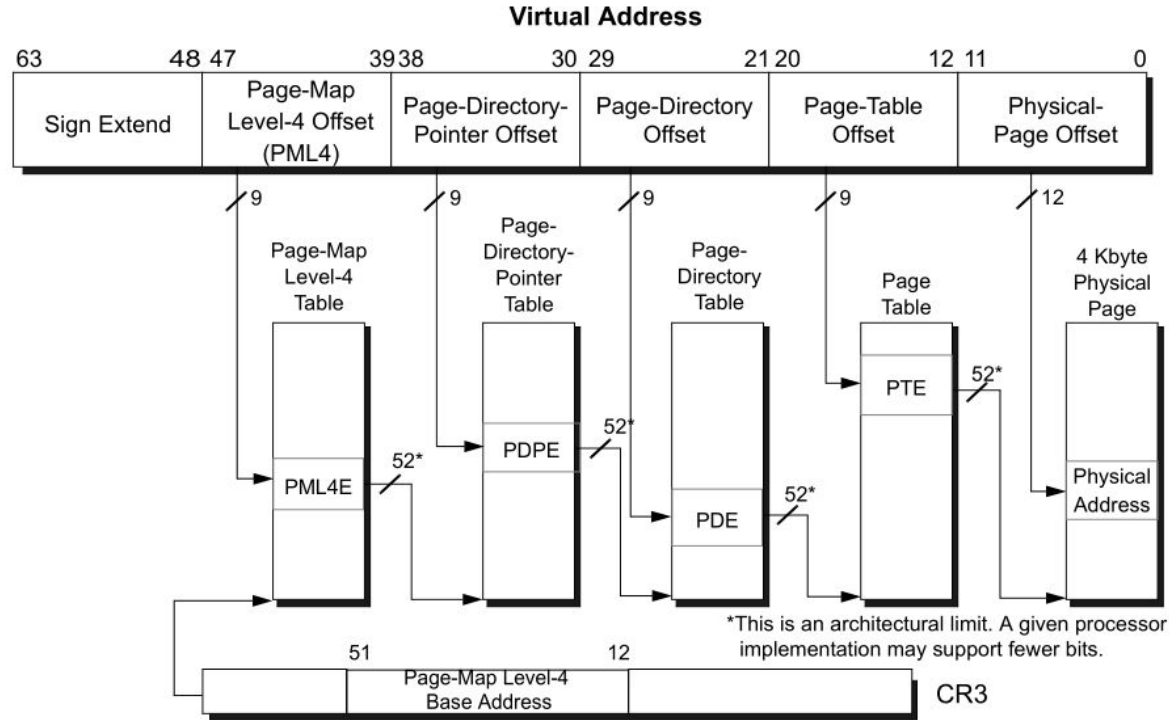
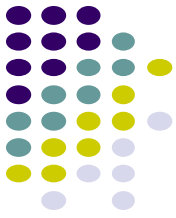
\*) 32 bits aligned to a 4-KByte boundary

# Recap: x86\_64: 48-bit address space



- Initial amd64 processors use only 48-bit virtual address space
  - $2^{48} = 256 \text{ TB}$
- Each level of table tree can process 9 bits (512 entries in table)
- We ignore lower 12 bits
  - $48 - 12 = 36$  (total  $2^{36}$  pages)
  - $36 / 9 = 4$  (each table can process 9 bits of address space)
- **We use 4-level page table**

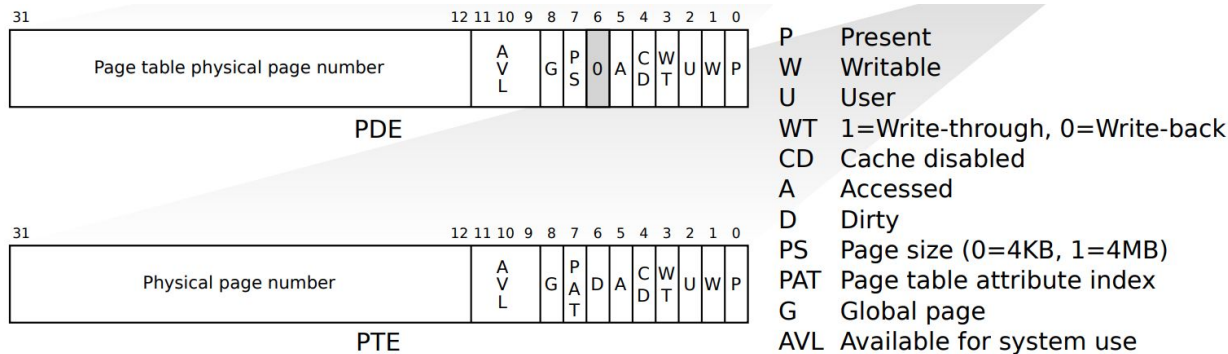
# Recap: Four-level paging (64-bit)



# Recap: Page Directory / Table Entry (PDE/PTE)



- Top 20 bits: physical page number
  - Physical page number of a page table (PDE)
  - Physical page number of the requested virtual address (PTE)
- Lower 12 bits: some flags
  - Permission
  - Etc.



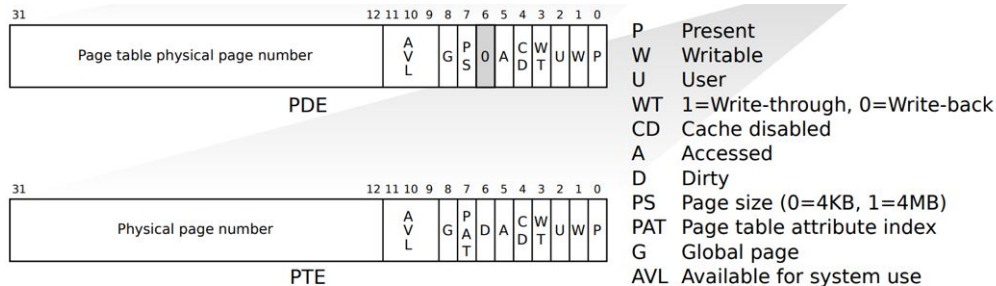


# Recap: Permission Flags

- PTE\_P (PRESENT)
  - 0: invalid entry
  - 1: valid entry
- PTE\_W (WRITABLE)
  - 0: read only
  - 1: writable

	Page Table Entry	
0	Addr PT	
0x48	0x10000 << 12   PTE_U   PTE_W	Invalid
0x49	0x11000 << 12   PTE_P   PTE_W	Kernel, writable
0x4a	0x50000 << 12   PTE_P   PTE_U	User, read-only

- PTE\_U (USER)
  - 0: kernel (only ring 0 can access)
  - 1: user (accessible by ring 3)



# Recap: When CPU Checks Permission Bits?

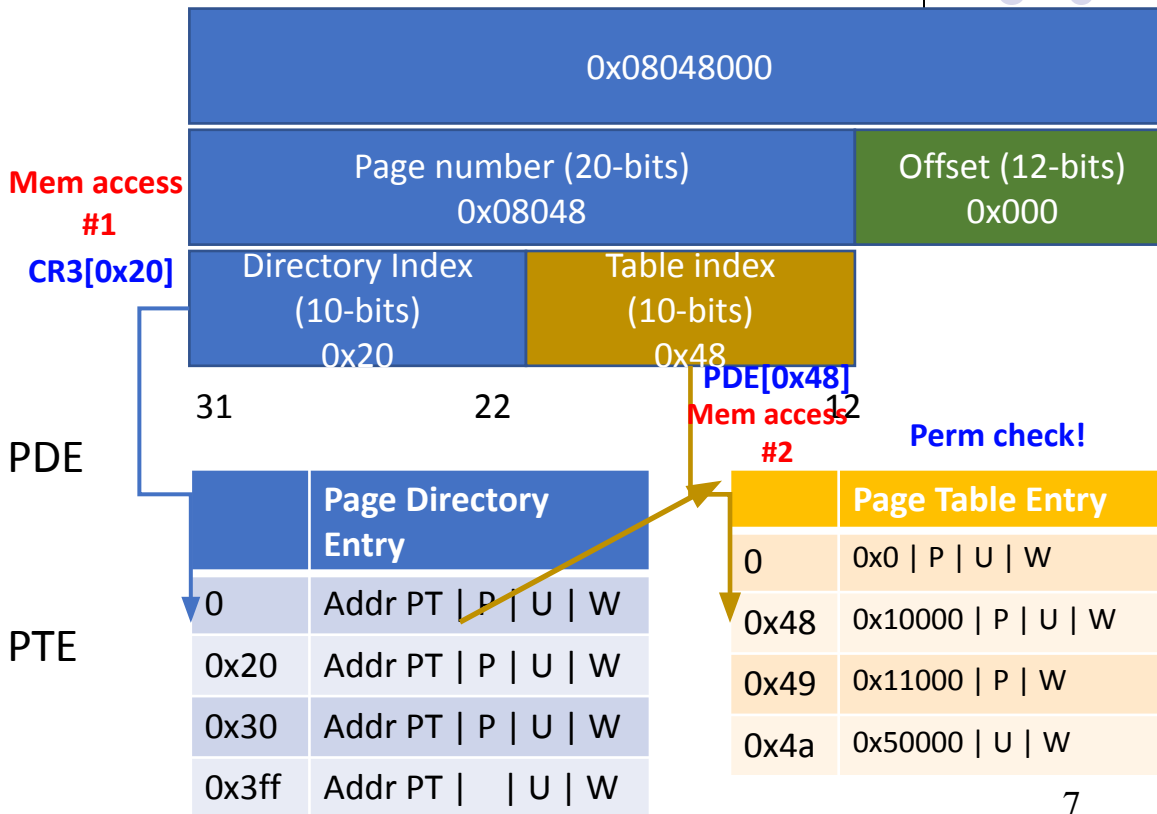


- A virtual memory address is inaccessible if **PDE disallows the access**
- A virtual memory address is inaccessible if **PTE disallows the access**
- Both **PDE and PTE should allow the access...**

# Recap: When CPU Checks Permission Bits?



- In address translation
- 1. Virtual address
- 2.  $PDE = CR3[PDX]$ 
  - Checks permission bits in PDE
- 3.  $PTE = PDE[PTX]$ 
  - Checks permission bits in PTE





# Recap: Allocating Virtual Memory

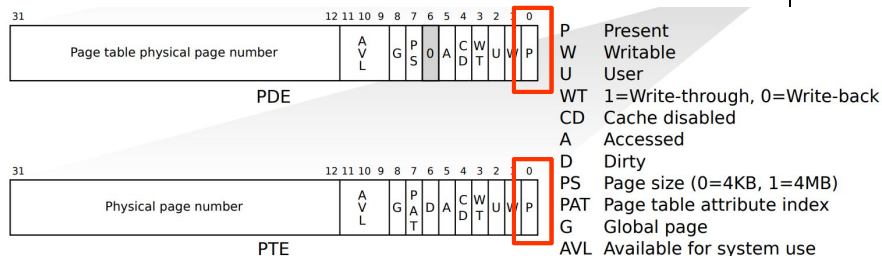


- Static allocation is inefficient:
  - Why don't we just allocate entire virtual address space to a process?
  - **Inefficient**: The process may not access entire virtual address space
- Solution: Dynamic, Request based



# What happens when we call malloc?

- Before malloc()?
  - No PTEs

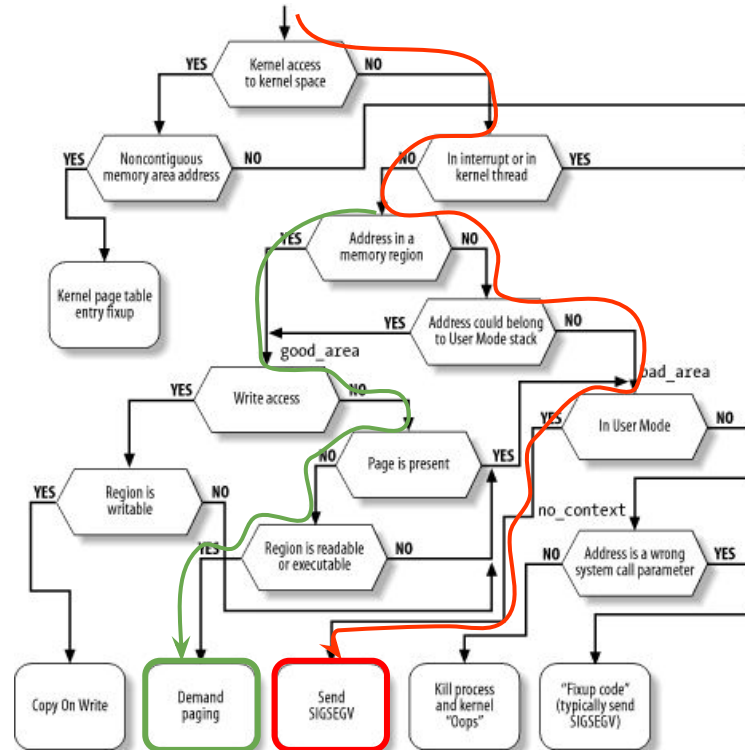


- After malloc()?
  - **PDE/PTE updated but present bit not-set**
- Upon first access?
  - Assign physical page (and page table) and set the valid bit.



# Handling Page faults in Kernel

- Accessing unallocated memory: Demand Paging



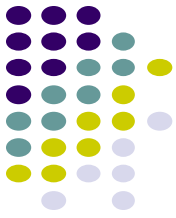
- Accessing Unassigned Virtual Memory: Segmentation Fault

# JOS Memory Management

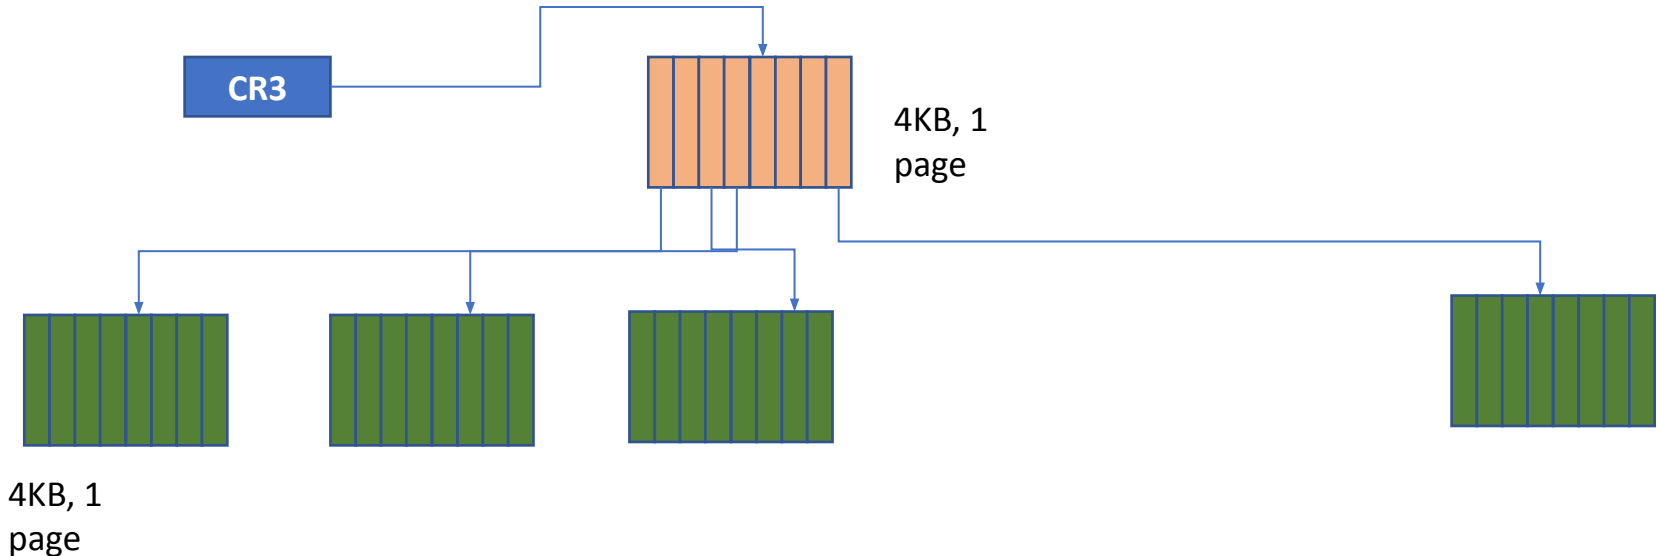


- Handling VA  $\leftrightarrow$  PA mapping
- Managing Physical pages

# Creating a Virtual Memory Space



- A page directory manages the entire virtual memory space
  - Of a process
- CR3 points to the Page directory, and each PDE entry points to PTs..





# Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
  - 0x800000 (RW from user)
- Allocation procedure
  - Check page directory entry (PDE)
    - If not set with PTE\_P, **allocate a physical page** for a new page table

PDE 0: EMPTY
PDE 1: EMPTY
PDE 2: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE 1022: EMPTY
PDE 1023: EMPTY



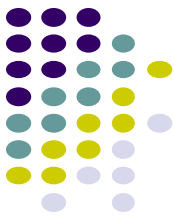
# Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
  - 0x800000 (RW from user)
- Allocation procedure
  - Check page directory entry (PDE)
    - If not set with PTE\_P, **allocate a physical page** for a new page table

PDE 0: EMPTY
PDE 1: <b>0x11223</b>
PDE 2: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE ...: EMPTY
PDE 1022: EMPTY
PDE 1023: EMPTY

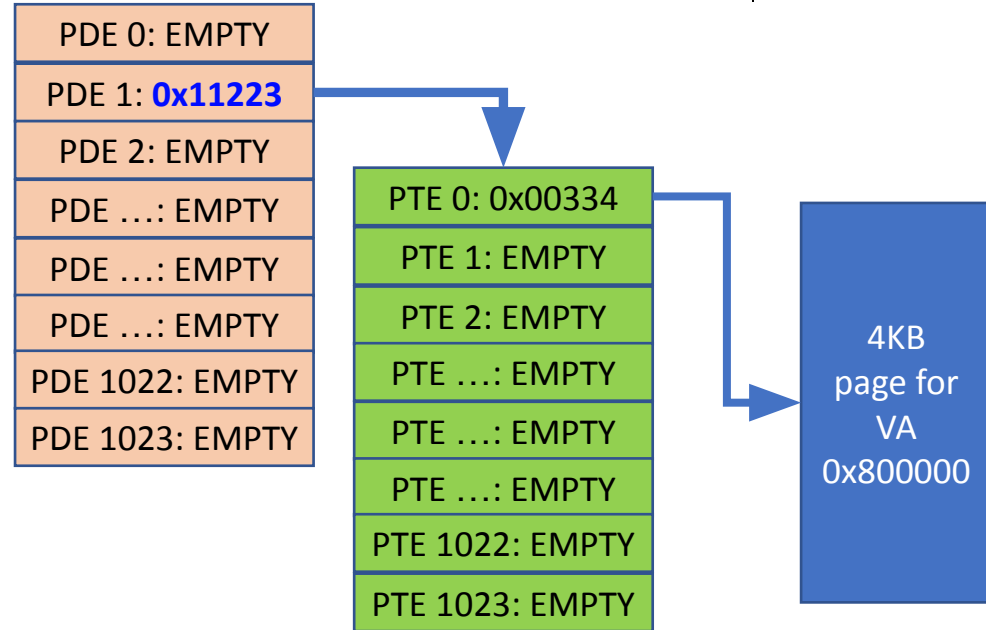


PTE 0: EMPTY
PTE 1: EMPTY
PTE 2: EMPTY
PTE ...: EMPTY
PTE ...: EMPTY
PTE ...: EMPTY
PTE 1022: EMPTY
PTE 1023: EMPTY



# Assigning VA -> PA mapping

- Suppose a process would like to use a virtual address
  - 0x800000 (RW from user)
- Allocation procedure
  - Check page directory entry (PDE)
    - If not set with PTE\_P, **allocate a physical page** for a new page table
  - Check page table entry (PTE)
    - If not set with PTE\_P, **allocate a physical page** to enable access





# Assigning VA -> PA mapping



- Suppose a process would like to use a virtual address
  - 0x800000 (RW from user)
- Allocation procedure
  - Check page directory entry (PDE)
    - If not set with PTE\_P, **allocate a physical page** for a new page table
  - Check page table entry (PTE)
    - If not set with PTE\_P, **allocate a physical page** to enable access
- We need to keep track of **'free'** physical pages...

# Managing Physical memory in JOS



- Struct PageInfo
  - A metadata type that counts number of 'references' of the page
  - NOT IN USE : pp\_ref == 0
- Struct PageInfo \* page\_free\_list
  - A linked list that contains free physical pages
- We will create Struct PageInfo per each Physical Page and then
  - Create a linked list of free pages...

# Struct PageInfo in JOS



- A **one-to-one** mapping from a **struct PageInfo** to a physical page
  - An 8 byte struct per each physical memory page
  - If we support 128MB memory, then we will create
    - Total number of physical pages:  $128 * 1048576 / 4096 = 32768$
  - Total size = **32768** \* 8 = 262,144 = 256KB
- A linked-list for managing free physical pages
  - Starting from `page_free_list->pp_link...`
- `pp_ref`
  - Count references
  - Non-zero – in-use
  - Zero – free

```
struct PageInfo {
    // Next page on the free list.
    struct PageInfo *pp_link;

    // pp_ref is the count of pointers (usually in page table entries)
    // to this page, for pages allocated using page_alloc.
    // Pages allocated at boot time using pmap.c's
    // boot_alloc do not have valid reference count fields.

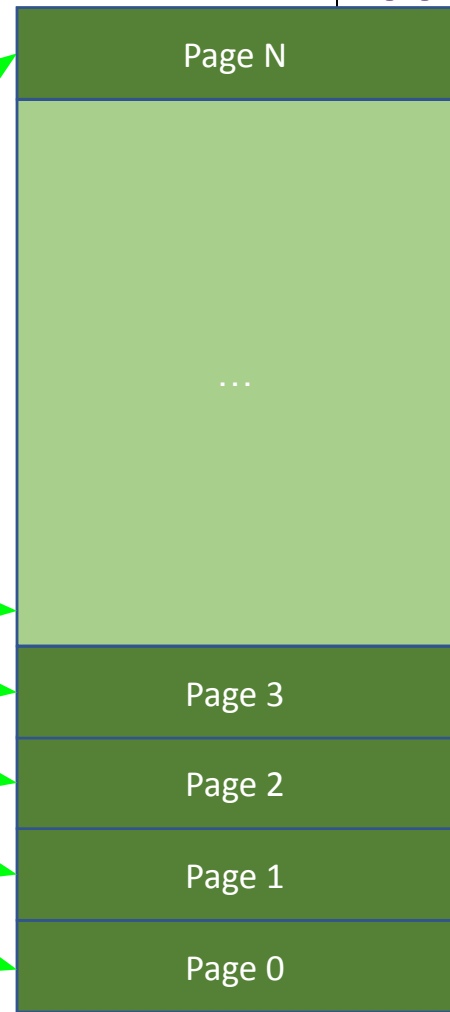
    uint16_t pp_ref;
};
```

# Struct PageInfo

```
struct PageInfo *pages; // Physical page state array
```

Struct PageInfo \* pages (array)

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	



# Struct PageInfo (128MB)

$128 * 1048576 / 4096 = 32768$  Pages

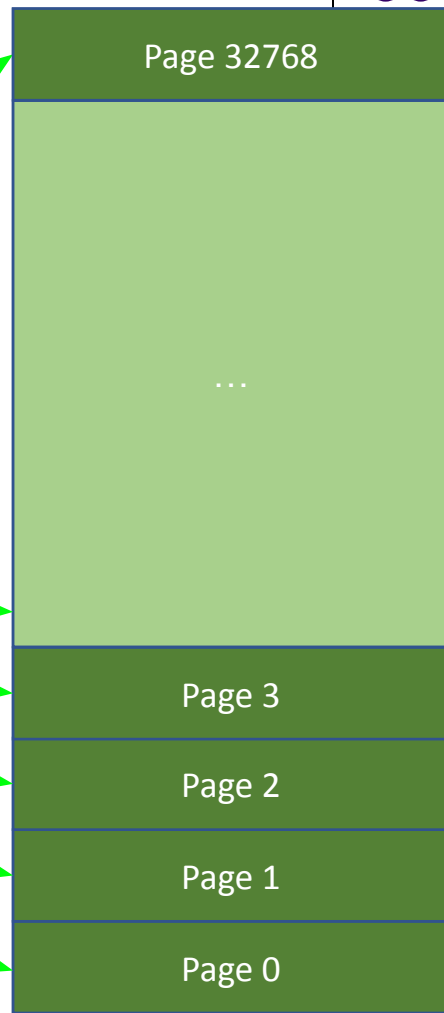
8 byte per each entry =  $32K * 8 = 256KB$

Struct PageInfo \* pages (array)

idx	pp_ref	pp_link
32K	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

We can put this array into our physical memory

Physical memory (128MB)



# Free Physical Memory (init)

In kern/pmap.c, boot\_alloc

```
static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did *not* assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }
}
```

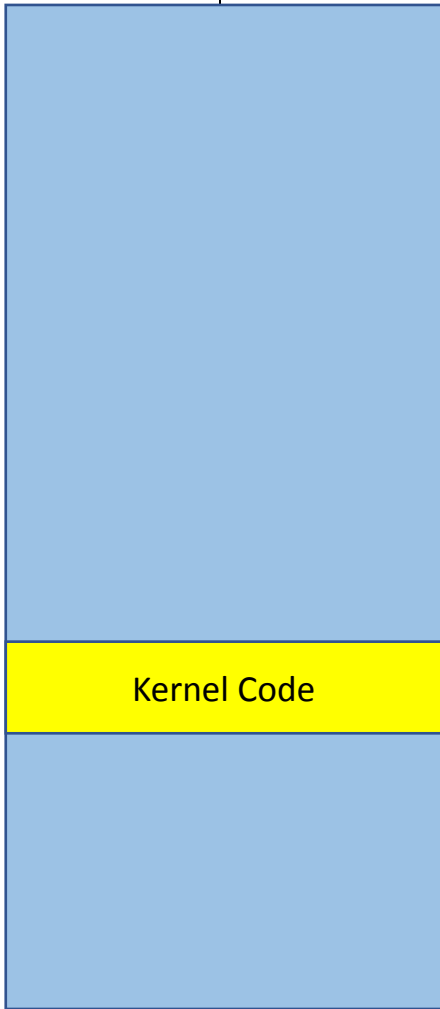
nextfree will point to the end of the kernel code/data

nextfree end

0x100000

Physical memory

Kernel Code



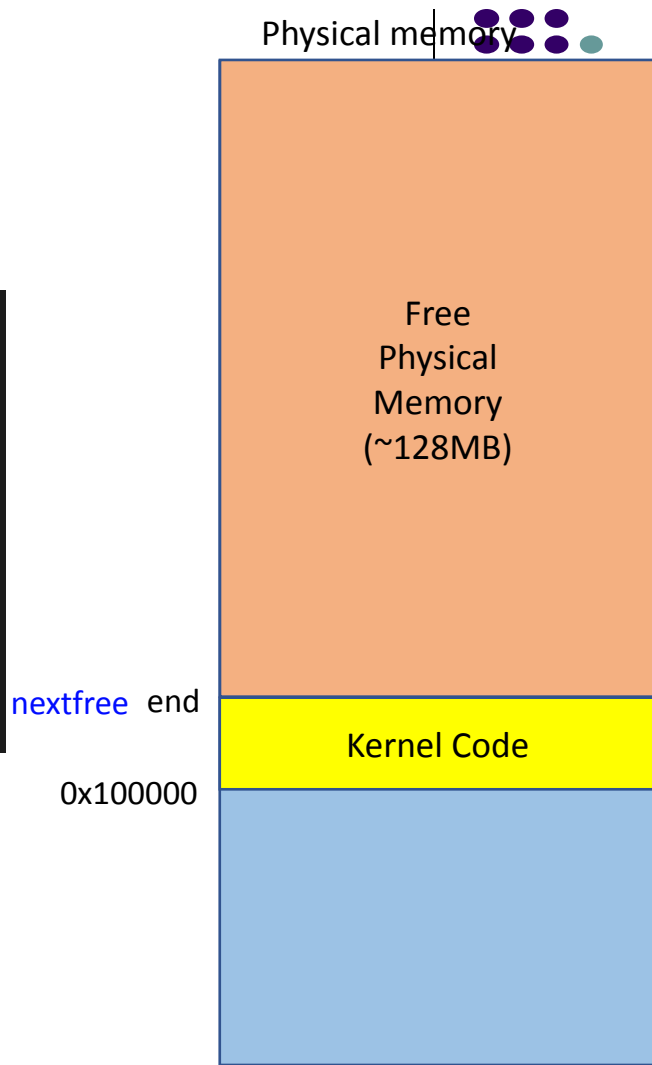
# Free Physical Memory (init)

In kern/pmap.c, boot\_alloc

```
static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual address of next byte of free memory
    char *result;

    // Initialize nextfree if this is the first time.
    // 'end' is a magic symbol automatically generated by the linker,
    // which points to the end of the kernel's bss segment:
    // the first virtual address that the linker did *not* assign
    // to any kernel code or global variables.
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *) end, PGSIZE);
    }
}
```

nextfree will point to the end of the kernel code/data



# Allocating struct PageInfo

```
// These variables are set in mem_init()
pde_t *kern_pgdir;    // Kernel's initial page directory
struct PageInfo *pages; // Physical page state array
static struct PageInfo *page_free_list; // Free list of physical pages
```

```
////////////////////////////////////
// Allocate an array of npages 'struct PageInfo's and store it in 'pages'.
// The kernel uses this array to keep track of physical pages: for
// each physical page, there is a corresponding struct PageInfo in this
// array. 'npages' is the number of physical pages in memory. Use memset
// to initialize all fields of each struct PageInfo to 0.
// Your code goes here:
```

```
pages =
boot_alloc(npages * sizeof(struct PageInfo));
```

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

Physical page N

Physical page 2

Physical page 1

Physical page 0

nextfree

end

0x100000

0x7fff000

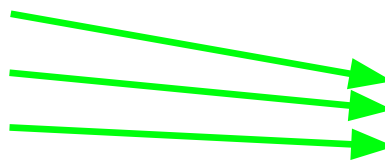
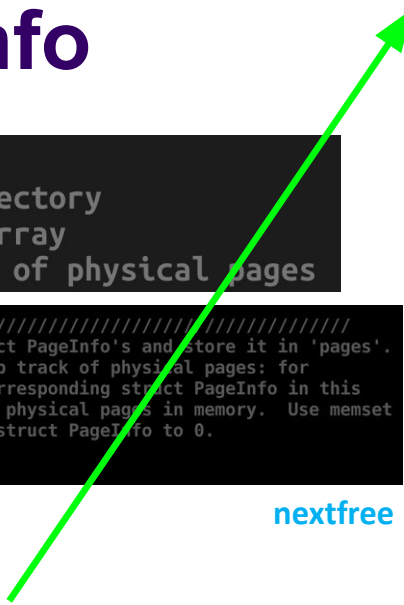
Free  
Physical  
Memory

struct PageInfo \* pages

Kernel Code

0x2000  
0x1000  
0x0000

Physical memory

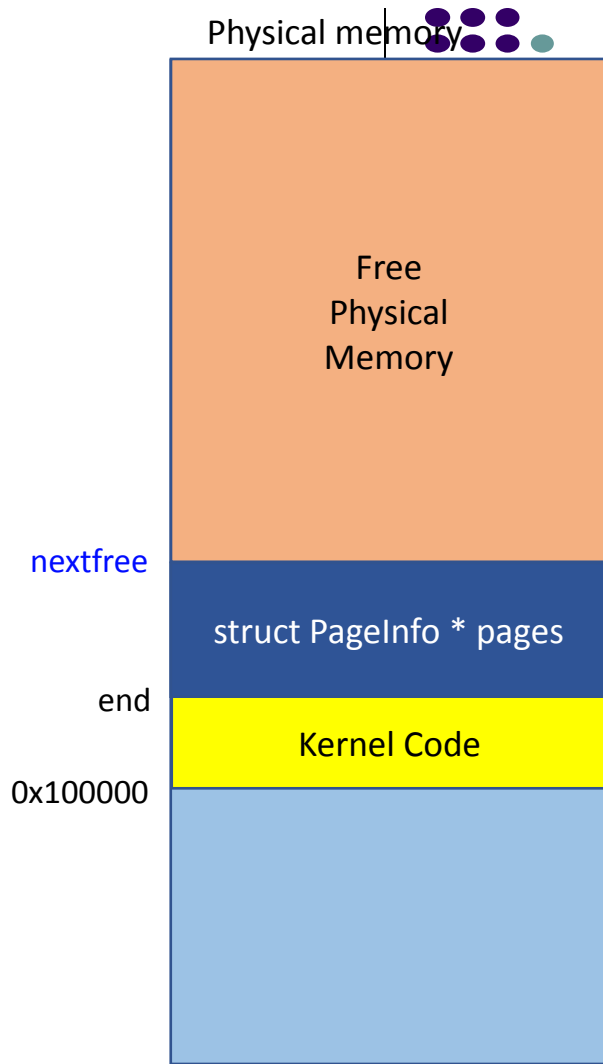




# free pages!

- in `page_init()`

```
// The example code here marks all physical pages as free.  
// However this is not truly the case. What memory is free?  
// 1) Mark physical page 0 as in use.  
// This way we preserve the real-mode IDT and BIOS structures  
// in case we ever need them. (Currently we don't, but...)  
// 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)  
// is free.  
// 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must  
// never be allocated.  
// 4) Then extended memory [EXTPHYSMEM, ...).  
// Some of it is in use, some is free. Where is the kernel  
// in physical memory? Which pages are already in use for  
// page tables and other data structures?  
//  
// Change the code to reflect this.  
// NB: DO NOT actually touch the physical memory corresponding to  
// free pages!
```

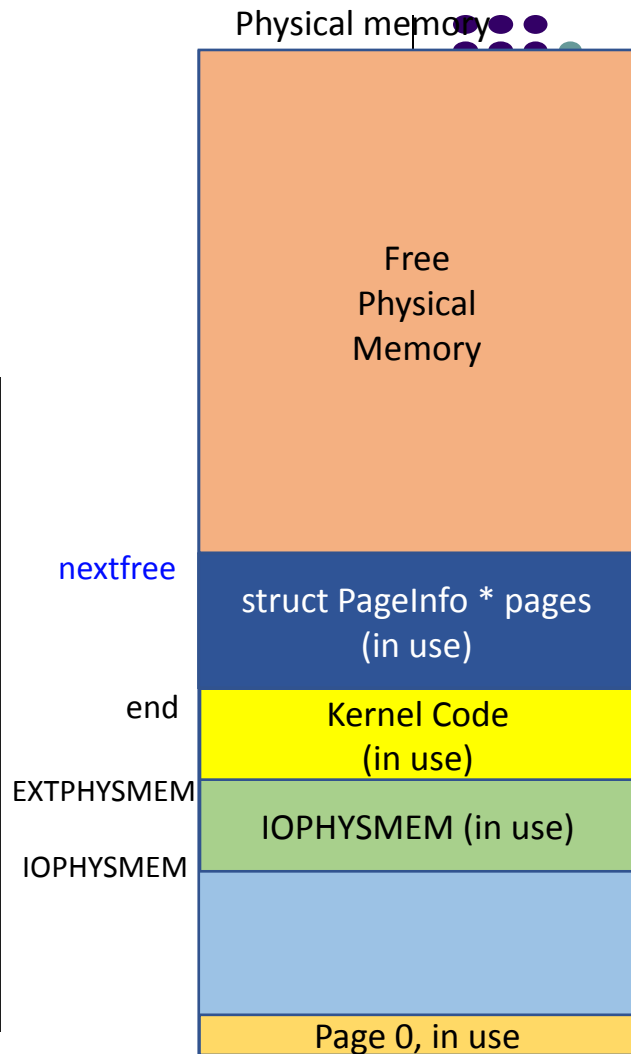


# free pages!

- in `page_init()`

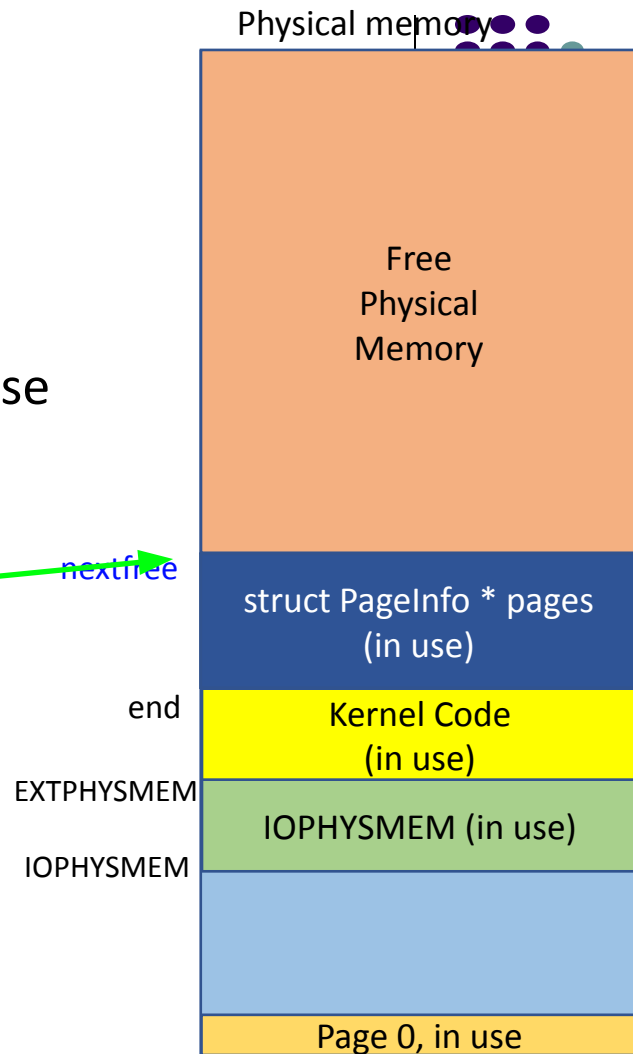
```
// The example code here marks all physical pages as free.
// However this is not truly the case. What memory is free?
// 1) Mark physical page 0 as in use.
// This way we preserve the real-mode IDT and BIOS structures
// in case we ever need them. (Currently we don't, but...)
// 2) The rest of base memory, [PGSIZE, npages_basemem * PGSIZE)
// is free.
// 3) Then comes the IO hole [IOPHYSMEM, EXTPHYSMEM), which must
// never be allocated.
// 4) Then extended memory [EXTPHYSMEM, ...).
// Some of it is in use, some is free. Where is the kernel
// in physical memory? Which pages are already in use for
// page tables and other data structures?
//
// Change the code to reflect this.
// NB: DO NOT actually touch the physical memory corresponding to
// free pages!
```

25



# free pages!

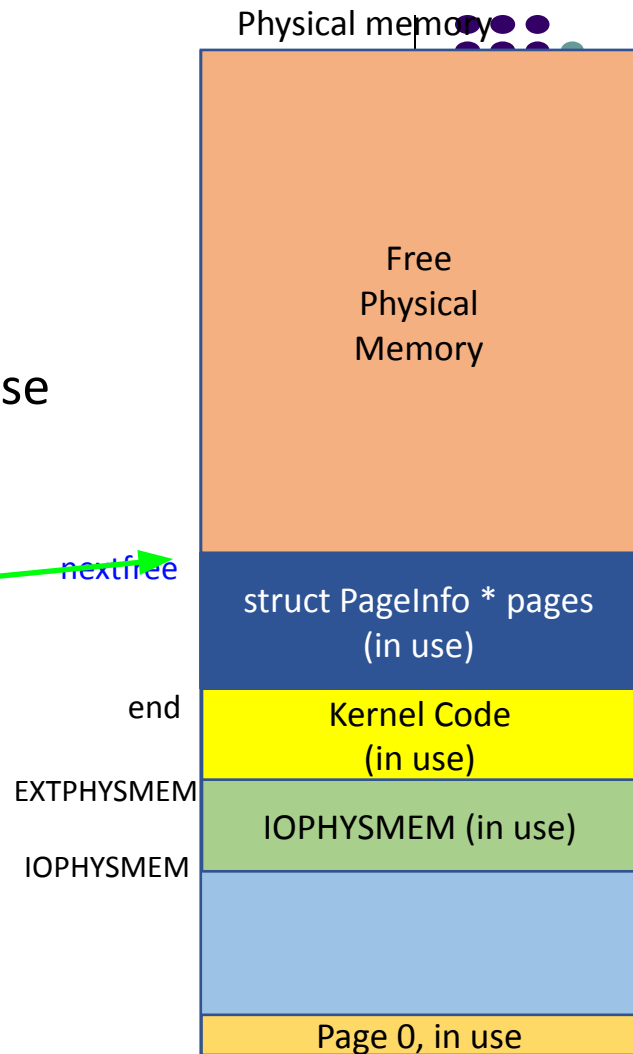
- Page 0 is in-use
- Pages in [IOPHYSMEM ~ EXTPHYSMEM] are in-use
- Pages for the kernel code are in-use
- Pages for struct PageInfo \*pages are in-use
- How can you point this?
  - pages + npages ?
  - boot\_alloc(0)?



# free pages!

- Page 0 is in-use
- Pages in [IOPHYSMEM ~ EXTPHYSMEM] are in-use
- Pages for the kernel code are in-use
- Pages for struct PageInfo \*pages are in-use
- How can you point this?
  - pages + npages ?
  - boot\_alloc(0)?

**boot\_alloc(0) is  
better...**



# Reference counting: Knowing when a page is free



- A typical mechanism for tracking free memory blocks
- Mechanism
  - Count up the value (`pp_ref++`) if the page is referenced by others (in use!)
  - Count down the value (`pp_ref--`) if not used for one of usages anymore
  - Free if `pp_ref == 0`
- In C++, `shared_ptr<T>`
  - When a pointer is assigned to a variable, count up!
  - When the variable no longer uses the variable, count down!
  - Free the memory when the count become 0



# Reference counting with `pp_ref`

- For in-use memory
  - Set `pp_ref = 1`
- For not-in-use memory
  - Invariant: `pp_ref == 0`
  - **Must be linked with `pages_free_list`**
- When assigning the page to a virtual address
  - `pp_ref++`
- When releasing the page from a virtual address
  - `pp_ref--`

# Allocating struct PageInfo

```
// These variables are set in mem_init()
pde_t *kern_pgdir;    // Kernel's initial page directory
struct PageInfo *pages; // Physical page state array
static struct PageInfo *page_free_list; // Free list of physical pages
```

```
////////////////////////////////////
// Allocate an array of npages 'struct PageInfo's and store it in 'pages'.
// The kernel uses this array to keep track of physical pages: for
// each physical page, there is a corresponding struct PageInfo in this
// array. 'npages' is the number of physical pages in memory. Use memset
// to initialize all fields of each struct PageInfo to 0.
// Your code goes here:
```

```
pages =
boot_alloc(npages * sizeof(struct PageInfo));
```

idx	pp_ref	pp_link
N	0	
...	0	
...	0	
3	0	
2	0	
1	0	
0	0	

Physical page N

Physical page 2

Physical page 1

Physical page 0

nextfree

end

0x100000

0x7fff000

Free  
Physical  
Memory

struct PageInfo \* pages

Kernel Code

0x2000  
0x1000  
0x0000

Physical memory

# Linked list of free pages



- Start with NULL at the head
  - `page_free_list = NULL;`
- After set `pp_ref` of all pages, do something like the following..

This will build a linked list...

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

page\_free\_list  NULL

# Building free list

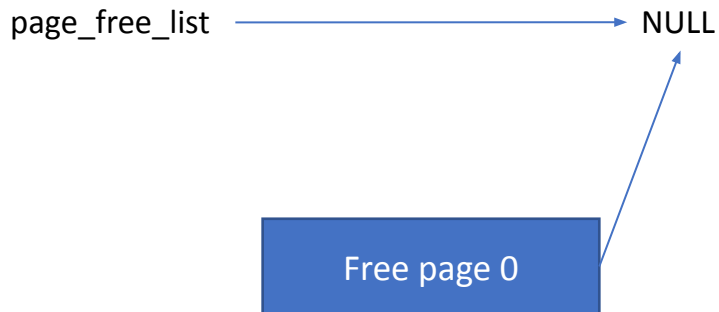
```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

page\_free\_list → NULL

Free page 0

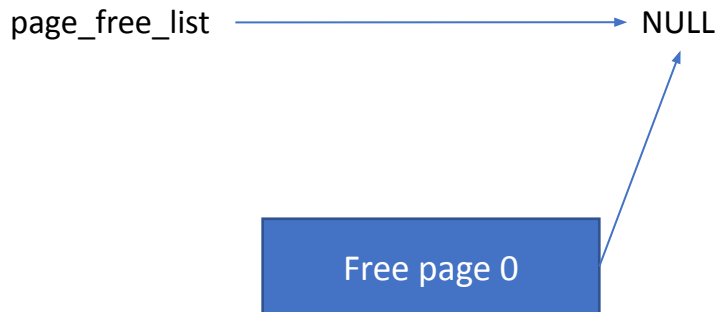
# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    → if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



Free page 1

# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```

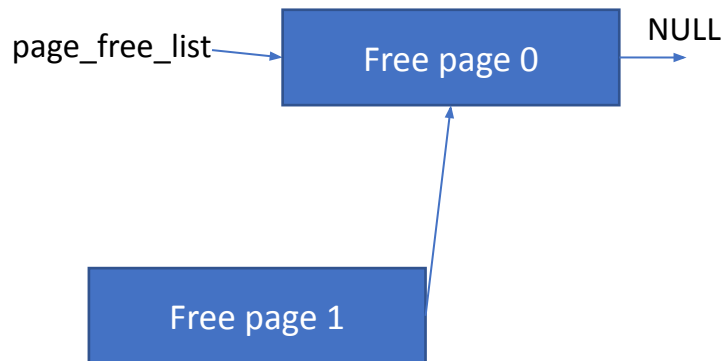


Free page 1



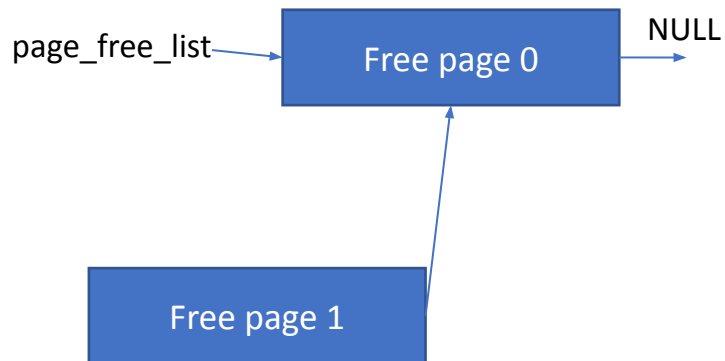
# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        → pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



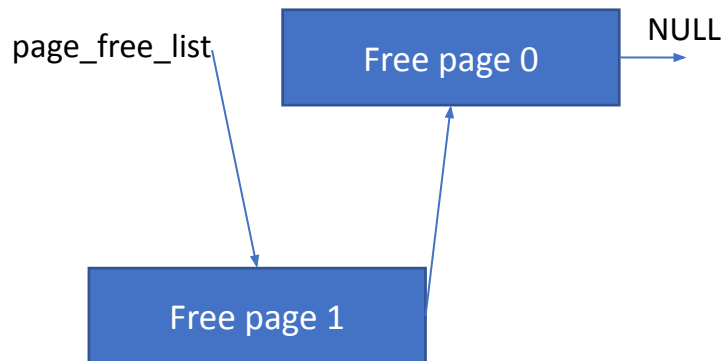
# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        → page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



# Building free list

```
for (int i=0; i < npages; ++i) {  
    if (pages[i].pp_ref == 0) {  
        pages[i].pp_link = page_free_list;  
        page_free_list = &pages[i];  
    }  
}
```



# page2pa(struct PageInfo \*pp)



- Changes a pointer to **struct PageInfo** to a physical address
- $idx = (pp - pages)$ 
  - Gets the index of pp in pages
  - E.g.,  $\&pages[idx] == pp$
- idx here is a physical page number

	idx	pp_ref	pp_link
pp →	4	0	
	3	0	
	2	0	
	1	0	
pages →	0	0	

```
static inline physaddr_t
page2pa(struct PageInfo *pp)
{
    return (pp - pages) << PGSHIFT;
}
```

$pp - pages = 4$   
 $0x4000$  □ physical page  
address!

# pa2page(physaddr\_t pa)



- PGNUM(pa)
  - Returns page number
- &pages[PGNUM(pa)]
  - Returns struct PageInfo \* of that pa..

```
static inline struct PageInfo*
pa2page(physaddr_t pa)
{
    if (PGNUM(pa) >= npages)
        panic("pa2page called with invalid pa");
    return &pages[PGNUM(pa)];
}
```

idx	pp_ref	pp_link
4	0	
3	0	
2	0	
1	0	
0	0	

# Sample Qs



- Which one of the following is not a job that JOS Bootloader does?
  - A. Enable protected mode
  - B. Enable paging
  - C. Load kernel image from disk
  - D. Enable A20



# Sample Qs



- Which one of the following is not a job that JOS Bootloader does?
  - A. Enable protected mode
  - B. Enable paging (is done in kernel, in kern/entry.S)
  - C. Load kernel image from disk
  - D. Enable A20

# Sample Qs



- In the x86 real mode, which address the following segment:offset pair points to?
- 0x8000:0x3131
  - A. 0xb131
  - B. 0x3131
  - C. 0x83131
  - D. 0x103131
  - E. 0x11131

# Sample Qs



- In the x86 real mode, which address the following segment:offset pair points to?
- 0x8000:0x3131
  - A. 0xb131
  - B. 0x3131
  - C. 0x83131 ( $0x8000 * 16 + 0x3131 = 0x80000 + 0x3131 = 0x83131$ )
  - D. 0x103131
  - E. 0x11131

# Sample Qs



- Which of the following x86 register stores the current privilege level?
  - A. ds
  - B. eip
  - C. ebp
  - D. esp
  - E. cs

# Sample Qs



- Which of the following x86 register stores the end of the current stack frame (and moves if the CPU runs push/pop) ?
  - A. ds
  - B. eip
  - C. ebp
  - D. esp
  - E. cs

# Sample Qs



- Which of the following x86 register stores the start of the current stack frame (also points to the address that stores previous frame's stack base pointer) ?
  - A. ds
  - B. eip
  - C. ebp
  - D. esp
  - E. cs

# Sample Qs



- What kind of benefit can we enjoy by enabling virtual memory?
- Choose all (no partial credits)
  - A. Performs faster execution than when using physical memory
  - B. Suffers less memory fragmentation than when using physical memory
  - C. Provides a better isolation / protection than when using physical memory
  - D. Provides memory transparency
  - E. Enables virtual reality