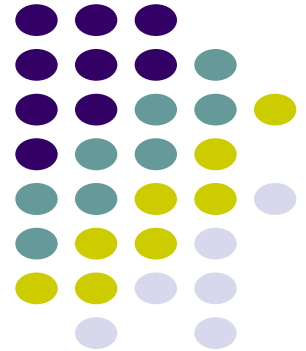


# System calls and Page faults

---

ECE 469, Feb 10

Aravind Machiry



# Recap: Interrupts

- Hardware Interrupts
- Software Interrupts



# Recap: Hardware Interrupts



- A way of hardware interacting with CPU
- Example: a network device
  - NIC: “Hey, CPU, I have a packet received for the OS, so please wake up the OS to handle the data”
  - CPU: call the interrupt handler for network device in ring 0 (set by the OS)
- Asynchronous (can happen at any time of execution)
  - It’s a request from a hardware, so it comes at any time of CPU’s execution
- Read
  - [https://en.wikipedia.org/wiki/Intel\\_8259](https://en.wikipedia.org/wiki/Intel_8259)
  - [https://en.wikipedia.org/wiki/Advanced\\_Programmable Interrupt\\_Controller](https://en.wikipedia.org/wiki/Advanced_Programmable Interrupt_Controller)

# Recap: Software Interrupts / exceptions



- A software mean to run code in ring 0 (e.g., int \$0x30 )
  - Telling CPU that "Please run the interrupt handler at 0x30"
- Synchronous (caused by running an instruction, e.g., int \$0x30)
- System call
  - int \$0x30 □ system call in JOS

# Recap: Types of exceptions

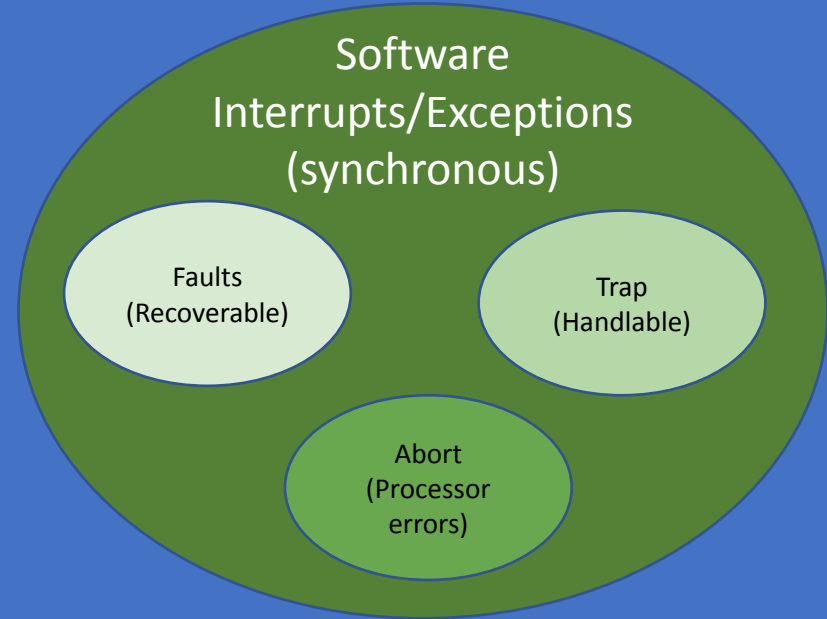


- Classification based on how they are handled:
  - Fault
    - Exception occurred but can be fixed
    - IP points to the current instruction
  - Trap
    - Exception occurred but the program could continue execution
    - IP points to next instruction
  - Abort
    - Unhandleable exception
    - Hardware failures in processor

# Recap: Interrupts classification



## Interrupts



# Recap: Handling Interrupts



- Setting an Interrupt Descriptor Table (IDT)

Interrupt Number	Code address
0 (Divide error)	0xf0130304
1 (Debug)	0xf0153333
2 (NMI, Non-maskable Interrupt)	0xf0183273
3 (Breakpoint)	0xf0223933
4 (Overflow)	0xf0333333
...	
8 (Double Fault)	0xf0222293
...	
14 (Page Fault)	0xf0133390
...	...
0x30 (syscall in JOS)	0xf0222222

# Recap: Handling Interrupts



- Setting an Interrupt Descriptor Table (IDT)

Interrupt Number	Code address
0 (Divide error)	0xf0130304
1 (Debug)	0xf0153333
2 (NMI, Non-maskable Interrupt)	0xf0183273
3 (Breakpoint)	0xf0223933
4 (Overflow)	0xf0333333
...	
8 (Double Fault)	0xf0222293
...	
14 (Page Fault)	0xf0133390
...	...
0x30 (syscall in JOS)	0xf0222222

Load the base address into IDTR




Figure 6-1. Relationship of the IDTR and IDT



# Recap: Handling Interrupts



- Setting an Interrupt Descriptor Table (IDT)

Interrupt Number	Code address
0 (Divide error)	t_divide
1 (Debug)	t_debug
2 (NMI, Non-maskable Interrupt)	t_nmi
3 (Breakpoint)	t_brkpt
4 (Overflow)	t_oflow
...	
8 (Double Fault)	t_dblflt
...	
14 (Page Fault)	t_pgflt
...	...
0x30 (syscall in JOS)	t_syscall

```
TRAPHANDLER_NOEC(t_divide, T_DIVIDE);    // 0
TRAPHANDLER_NOEC(t_debug, T_DEBUG);      // 1
TRAPHANDLER_NOEC(t_nmi, T_NMI);          // 2
TRAPHANDLER_NOEC(t_brkpt, T_BRKPT);      // 3
TRAPHANDLER_NOEC(t_oflow, T_OFLOW);      // 4
TRAPHANDLER_NOEC(t_bound, T_BOUND);      // 5
TRAPHANDLER_NOEC(t_illop, T_ILLOP);      // 6
TRAPHANDLER_NOEC(t_device, T_DEVICE);    // 7

TRAPHANDLER(t_dblflt, T_DBLFLT);          // 8

TRAPHANDLER(t_tss, T_TSS);                // 10
TRAPHANDLER(t_segnp, T_SEGNP);            // 11
TRAPHANDLER(t_stack, T_STACK);            // 12
TRAPHANDLER(t_gpflt, T_GPFLT);            // 13
TRAPHANDLER(t_pgflt, T_PGFLT);            // 14

TRAPHANDLER_NOEC(t_fperr, T_FPERR);       // 16

TRAPHANDLER(t_align, T_ALIGN);            // 17

TRAPHANDLER_NOEC(t_mchk, T_MCHK);          // 18
TRAPHANDLER_NOEC(t_simderr, T_SIMDERR);   // 19

TRAPHANDLER_NOEC(t_syscall, T_SYSCALL);   // 48, 0x30
```

# Recap: JOS Interrupt Handling

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

- Setup the IDT at trap\_init() in kern/trap.c

# Recap: JOS Interrupt Handling

- Setup the IDT at trap\_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call \_alltraps (in kern/trapentry.S)

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}

#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

# Recap: JOS Interrupt Handling

- Setup the IDT at trap\_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call \_alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds
    pushl %es
    pushal
```

**Build a Trapframe!**

# Recap: JOS Interrupt Handling

```
struct Trapframe {
    struct PushRegs tf_regs;
    uint16_t tf_es;
    uint16_t tf_padding1;
    uint16_t tf_ds;
    uint16_t tf_padding2;
    uint32_t tf_trapno;
    /* below here defined by x86 hardware */
    uint32_t tf_err;
    uintptr_t tf_eip;
    uint16_t tf_cs;
    uint16_t tf_padding3;
    uint32_t tf_eflags;
    /* below here only when crossing rings, such as from user to kernel */
    uintptr_t tf_esp;
    uint16_t tf_ss;
    uint16_t tf_padding4;
} __attribute__((packed));
```

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds
    pushl %es
    pushal
```

**Build a  
Trapframe!**

# Recap: JOS Interrupt Handling

- Setup the IDT at trap\_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call \_alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds    Build a
    pushl %es    Trapframe!
    pushal
```

```
void
trap(struct Trapframe *tf)
{
```

# Recap: JOS Interrupt Handling

- Setup the IDT at trap\_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call \_alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c
- Call trap\_dispatch() in kern/trap.c

```
static void
trap_dispatch(struct Trapframe *tf)
{
    // Handle processor exceptions.
    // LAB 3: Your code here.
```

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);

#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds    Build a
    pushl %es    Trapframe!
    pushal
```

```
void
trap(struct Trapframe *tf)
{
```


# Today

- Syscalls
- Page fault





# Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```

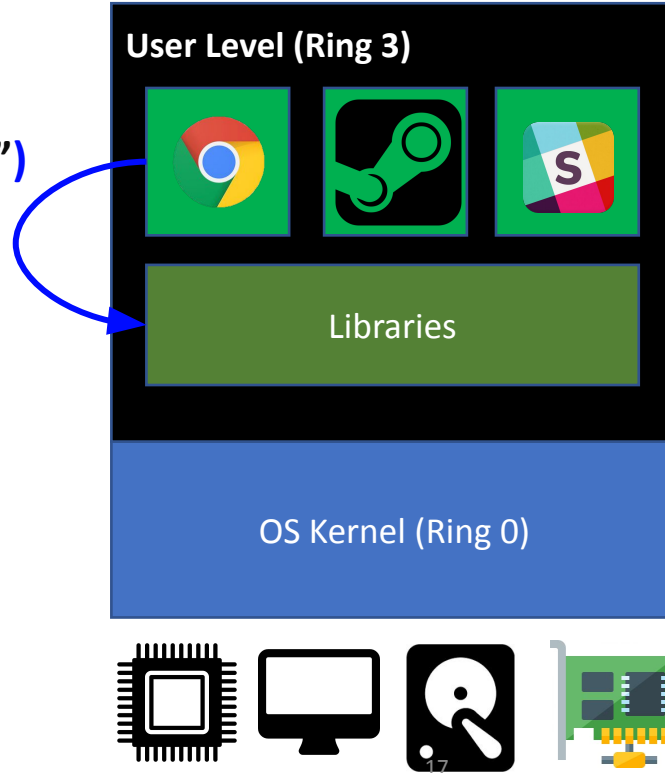


# Syscall: User/Kernel communication



`printf("ECE469")`

A library call in ring 3



```
int main() {  
    printf("ECE469");  
}
```

# Syscall: User/Kernel communication




`printf("ECE469")`

A library call in ring 3

`sys_write(1, "ECE469", 6);`

A system call, **From ring 3**



```
int main() {  
    printf("ECE469");  
}
```



# Syscall: User/Kernel communication



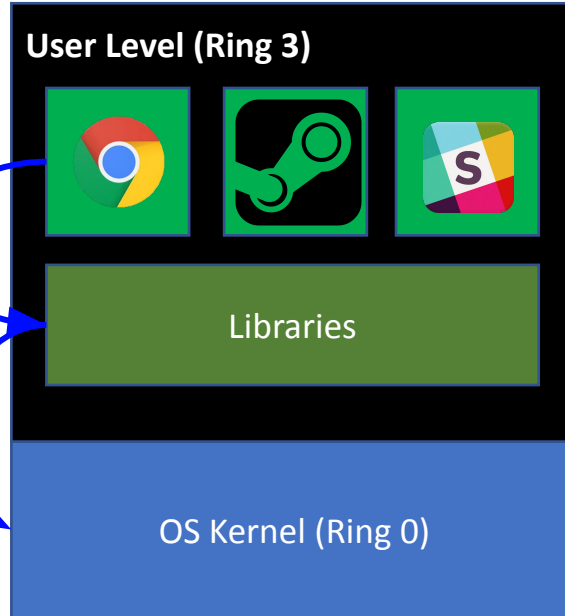
`printf("ECE469")`

A library call in ring 3

`sys_write(1, "ECE469", 6);`

A system call, **From ring 3**

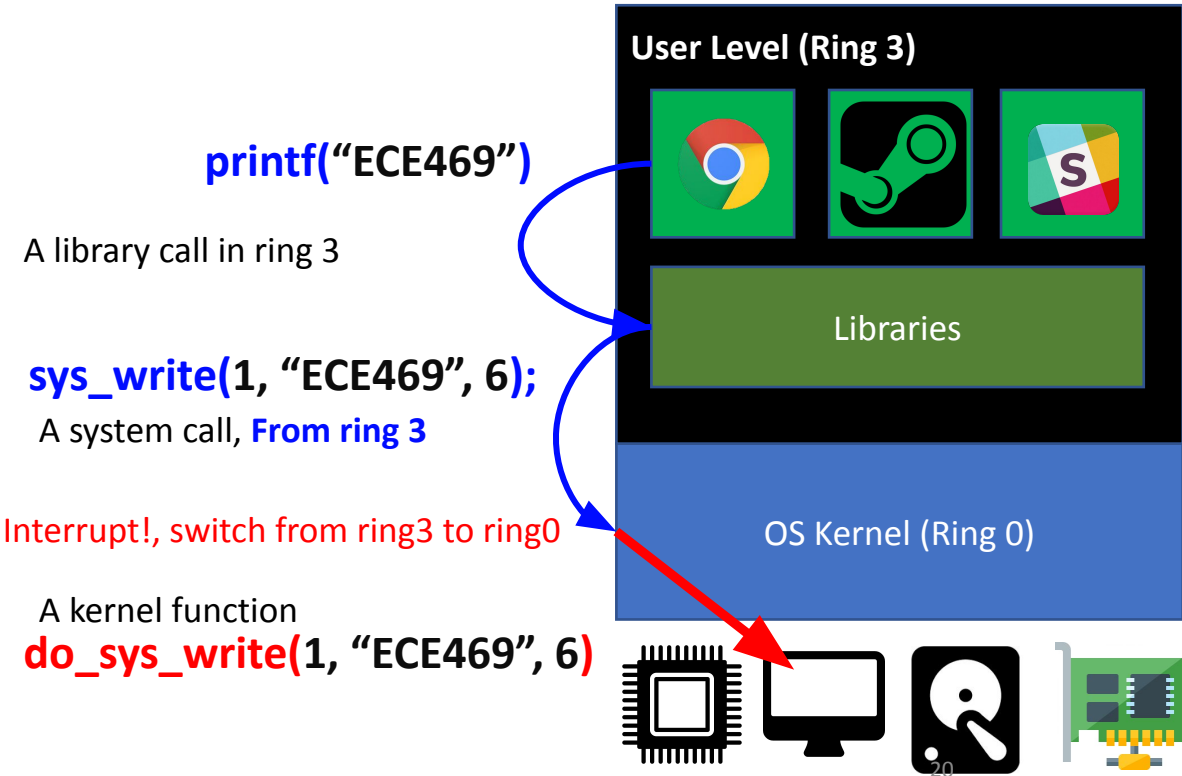
Interrupt!, switch from ring3 to ring0



```
int main() {  
    printf("ECE469");  
}
```

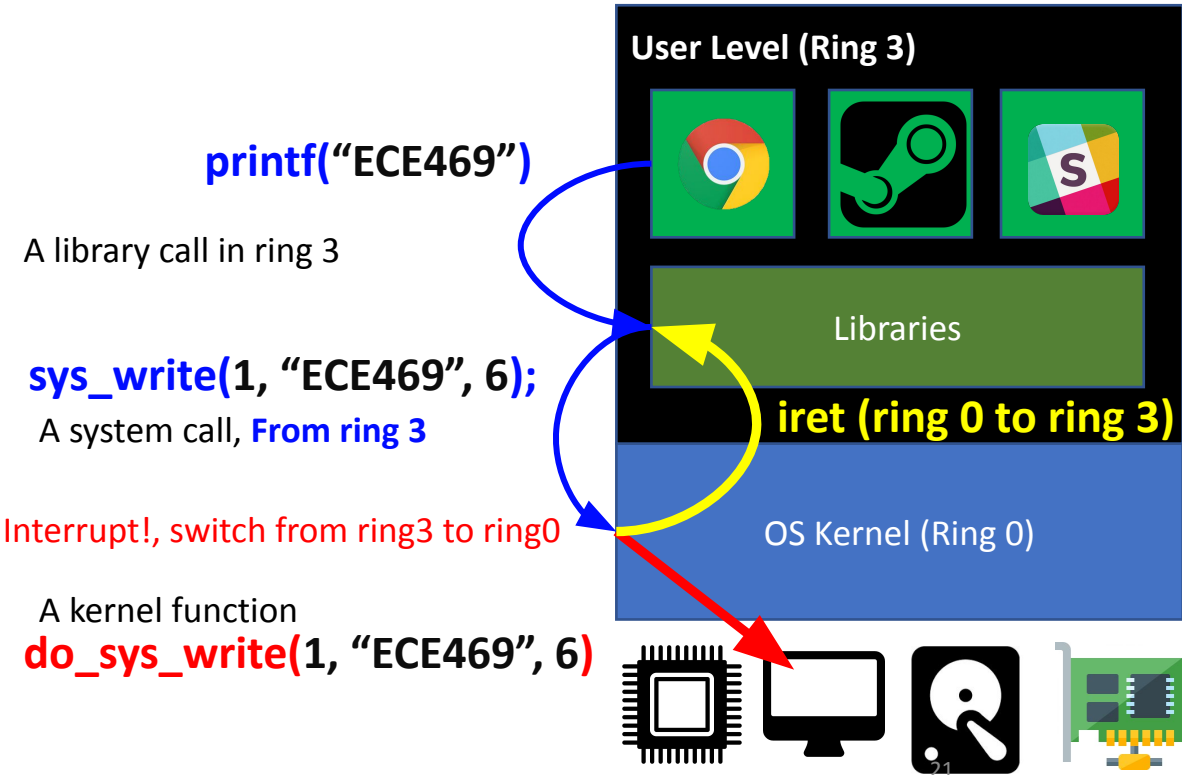


# Syscall: User/Kernel communication



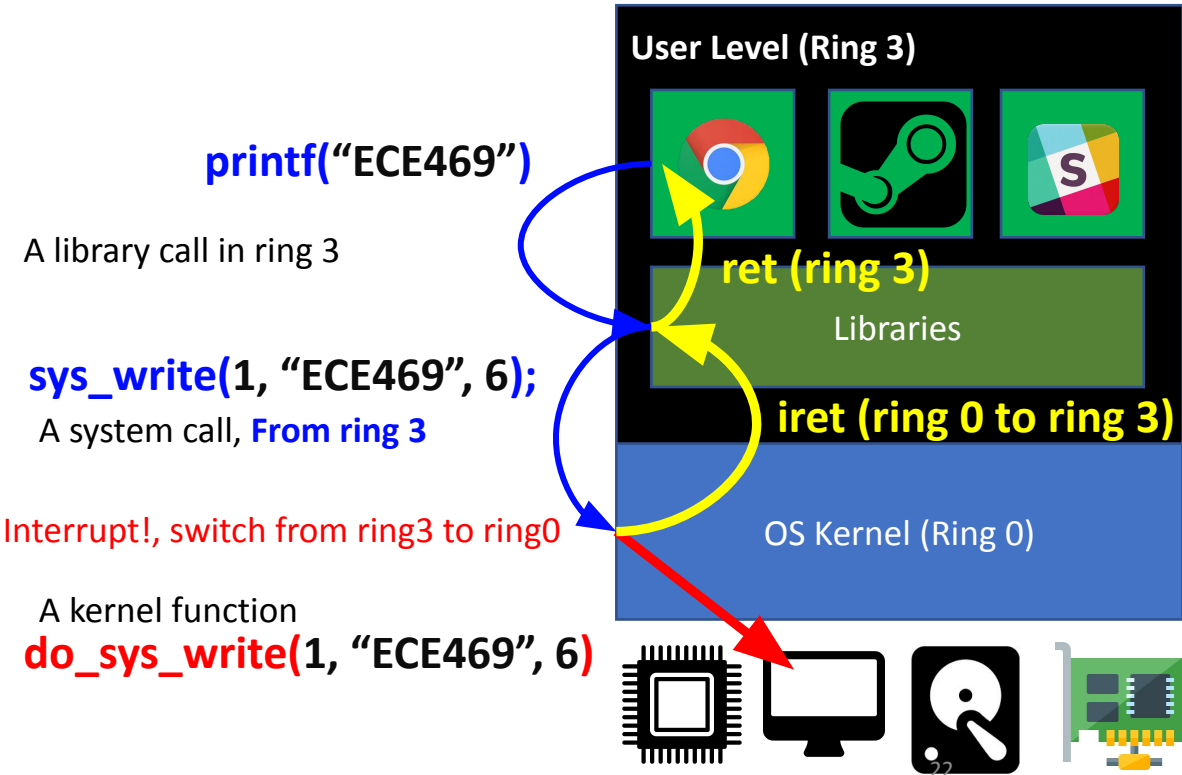
```
int main() {  
    printf("ECE469");  
}
```

# Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```


# Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```

# The need for syscalls?


- We cannot let a process access peripherals.





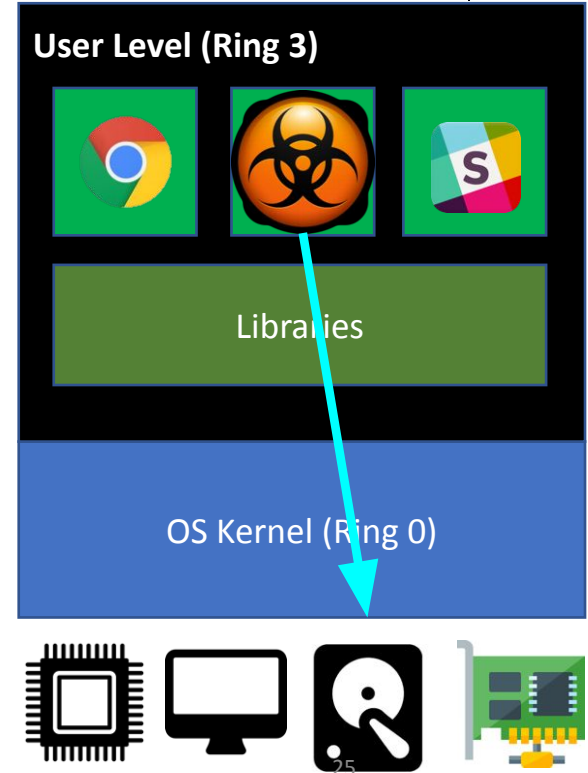
# The need for syscalls?

- We cannot let a process access peripherals.




# The need for syscalls?

- We cannot let a process access peripherals.



# The need for syscalls?


- We cannot let a process access peripherals.
- Why do we have privilege separation?
  - Security!
- We do not know what application will do
  - Do not allow dangerous operations to system
    - Flash BIOS, format disk, deleting system files, etc.
  - Let only the OS can access hardware
    - Apply access control on accessing hardware resources!
    - E.g., only the administrator can format disk



# The need for syscalls?

- We cannot let a process access peripherals.
- Why do we have privilege separation?
  - Security!
- We do not know what application will do
  - Do not allow dangerous operations to system
    - Flash BIOS, format disk, deleting system files, etc.
  - Let only the OS can access hardware
    - Apply access control on accessing hardware resources!
    - E.g. only the administrator can format disk

OS must mediate hardware access request from userspace, and we handle this via system calls




# Library Calls v/s System calls

- Library Calls

- APIs available in Ring 3
- DO NOT include operations in Ring 0
  - Cannot access hardware directly
- Could be a wrapper for some computation or
- Could be a wrapper for system calls
  - E.g., printf() internally uses write(), which is a system call

- Some system calls are available as library calls

- As wrappers in Ring 3



## NAME

read - read from a file descriptor

## SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

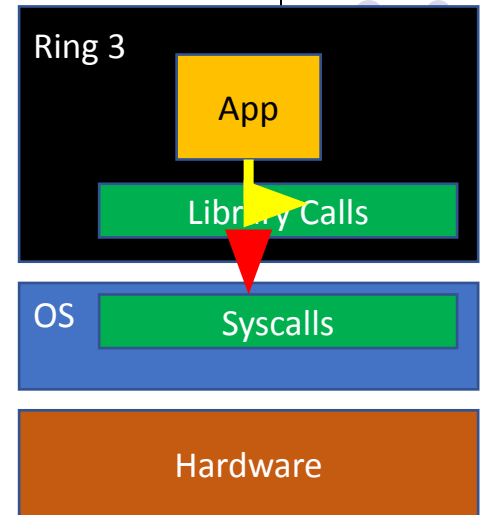
# Library Calls v/s System calls

- Library Calls

- APIs available in Ring 3
- DO NOT include operations in Ring 0
  - **Cannot access hardware directly**
- Could be a wrapper for some computation or
- Could be a wrapper for system calls
  - E.g., `printf()` internally uses `write()`, which is a system call

- Some system calls are available as library calls

- As wrappers in Ring 3



## NAME

read - read from a file descriptor

## SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

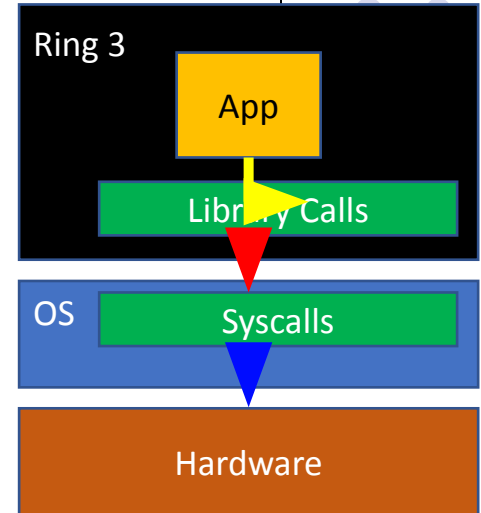
# Library Calls v/s System calls

- Library Calls

- APIs available in Ring 3
- DO NOT include operations in Ring 0
  - **Cannot access hardware directly**
- Could be a wrapper for some computation or
- Could be a wrapper for system calls
  - E.g., `printf()` internally uses `write()`, which is a system call

- Some system calls are available as library calls

- As wrappers in Ring 3



## NAME

read - read from a file descriptor

## SYNOPSIS


```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

# Library Calls v/s System calls

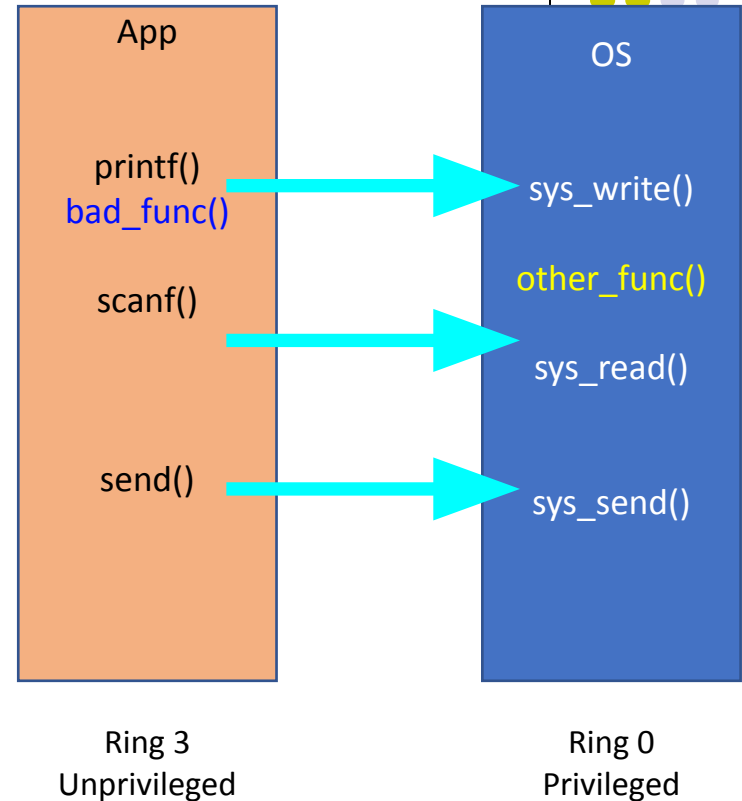
- System Calls

- APIs available in Ring 0
- OS's abstraction for hardware interface for userspace
- Called when Ring 3 application need to perform Ring 0 operations

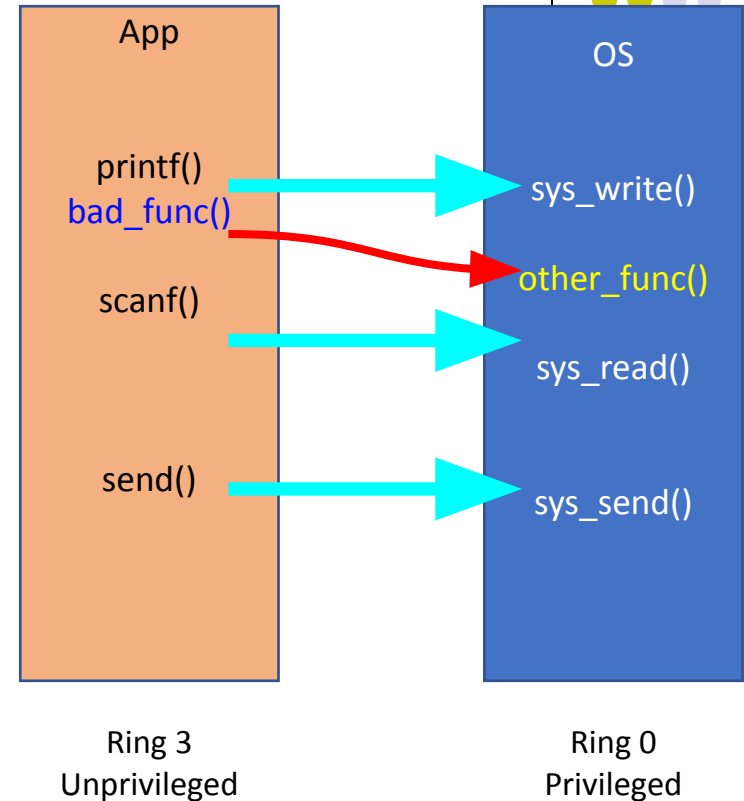




# System calls are not function calls!




# System calls are not function calls!



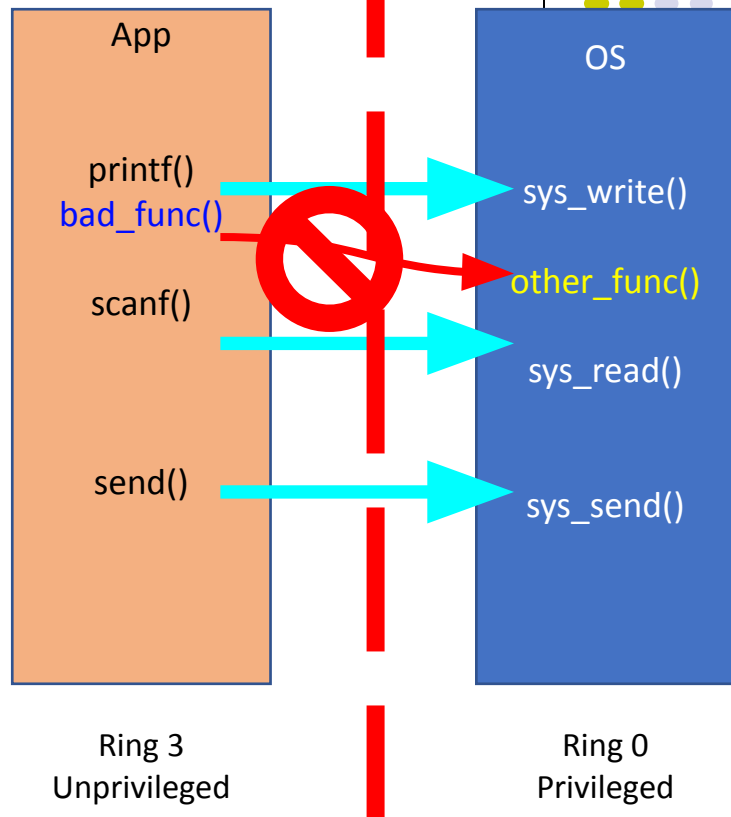
# System calls are not function calls!

- Application should not call arbitrary OS functions
  - If so, app can do all operations that OS can do; privilege separation is meaningless!



# System calls are not function calls!

- Application should not call arbitrary OS functions
  - If so, app can do all operations that OS can do; privilege separation is meaningless!
- How can we protect this, in other words, how can we **let apps invoke system calls** but **no other OS functions**?



# System call Design via call gate

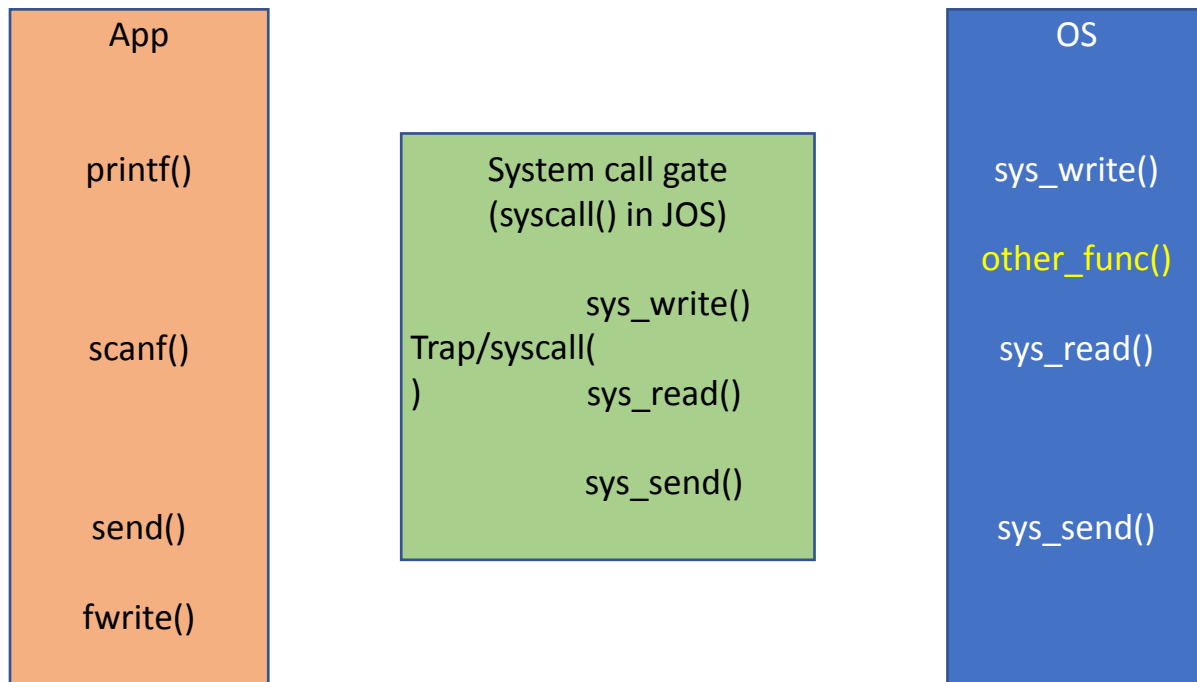


- Call gate: a secure method to control access to Ring 0!

# System call Design via call gate



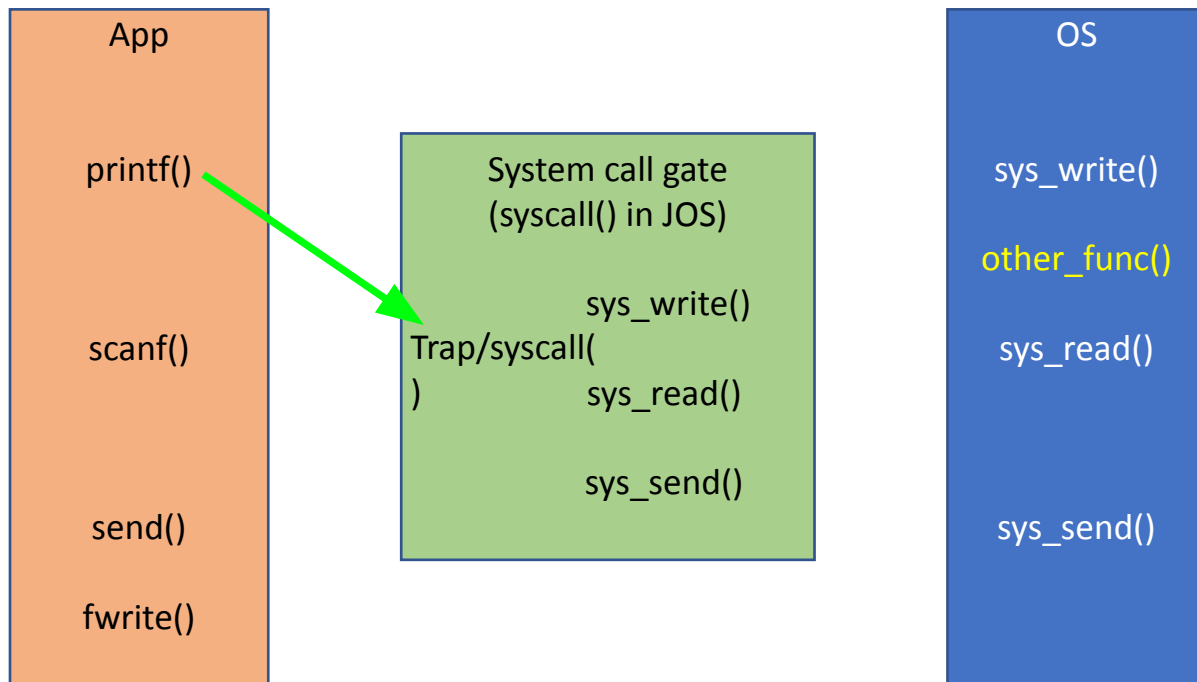
- Call gate: a secure method to control access to Ring 0!



# System call Design via call gate



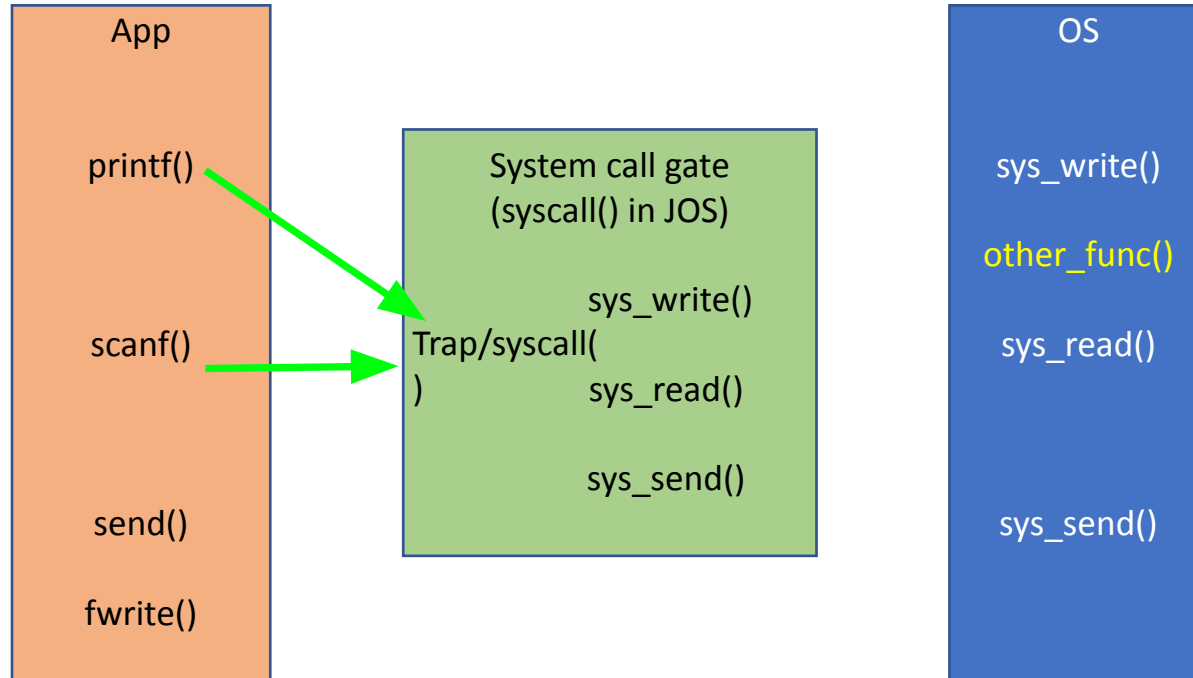
- Call gate: a secure method to control access to Ring 0!



# System call Design via call gate



- Call gate: a secure method to control access to Ring 0!






# System call Design via call gate



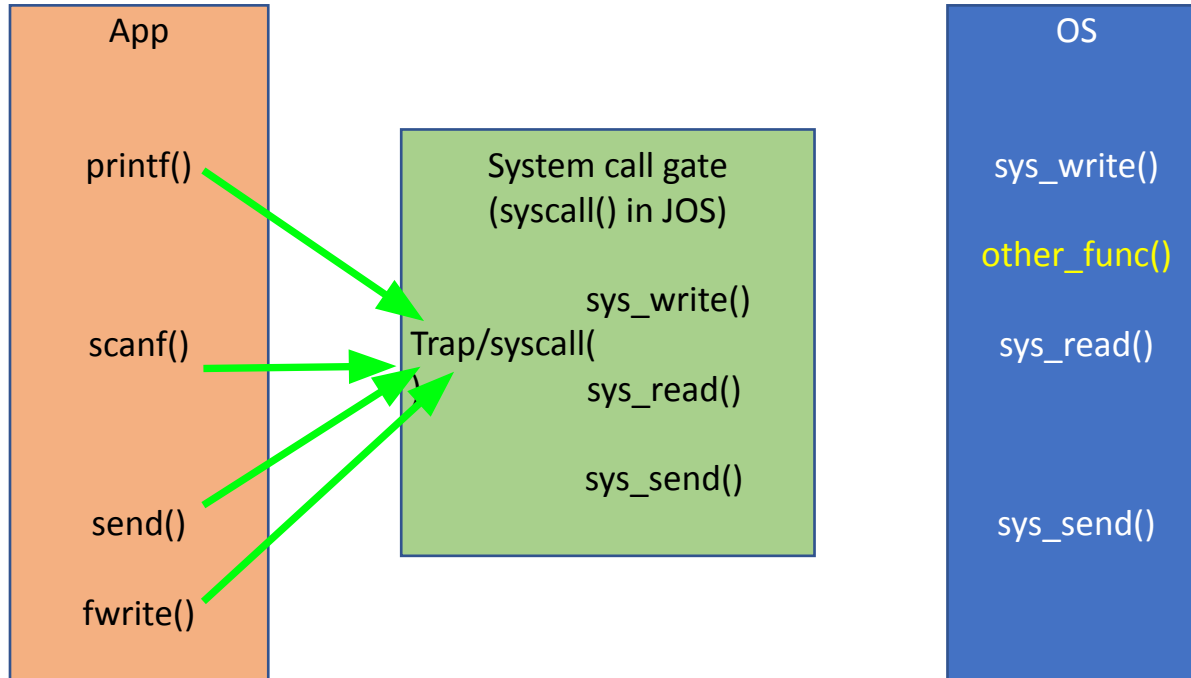
- Call gate: a secure method to control access to Ring 0!



# System call Design via call gate



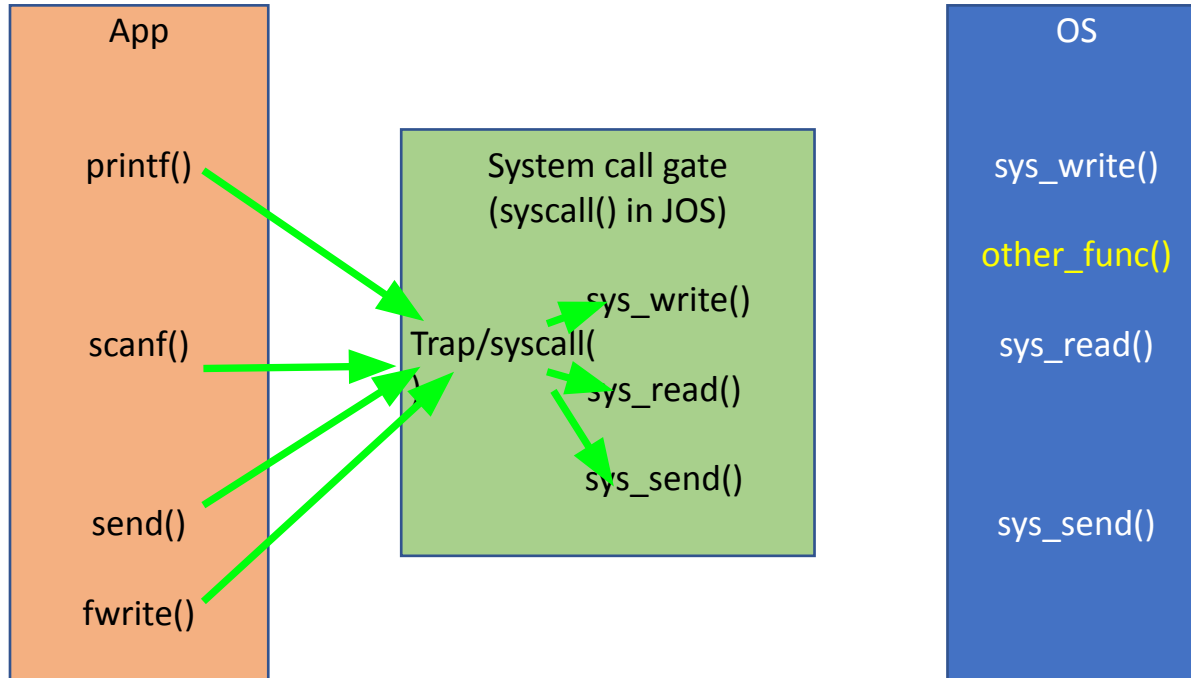
- Call gate: a secure method to control access to Ring 0!



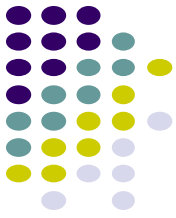
# System call Design via call gate



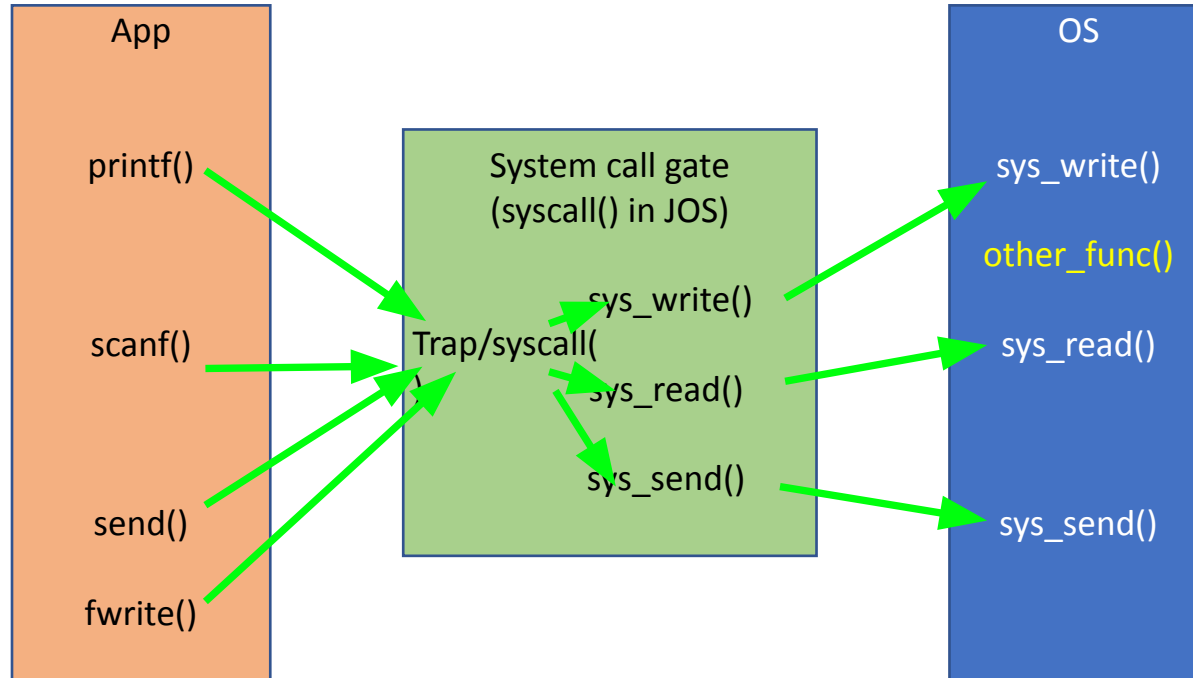
- Call gate: a secure method to control access to Ring 0!



# System call Design via call gate




- Call gate: a secure method to control access to Ring 0!



# Call gate via Interrupt Handler



- Call gate
  - System call can be invoked only with trap handler
    - `int $0x30` – in JOS
    - `int $0x80` – in Linux (32-bit)
    - `int $0x2e` – in Windows (32-bit)
    - `sysenter/sysexit` (32-bit)
    - `syscall/sysret` (64-bit)
- OS performs checks if userspace is doing a right thing
  - Before performing important ring 0 operations
  - E.g., accessing hardware..




# Call gate via Interrupt Handler



- Call gate
  - System call can be invoked only with trap handler
    - `int $0x30` – in JOS
    - `int $0x80` – in Linux (32-bit)
    - `int $0x2e` – in Windows (32-bit)
    - `sysenter/sysexit` (32-bit)
    - `syscall/sysret` (64-bit)
- OS performs checks if userspace is doing a right thing
  - Before performing important ring 0 operations
  - E.g., accessing hardware..

**int \$0x30**  
**CHECK!!**



# Why should we check arguments?



- How can we protect 'read()' system call?
  - `read(int fd, void *buf, size_t count)`
  - Read `count` bytes from a file pointed by `fd` and store those in `buf`
- Usage

```
// buffer at the stack
char buf[512];
// read 512 bytes from standard input
read(0, buf, 512);
```

# Why should we check arguments?



- Problem: what will happen if we call...

```
// kernel address will points to a dirmap of  
// the physical address at 0x100000  
char kernel_address = KERNBASE + 0x100000;  
// read 512 bytes from standard input  
read(0, buf, 512);
```

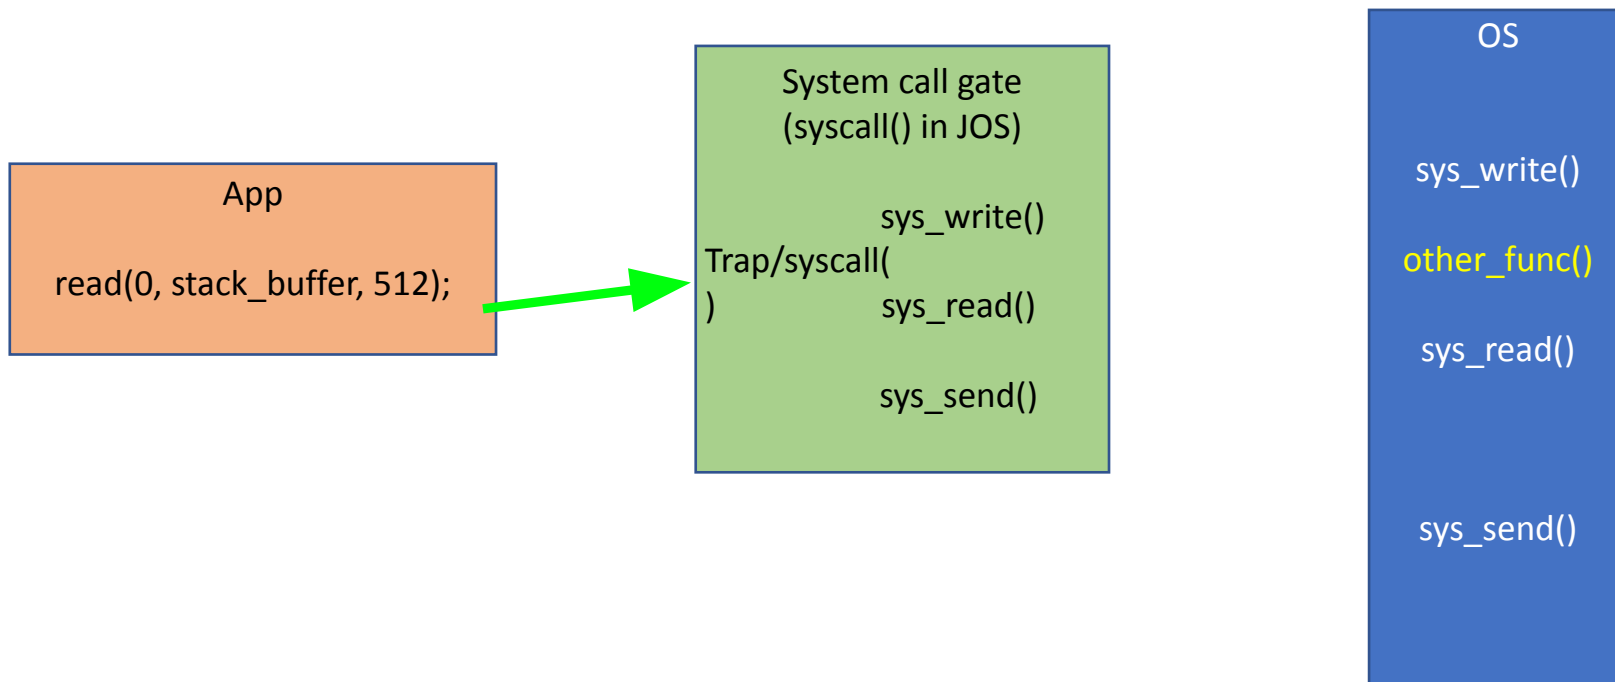
- This will **overwrite kernel code** with your keystroke typing..
  - Changing kernel code from Ring 3 is possible!



# Checking arguments for syscalls



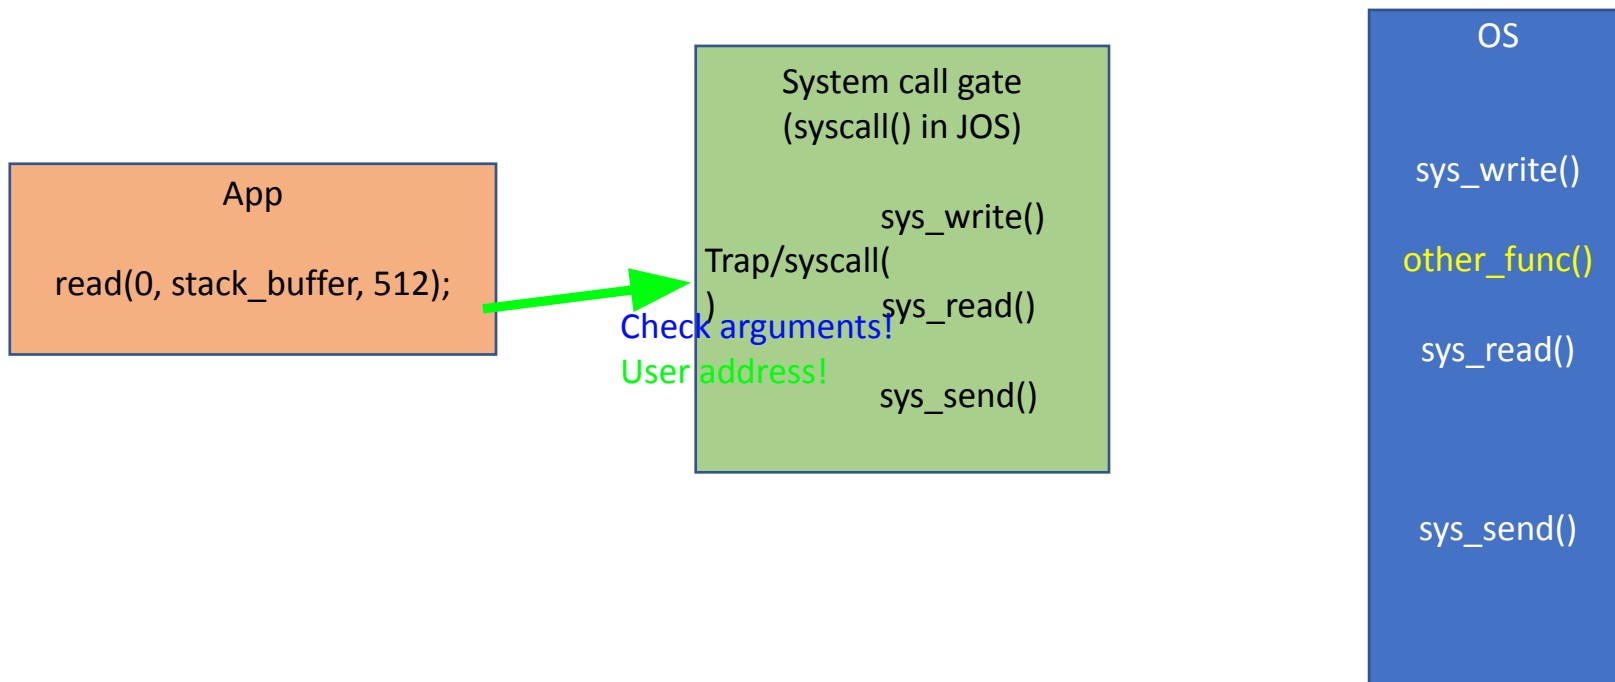
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



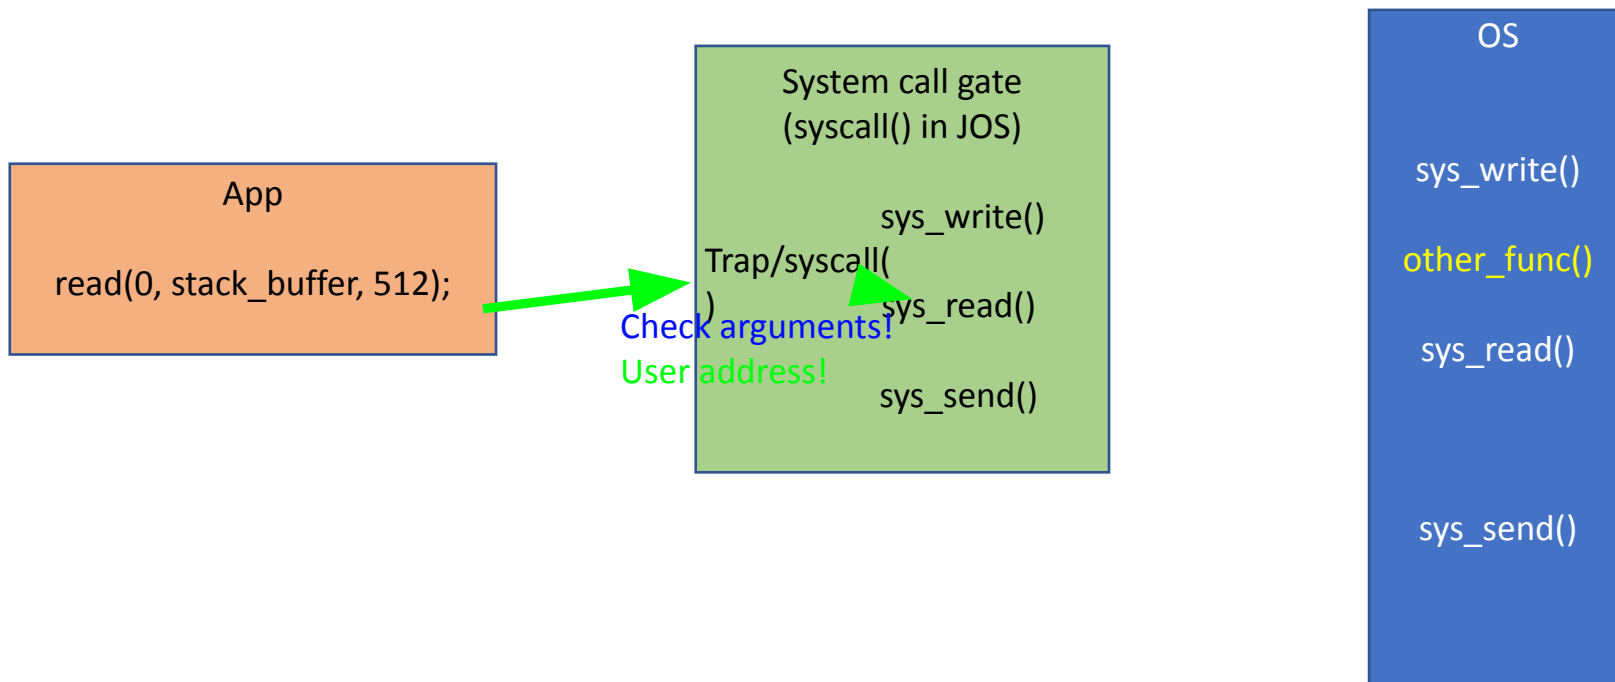
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



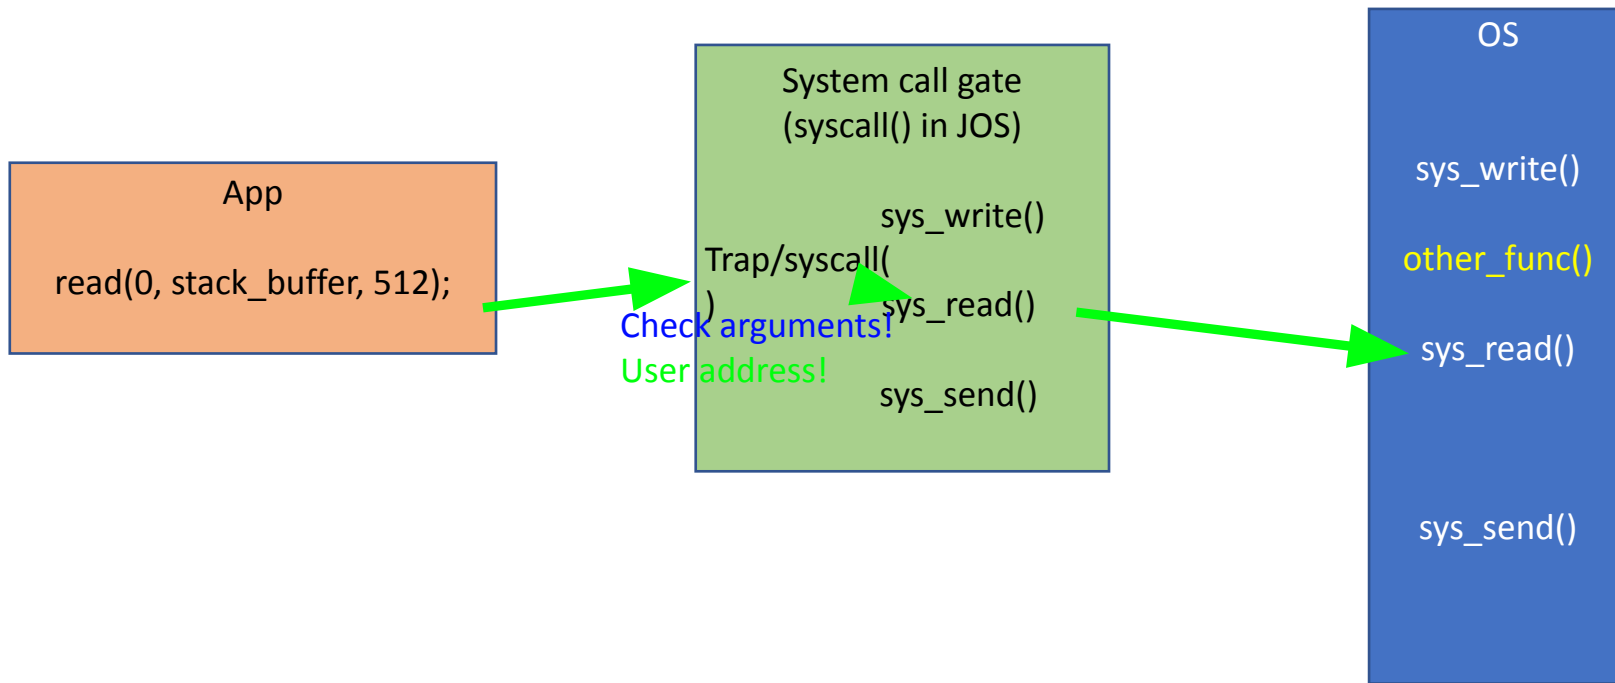
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



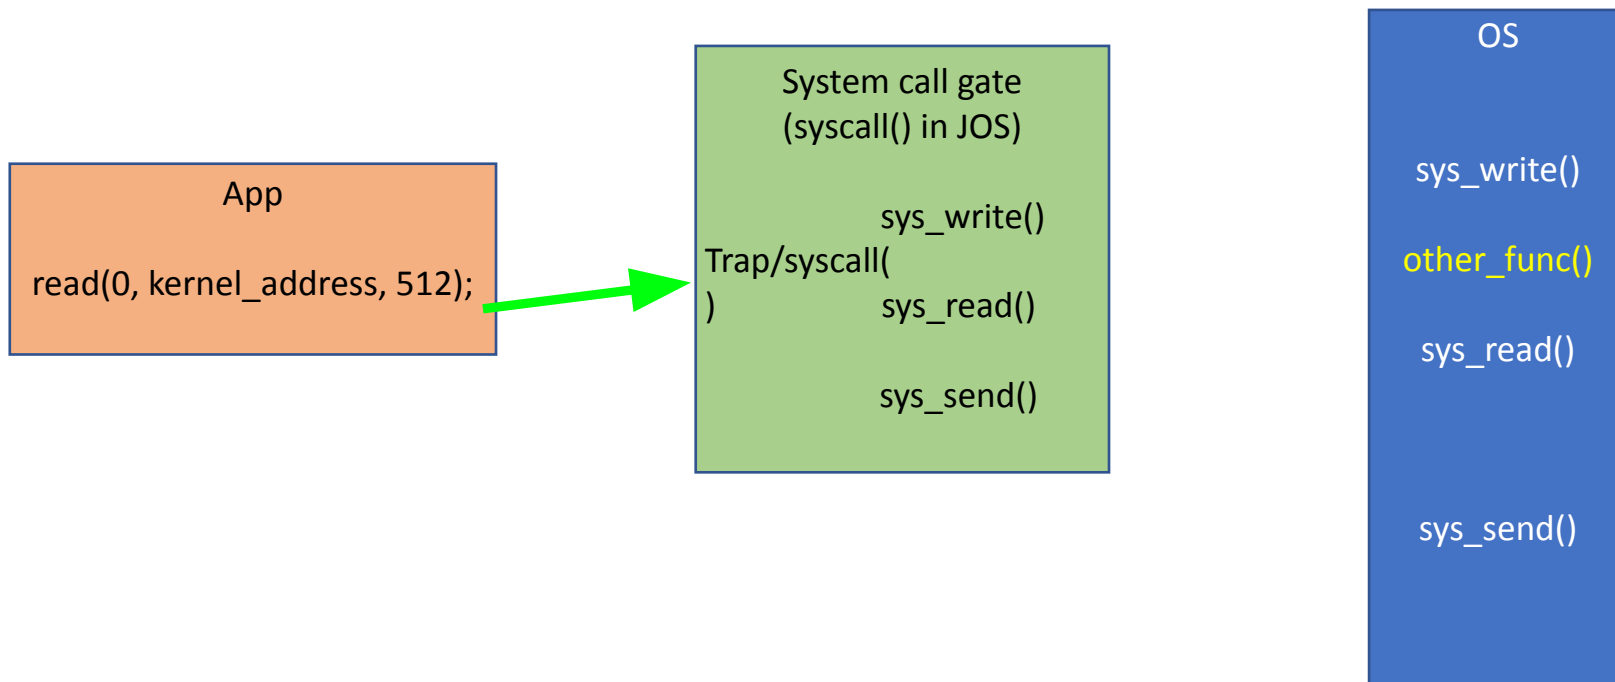
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



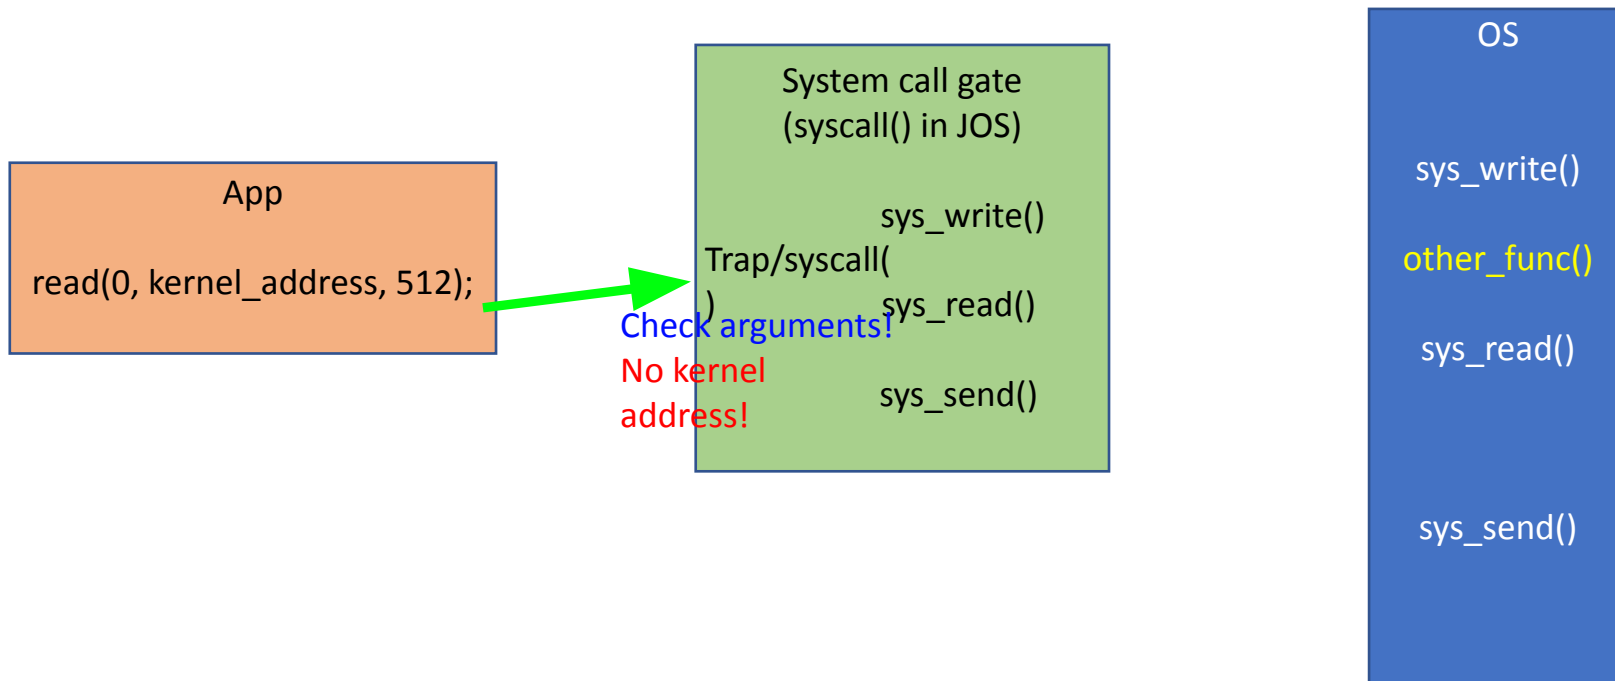
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



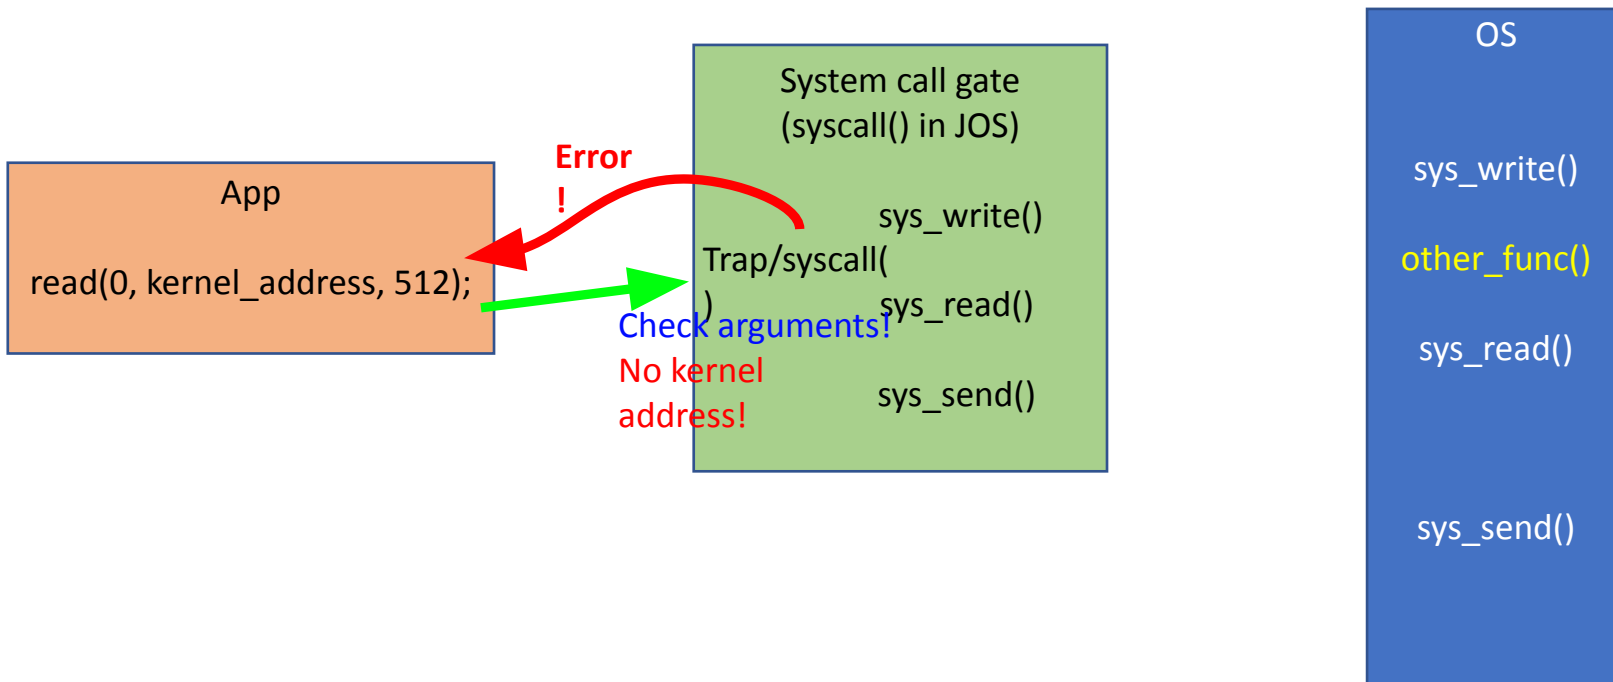
- We can hook all syscalls from Ring 3 at our syscall trap handler



# Checking arguments for syscalls



- We can hook all syscalls from Ring 3 at our syscall trap handler



# Test: using ltrace and strace



```
// buffer at the stack
char buf[512];
// read 512 bytes from stdin to stack.
int ret = read(0, buf, 512);

printf("Read to stack memory returns: %d\n", ret);

// read 512 bytes from stdin to kernel.
ret = read(0, (void*)0xffffffff01000000, 512);

printf("Read to kernel memory returns: %d\n", ret);
perror("Reason for the error:");
```



# Summary: Syscalls



- Prevent Ring 3 from accessing hardware directly
  - Security reasons!
  - OS mediates hardware access via system calls
- You may regard system calls as APIs of an OS
- How to prevent an application from running arbitrary ring 0 operation?
  - Call gate
- Modern OS use call gate to protect system calls
  - At trap handler, an OS can apply access control to system call request

# Faults



- Faults
  - Faulting instruction has not executed (e.g., page fault)
  - Resume the execution after handling the fault
- Resume the execution after handling the fault

# Page faults



- Occurs when paging (address translation) fails

# Page faults



- Occurs when paging (address translation) fails
  - Access from user but ! (pte & PTE\_U) : protection violation



# Page faults

- Occurs when paging (address translation) fails
  - Access from user but ! (pte & PTE\_U) : protection violation

```
int main() {  
    char *kernel_memory = (char*)0xf0100000;  
    // I am a bad guy, and I would like to change  
    // some contents in kernel memory  
    kernel_memory[100] = '!';  
}
```

```
0x00800039 ? movb    $0x21,0xf0100064
```



# Page faults

- Occurs when paging (address translation)
  - Access from user but ! (pte & PTE\_U):

```
int main() {  
    char *kernel_memory = (char*)0xf01  
    // I am a bad guy, and I would lik  
    // some contents in kernel memory  
    kernel_memory[100] = '!';  
}
```

```
0x00800039 ? movb $0x21,0x
```

TRAP frame at 0xf01c0000

edi 0x00000000

esi 0x00000000

ebp 0xeebdfd0

oesp 0xefffffff

ebx 0x00000000

edx 0x00000000

ecx 0x00000000

eax 0xeec00000

es 0x----0023

ds 0x----0023

trap 0x0000000e Page Fault

cr2 0xf0100064

err 0x00000007 [user, write, protection]

eip 0x00800039


cs 0x----001b

flag 0x00000096

esp 0xeebdfdb8

ss 0x----0023

[00001000] free env 00001000



# What does CPU do on a page fault?

- CPU let OS know why and where such a page fault happened

```
TRAP frame at 0xf01c0000
edi  0x00000000
esi  0x00000000
ebp  0xeebdfd0
oesp 0xefffffff
ebx  0x00000000
edx  0x00000000
ecx  0x00000000
eax  0xeec00000
es   0x---0023
ds   0x---0023
trap 0x0000000e Page Fault
cr2  0xf0100064
err  0x00000007 [user, write, protection]
eip  0x00800039
cs   0x---001b
flag 0x00000096
esp  0xeebdfb8
```


```
kernel_memory[100] = '!'; 00001000
```



# What does CPU do on a page fault?

- CPU let OS know why and where such a page fault happened
  - CR2: stores the address of the fault

```
TRAP frame at 0xf01c0000
edi  0x00000000
esi  0x00000000
ebp  0xeebfdfd0
oesp 0xefffffffdc
ebx  0x00000000
edx  0x00000000
ecx  0x00000000
eax  0xeec00000
es   0x---0023
ds   0x---0023
trap 0x0000000e Page Fault
cr2  0xf0100064
err  0x00000007 [user, write, protection]
eip  0x00800039
cs   0x---001b
flag 0x00000096
esp  0xeebdfdb8
```



```
kernel_memory[100] = '!'; 00001000
```





# How does OS handle page fault?



- User program accesses 0xf0100064



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault (pte & PTE\_U == 0)
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault ( $\text{pte} \& \text{PTE\_U} == 0$ )
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault (`pte & PTE_U == 0`)
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`
  - Read CR2 (address of the fault, 0xf0100064)



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault ( $\text{pte} \& \text{PTE\_U} == 0$ )
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`
  - Read CR2 (address of the fault, 0xf0100064)
  - Read error code (contains the reason of the fault)



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault (`pte & PTE_U == 0`)
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`
  - Read CR2 (address of the fault, 0xf0100064)
  - Read error code (contains the reason of the fault)
  - Resolve error (if not, destroy the environment)



# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault (`pte & PTE_U == 0`)
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`
  - Read CR2 (address of the fault, 0xf0100064)
  - Read error code (contains the reason of the fault)
  - Resolve error (if not, destroy the environment)
  - Continue user execution






# How does OS handle page fault?

- User program accesses 0xf0100064
- CPU generates page fault (`pte & PTE_U == 0`)
  - Put the faulting address on CR2
  - Put an error code
  - Calls page fault handler in IDT
- OS: `page_fault_handler`
  - Read CR2 (address of the fault, 0xf0100064)
  - Read error code (contains the reason of the fault)
  - Resolve error (if not, destroy the environment)
  - Continue user execution
- User: resume on that instruction (or destroyed by the OS)



## Page fault example (2): Handling call stack


- inc/memlayout.h
- We allocate one (1) page for the user stack





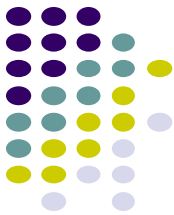
# Page fault example: Handling call stack

- inc/memlayout.h
- We allocate one (1) page for the user stack



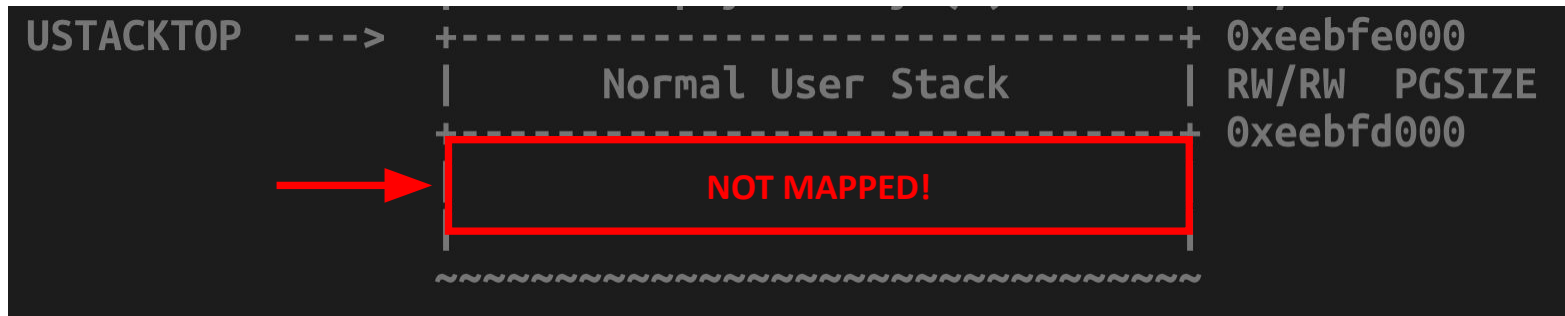
- If you use a large local variable on the stack
  - Stack overflow (stack grows down...)

```
int func() {  
    char buf[8192];  
    buf[0] = '1';  
}
```



# Page fault example: Handling call stack

- inc/memlayout.h
- We allocate one (1) page for the user stack



- If you use a large local variable on the stack
  - Stack overflow (stack grows down...)

```
int func() {  
    char buf[8192];  
    buf[0] = '1';  
}
```

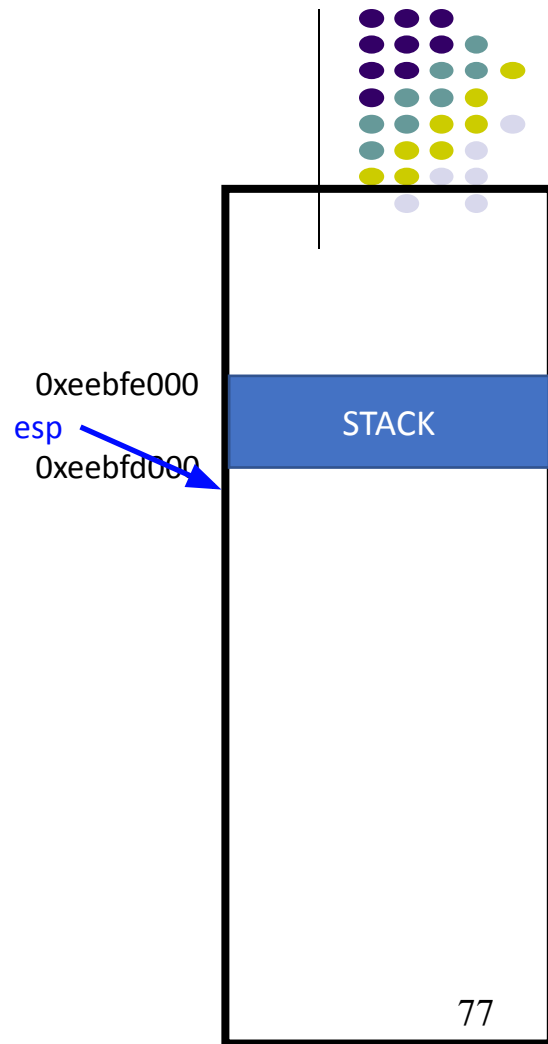
# Expand stack automatically



- Can we detect such an access and allocate a new page for the stack automatically?
  - Yes
- We will utilize 'Page Fault'
- Observations
  - Stack overflow would be sequential (access pages adjacent to the stack)
  - We should catch both read/write access (both should fault)

# Expand stack automatically

- Stack ends at 0xeebfd000
- Suppose the current value of `esp` (stack) is
  - 0xeebfd010




# Expand stack automatically

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        buf[i] = '1' + i;  
    }  
}
```

- Stack ends at 0xeebfd000
- Suppose the current value of esp (stack) is
  - 0xeebfd010
- User program creates a new variable: char buf[32]
  - buf = 0xeebfcff0
  - Buffer range: 0xeebfcff0 ~ 0xeebfd010

0xeebfe000  
esp  
0xeebfd000  
buf  
0xeebfc000




# Expand stack automatically

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        buf[i] = '1' + i;  
    }  
}
```

- Stack ends at 0xeebfd000
- Suppose the current value of esp (stack) is
  - 0xeebfd010
- User program creates a new variable: char buf[32]
  - buf = 0xeebfcff0
  - Buffer range: 0xeebfcff0 ~ 0xeebfd010
- On accessing buf[0] = '1';
  - movb \$0x31, (%eax)

0xeebfe000  
esp  
0xeebfd000  
buf  
0xeebfc000





# Expand stack automatically

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        buf[i] = '1' + i;  
    }  
}
```

- Stack ends at 0xeebfd000
- Suppose the current value of esp (stack) is
  - 0xeebfd010
- User program creates a new variable: char buf[32]
  - buf = 0xeebfcff0
  - Buffer range: 0xeebfcff0 ~ 0xeebfd010
- On accessing buf[0] = '1';
  - movb \$0x31, (%eax)
  - eax = 0xeebfcff0 No translation for 0xeebfc000



# Expand stack automatically

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        buf[i] = '1' + i;  
    }  
}
```

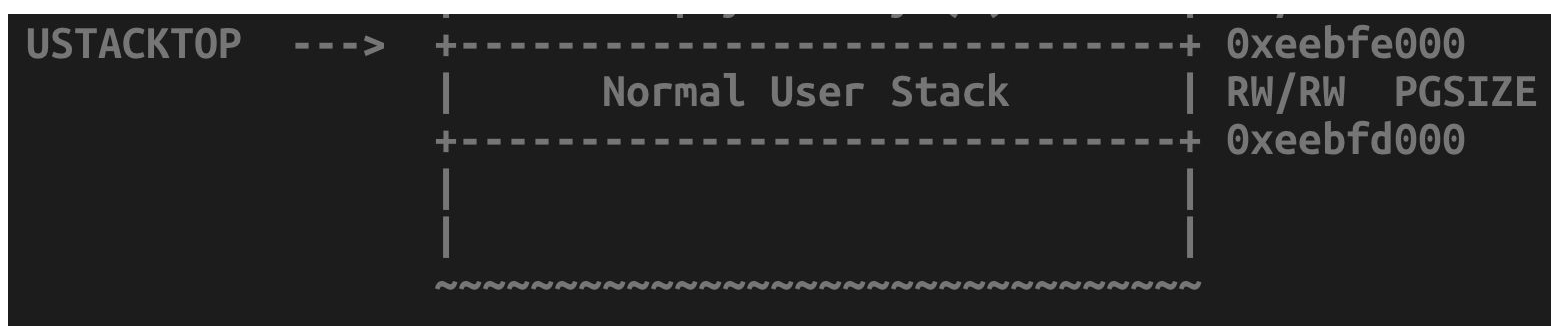
- Stack ends at 0xeebfd000
- Suppose the current value of esp (stack) is
  - 0xeebfd010
- User program creates a new variable: char buf[32]
  - buf = 0xeebfcff0
  - Buffer range: 0xeebfcff0 ~ 0xeebfd010
- On accessing buf[0] = '1';
  - movb \$0x31, (%eax)
  - eax = 0xeebfcff0 No translation for 0xeebfc000
  - **Need to allocate 0xeebfc000 ~ 0xeebfd000**





# What does processor do?

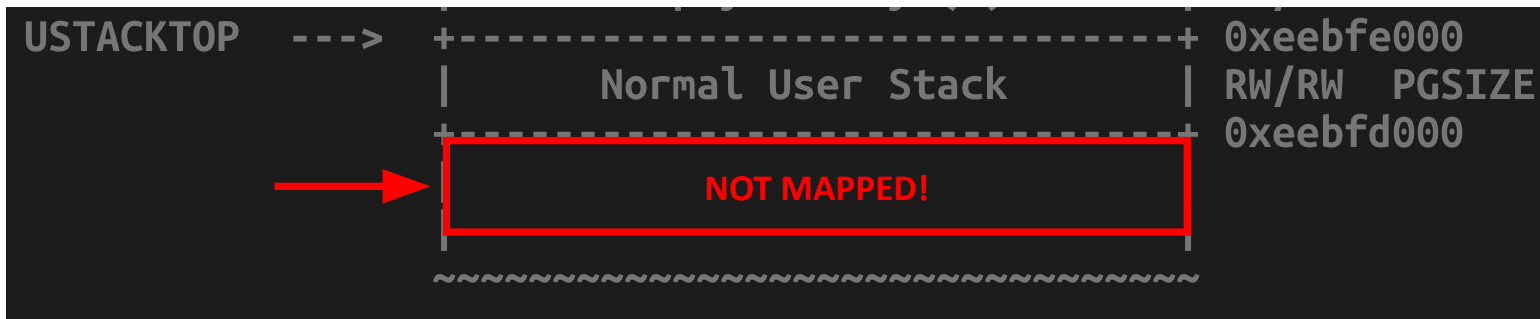
- Lookup page table
  - No translation!




# What does processor do?




- Lookup page table
  - No translation!





# What does processor do?

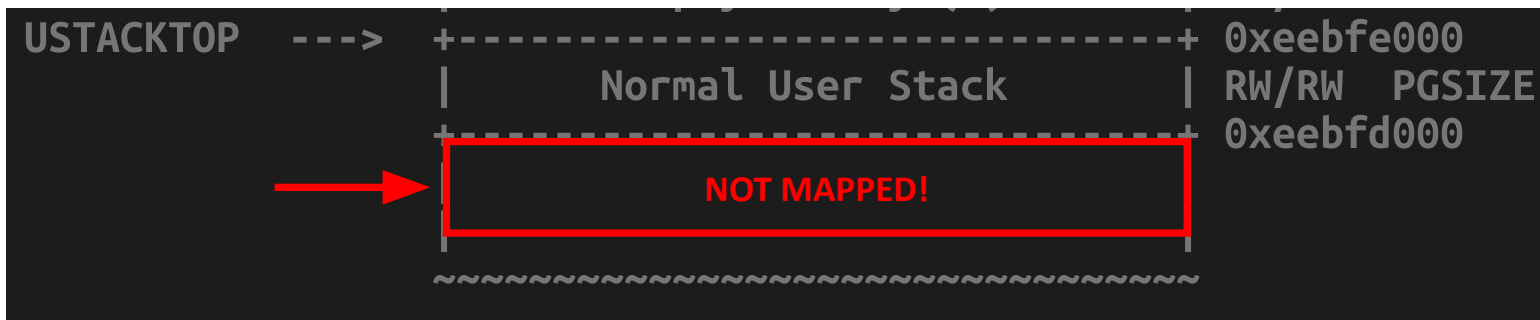
- Lookup page table
  - No translation!
- Store `0xeebfcff0` to CR2






# What does processor do?

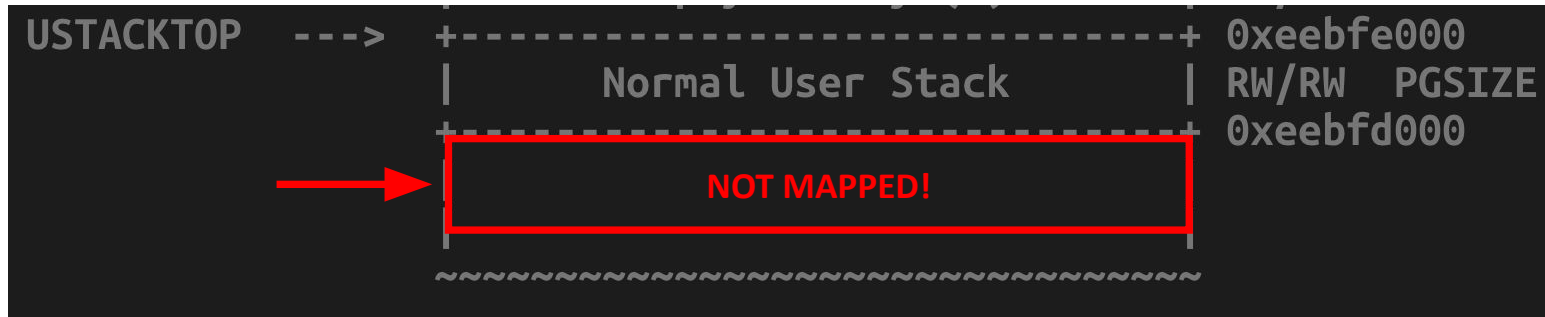
- Lookup page table
  - No translation!
- Store `0xeebfcff0` to CR2
- Set error code
  - “The fault was caused by a non-present page!”





# What does processor do?

- Lookup page table
  - No translation!
- Store `0xeebfcff0` to CR2
- Set error code
  - “The fault was caused by a non-present page!”
- Raise page fault exception (interrupt #14) -> call page fault handler



# Handling page fault on Stack access


- Interrupt will make CPU invoke the `page_fault_handler()`

0xeebfe000

0xeebfd000

STACK

No mapping!





# Handling page fault on Stack access

- Interrupt will make CPU invoke the `page_fault_handler()`
- Read CR2
  - `0xeebfcff0`

0xeebfe000

0xeebfd000

STACK


No mapping!

# Handling page fault on Stack access

- Interrupt will make CPU invoke the `page_fault_handler()`
- Read CR2
  - `0xeebfcff0`, it seems like the page right next to current stack end
  - The current stack end is: `0xeebfd000`

0xeebfe000

0xeebfd000




# Handling page fault on Stack access

- Interrupt will make CPU invoke the `page_fault_handler()`
- Read CR2
  - `0xeebfcff0`, it seems like the page right next to current stack end
  - The current stack end is: `0xeebfd000`
- Read error code
  - “The fault was caused by a non-present page!”

0xeebfe000

0xeebfd000




# Handling page fault on Stack access

- Interrupt will make CPU invoke the `page_fault_handler()`
- Read CR2
  - `0xeebfcff0`, it seems like the page right next to current stack end
  - The current stack end is: `0xeebfd000`
- Read error code
  - “The fault was caused by a non-present page!”
- Let’s allocate a new page for the stack!

0xeebfe000

0xeebfd000



# Adding new page for stack

- Allocate a new page for the stack
  - `Struct PageInfo *pp = page_alloc(ALLOC_ZERO);`
    - Get a new page, and wipe it to have all zero as its contents



0xeebfe000

0xeebfd000

0xeebfc000

STACK

# Adding new page for stack

- Allocate a new page for the stack

- `Struct PageInfo *pp = page_alloc(ALLOC_ZERO);`

- Get a new page, and wipe it to have all zero as its contents

- `page_insert(env_pgdir, pp, 0xeebfc000, PTE_U|PTE_W);`

- Map a new page to that address!

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

# Adding new page for stack

- Allocate a new page for the stack

- `Struct PageInfo *pp = page_alloc(ALLOC_ZERO);`

- Get a new page, and wipe it to have all zero as its contents

- `page_insert(env_pgdir, pp, 0xeebfc000, PTE_U|PTE_W);`

- Map a new page to that address!

- `iret!`

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

# Resuming execution of user process

- On accessing `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` No translation for `0xeebfc000`

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000



# Resuming execution of user process

- On accessing `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` No translation for `0xeebfc000`
- Execute the faulting instruction again: `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0`

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        → buf[i] = '1' + i;  
    }  
}
```

# Resuming execution of user process

- On accessing `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` No translation for `0xeebfc000`
- Execute the faulting instruction again: `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` Now translation is valid!

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        → buf[i] = '1' + i;  
    }  
}
```

# Resuming execution of user process

- On accessing `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` No translation for `0xeebfc000`
- Execute the faulting instruction again: `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` Now translation is valid!
- Continue to execute the loop..

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        → buf[i] = '1' + i;  
    }  
}
```

# Resuming execution of user process

- On accessing `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` No translation for `0xeebfc000`
- Execute the faulting instruction again: `buf[0] = '1';`
  - `movb $0x31, (%eax)`
  - `eax = 0xeebfcff0` Now translation is valid!
- Continue to execute the loop..

0xeebfe000

STACK

0xeebfd000

STACK

0xeebfc000

By exploiting **page fault and its handler**, we can implement **automatic allocation of user stack!**

```
int func() {  
    char buf[32];  
    for(int i=0; i<32; ++i) {  
        → buf[i] = '1' + i;  
    }  
}
```