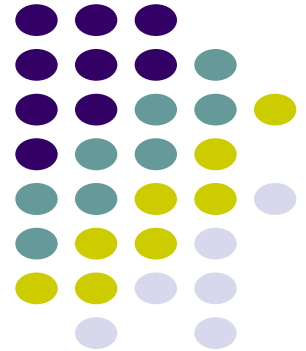


# Multi Processes and Scheduling

---

ECE 469, Feb 17

Aravind Machiry



# Recap: Users, Programs, Processes



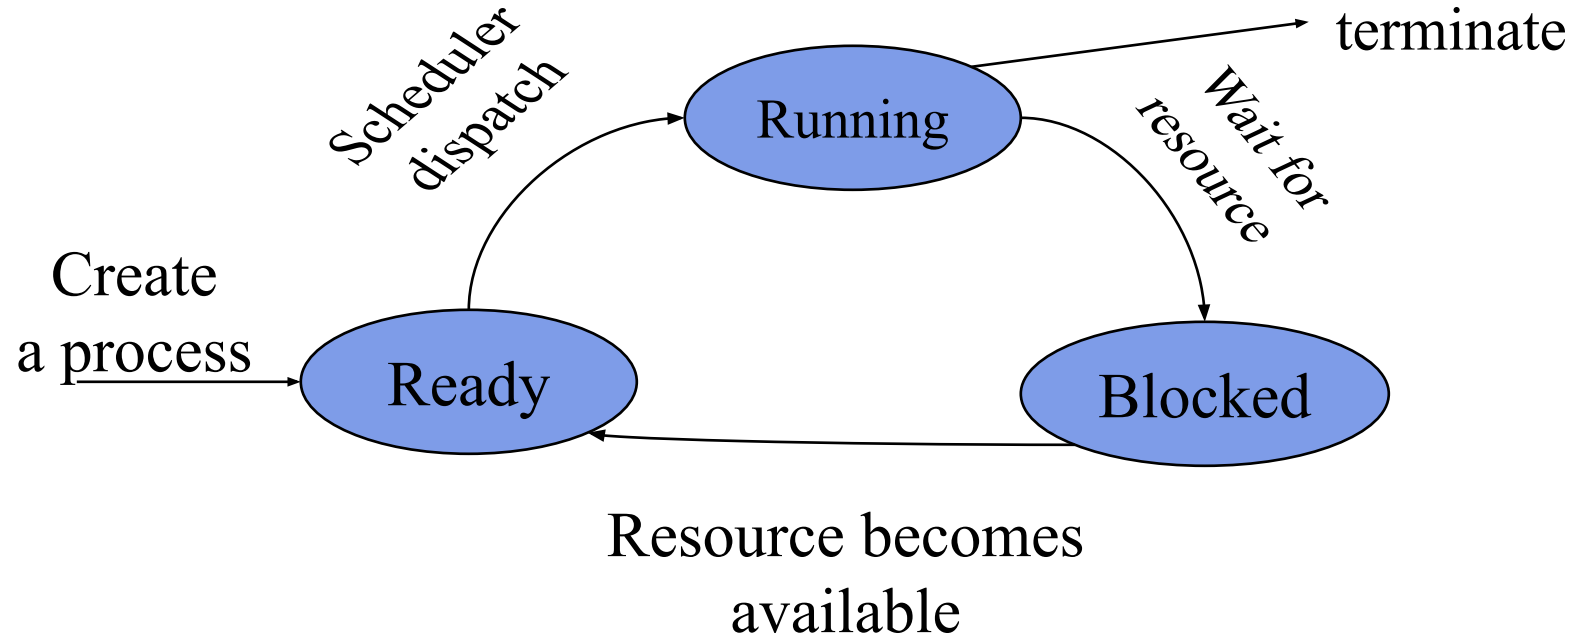
- Users have accounts on the system
- Users launch programs
- There can be multiple programs (i.e., processes), which want to run at the same time

# Sequential execution of each process

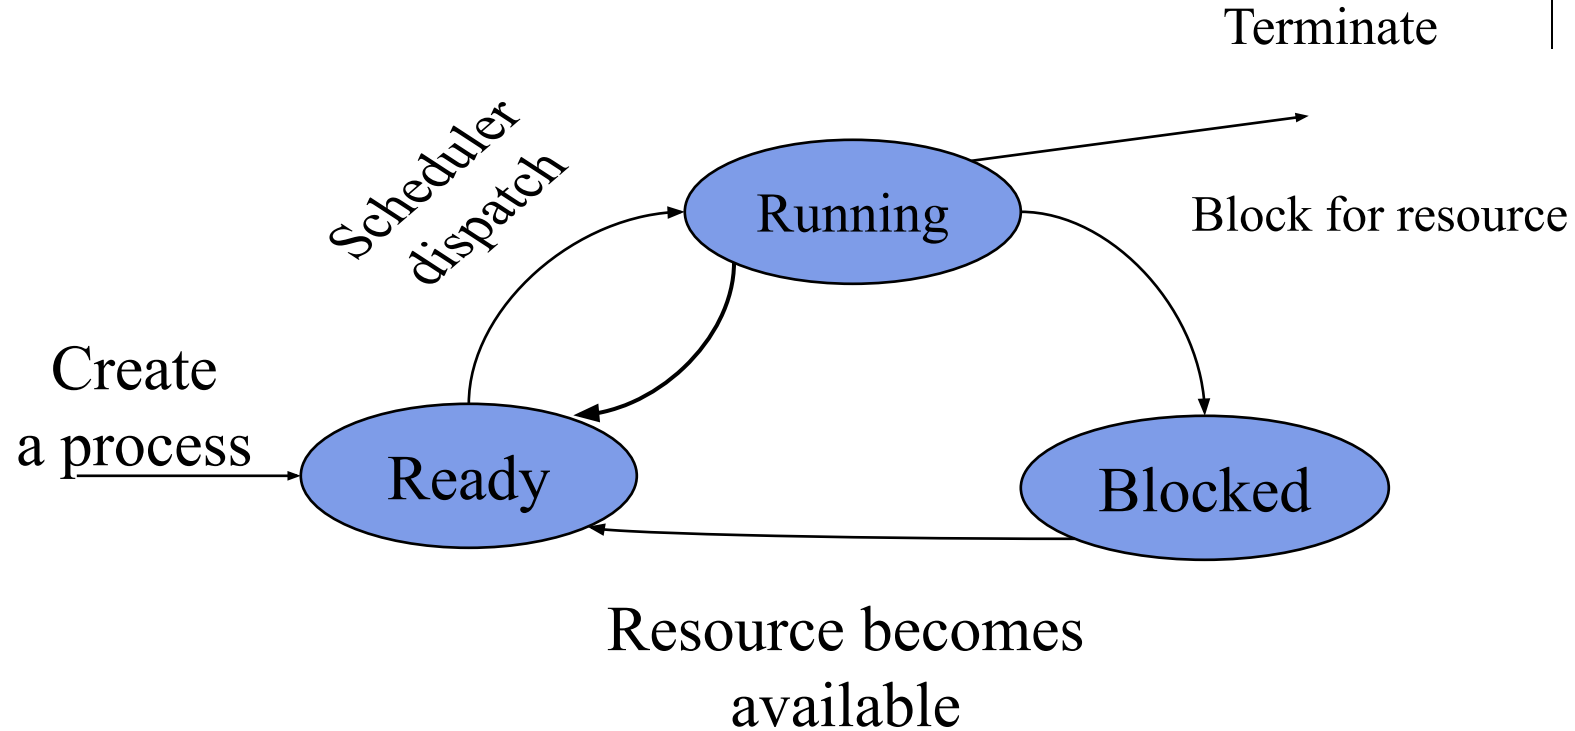


- Assuming single-threaded program
- No concurrency inside a process
- Everything happens sequentially
- Often with interleaved CPU/IO operations

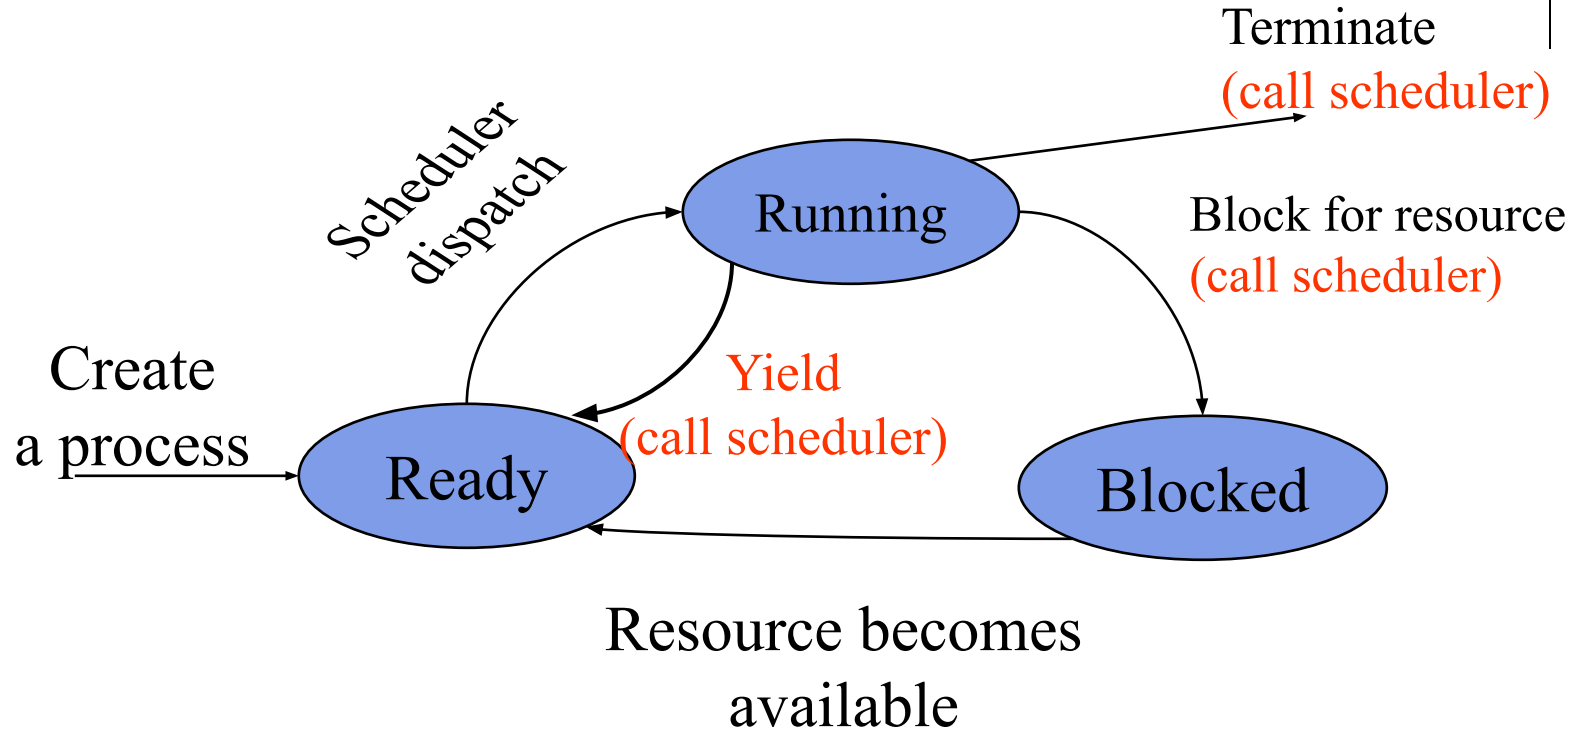
# Process Life Cycle



# Non-Preemptive Scheduling



# Non-Preemptive Scheduling




# Non-Preemptive Scheduling



- Any issues?
- What if a process runs:

```
int main() {  
    while(1);  
}
```



# Concurrent Processes



- Processes in a system can execute concurrently (multitasking)
- Motivations for allowing concurrent execution
  - Physical resource sharing (system utilization)
  - Computational speedup – with several CPUs
  - Modularity (chrome)
  - Convenience (desktop: chrome, google drive, clock, weather)
- Logical resource sharing (eg password files)



# Time Sharing Systems



- Timesharing systems support interactive use each user feels he/she has the entire machine
- How?
  - optimize response time
  - based on time-slicing

# Preemptive Scheduling



- Basic idea
  - before moving process to running, OS sets timer
  - if process yields/blocks, clear timer, go to scheduler
  - If timer expires, go to scheduler

# Preemptive Scheduling



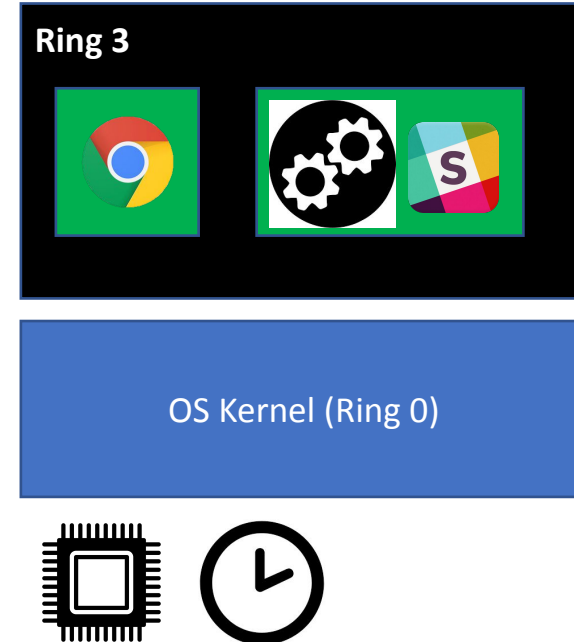
- How does the OS know that the timer expired?

```
int main() {  
    while(1);  
}
```

# Preemptive Scheduling



- Preemptive Multitasking (Lab 4)
- CPU generates an interrupt to force execution at kernel after some time quantum
  - E.g., 1000Hz, on each 1ms..





# Preemptive Scheduling



- Preemptive Multitasking (Lab 4)
- CPU generates an interrupt to force execution at kernel after some time quantum
  - E.g., 1000Hz, on each 1ms..

After  
1ms

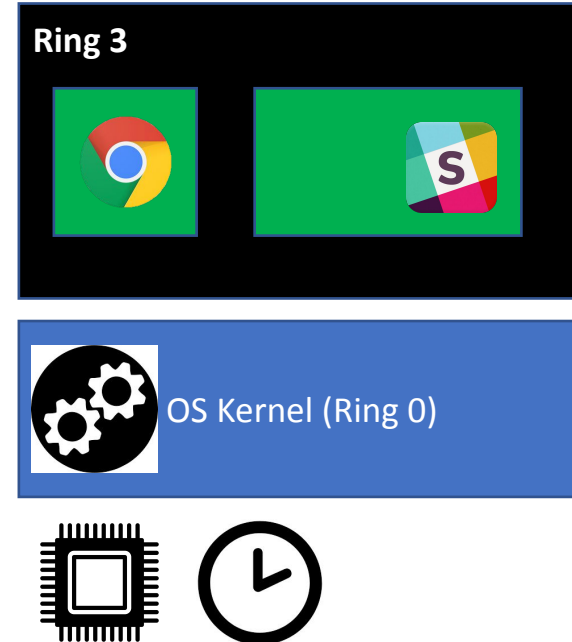


**Timer interrupt!**

# Preemptive Scheduling



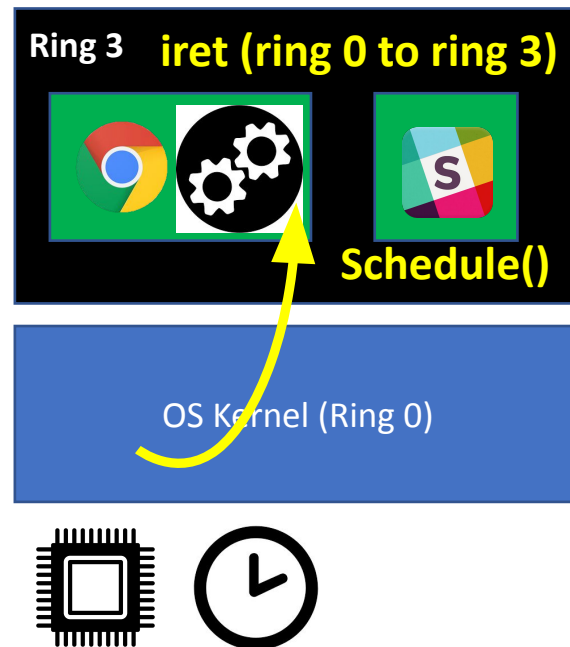
- Preemptive Multitasking (Lab 4)
- CPU generates an interrupt to force execution at kernel after some time quantum
  - E.g., 1000Hz, on each 1ms..
- Guaranteed execution in kernel
  - Let kernel mediate resource contention



# Preemptive Scheduling



- Preemptive Multitasking (Lab 4)
- CPU generates an interrupt to force execution at kernel after some time quantum
  - E.g., 1000Hz, on each 1ms..
- Guaranteed execution in kernel
  - Let kernel mediate resource contention



# Context Switch



- **Definition:** Switching the CPU to another process, which involves saving the state of the old process and loading the state of the new process
- **What state?**
- **Where to store them?**

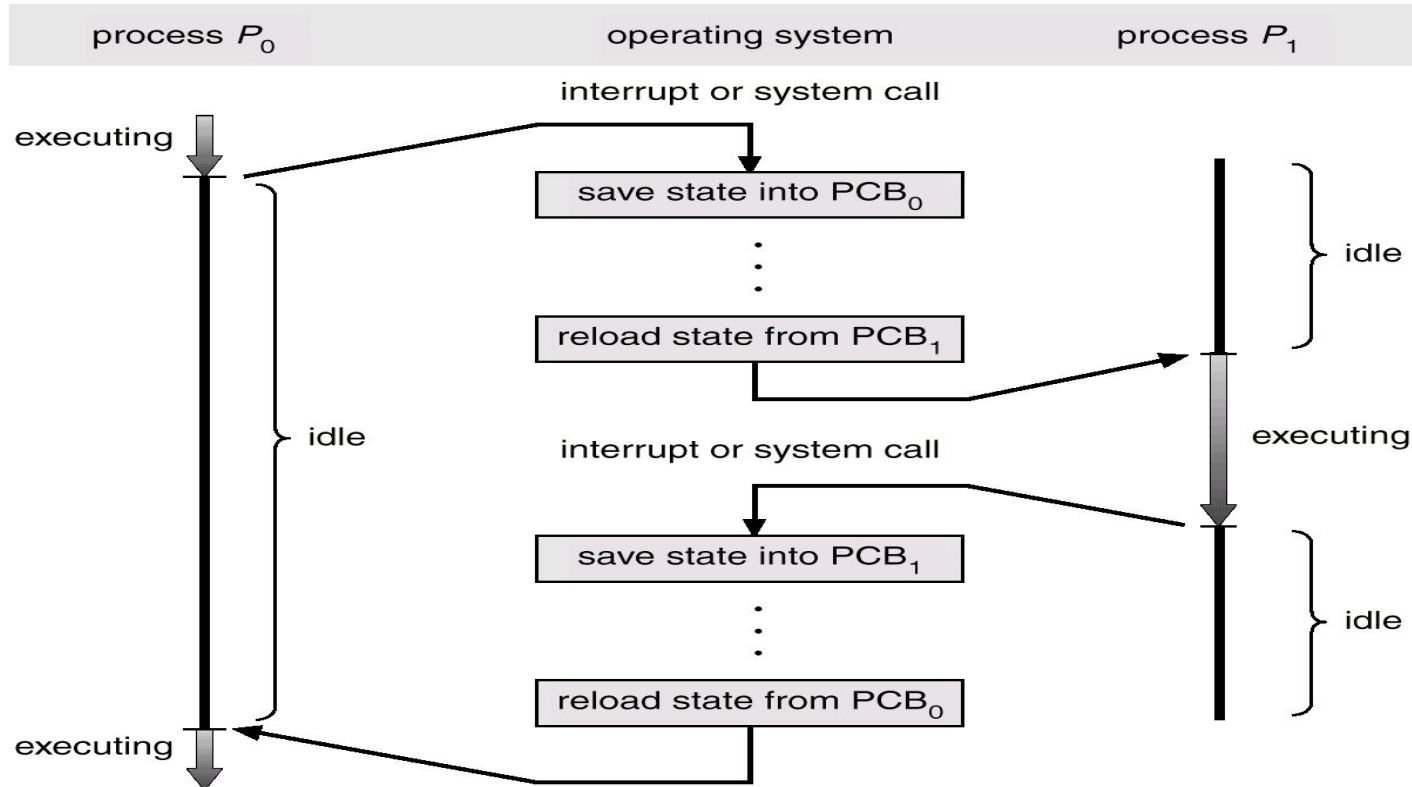


# Process State: Process Control Block (PCB)

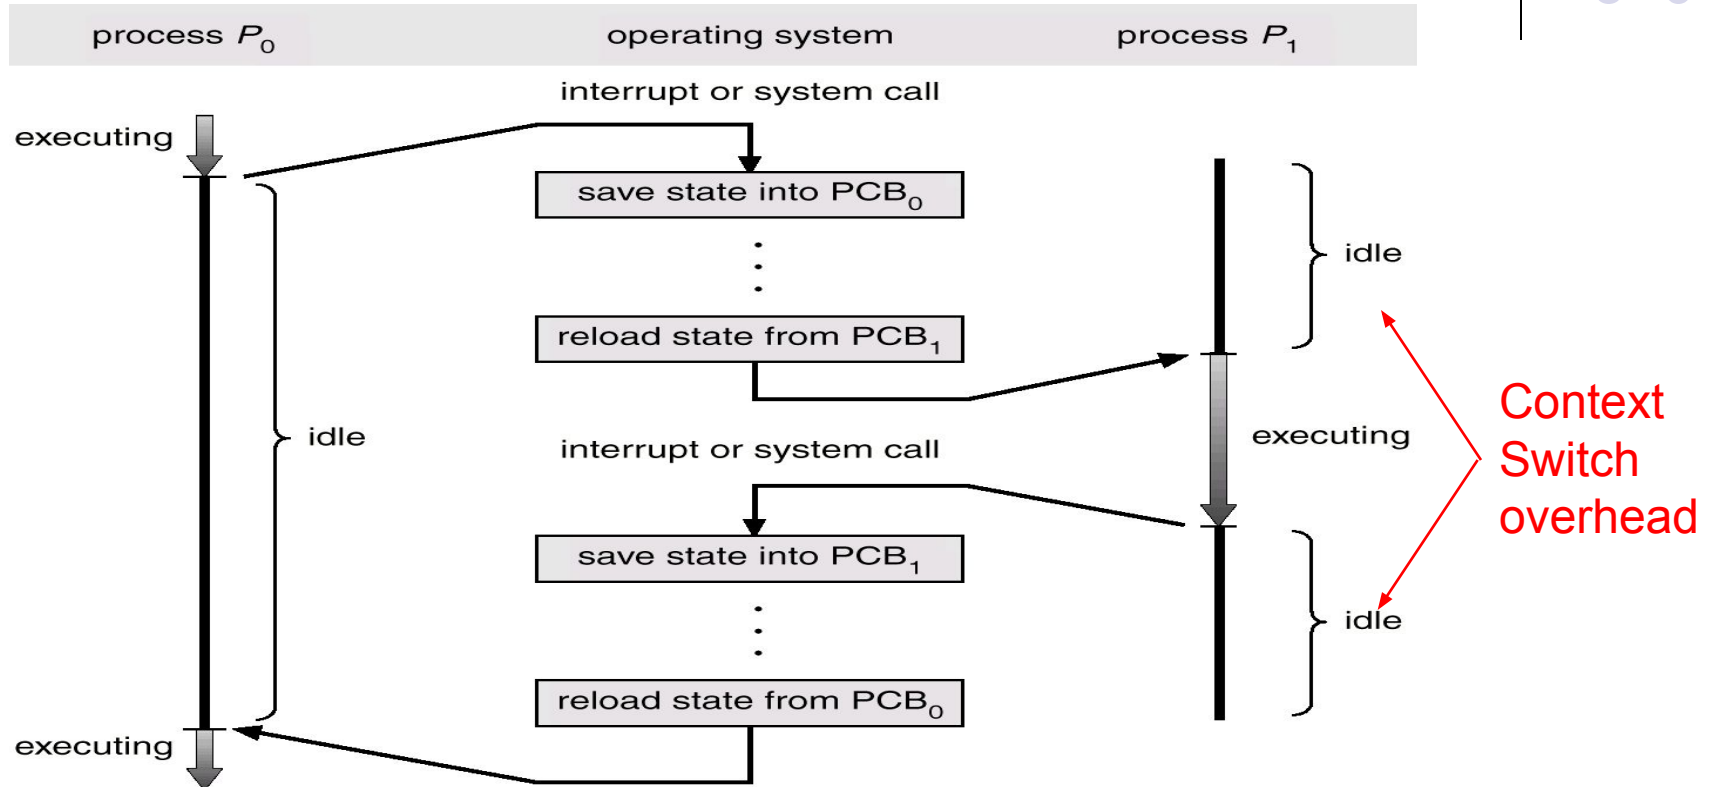


- A.K.A User Environment (JOS)
- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
- I/O and file management
  - Communication ports, directories, file descriptors, etc

# Context Switch



# Context Switch



# Preemptive Scheduling Considerations



- Timer granularity
  - **Finer timers = more responsive, high overhead**
  - **Coarser timers = less responsive, more efficient**
- CPU Accounting (CPU running stats)
  - Used by the scheduler
  - Useful for the programmer

# Preemptive Scheduling Considerations



- Mechanism + policy
- Mechanisms fairly simple:
  - Save state into a PCB and Restore state from another PCB

# Preemptive Scheduling Considerations



- Mechanism + policy
- Mechanisms fairly simple:
  - Save state into a PCB and Restore state from another PCB
- Policy choices harder:
  - When should we switch?

# Challenges in Policy



- Flexibility - variability in job types
  - Long vs. short
  - Interactive vs. non-interactive
  - I/O-bound vs. compute-bound
- Issues
  - Short jobs shouldn't suffer
  - (Interactive) Users shouldn't be annoyed

# Challenges in Policy (2)



- Fairness
  - All users should get access to CPU
  - Amount of CPU should be roughly even?
- Issue
  - Short-term vs. long-term fairness



# Goals and Assumptions



- Goals (Performance metrics)
  - Minimize turnaround time
    - avg time to complete a job
    - $T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$
  - Maximize throughput
    - operations (jobs) per second
    - Minimize overhead of context switches: large quanta
    - Efficient utilization (CPU, memory, disk etc)
  - Short response time
    - $T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$
    - type on a keyboard
    - Small quanta
  - Fairness
    - fair, no starvation, no deadlock

# Goals and Assumptions




- Goals often conflict
  - Response time vs. throughput
  - fairness vs. avg turnaround time?
- Assumptions
  - One process/program per user
  - Programs are independent

# Goals and Assumptions



- Goals (Performance metrics)
  - Minimize turnaround time
    - avg time to complete a job
    - $T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$
  - Maximize throughput
    - operations (jobs) per second
    - Minimize overhead of context switches: large quanta
    - Efficient utilization (CPU, memory, disk etc)
  - Short response time
    - $T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$
    - type on a keyboard
    - Small quanta
  - Fairness
    - fair, no starvation, no deadlock



# Scheduling Policies




- Is there an optimal scheduling policy?
  - Even if we narrow down to one goal?
- But we don't know about future
  - Offline vs. online

# Scheduling Policies

- Round Robin
- SJCF
- SRTCF



# Round Robin



- Each runs a time slice or quantum: Fair
- How do you choose time slice?
  - Overhead vs. response time
  - Overhead is typically about 1% or less
  - Quantum typically between 10 ~ 100 millisec

# Is Fairness always good?



- Assume 10 jobs waiting to be scheduled, each takes 100 seconds
  - Assume no other overhead
  - Total CPU time? 1000 seconds, always
- Implications?
  - Last job always finishes at 1000 seconds
  - So what's the point of scheduling?

# Non-Preemptive Scheduling



- Job 1 – start 0, end 100
  - Job 2 – start 100, end 200
  - ...
  - Job 10 – start 900, end 1000
- 
- Average turnaround time =  $100 + 200 + \dots / 10 = 550$  sec



# Round Robin



- Assume each quantum is 1 second
- Job 0 – 0, 10, 20, 30, 40,..., 990
- Job 1 – 1, 11, 21, 31,..., 991
- Job 2 – 2, 12, 22, 32,..., 992
- ...
- Avg turnaround time =  $990 + 991 + \dots / 10 = 995$  sec

# Is Fairness always good?



- Unfair policy was faster!
- Job 10 always ended at the same time
- Round-Robin just hurt jobs 1-9 with no gain

# Why use Round Robin?



- Imagine 10 jobs
  - Jobs 1-9 are 100 seconds
  - Job 10 is 10 seconds
- Which policy is better now?

# Non-preemptive scheduling

- Jobs 1-9 are 100 seconds
- Job 10 is 10 seconds



# Non-preemptive scheduling



- Jobs 1-9 are 100 seconds
  - Job 10 is 10 seconds
- 
- Job 0 – start 0, end 100
  - Job 1 – start 100, end 200
  - Job 10 – start 900, end 910
- 
- Avg turnaround time =  $100+200+\dots+910/N = 541$

# Round Robin scheduling



- Jobs 1-9 are 100 seconds
  - Job 10 is 10 seconds
- 
- Job 0 – 0, 10, 20, ..., 900
  - Job 1 – 1, 11, 21, ..., 901
  - Job 10 – 9, 19, 29, ..., 99
- 
- Avg turnaround time =  $900 + 901 + 908 + 99 / 10 = 824$

# Round Robin scheduling



- Jobs 1-9 are 100 seconds
- Job 10 is 10 seconds

**9% work drop**

**2%** avg turnaround drop for  
FIFO

- Job 0 – 0, 10, 20, ..., 900
- Job 1 – 1, 11, 21, ..., 901
- Job 10 – 9, 19, 29, ..., 99

**17%** avg turnaround drop for  
RR

- Avg turnaround time =  $900 + 901 + 908 + 99 / 10 = 824$

# Why use Round Robin?




- Imagine 10 jobs
  - Job 1 is 100 seconds
  - Job 2-10 is 10 seconds
- Which policy is better now?
  - FIFO: average turnaround 145
  - RR: average turnaround 105



# STCF (SJF) – Shortest Job First



- What shall we do if we care about turn-around time?
  - FIFO can be bad



- STCF/SJF
  - schedule shortest (total completion time) job first





# SJF: Pros and Cons



- Can we do better than Shortest Job First in terms of average turnaround time?
  - Assume all jobs arrive at the beginning
- In fact, SJF can be proved to be the optimal scheduling algorithm with the above assumption
  - But we are not going to prove it, since this is not a theory class 😊
- SJF Advantage
  - Minimal average turnaround time
- Disadvantage
  - Difficult to know the future, has to run until finish

# STCF


- Shortest time to completion first (shortest job first)
  - Non-preemptive



# SRTCF



- Shortest remaining time to completion first
  - Preemptive



Any potential problems?  
- Can cause **starvation!**

# Policy Decisions



- Need to accommodate interactive jobs
  - Need some kind of RR
- Diversity in jobs – job length, I/O mix
  - RR also appears to help
- SJF also has virtue
  - Reduce avg. turnaround time
- Can we accommodate all?

# Scheduling Policies Advantages



FIFO

Response time

RR

Throughput

SJF

Avg. turnaround time

Fairness

# Scheduling Policies Advantages



FIFO

Response time

RR

Throughput

SJF

Avg. turnaround time

Fairness

# Scheduling Policies Advantages



FIFO

Response time

RR

Throughput

Avg. turnaround time

SJF

Fairness



# Scheduling Policies Advantages



FIFO

Response time


RR

Throughput


SJF

Avg. turnaround time

Fairness



# Scheduling Policies Advantages



# Scheduling Policy Issues



- Fairness
- Flexibility
- **High utilization (efficiency)**
- Good response time
- Good turnaround time

# Scheduling Policy Issues



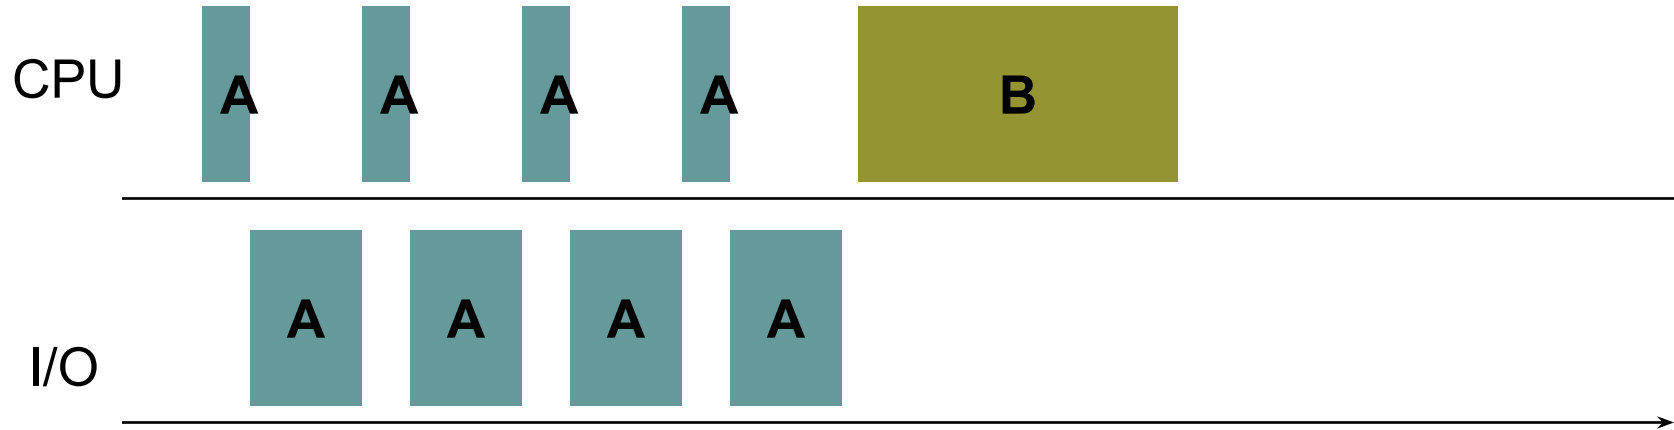
- High utilization (efficiency)
  - Lots of processes (want diff resources)
  - Lots of resources (want full parallelism)
- Issue?
  - How do you get the most useful work out of the system? (job throughput)

# Adding I/O into mix



- Resource utilization example
  - A and B each uses 100% CPU
  - C loops forever (1ms CPU and 10ms disk)
  - Time slice 99ms: roughly 30% of disk utilization with Round Robin and roughly 70% of CPU utilization
  - Time slice 1ms: roughly 90% of disk utilization with Round Robin and nearly 100% of CPU utilization
- What do we learn from this example?
  - Small time slice can improve utilization / fairness to I/O jobs

# Handling I/Os



# Handling I/Os

