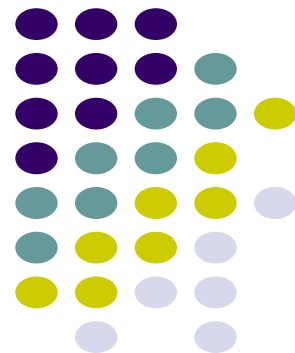


# Multi Threading and Synchronization

---

ECE 469, Feb 24

Aravind Machiry



# Web Server Example



- How does a web server handle 1 request?
- A web server needs to handle many concurrent requests
- Solution 1:
  - Have the parent process fork as many processes as needed
  - Processes communicate with each other via inter-process communication

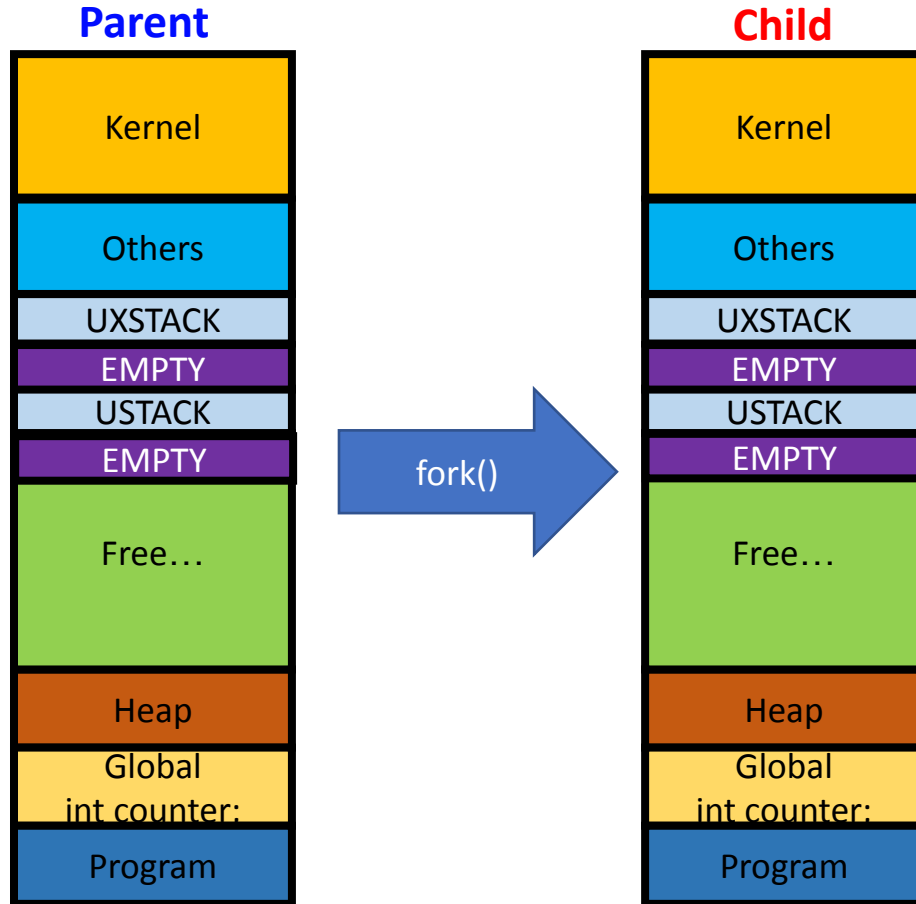
# Multiple Processes



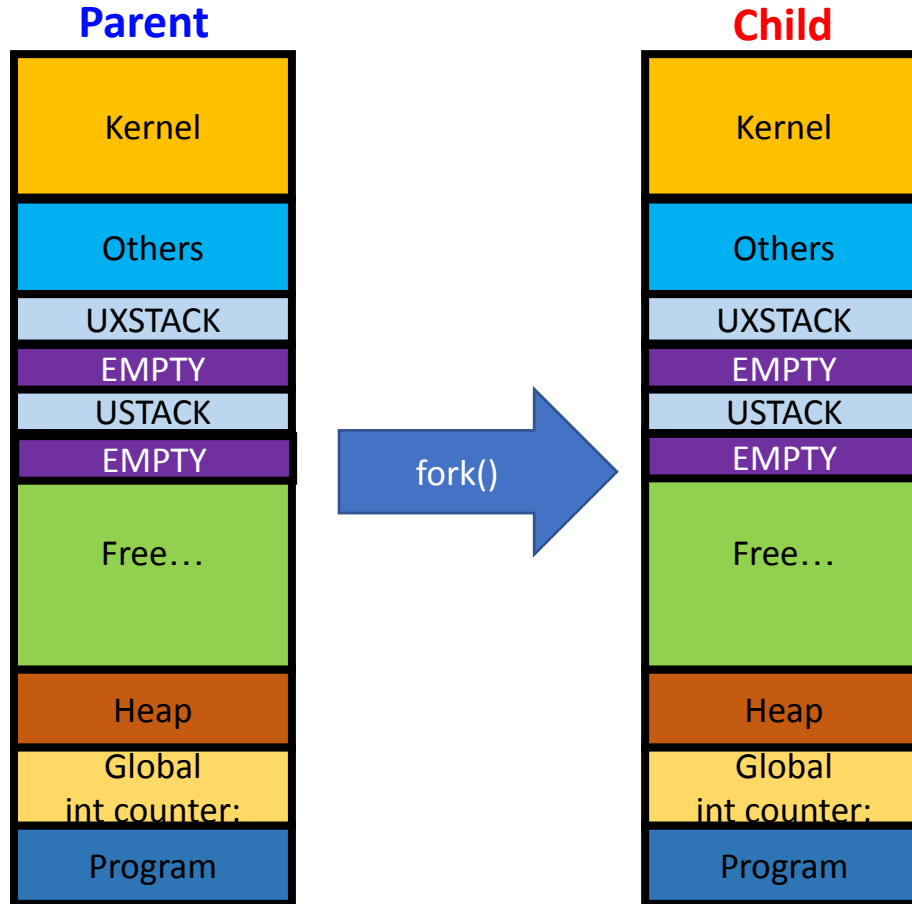
Parent



# Multiple Processes



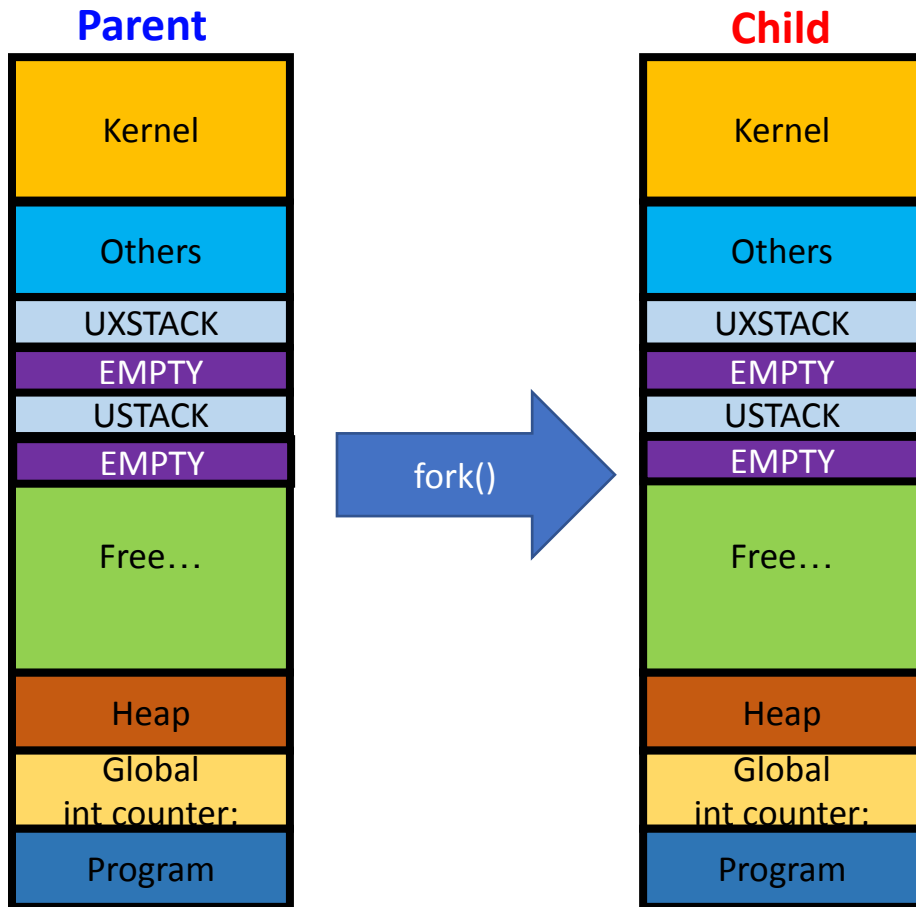
# Multiple Processes



Fork() creates new process by copying memory space

Process creates a new PRIVATE memory space

# Multiple Processes



**Fork() creates new process by copying memory space**

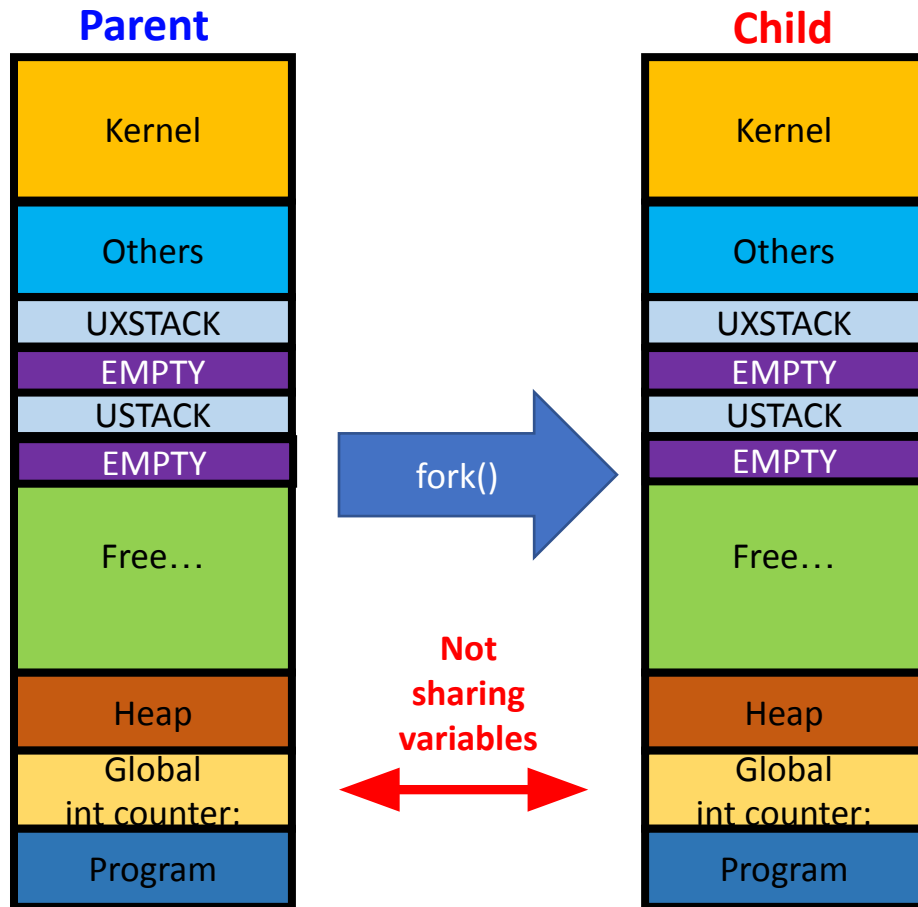
```
#include <stdio.h>
#include <unistd.h>

int counter;
volatile int value = 1;

void countup() {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
}

int main() {
    pid_t pid = fork();
    countup();
    printf("%s: %d\n", pid ? "Parent" : "Child", counter);
}
```

# Multiple Processes



Fork() creates new process by copying memory space

```
#include <stdio.h>
#include <unistd.h>

int counter;
volatile int value = 1;

void countup() {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
}

int main() {
    pid_t pid = fork();
    countup();
    printf("%s: %d\n", pid ? "Parent" : "Child", counter);
}
```

**Parent: 1000000**  
**Child: 1000000**

# How do Process communicate?



- At process creation time
  - Parents get one chance to pass everything at `fork()`
- OS provides generic mechanisms to communicate
  - Shared Memory: multiple processes can read/write same physical portion of memory; implicit channel
    - System call to declare shared region
    - No OS mediation required once memory is mapped
  - Message Passing: explicit communication channel provided through `send()/receive()` system calls
    - A system call is required



# How do Process communicate?



- IPC is, in general, expensive due to the need for system calls
  - Although many OSes have various forms of lightweight IPC

# The Soul of a Process



- But all the processes in the web-server are **cooperating!**
  - They all share the same code and data (address space)
  - They all share the same privileges
  - They all share the same resources (files, sockets, etc.)
- What don't they share?
  - Each has its own execution state: PC, SP, and registers

# The Soul of a Process



- Key idea: Why don't we **separate the concept of a process from its execution state**?
  - Process: address space, privileges, resources, etc.
  - Execution state: PC, SP, registers
- Exec state also called thread of control, or thread

# Threads




- **Separate** the concepts of a “thread of control” (PC, SP, registers) from the rest of the process (address space, resources, accounting, etc.)
- Modern OSes support two entities:
  - the **task** (process), which defines an address space, a resource container, accounting info
  - the **thread** (lightweight process), which defines a single sequential execution stream within a task (process)

# Threads vs. Process



- There can be several threads in a single address space
- Threads are the unit of scheduling; tasks are containers (address space, other shared resources) in which threads execute

# Single threaded v/s multithreaded



# What differs in threads of a process?



- A.K.A User Environment (JOS)
- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
  - Code, data, heap, execution stack
- I/O and file management
  - Communication ports, directories, file descriptors, etc

# What differs in threads of a process?



- A.K.A User Environment (JOS)
- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
  - Code, data, heap, execution stack
- I/O and file management
  - Communication ports, directories, file descriptors, etc



# Thread Control Block

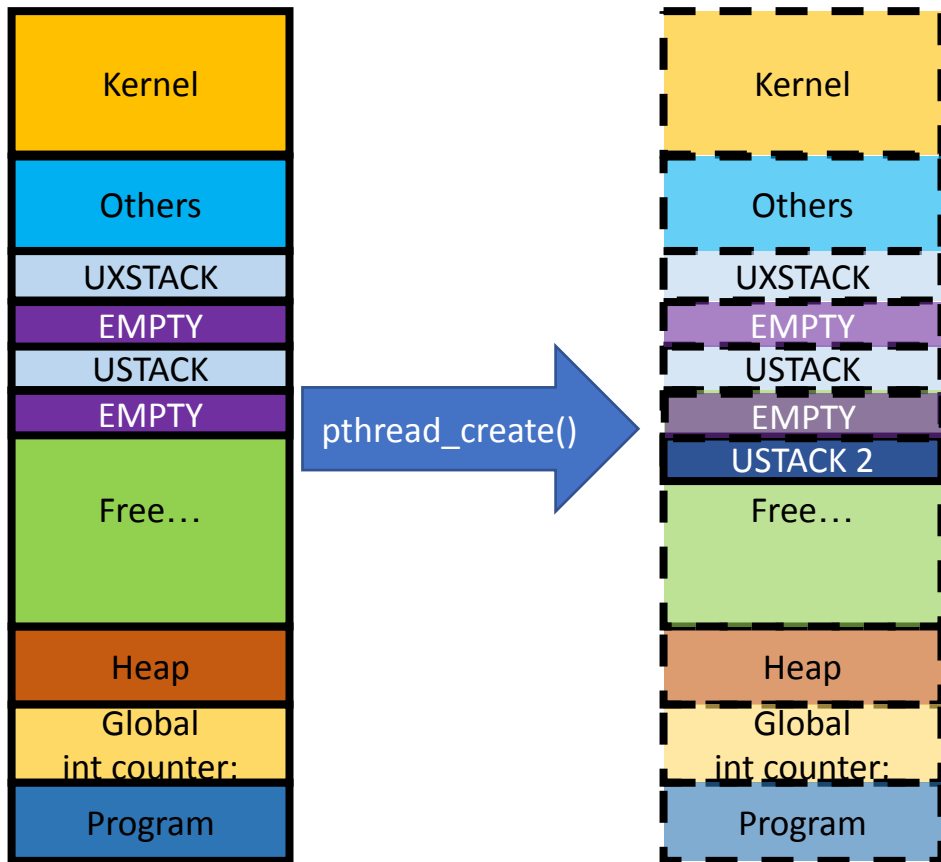


- Shared information
  - Process info: parent process
  - Memory: code/data segments, page table, and stats
  - I/O and file: comm ports, open file descriptors
- Private state
  - State (ready, running and blocked)
  - PC, Registers
  - Execution stack

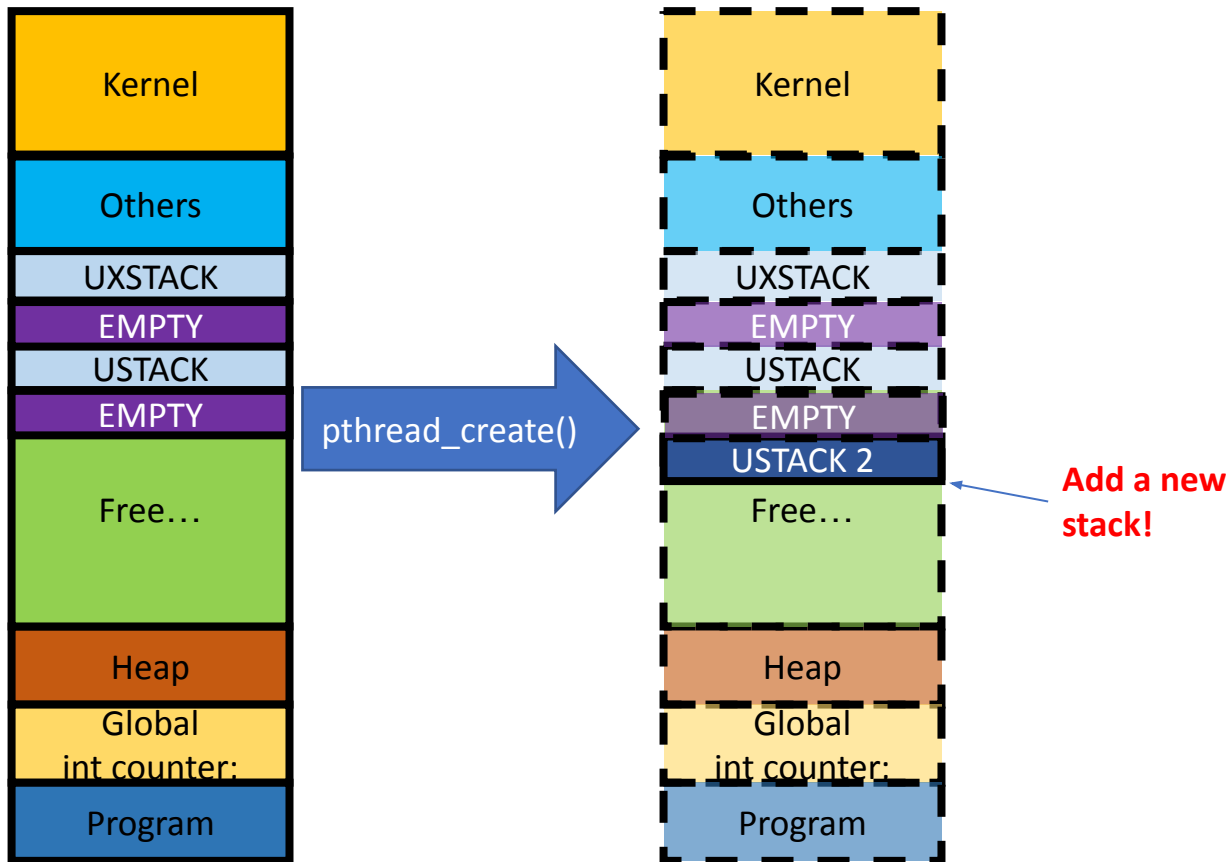
# Threads



# Threads



# Threads



# Threads



`pthread_create()`



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

**Add a new  
stack!**

# Threads



pthread\_create()



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

Add a new  
stack!

Adding  
value..

# Threads



pthread\_create()



```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int counter;
volatile int value = 1;

void * countup(void *arg) {
    for(int i=0; i<1000000; ++i) {
        counter += value;
    }
    printf("%s: %d\n", arg ? "Parent" : "Child", counter);
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, countup, NULL);
    countup((void*) 1);
    pthread_join(thread, NULL);
}
```

Add a new  
stack!

Adding  
value..

The same  
variable..

# Programming with Threads



- Flexible, but error-prone, since there no protection between threads
  - In C/C++,
    - automatic variables are private to each thread
    - global variables and dynamically allocated memory (malloc) are **shared**
- Need synchronization!



# The need for synchronization!



- Cooperating processes may share data via
  - shared address space (code, data, heap) by using threads
  - Files
  - (Sending messages)
- What can happen if processes try to access shared data (address) concurrently?
  - Sharing bank account with sibling:  
At 3pm: If (balance > \$10) withdraw \$10
- How hard is the solution?

# “Too much milk” Problem



## Person A

1. Look in fridge: out of milk
2. Leave for Walmart
5. Arrive at Walmart
6. Buy milk
7. Arrive home

## Person B

3. Look in fridge: out of milk
4. Leave for Walmart
8. Arrive at Walmart
9. Buy milk
10. Arrive home

- How to put in a locking mechanism?

# Possible Solution 1



Person A

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Person B

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

# Will this work?



Person A

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Person B

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

# Will this work?



Person A

```
1.if ( noMilk ) {  
  2.if (noNote) {  
    5.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

Person B

```
3.if ( noMilk ) {  
  4.if (noNote) {  
    6.leave note;  
    buy milk;  
    remove note;  
  }  
}
```

- Process can get context switched after checking milk and note, but before leaving note

# Why does this work for humans?



- Human can perform *test* (look for other person & milk) and *set* (leave note) at the same time.

# Possible Solution 2



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

# Will this work?



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```



# Will this work?



Person A

```
leave noteA
if (no noteB) {
    if (noMilk) {
        buy milk
    }
}
remove noteA
```

Person B

```
leave noteB
if (no noteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

- We may not have Milk: Both process can leave note and skip buying milk

# Possible Solution 3



## Process A

```
leave noteA
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
remove noteA
```

## Process B

```
leave noteB
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

# Will this work?



## Process A

```
leave noteA
while (noteB)
    do nothing;
if (noMilk)
    buy milk;
remove noteA
```

## Process B

```
leave noteB
if (noNoteA) {
    if (noMilk) {
        buy milk
    }
}
remove noteB
```

# Works, but complicated!



## Process A

leave noteA

```
while (noteB)
  do nothing;
```

```
if (noMilk)
```

```
  buy milk;
```

remove noteA

## Process B

leave noteB

```
if (noNoteA) {
  if (noMilk) {
    buy milk
  }
}
```

remove noteB

- A's code is **different** from B's
- **busy waiting** is a waste

# How can we solve this?





- Root cause: **Data Race**
- A thread's execution result could be inconsistent if other threads intervene its execution...

- counter += value

|                                 |                  |                                  |   |
|---------------------------------|------------------|----------------------------------|---|
| • <code>edx = value;</code>     | <code>mov</code> | <code>0x20087b(%rip),%edx</code> | <code># 0x201010 &lt;value&gt;</code>   |
| • <code>eax = counter;</code>   | <code>mov</code> | <code>0x20087d(%rip),%eax</code> | <code># 0x201018 &lt;counter&gt;</code> |
| • <code>eax = edx + eax;</code> | <code>add</code> | <code>%edx,%eax</code>           |   |
|                                 | <code>mov</code> | <code>%eax,0x200875(%rip)</code> | <code># 0x201018 &lt;counter&gt;</code> |
| • <code>counter = eax;</code>   |                  |                                  |   |


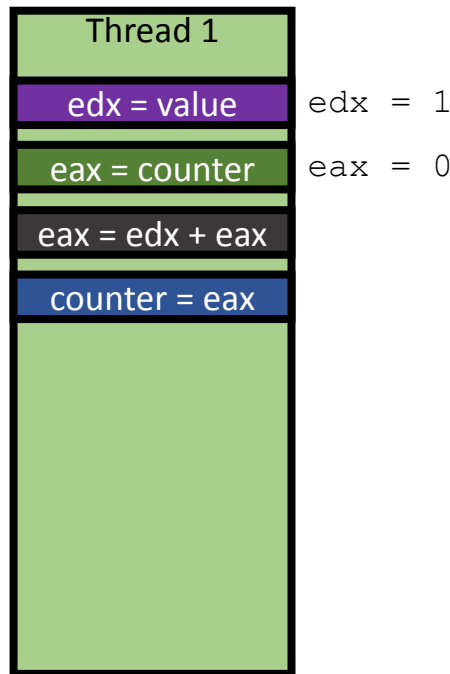
# Shared variable: No race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;




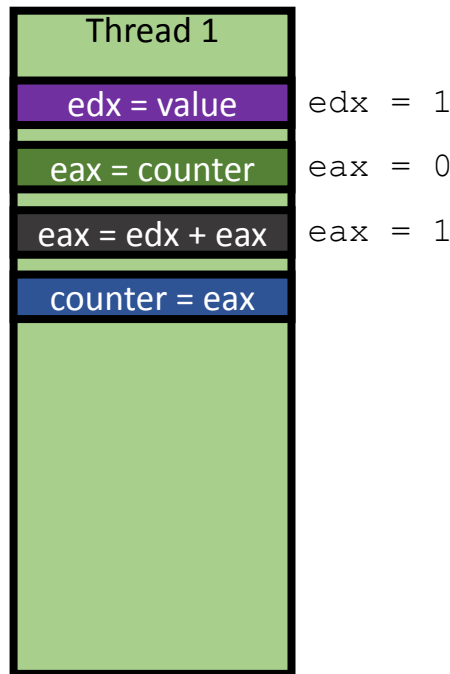
# Shared variable: No race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: No race


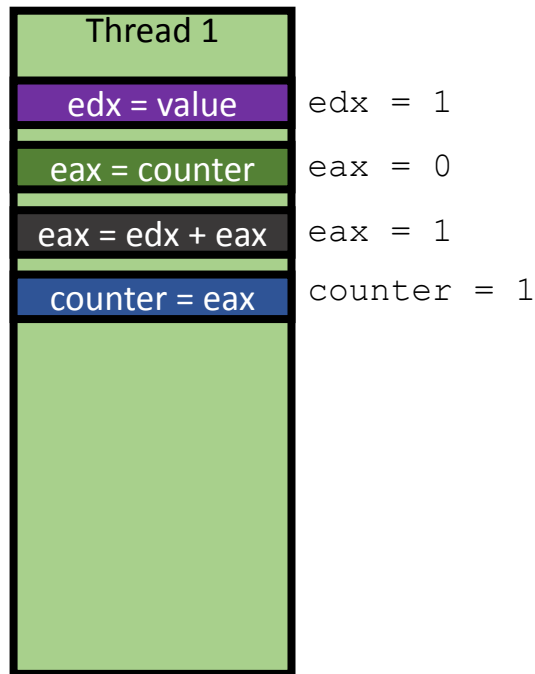
- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;





# Shared variable: No race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;


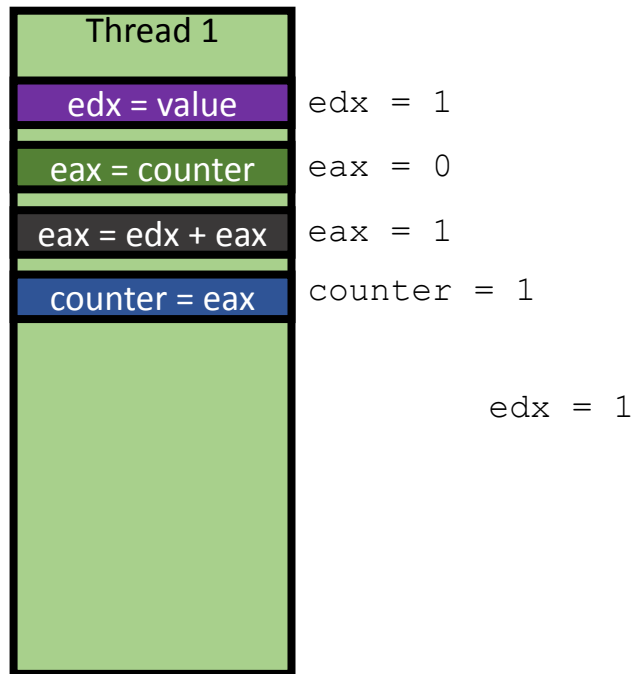


# Shared variable: No race




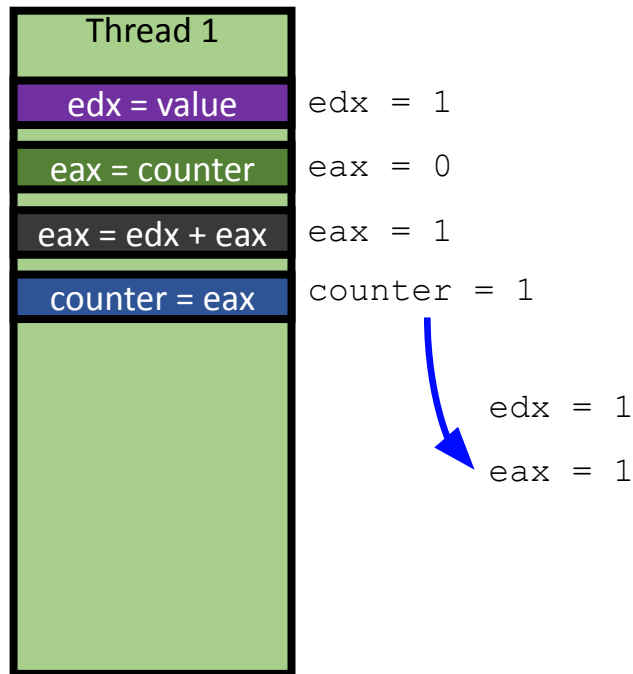
- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`

- Assume counter = 0 at start, and value = 1;




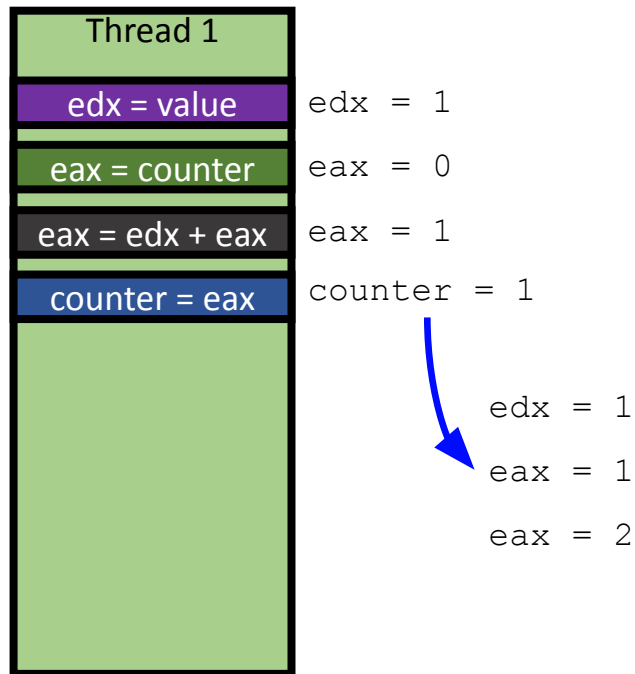
# Shared variable: No race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: No race

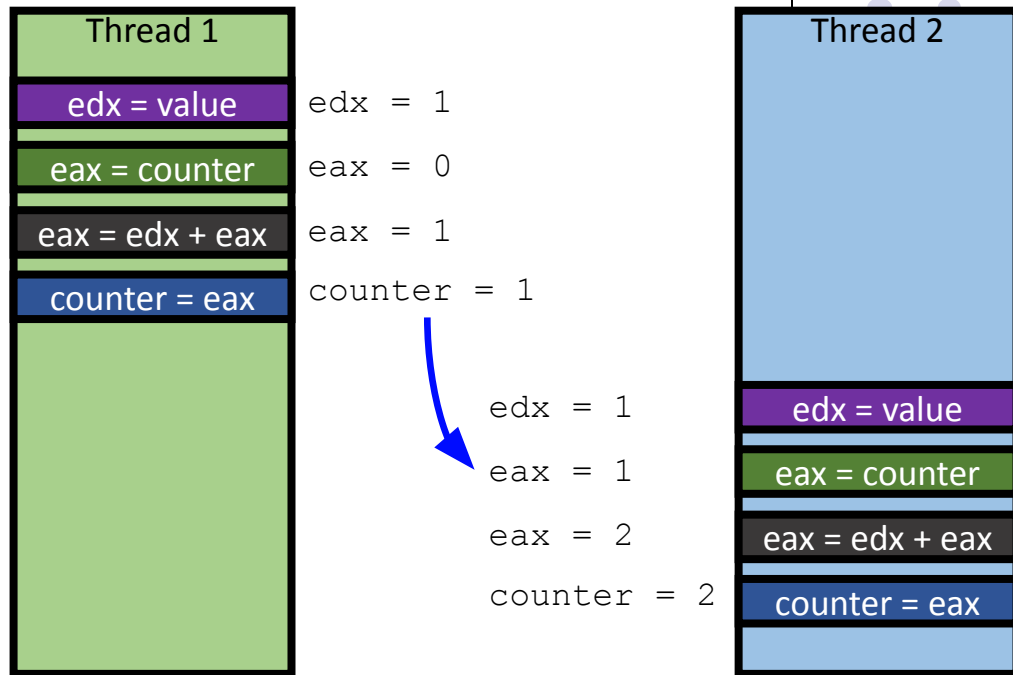
- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: No race



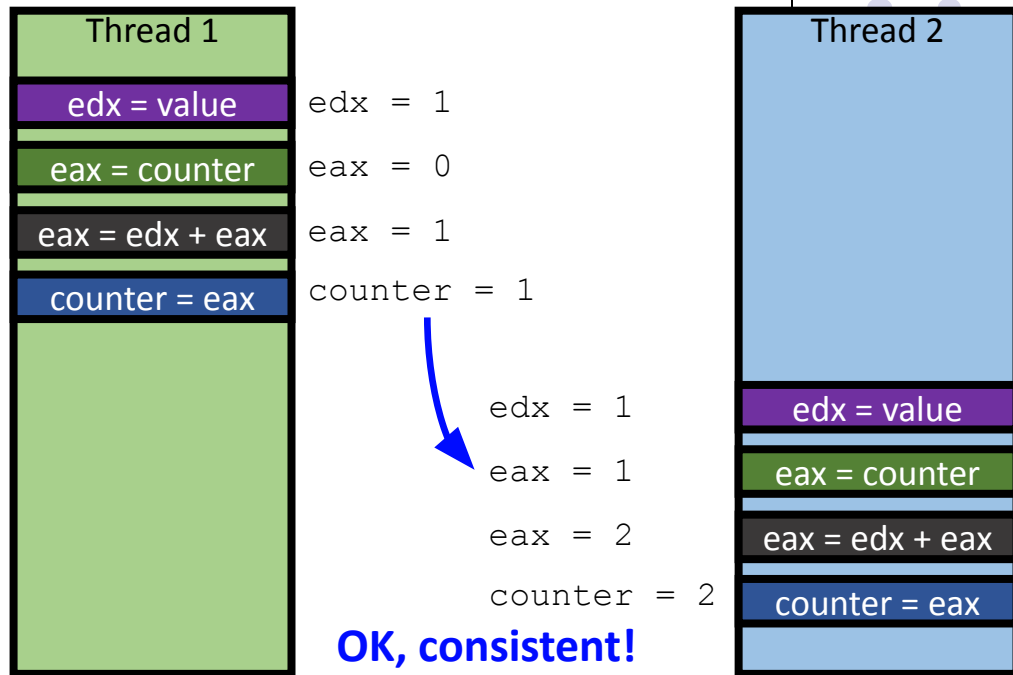
- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: No race





- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: Data race





- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: Data race

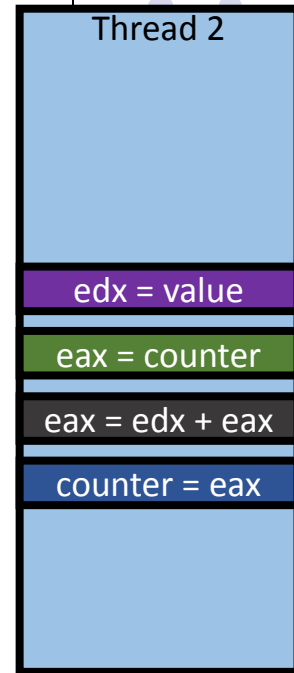


- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



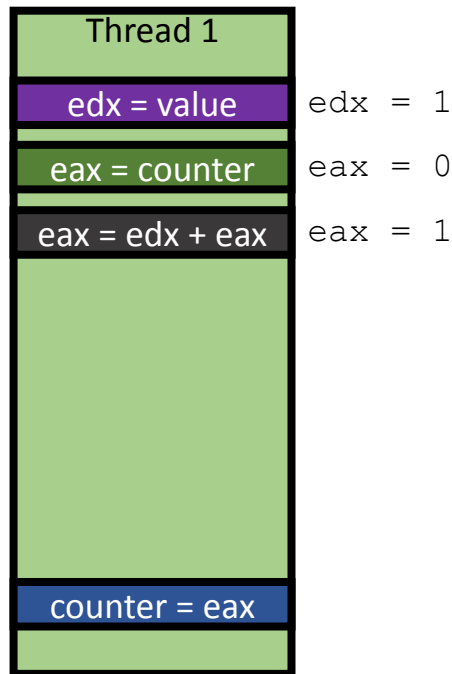


- 
- The diagram illustrates the execution flow of Thread 1. It consists of a vertical stack of colored boxes representing instructions or states, with corresponding register values shown to the right.
- Thread 1** (Green box): The initial state of the thread.
  - edx = value** (Purple box): The register `edx` is assigned the value 1. To the right, `edx = 1` is shown.
  - eax = counter** (Green box): The register `eax` is assigned the value 0. To the right, `eax = 0` is shown.
  - eax = edx + eax** (Dark Grey box): The register `eax` is updated to the sum of `edx` and `eax`, resulting in 1. To the right, `eax = 1` is shown.
  - counter = eax** (Blue box): The variable `counter` is updated with the value of `eax`, which is 1.
- The diagram shows a sequence of operations where the register `edx` is initialized to 1, `eax` is initialized to 0, and then `eax` is updated to 1. Finally, the variable `counter` is set to the value of `eax`.




# Shared variable: Data race

- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



`edx = 1`




- 
- The diagram illustrates two threads, Thread 1 and Thread 2, and their execution flow. Thread 1 is represented by a green bar, and Thread 2 is represented by a blue bar. The execution flow is shown as a sequence of steps, with the state of variables `edx` and `eax` indicated next to each step.
- Thread 1 (Green Bar):**
- Initial state: `edx = value`
  - Step 1: `eax = counter` (State: `eax = 0`)
  - Step 2: `eax = edx + eax` (State: `eax = 1`)
  - Step 3: `counter = eax` (State: `counter = 1`)
- Thread 2 (Blue Bar):**
- Initial state: `edx = value`
  - Step 1: `eax = counter` (State: `eax = 0`)
  - Step 2: `eax = edx + eax` (State: `edx = 1`, `eax = 0`)
  - Step 3: `counter = eax` (State: `counter = 0`)
- The diagram shows that Thread 2's update to `counter` occurs before Thread 1's, leading to an incorrect final value of `counter`.

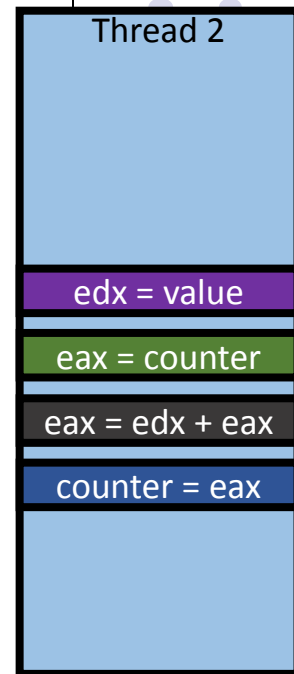
# Shared variable: Data race



- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



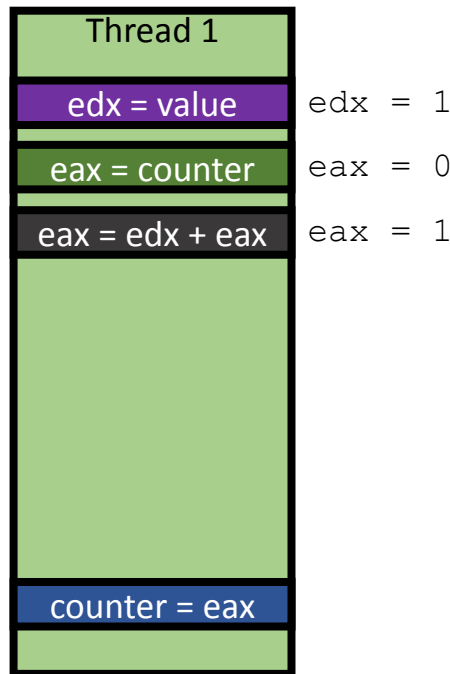
edx = 1  
eax = 0  
eax = 1



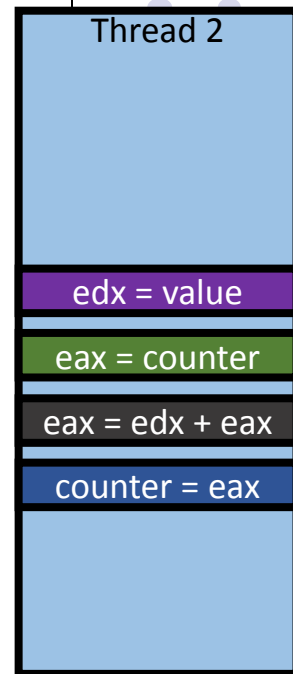
# Shared variable: Data race



- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



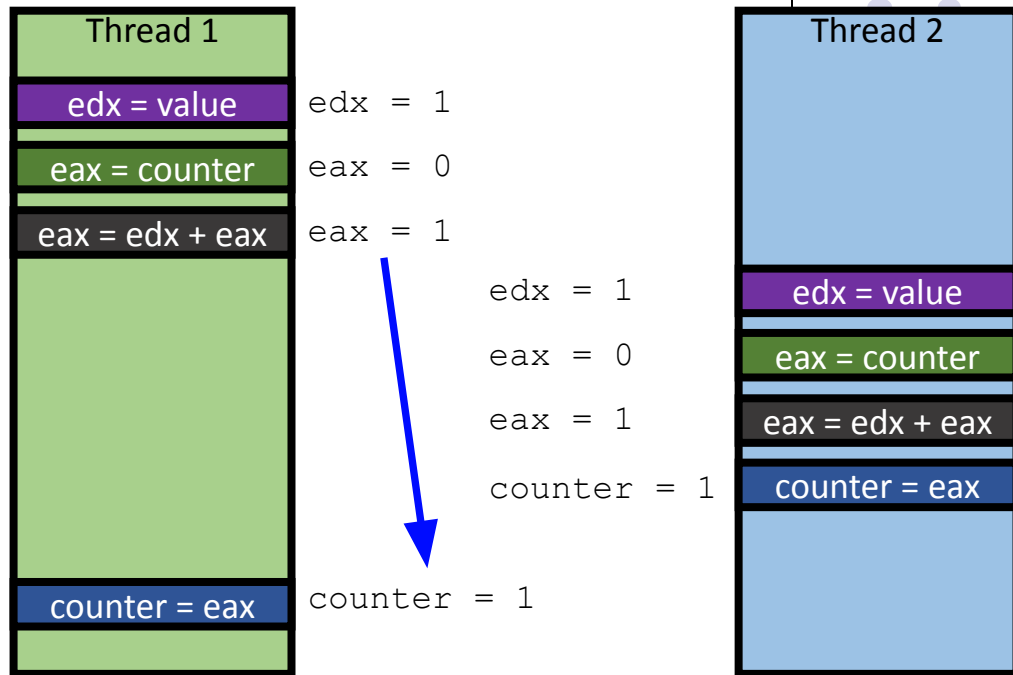
edx = 1  
eax = 0  
eax = 1  
counter = 1



# Shared variable: Data race

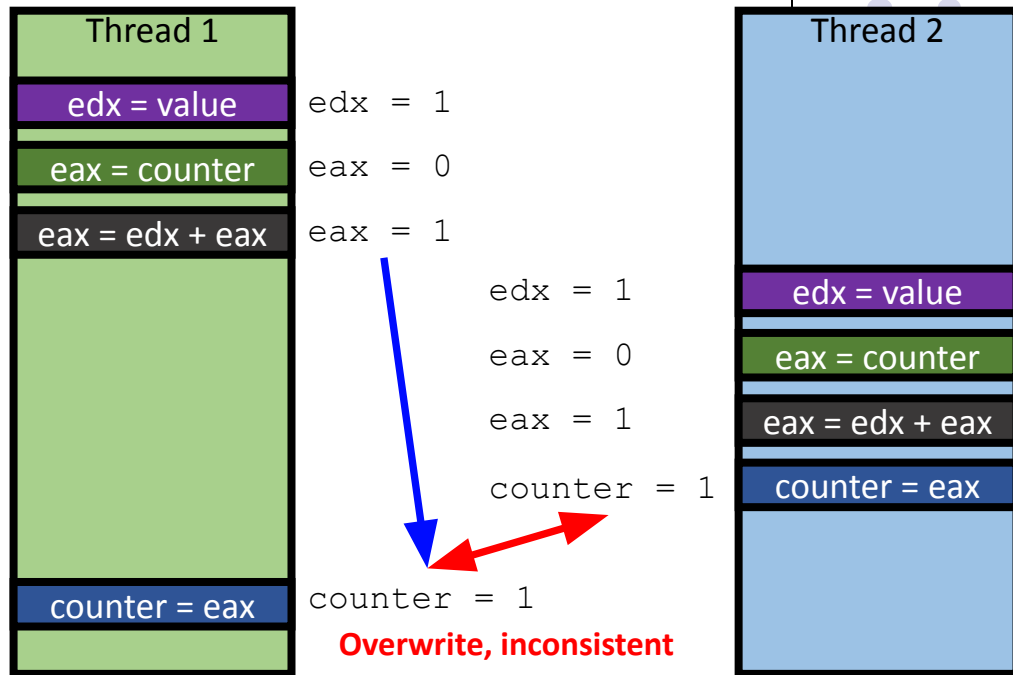


- counter += value
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume counter = 0 at start, and value = 1;



# Shared variable: Data race

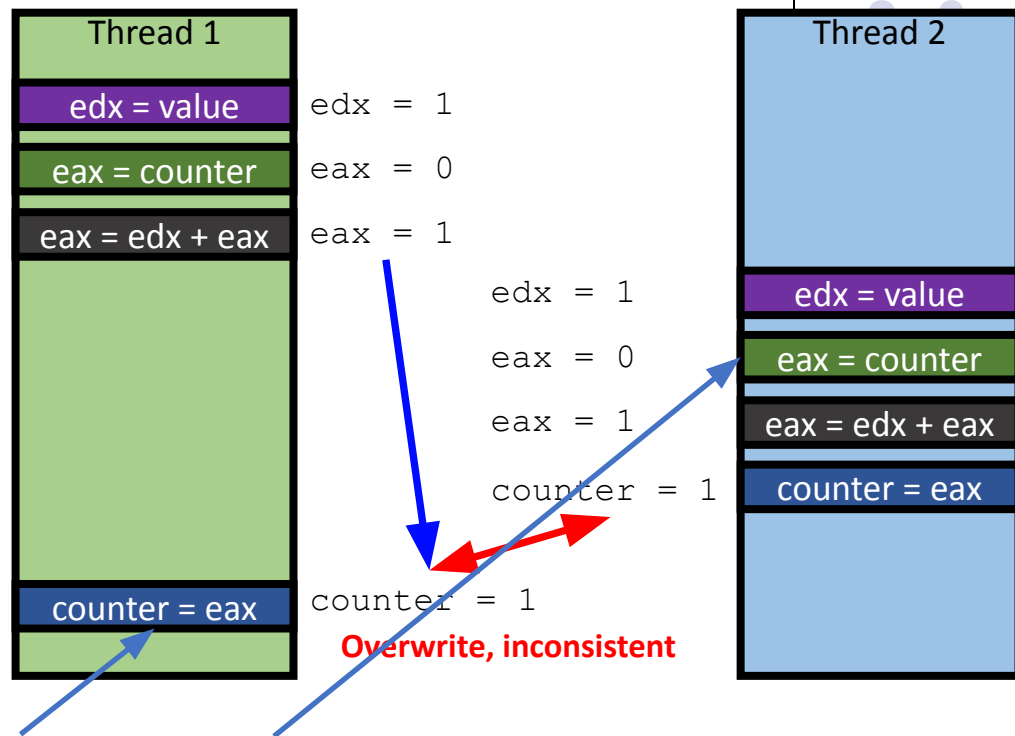
- `counter += value`
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume `counter = 0` at start, and `value = 1`;



# Shared variable: Data race



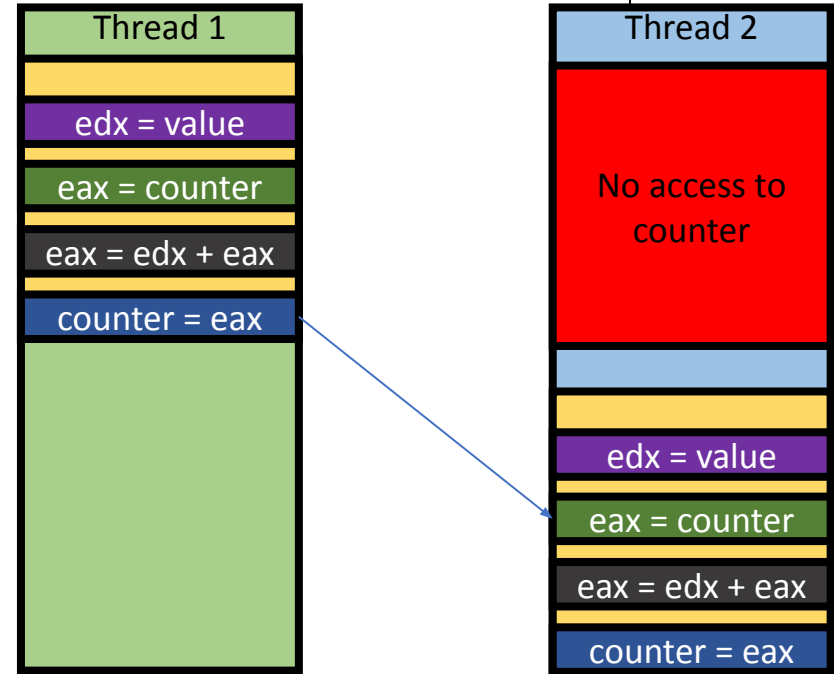
- `counter += value`
  - `edx = value;`
  - `eax = counter;`
  - `eax = edx + eax;`
  - `counter = eax;`
- Assume `counter = 0` at start, and `value = 1`;





# How can we prevent data races?

- What we need?
  - **Exclusive access** to counter (shared variable)




# How can we prevent data races?



- *Critical section* – a section of code, or collection of operations, in which only one process shall be executing at a given time
- *Mutual exclusion (Mutex)* - mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded)

# How can we prevent data races?

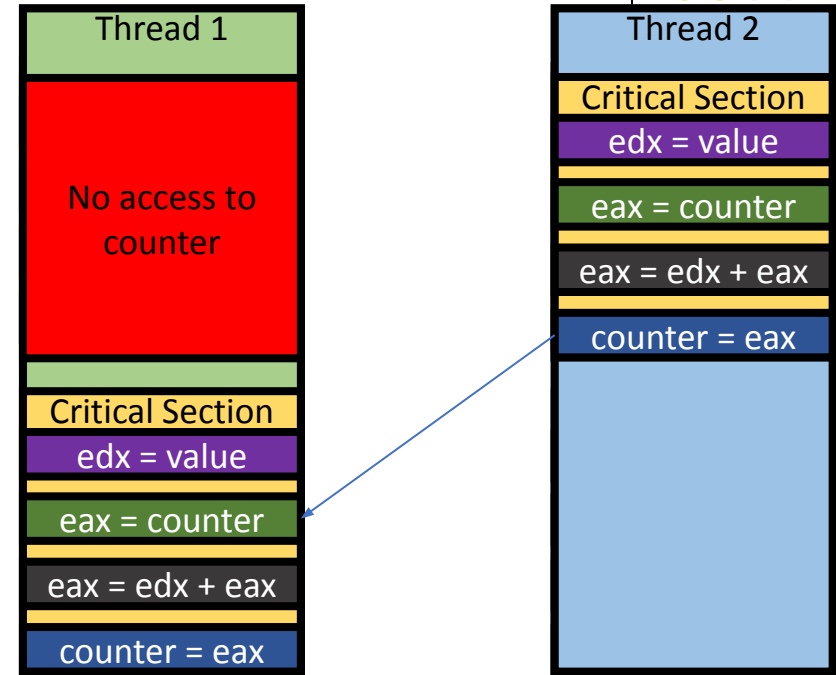
- **Mutual Exclusion / Critical Section**
  - Combine multiple instructions as a chunk
  - Let only one chunk execution runs
  - Block other executions




# How can we prevent data races?



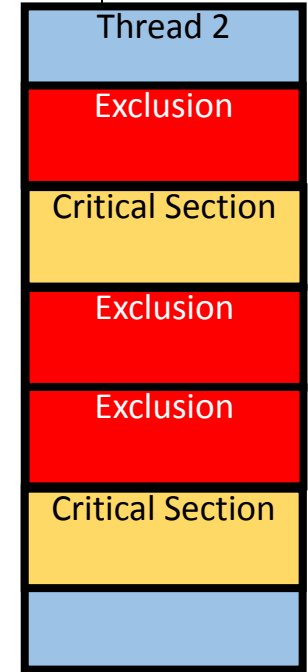
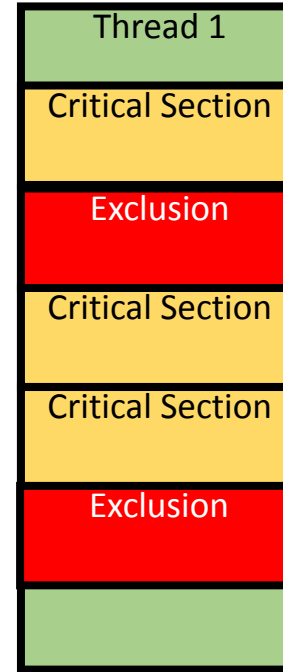
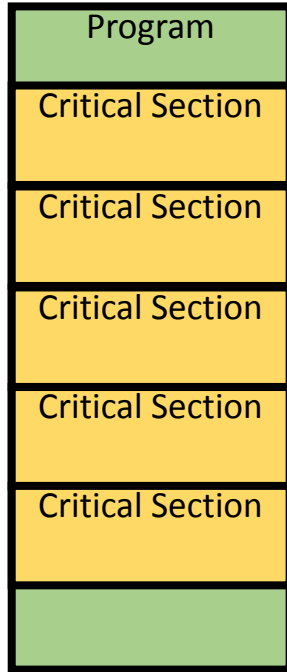
- **Mutual Exclusion / Critical Section**
  - Combine multiple instructions as a chunk
  - Let only one chunk execution runs
  - Block other executions






# Does mutex renders threading useless?



# Does mutex renders threading useless?



# Does mutex renders threading useless?



# Mutex Considerations



- Mutex can synchronize multiple threads and yield consistent result
  - No read before previous thread store the shared data
- Making the **entire program as critical section is meaningless**
  - Running time will be the same as single-threaded execution
- Apply critical section **as short as possible** to maximize benefit of having concurrency
  - Non-critical sections will run concurrently!