


Midterm Review

ECE 469, Mar 08

Aravind Machiry



Midterm

- Online on Brightspace
- Thursday, 14 March
- Time: 75 min
- Available Period: 7:30 am - 10:00 am
- Open notes

What happens, when we turn on the machine?

1. Power up.
2. BIOS initializes basic devices.
3. After initializing peripheral devices, it will put some initialization code to
 - a. DRAM physical address 0xffff0 ([f000:fff0])
 - b. Copy the code from ROM to RAM
 - c. Run from RAM
4. What does the code do? Load and run the boot sector from disk
 - a. Read the 1st sector from the boot disk (512 bytes)
 - b. Put the sector at 0x7c00
 - c. **Run it!** (set the instruction pointer = 0x7c00)

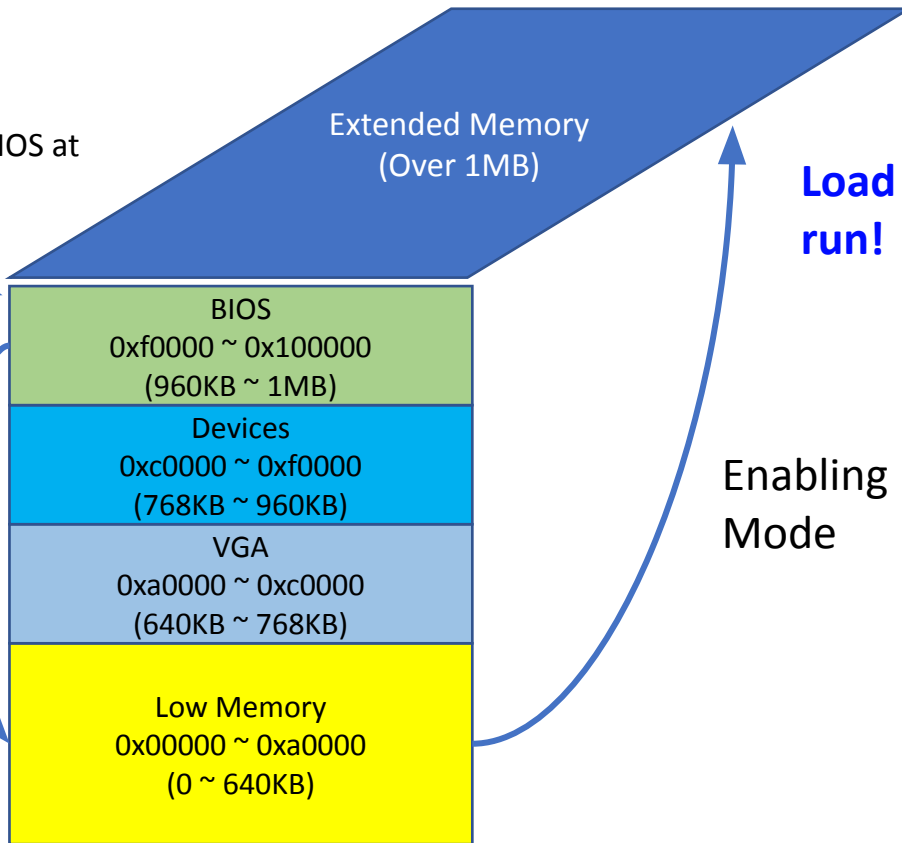
```
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: jmp $0xf000,$0xe05b
0x0000fff0 in ?? ()
```

Summary!



Read Master Boot Record
(MBR)
from the boot disk
and load it at 0x7c00

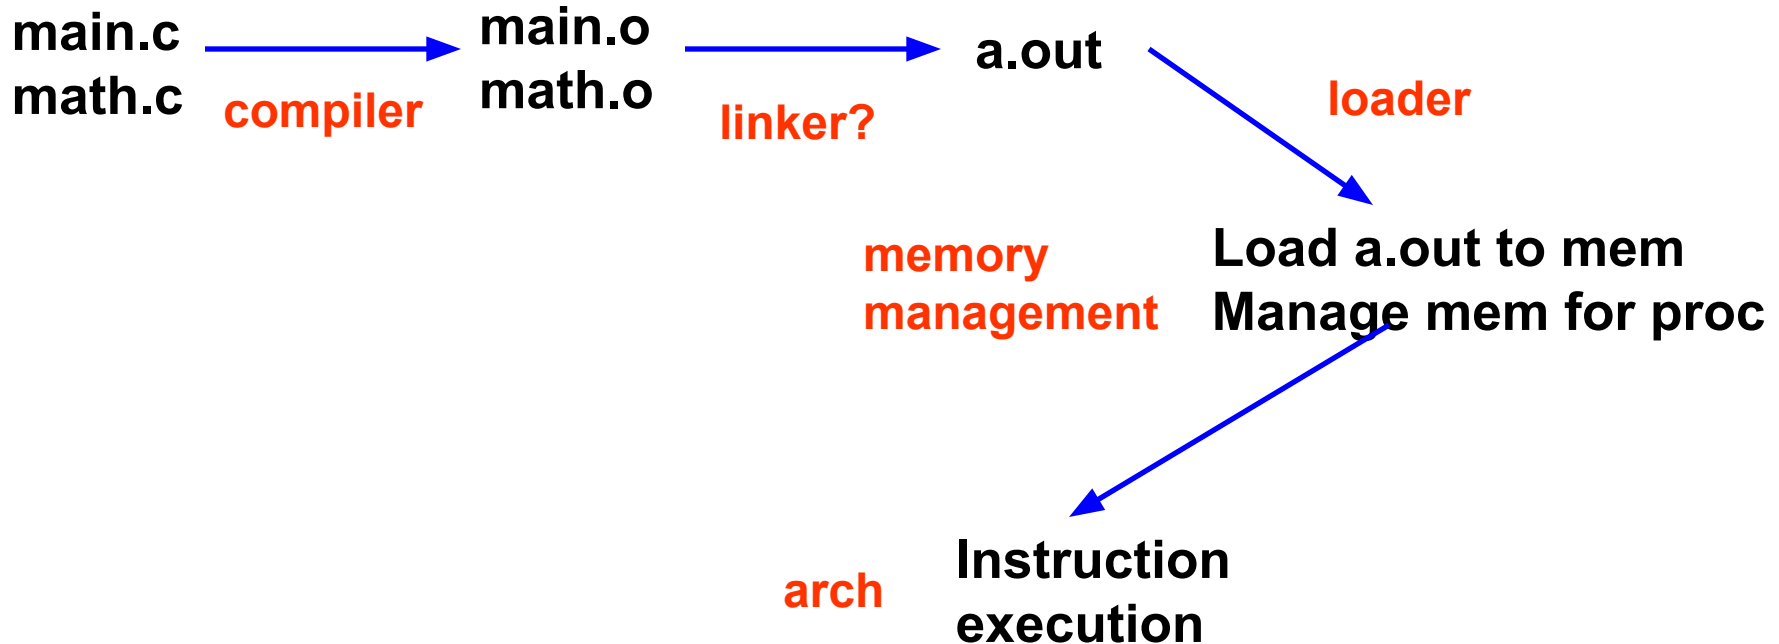
Map code in BIOS at
f000:ffff



Load kernel and
run!

Enabling Protected
Mode

A gap among Architecture, Compiler and OS courses



Virtual Memory


- Three goals
 - Transparency: does not need to know system's internal state
 - Program A is loaded at 0x8048000. Can Program B be loaded at 0x8048000?
 - Efficiency: do not waste memory; manage memory fragmentation
 - Can Program B (288KB) be loaded if 288 KB of memory is free, regardless of its allocation?
 - Protection: isolate program's execution environment
 - Can we prevent an overflow from Program A from overwriting Program B's data?

Paging!

- Idea: Make all chunks of memory the same size, called **pages**
 - Both virtual and physical memory divided into same size chunks.
- For each process, a **page table** defines the base address of each of that process' pages along with existence and read/write bits

Paging!


Virtual address



Virtual pages

physical pages

Physical memory



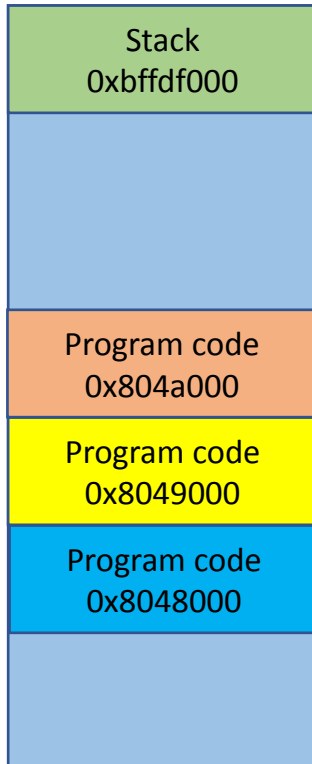
Page size / fragmentation

- If a page size is too small, it requires a big page table
 - 1B, 4GB
 - 4KB, 4MB
 - 4MB, 4KB
 - 1G, 16B
- If a page size is too big, unused memory in a page will be wasted
 - 1B - 1B (no waste)
 - 4KB - 1B
 - 4MB - 1B
 - 1G - 1B

Design consideration:
Memory fragmentation matters!

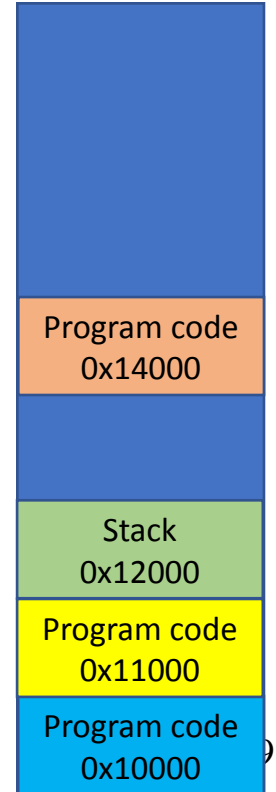
Virtual Memory - Paging

Having an indirect table that maps virt-addr to phys-addr




Virtual	Physical
0x8048000	0x10000
0x8049000	0x11000
0x804a000	0x14000
0xbffdf000	0x12000
...	...

Physical Memory

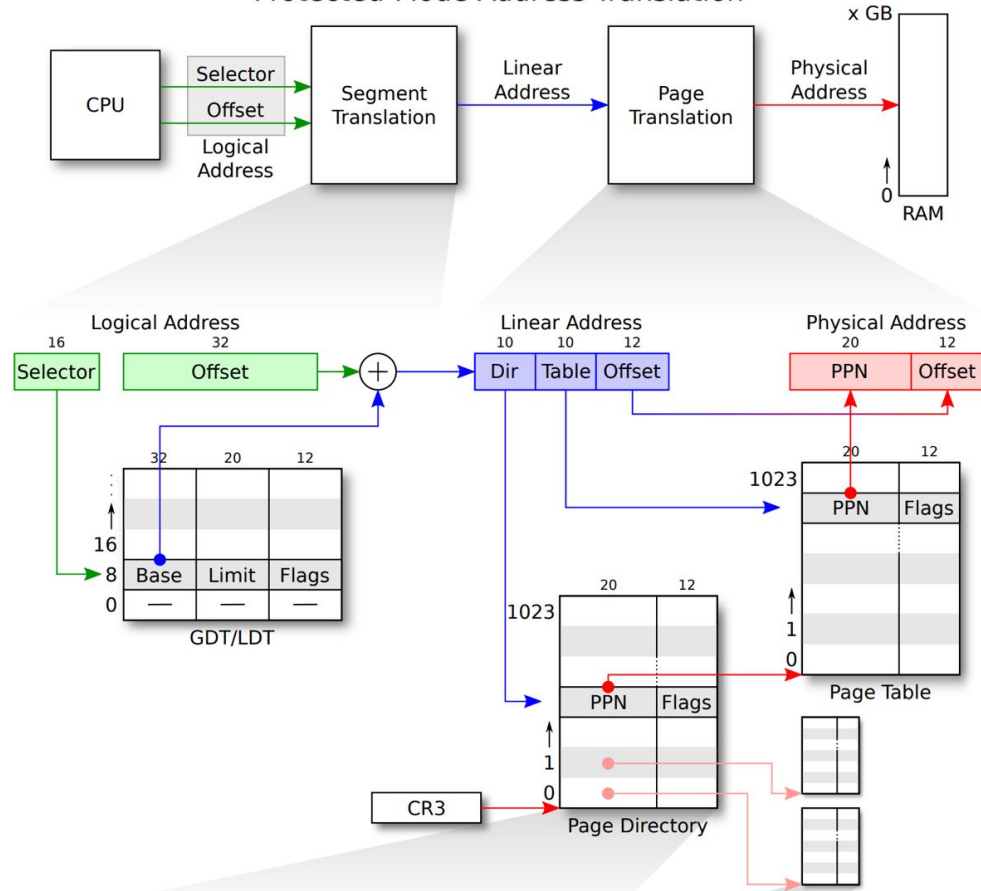


Page Table

- We access page table by virtual address
- Page size: 4 KB (12bits)
- Page number: 20 bits
- What is the page number and offset of
 - 0x8048000
 - 0xb7ff3100



Protected-Mode Address Translation



Caching!

- Cache the frequently used page table entries:
 - Exploit locality
 - Translation Lookaside Buffer (TLB)
 - Cache for translation entries.

Translation Lookaside Buffer (TLB)

- Stores VA-PA mappings and caches them!

VPN (Virt Page Number)	PPN (Phy Page Number)	Valid
0x12345	0x0	1
0x12346	0x5	1
0x12347	0xff	1
0x12348	0xfff	1


0x12345678 -> 0x678

0x12346678 -> 0x5678

0x12347678 -> 0xff678

0x12348678 -> 0xfff678

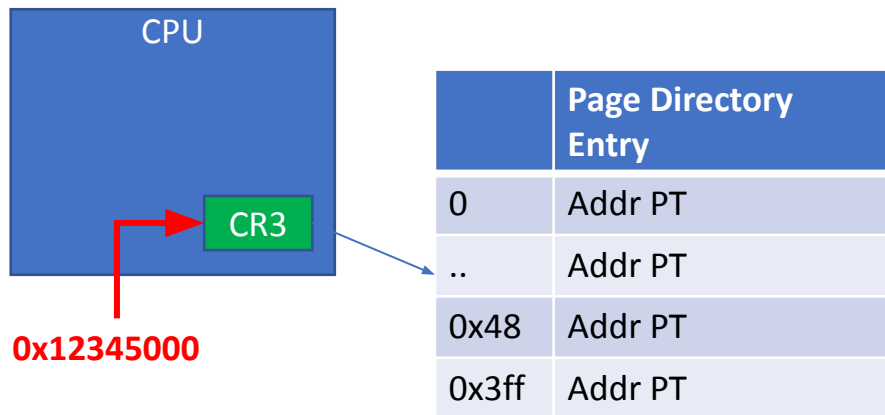
Without TLB



	Page Directory Entry
0	Addr PT
..	Addr PT
0x48	Addr PT
0x3ff	Addr PT

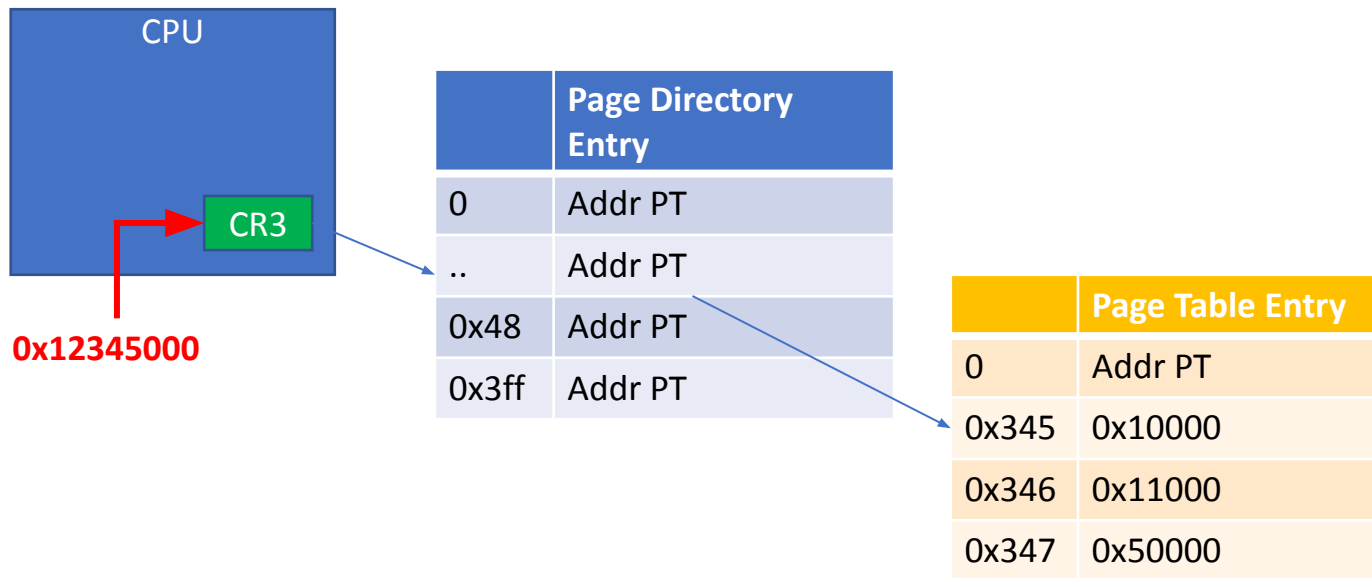
	Page Table Entry
0	Addr PT
0x345	0x10000
0x346	0x11000
0x347	0x50000

Without TLB

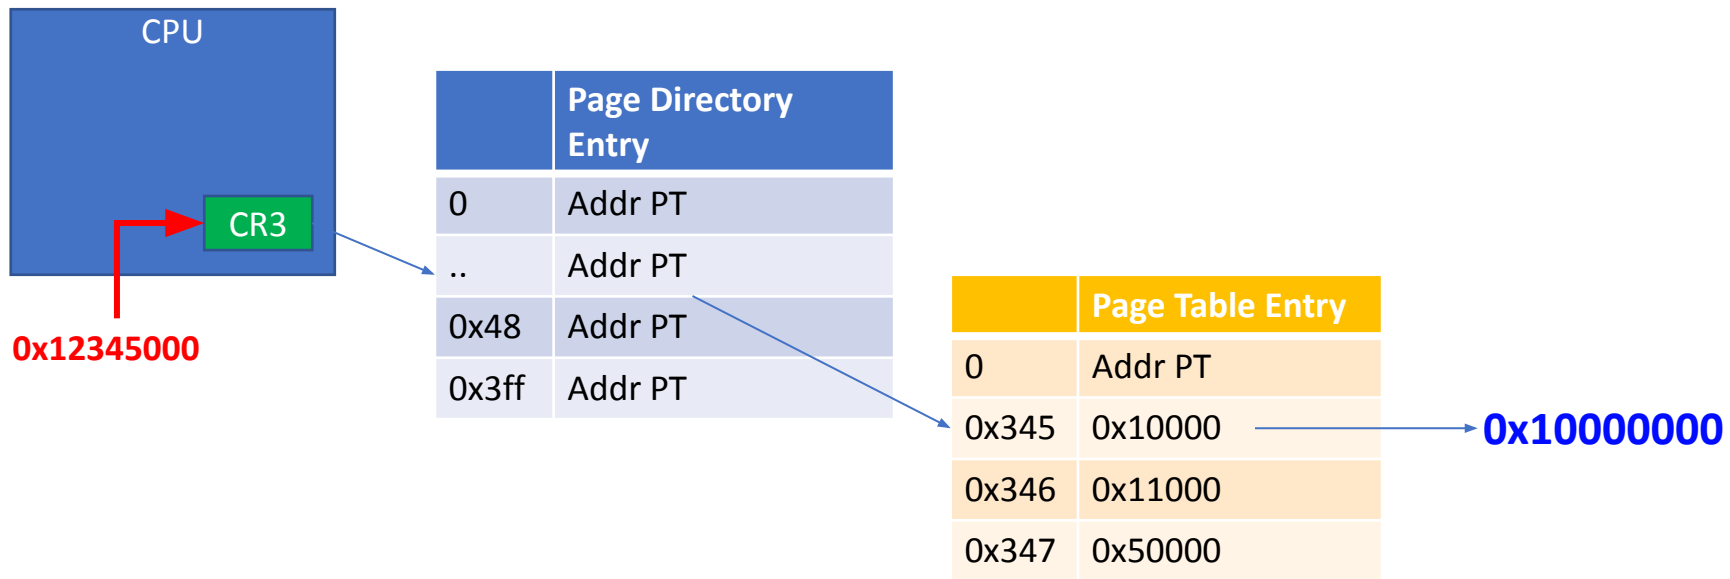


	Page Table Entry
0	Addr PT
0x345	0x10000
0x346	0x11000
0x347	0x50000

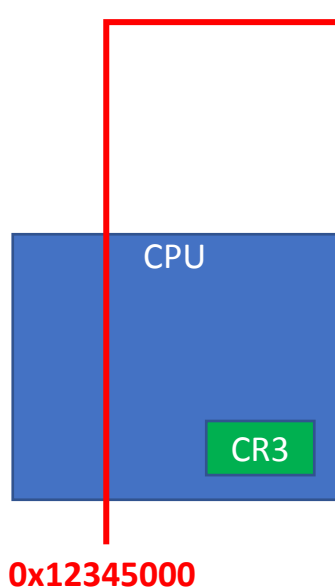
Without TLB



Without TLB



With TLB

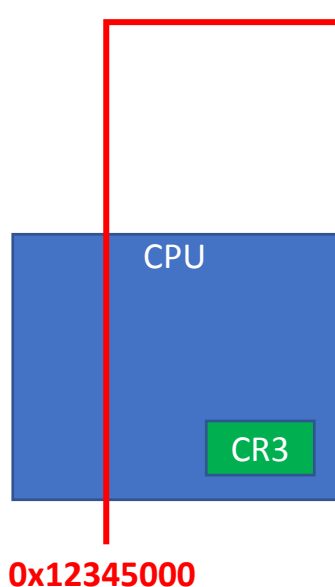


VPN (Virt Page Number)	PPN (Phy Page Number)	Valid
0x12345	0x10000	1
0x12346	0x5	1
0x12347	0xff	1
0x12348	0xfff	1

	Page Directory Entry
0	Addr PT
..	Addr PT
0x20	Addr PT
0x3ff	Addr PT

	Page Table Entry
0	Addr PT
0x48	0x10000
0x49	0x11000
0x4a	0x50000

With TLB



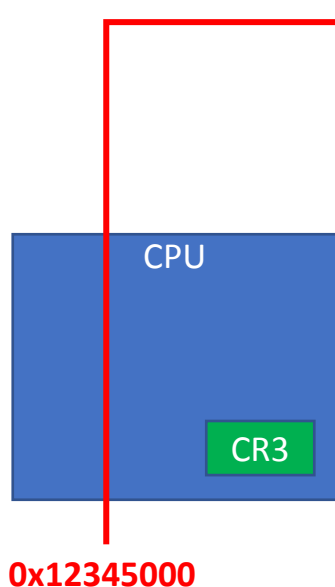
VPN (Virt Page Number)	PPN (Phy Page Number)	Valid
0x12345	0x10000	1
0x12346	0x5	1
0x12347	0xff	1
0x12348	0xfff	1

	Page Directory Entry
0	Addr PT
..	Addr PT
0x20	Addr PT
0x3ff	Addr PT

	Page Table Entry
0	Addr PT
0x48	0x10000
0x49	0x11000
0x4a	0x50000

0x10000000

With TLB



VPN (Virt Page Number)	PPN (Phy Page Number)	Valid
0x12345	0x10000	1
0x12346	0x5	1
0x12347	0xff	1
0x12348	0xfff	1

	Page Directory Entry
0	Addr PT
..	Addr PT
0x20	Addr PT
0x3ff	Addr PT

	Page Table Entry
0	Addr PT
0x48	0x10000
0x49	0x11000
0x4a	0x50000

0x10000000


No page table access..

TLB Performance

- TLB hit requires 4 cycles, 1ns!
- Page table walk requires 2 memory access for translation
 - Uncached: 9 cycles + (42 cycles + 51ns) * 2
 - [TLB miss] [RAM latency]
 $2\text{ns} + (10\text{ns} + 51\text{ns}) * 2 = 124\text{ns}$ (**124 times slower...**)
 - Cached: $9 + 4 * 2 = 17$ cycles if all blocks cached in L1 (4 ns, **4 times slower!**)
 - Data TLB L1: 64 items. 4-way. Miss penalty = 9 cycles. Parallel miss: 1 cycle per access
 - Data TLB L2 (STLB): 1536 items. 12-way. Miss penalty = 17 ? cycles. Parallel miss: 14 cycles per access
 - PDE cache = ? items. Miss penalty = ? cycles.
 - L1 Data Cache Latency = 4 cycles for simple access via pointer
 - L1 Data Cache Latency = 5 cycles for access with complex address calculation (`size_t n, *p; n = p[n];`).
 - L2 Cache Latency = 12 cycles
 - L3 Cache Latency = 42 cycles (core 0) (i7-6700 Skylake 4.0 GHz)
 - L3 Cache Latency = 38 cycles (i7-7700K 4 GHz, Kaby Lake)
 - RAM Latency = 42 cycles + 51 ns (i7-6700 Skylake)


Page Directory / Table

- In x86 (32-bit), CPU uses 2-level page table
- 10-bit directory index
- 10-bit page table index
- 12-bit offset
- **2-level paging**



Size of Page Directory!

- Page Size = 4 KB



$$4096 / 4 = 1024 \text{ entries}$$

$$1024 == 2^{10}$$


10-bit index for PD

Each entry is 4-byte (32 bits)

Size of Page Table!

- Page Size = 4 KB

One page,
4KB



Each entry is 4-byte (32 bits)

$$4096 / 4 = 1024 \text{ entries}$$

$$1024 == 2^{10}$$


10-bit index for PT

Permission Flags

- PTE_P (PRESENT)
 - 0: invalid entry
 - 1: valid entry
- PTE_W (WRITABLE)
 - 0: read only
 - 1: writable

- PTE_U (USER)
 - 0: kernel (only ring 0 can access)
 - 1: user (accessible by ring 3)

	Page Table Entry	
0	Addr PT	
0x48	0x10000 << 12 PTE_U PTE_W	Invalid
0x49	0x11000 << 12 PTE_P PTE_W	Kernel, writable
0x4a	0x50000 << 12 PTE_P PTE_U	User, read-only



Cannot have permissions such as ...

- Kernel: RW, User: R
 - PTE_P | PTE_W | PTE_U -> User RW...
 - PTE_P | PTE_W -> User --
- Kernel: R, User: RW
 - PTE_P | PTE_U | PTE_W -> Kernel RW...
 - PTE_P | PTE_U -> User R...
- Kernel: --, User: RW
 - PTE_P | PTE_U | PTE_W -> Kernel RW...

Struct PageInfo in JOS

- A **one-to-one** mapping from a **struct PageInfo** to a physical page
 - An 8 byte struct per each physical memory page
 - If we support 128MB memory, then we will create
 - Total number of physical pages: $128 * 1048576 / 4096 = 32768$
 - Total size = **32768** * **8** = 262,144 = 256KB
- A linked-list for managing free physical pages
 - Starting from `page_free_list->pp_link...`
- `pp_ref`
 - Count references
 - Non-zero – in-use
 - Zero – free

```
struct PageInfo {
    // Next page on the free list.
    struct PageInfo *pp_link;

    // pp_ref is the count of pointers (usually in page table entries)
    // to this page, for pages allocated using page_alloc.
    // Pages allocated at boot time using pmap.c's
    // boot_alloc do not have valid reference count fields.

    uint16_t pp_ref;
};
```

Users, Programs, Processes

- Users have accounts on the system
- Users launch programs
 - Can many users launch the same program?
 - Can one user launch many instances of the same program?

□ A process is an “instance” of a program

Program vs. Process

```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Program

```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Process

Code

Data

heap

stack


main

foo

registers

PC

Process State Transition



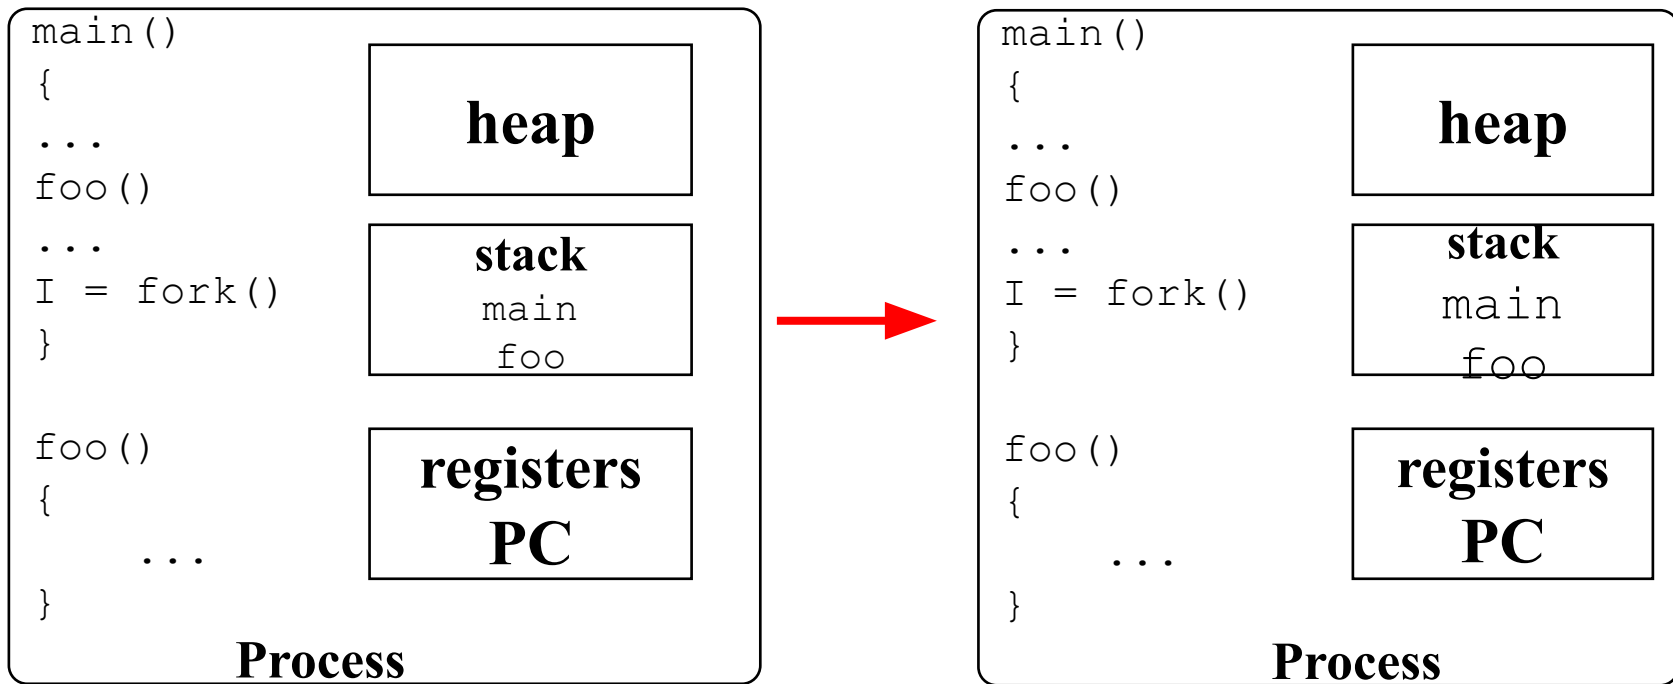
OS Process API

- 4 system calls related to process creation/termination:
 - Process Creation:
 - fork/clone – create a copy of this process
 - exec – replace this process with this program
 - Wait for completion:
 - wait – wait for child process to finish
 - Terminate a process:
 - kill - send a signal (to terminate) a process

fork


fork causes OS creates a copy of the calling process:

- Why?
- How can we disambiguate between new process and the calling process?




exec

Replaces current process with the content from new program.



wait

What happens when the parent process dies? what happens to child process?




Interrupts

- Hardware Interrupts
- Software Interrupts

Interrupts classification

Interrupts



Handling Interrupts

- Interrupts are numbered
- We need to define “what to do” (i.e., code to run) when an interrupt with corresponding number occurs

Handling Interrupts

- Setting an Interrupt Descriptor Table (IDT)

Interrupt Number	Code address
0 (Divide error)	t_divide
1 (Debug)	t_debug
2 (NMI, Non-maskable Interrupt)	t_nmi
3 (Breakpoint)	t_brkpt
4 (Overflow)	t_oflow
...	
8 (Double Fault)	t_dblflt
...	
14 (Page Fault)	t_pgflt
...	...
0x30 (syscall in JOS)	t_syscall

```
TRAPHANDLER_NOEC(t_divide, T_DIVIDE);    // 0
TRAPHANDLER_NOEC(t_debug, T_DEBUG);      // 1
TRAPHANDLER_NOEC(t_nmi, T_NMI);          // 2
TRAPHANDLER_NOEC(t_brkpt, T_BRKPT);      // 3
TRAPHANDLER_NOEC(t_oflow, T_OFLOW);      // 4
TRAPHANDLER_NOEC(t_bound, T_BOUND);      // 5
TRAPHANDLER_NOEC(t_illop, T_ILLOP);      // 6
TRAPHANDLER_NOEC(t_device, T_DEVICE);    // 7

TRAPHANDLER(t_dblflt, T_DBLFLT);          // 8

TRAPHANDLER(t_tss, T_TSS);                // 10
TRAPHANDLER(t_segnp, T_SEGNP);            // 11
TRAPHANDLER(t_stack, T_STACK);            // 12
TRAPHANDLER(t_gpflt, T_GPFLT);            // 13
TRAPHANDLER(t_pgflt, T_PGFLT);            // 14

TRAPHANDLER_NOEC(t_fperr, T_FPERR);        // 16

TRAPHANDLER(t_align, T_ALIGN);            // 17

TRAPHANDLER_NOEC(t_mchk, T_MCHK);          // 18
TRAPHANDLER_NOEC(t_simderr, T_SIMDERR);    // 19

TRAPHANDLER_NOEC(t_syscall, T_SYSCALL);    // 48, 0x30
```

JOS Interrupt Handling

- Setup the IDT at trap_init() in kern/trap.c

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```


JOS Interrupt Handling

- Setup the IDT at trap_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call _alltraps (in kern/trapentry.S)

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}

#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

JOS Interrupt Handling

- Setup the IDT at trap_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call _alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds
    pushl %es
    pushal
```

Build a Trapframe!

JOS Interrupt Handling

```
struct Trapframe {
    struct PushRegs tf_regs;
    uint16_t tf_es;
    uint16_t tf_padding1;
    uint16_t tf_ds;
    uint16_t tf_padding2;
    uint32_t tf_trapno;
    /* below here defined by x86 hardware */
    uint32_t tf_err;
    uintptr_t tf_eip;
    uint16_t tf_cs;
    uint16_t tf_padding3;
    uint32_t tf_eflags;
    /* below here only when crossing rings, such as from user to kernel */
    uintptr_t tf_esp;
    uint16_t tf_ss;
    uint16_t tf_padding4;
} __attribute__((packed));
```

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds
    pushl %es
    pushal
```

**Build a
Trapframe!**

JOS Interrupt Handling

- Setup the IDT at trap_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call _alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);
}
```

```
#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds    Build a
    pushl %es    Trapframe!
    pushal
```

```
void
trap(struct Trapframe *tf)
{
```

JOS Interrupt Handling

- Setup the IDT at trap_init() in kern/trap.c
- Interrupt arrives to CPU!
- Call interrupt handler in IDT
- Call _alltraps (in kern/trapentry.S)
- Call trap() in kern/trap.c
- Call trap_dispatch() in kern/trap.c

```
static void
trap_dispatch(struct Trapframe *tf)
{
    // Handle processor exceptions.
    // LAB 3: Your code here.
```

```
void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE], 0, GD_KT, t_divide, 0);
    SETGATE(idt[T_DEBUG], 0, GD_KT, t_debug, 0);


#define TRAPHANDLER_NOEC(name, num)
    .globl name;
    .type name, @function;
    .align 2;
    name:
    pushl $0;
    pushl $(num);
    jmp _alltraps
```

```
/*
 * Lab 3: Your code here for _alltraps
 */

_alltraps:
    pushl %ds    Build a
    pushl %es    Trapframe!
    pushal
```

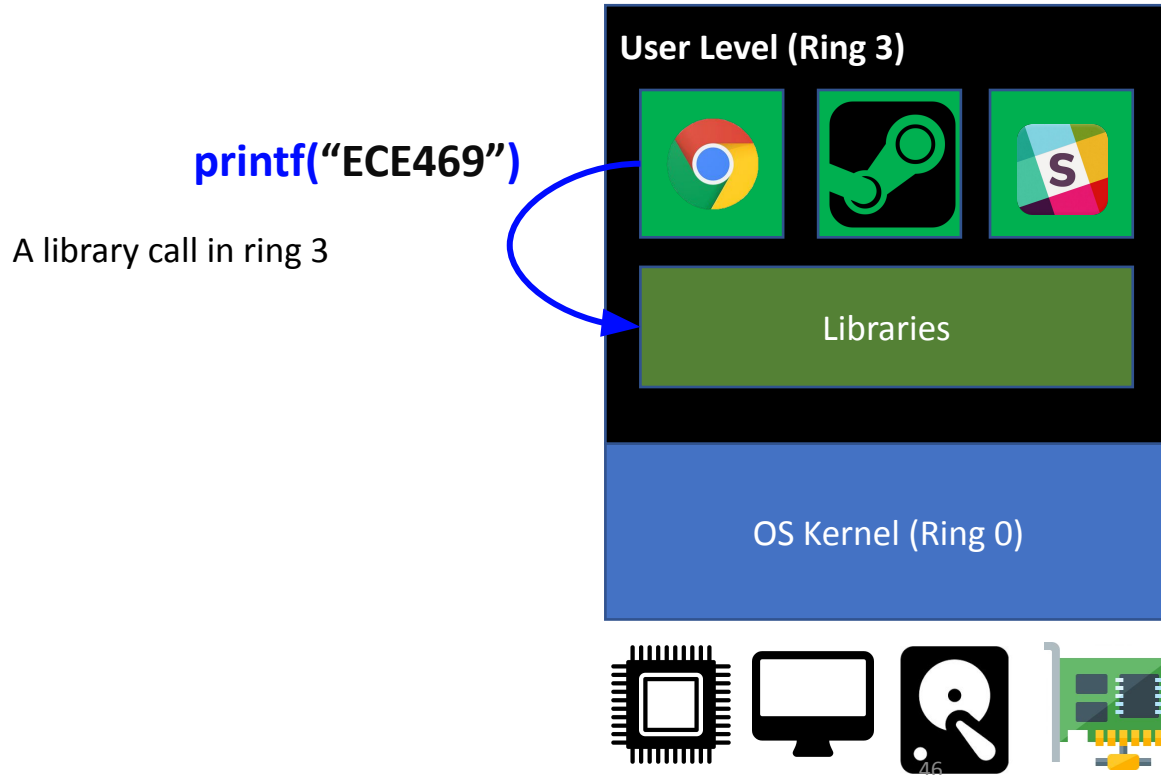
```
void
trap(struct Trapframe *tf)
{
```

Syscall: User/Kernel communication



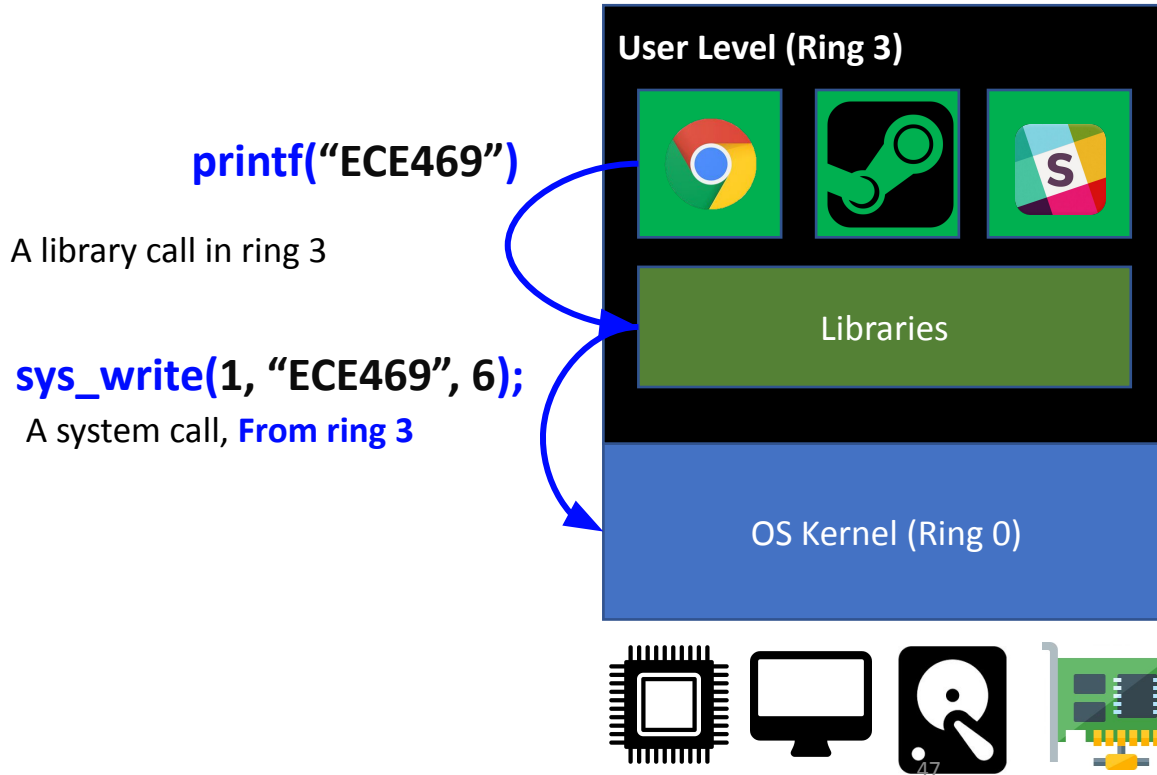
```
int main() {  
    printf("ECE469");  
}
```

Syscall: User/Kernel communication



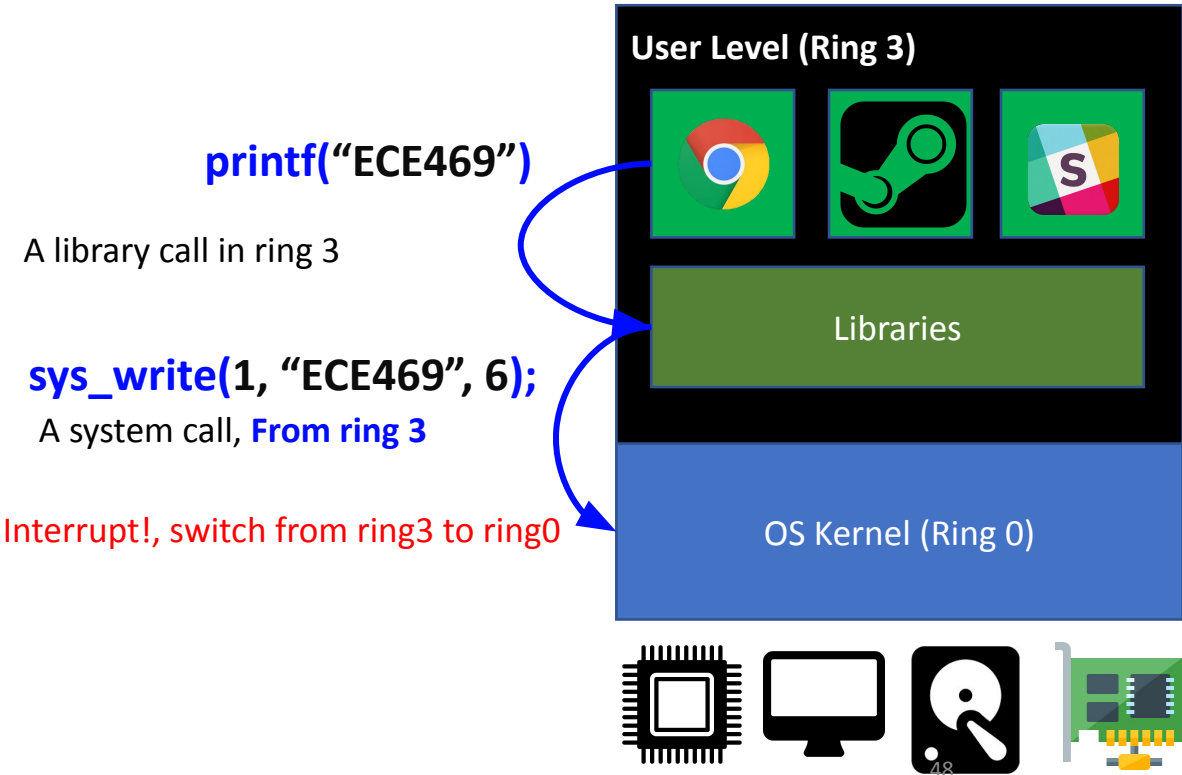
```
int main() {  
    printf("ECE469");  
}
```

Syscall: User/Kernel communication



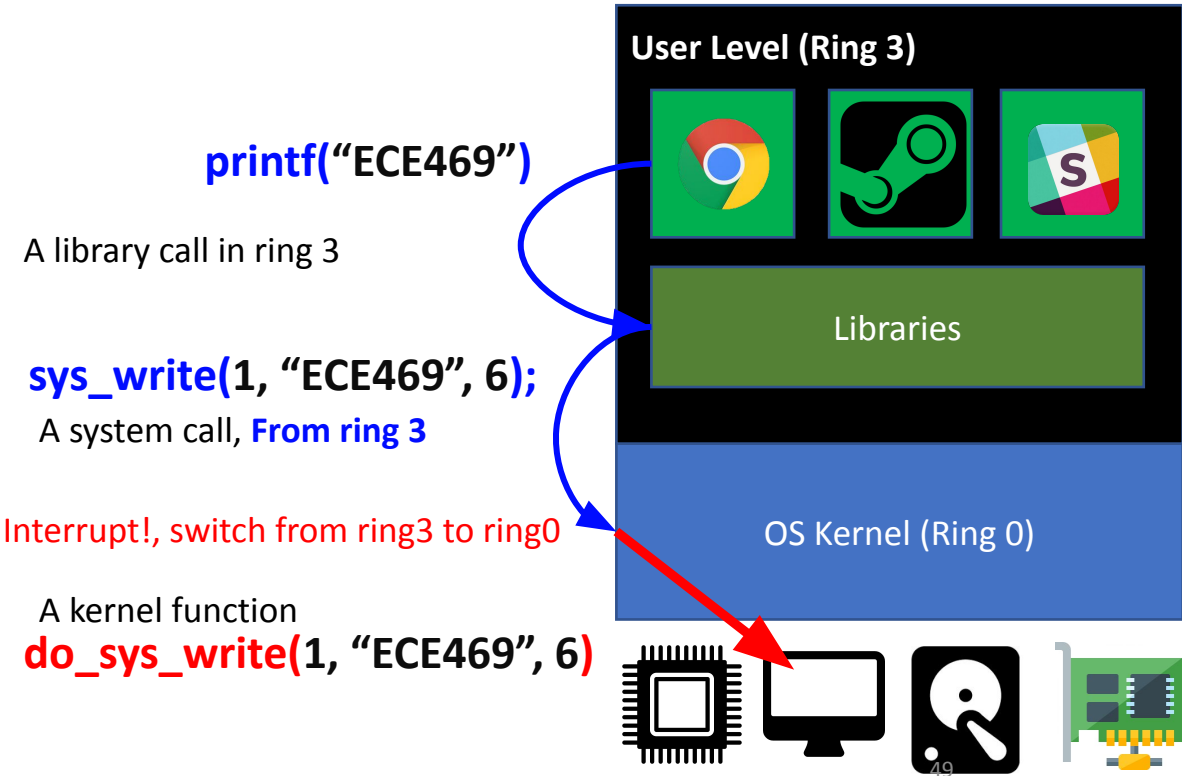
```
int main() {  
    printf("ECE469");  
}
```


Syscall: User/Kernel communication



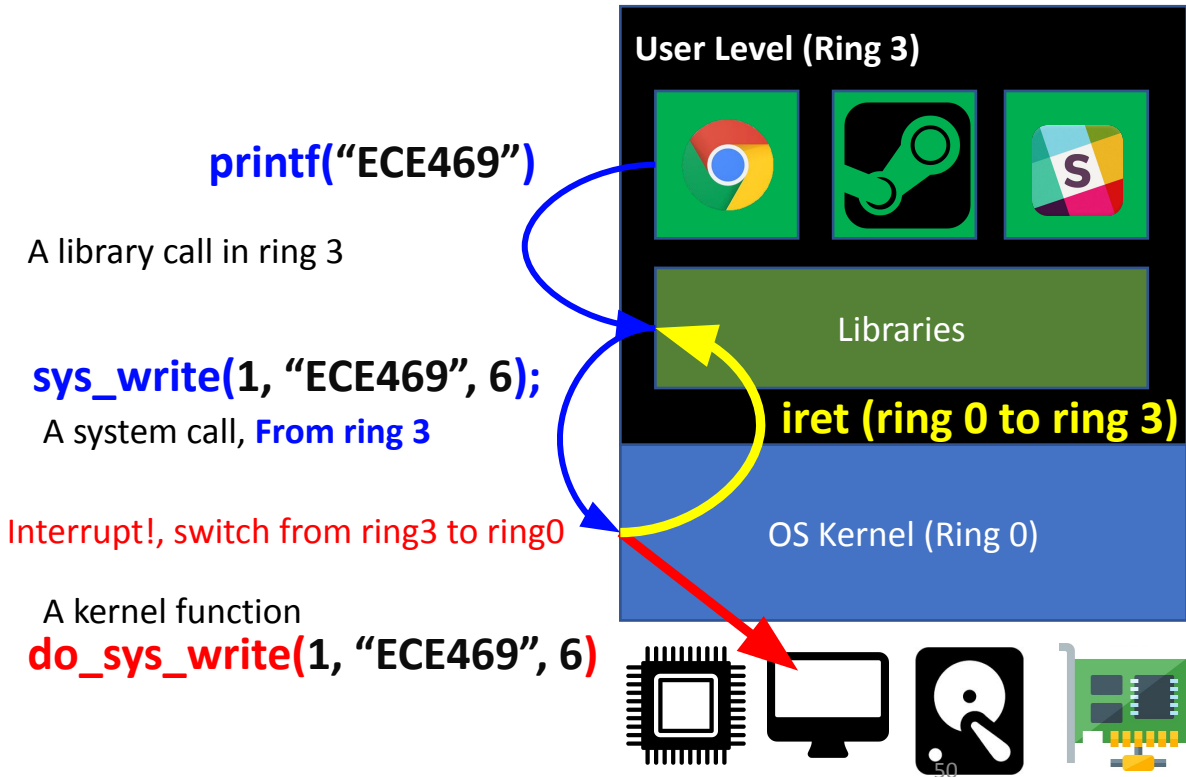
```
int main() {  
    printf("ECE469");  
}
```

Syscall: User/Kernel communication



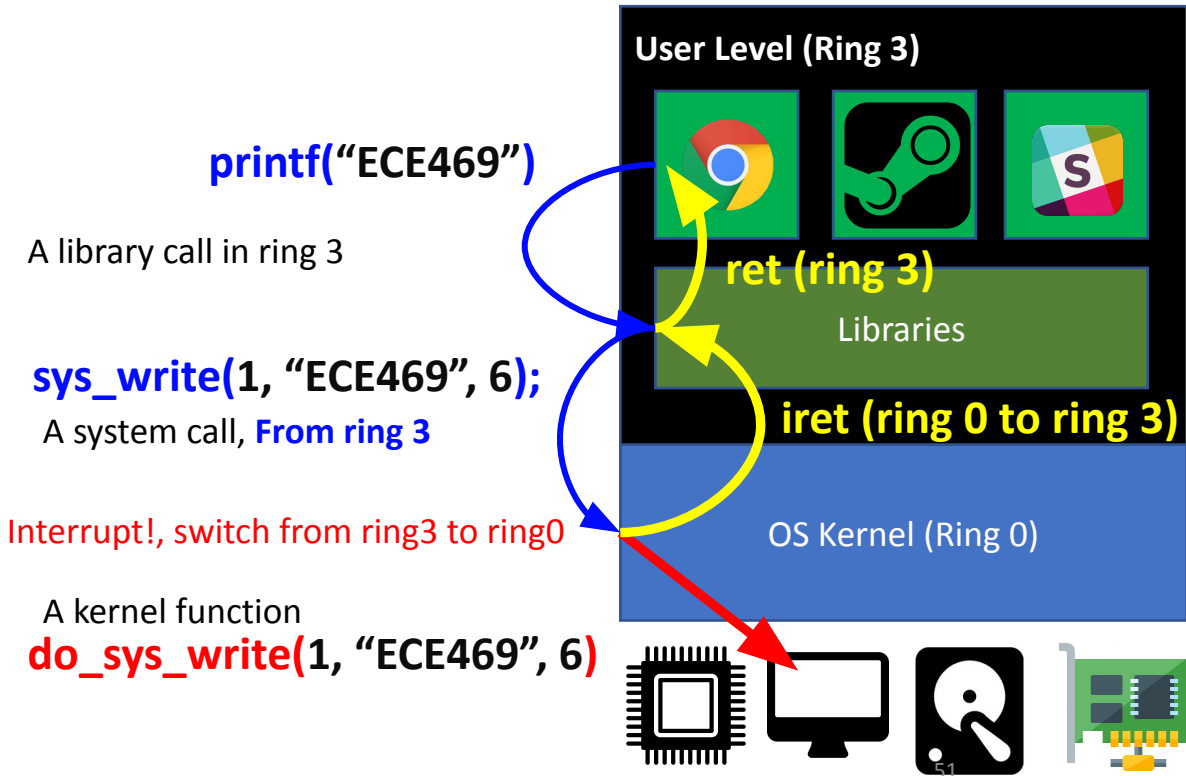
```
int main() {  
    printf("ECE469");  
}
```

Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```

Syscall: User/Kernel communication



```
int main() {  
    printf("ECE469");  
}
```


Invoking Syscalls

- Set all arguments in the registers
 - Order: edx ecx ebx edi esi
- int \$0x30 (in JOS)
 - Software interrupt 48
- int \$0x80 (in 32bit Linux)
 - Software interrupt 128

Invoking Syscalls in User mode

- User calls a function
 - `cprintf` -> calls `sys_cputs()`
- `sys_cputs()` at user code will call `syscall()` (`lib/syscall.c`)
 - This `syscall()` is at `lib/syscall.c`
 - Set args in the register and then
- `int $0x30`
- Now kernel execution starts...


Non-Preemptive Scheduling



Non-Preemptive Scheduling

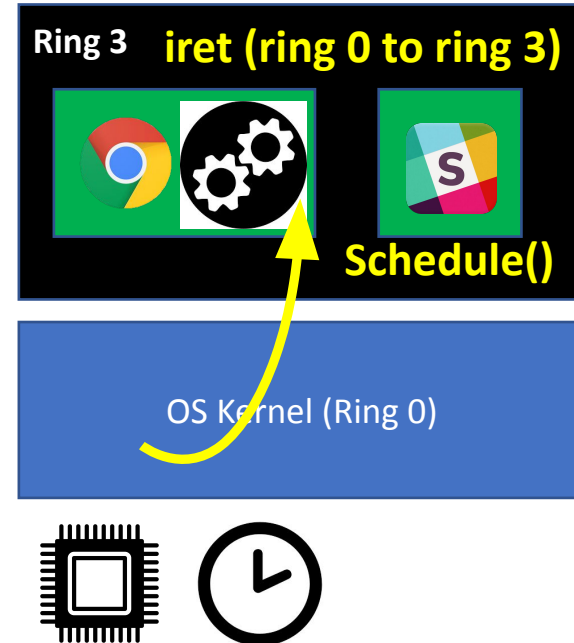
- Any issues?
- What if a process runs:

```
int main() {  
    while(1);  
}
```



Preemptive Scheduling


- Preemptive Multitasking (Lab 4)
- CPU generates an interrupt to force execution at kernel after some time quantum
 - E.g., 1000Hz, on each 1ms..
- Guaranteed execution in kernel
 - Let kernel mediate resource contention



Goals and Assumptions

- Goals (Performance metrics)
 - Minimize turnaround time
 - avg time to complete a job
 - $T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$
 - Maximize throughput
 - operations (jobs) per second
 - Minimize overhead of context switches: large quanta
 - Efficient utilization (CPU, memory, disk etc)
 - Short response time
 - $T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$
 - type on a keyboard
 - Small quanta
 - Fairness
 - fair, no starvation, no deadlock


Round Robin



- Each runs a time slice or quantum: Fair
- How do you choose time slice?
 - Overhead vs. response time
 - Overhead is typically about 1% or less
 - Quantum typically between 10 ~ 100 millisec


STCF

- Shortest time to completion first (shortest job first)
 - Non-preemptive



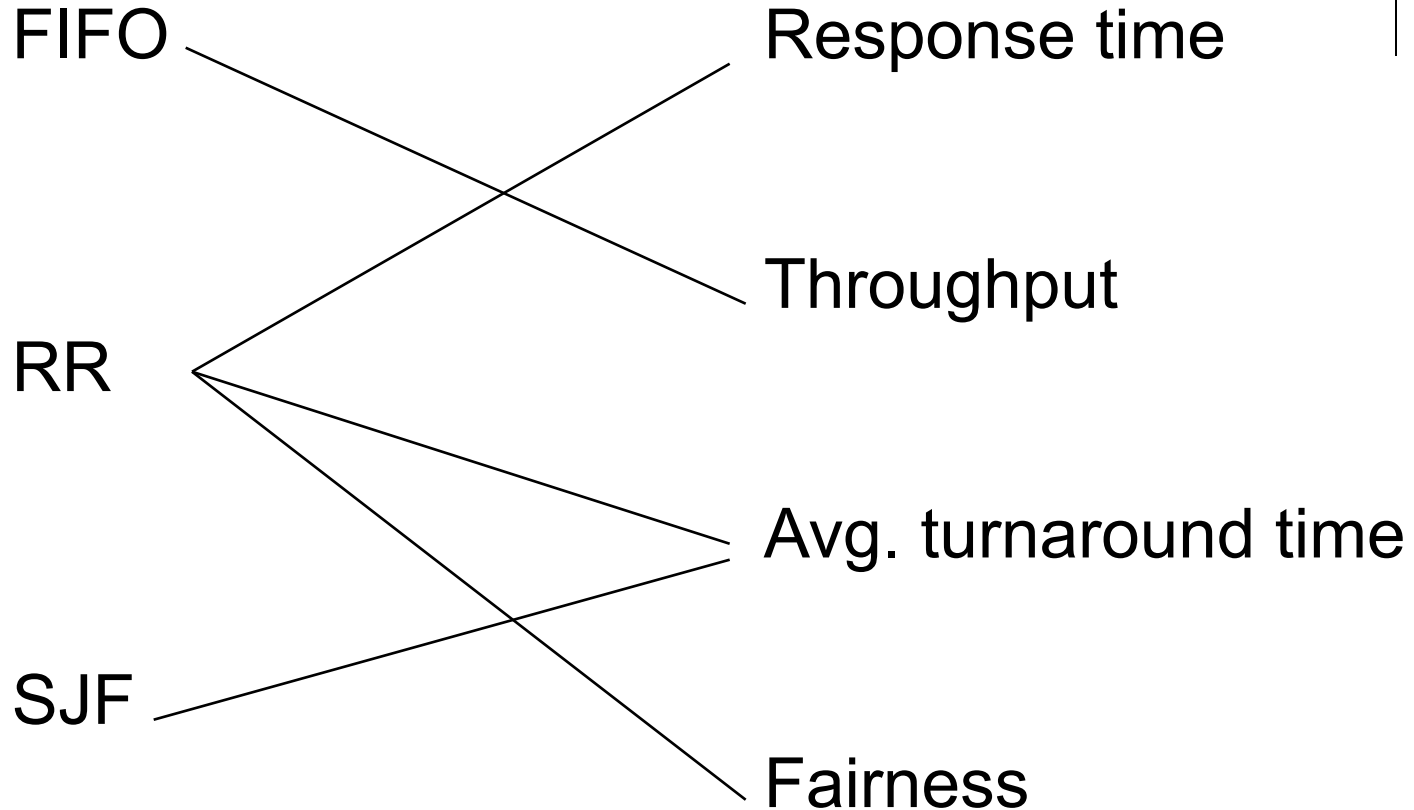
SRTCF

- Shortest remaining time to completion first
 - Preemptive



Any potential problems?
- Can cause **starvation!**

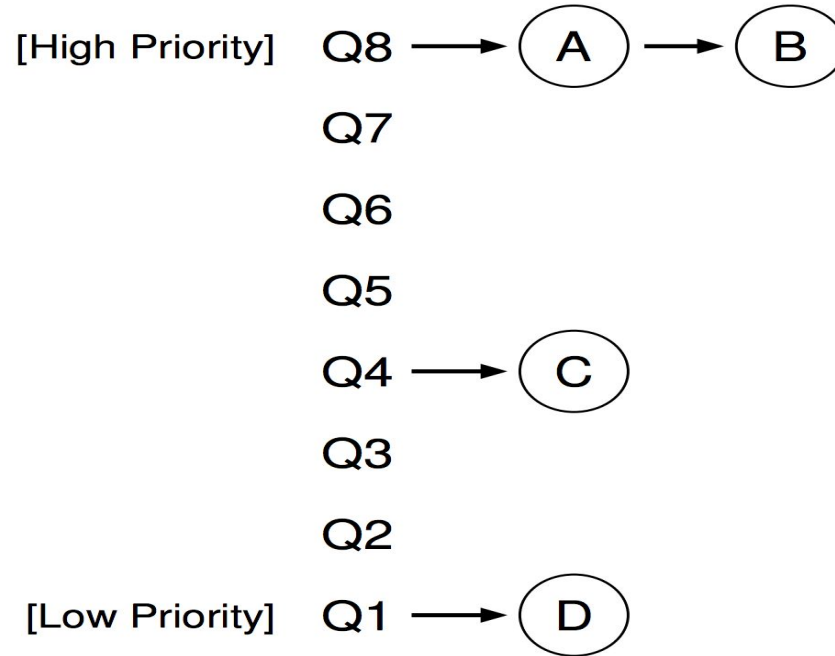
Scheduling Policies Advantages



Priority Scheduling

- To accommodate the spirits of SJF/RR/FIFO
- The method
 - Assign each process a *priority*
 - Run the process with highest priority in ready queue first
 - Use FIFO for processes with equal priority
 - Adjust priority dynamically
 - To deal with *all* issues: e.g. aging, I/O wait raises priority

Multiple Queue Scheduling



Multilevel Feedback Queue (MLFQ)

- Problem: how to change priority?
- Jobs start at highest priority queue
- Feedback
 - Priority Decreases: If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue).
 - Priority Unchanged: If a job gives up the CPU before the time slice is up, it stays at the same priority level.
 - Priority Increases: After a long time period, move all the jobs in the system to the topmost queue (aging)